

Multivariable Calculus

Jason Siefken



May 7, 2017



Contents

Introduction	iii
Licensing	v
Contributors	vii
1 Preliminaries	1
1.1 Mathematical Notation	1
1.2 Proof	6
2 Vectors	9
2.1 Vector Notation	10
2.2 Vector Arithmetic	12
2.3 Coordinates	14
2.4 Dot Products & Projections	19
2.5 The Cross Product	23
2.6 Lines and Planes	26
3 Parameterization	35
3.1 Speed and Velocity of a Parameterization	36
3.2 Arc-length Parameterization	41
3.3 Acceleration and Curvature	48
3.4 Line Integrals	53
3.5 Multi-dimensional Parameterizations	58
3.6 Coordinate Systems	62
4 Multi-variable Functions	69
4.1 Graphing in Many Variables	69
4.2 Rates of Change and the Directional Derivative	78
4.3 Tangent Planes and Differentiability	82
4.4 Multidimensional Limits and Continuity	87
4.5 Polynomial Approximation	90
4.6 Optimization	94
5 Integrals	101
5.1 Multivariable Integrals	101

5.2	The Volume Form	110
5.3	Surface Integrals	118
6	Vector Fields	123
6.1	Graphing Vector Fields	124
6.2	The Gradient	125
6.3	Flux and Divergence	129
6.4	Circulation and Curl	136
7	Differential Forms	141
7.1	Abstract Foundation for Differential Forms	141
7.2	1-forms and Line Integrals	146
7.3	2-forms and Surface Integrals	149
7.4	Higher-order Forms	154
7.5	The Exterior Derivative	156
7.6	The Hodge Dual	161
7.7	Stokes's Theorem	164
7.8	de Rham Cohomology	165
A	Proofs	169
B	Creative Commons License Fulltext	173

Introduction

Multivariable Calculus approaches the subject from a mathematical, but not overly technical, perspective. The key idea of calculus—chop things into little pieces and put them together again—is emphasized throughout.

Licensing

This book would not be possible without the long tradition of mathematical inquiry that came before. And like the ideas of mathematics, which are free for all to re-imagine, re-use, and re-purpose, so too is this book.

This book is licensed under the Creative Commons By-Attribution Share Alike 4.0 International license. This gives you permission to reuse, redistribute, and modify the contents of this book provided you attribute a derived work appropriately and that you license a derived work under the same terms. For the full text of the license, see Appendix B.

This book is a derived work of the Creative-Commons licensed *ISP Mathematics 281*¹ multivariable calculus textbook by Leonard Evans of Northwestern University. The problems from this textbook come from a variety of sources. Unannotated problems come from *ISP Mathematics 281*. Problems with a superscript OS come from *Calculus Volume 3*² from the Openstax project. The Openstax problems are licensed under a Creative Commons By-Attribution Non-Commercial Share Alike 4.0 license. If you require this textbook to be fully open-source, remove these problems.

¹ <https://github.com/siefkenj/ISPMathematics281>

² <https://openstax.org/details/calculus-volume-3>

Contributors

This book is a collaborative effort. The following people have contributed to its creation:

◦ Bailey Bjornstad ◦ Dylaan Cornish ◦ Julia Dierksheide ◦ Kathryn Riedel ◦ Maxwell
Sigal ◦ Kim Soram ◦ Ieva Stakvileviciute ◦ Paul Strauss ◦

Chapter 1

Preliminaries

1.1 Mathematical Notation

Mathematics is a sophisticated and precise language, and we best not adventure into calculus without learning some basic words.

The most basic mathematical word is that of a *set*. A set is an unordered collection of distinct objects. We won't try and pin it down more exactly than this—our intuition about collections of objects will suffice¹. We write a set with curly-braces { and } and list the objects inside. For instance

$$\{1, 2, 3\}.$$

This would be read aloud as “the set containing the elements 1, 2, and 3.” The symbol \in is used to specify that some object is an element of a set, and \notin is used to specify it is not. For example,

$$3 \in \{1, 2, 3\} \quad 4 \notin \{1, 2, 3\}.$$

Sets can contain mixtures of objects, including other sets. For example,

$$\{1, 2, a, \{-70, \infty\}, x\}$$

is a perfectly valid set.

It is tradition to use capital letters to name sets. So we might say $A = \{6, 7, 12\}$ or $X = \{7\}$. There is, however, a special set with a special name—the empty set. The *empty set* is the set containing no elements and is written \emptyset or $\{\}$. Note that $\{\emptyset\}$ is *not* the empty set. It is the set containing the empty set! It is also traditional to call elements of a set *points* regardless of whether you consider them “point-like” objects.

Operations on Sets

If the set A contains all the elements that the set B does, we call A a *superset* and B a *subset*. We'll give this a formal definition.

¹ When you pursue more rigorous math, you rely on definitions to get yourself out of philosophical jams. For instance, with our definition of set, consider “the set of all sets that don't contain themselves.” Such a set cannot exist! This is called *Russel's Paradox*, and shows that if we start talking about sets of sets, we may need more than intuition.

Definition 1.1.1 — Subset & Superset. The set B is a *subset* of the set A , written $B \subseteq A$, if for all $b \in B$ we also have $b \in A$. In this case, A is called a *superset* of B .^a

^a Some mathematicians use the symbol \subset instead of \subseteq .

Some simple examples are $\{1, 2, 3\} \subseteq \{1, 2, 3, 4\}$ and $\{1, 2, 3\} \subseteq \{1, 2, 3\}$. There's something funny about that last example, though. Those two sets are not only subsets/supersets of each other, they're *equal*. As surprising as it seems, we actually need to define what it means for two sets to be equal.

Definition 1.1.2 — Set Equality. The sets A and B are *equal*, written $A = B$, if $A \subseteq B$ and $B \subseteq A$.

Having a definition of equality to lean on will help us when we need to prove things about sets.

■ **Example 1.1** Let A be the set of numbers that can be expressed as $2n$ for some whole number n , and let B be the set of numbers that can be expressed as $m + 1$ where m is an odd whole number. We will show $A = B$.

First, let us show $A \subseteq B$. If $x \in A$ then $x = 2n$ for some whole number n . Therefore $x = 2n = 2(n - 1) + 1 + 1 = m + 1$ where $m = 2(n - 1) + 1$ is, by definition, an odd number. Therefore $x \in B$.

Now we will show $B \subseteq A$. Let $x \in B$. By definition, $x = m + 1$ for some odd m and so by the definition of oddness, $m = 2k + 1$ for some whole number k . Thus

$$\begin{aligned} x &= m + 1 = (2k + 1) + 1 = 2k + 2 \\ &= 2(k + 1) = 2n, \end{aligned}$$

where $n = k + 1$, and so $x \in A$. Since $A \subseteq B$ and $B \subseteq A$, by definition $A = B$. ■

Set-builder Notation

Specifying sets by listing all their elements can be a hassle, and if there are an infinite number of elements, it's impossible! Fortunately, *set-builder notation* solves these problems. If X is a set, we can define a subset

$$Y = \{a \in X : \text{some rule involving } a\},$$

which is read “ Y is the set of a in X such that some rule involving a is true.” If X is intuitive, we may omit it and simply write $Y = \{a : \text{some rule involving } a\}$ ². You may equivalently use “|” instead of “:”, writing $Y = \{a \mid \text{some rule involving } a\}$.

■ **Example 1.2** The set \mathbb{Z} is the set of integers (positive, negative, and zero whole numbers). To define E as the even integers, we could write

$$E = \{n \in \mathbb{Z} : n = 2k \text{ for some } k \in \mathbb{Z}\}.$$

² If you want to get technical, to make this notation unambiguous, you define a *universe of discourse*. That is, a set \mathcal{U} containing every object you might want to talk about. Then $\{a : \text{some rule involving } a\}$ is short for $\{a \in \mathcal{U} : \text{some rule involving } a\}$

To define P as the set of positive integers, we could write

$$P = \{n \in \mathbb{Z} : n > 0\}.$$

■

There are also some common operations we can do with two sets.

Definition 1.1.3 — Intersections & Unions. Let A and B be sets. Then the *intersection* of A and B , written $A \cap B$, is defined by

$$A \cap B = \{x : x \in A \text{ and } x \in B\}.$$

The *union* of A and B , written $A \cup B$, is defined by

$$A \cup B = \{x : x \in A \text{ or } x \in B\}.$$

For example, if $A = \{1, 2, 3\}$ and $B = \{-1, 0, 1, 2\}$, then $A \cap B = \{1, 2\}$ and $A \cup B = \{-1, 0, 1, 2, 3\}$. Set unions and intersections are *associative*, which means it doesn't matter how you apply parentheses to an expression involving just unions or just intersections. For example $(A \cup B) \cup C = A \cup (B \cup C)$, which means we can give an unambiguous meaning to an expression like $A \cup B \cup C$ (just put the parenthesis wherever you like). But watch out, $(A \cup B) \cap C$ means something different than $A \cup (B \cap C)$!

Definition 1.1.4 — Set Subtraction. For sets A and B , the *set-wise difference* between A and B , written $A \setminus B$, is the set

$$A \setminus B = \{x : x \in A \text{ and } x \notin B\}.$$

Definition 1.1.5 — Cardinality. For a set A , the *cardinality* of A , written $|A|$ is the number of elements in A . If A contains infinitely many elements, we write $|A| = \infty$.

Let's define some notation for common sets.

$$\emptyset = \{\}, \text{ the empty set}$$

$$\mathbb{N} = \{0, 1, 2, 3, \dots\} = \{\text{natural numbers}\}$$

$$\mathbb{Z} = \{\dots, -3, -2, -1, 0, 1, 2, 3, \dots\} = \{\text{integers}\}$$

$$\mathbb{Q} = \{\text{rational numbers}\}$$

$$\mathbb{R} = \{\text{real numbers}\}$$

$$\mathbb{R}^n = \{\text{vectors in } n\text{-dimensional Euclidean space}\}$$

Besides unions, there's another way to join sets together: *products*.

Definition 1.1.6 — Cartesian Product. Given two sets A and B , the *Cartesian product* (sometimes shortened to *product*) of the sets A and B is written $A \times B$ and defined to be

$$A \times B = \{(a, b) : a \in A \text{ and } b \in B\}.$$

The Cartesian product of two sets is the set of all ordered pairs of elements from those sets. For example,

$$\{1, 2\} \times \{1, 2, 3\} = \{(1, 1), (1, 2), (1, 3), (2, 1), (2, 2), (2, 3)\}.$$

You can repeat this operation more than once. $\mathbb{R} \times \mathbb{R} \times \mathbb{R}$ is the set of all triples of real numbers. Extending power notation notation, if you take the Cartesian product of a set with itself some number of times, you can represent it with an exponent. Thus, $\mathbb{R} \times \mathbb{R} \times \mathbb{R}$ can be written as \mathbb{R}^3 , which is a set we've seen before³.

Functions

You're probably used to seeing functions like $f(x) = x^2$, but it's worth reviewing some of the concepts and terminology associated with functions.

Definition 1.1.7 — Function. A *function* with *domain* the set A and *co-domain* the set B is an object that associates every point in the set A with *exactly one* point in the set B .

If a function f has domain A and co-domain⁴ B , we notate this by writing $f : A \rightarrow B$. If we want to further specify what the function f actually is, we need to express how f associates each point in A to a point in B . This can be done with an equation. For example, we could define the function $f : \mathbb{R} \rightarrow \mathbb{R}$ by

$$f(x) = 2x,$$

which says that each real number gets associated to its double. We can notate the same thing using a special type of arrow: " \mapsto ". Now we might write

$$f : \mathbb{R} \rightarrow \mathbb{R} \text{ where } x \mapsto 2x,$$

which is read " f is a function from \mathbb{R} to \mathbb{R} where $x \in \mathbb{R}$ gets mapped to $2x$."

Note that every point in the co-domain of a function doesn't need to get mapped to. For example $g : \mathbb{R} \rightarrow \mathbb{R}$ given by $g(x) = x^2$ outputs only non-negative numbers, but it is still valid to specify \mathbb{R} as the co-domain. However, if we wanted to make a point of it, we are perfectly justified in writing $g : \mathbb{R} \rightarrow [0, \infty)$ when defining g .

Many common math operations give rise to functions. For example, $f(x) = \sqrt{x}$ is the familiar square root function. Sometimes, when we wish to talk about a function for which notation already exists, we will put a " \cdot " where we would normally put a variable. Thus, we might say, " $\sqrt{\cdot}$ is the square root function."⁵

Definition 1.1.8 — Range. The *range* of a function $f : A \rightarrow B$ is the set of all outputs of f . That is

$$\text{range } f = \{y \in B : y = f(x) \text{ for some } x \in A\}.$$

Definition 1.1.9 — Image. Let $f : A \rightarrow B$ be a function. The *image* of a set $X \subseteq A$, written $f(X)$ is defined by

$$f(X) = \{y \in B : y = f(x) \text{ for some } x \in X\}.$$

We see that if $f : A \rightarrow B$, $\text{range } f = f(A)$. In words, the range of f is the image of its domain. This language will become useful when we think of functions as transformations that

³ If you're scratching your head saying, "I thought \mathbb{R}^3 was vectors in 3-dimensional space. How do we know that's the same thing as triples of real numbers?" your mind is keen. This is a theorem of linear algebra.

⁴ Some people use the word *range* interchangeably with co-domain.

⁵ Since \sqrt{x} is "the square root of the quantity x ," it is technically a quantity and not a function. This is why we write $\sqrt{\cdot}$ instead of x when we want to refer to the square root *function*.

move or bend space. If $f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ is a function that warps the Cartesian plane, then the image of X under f could be visualized by painting X on the Cartesian plane, warping the whole plane, and then looking at the resulting, painted shape.

Closely related to images, we have the idea of *restriction*. Suppose $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ is defined by $f(x, y) = xy$, but we were only really interested in f on the unit circle, C . In this case, we might say f attains a maximum on C , or f *restricted to* C attains a maximum, even though f itself is unbounded. This idea comes up often enough to deserve its own notation.

Definition 1.1.10 — Restriction. If $f : A \rightarrow B$ and $X \subseteq A$, the *restriction* of f to X is written $f|_X$ and is defined to be the function $g : X \rightarrow B$ where $x \mapsto f(x)$.

The last important function-related ideas for us are function composition and inverses. Given two functions $f : A \rightarrow B$ and $g : B \rightarrow C$, we can *compose* g and f to get a new function.

Definition 1.1.11 — Composition. Given two functions $f : A \rightarrow B$ and $g : B \rightarrow C$, the *composition* of g and f , written $g \circ f$, is the function $h : A \rightarrow C$ where $x \mapsto g(f(x))$.

Note that the composition $g \circ f$ has the domain of f and the co-domain of g . When a point is fed into $g \circ f$, it moves from $A \rightarrow B \rightarrow C$. The composition $g \circ f$ only makes sense because the outputs of f are allowed as inputs to g . If we wrote $f \circ g$, it wouldn't mean much, because g outputs points in C and f has no idea what to do with points in C .⁶

Inverses relate to composition and the *identity function*, the function that does nothing to its inputs.

Definition 1.1.12 — Identity Function. The *identity function* $\text{id} : A \rightarrow A$ is defined by the relation

$$\text{id}(x) = x$$

for all $x \in A$.

Notice that for every set, that set is the domain of an identity function. Since the domain and co-domain of a function are part of its definition, we don't want to confuse them. After all, $f : \{0, 1\} \rightarrow \{0, 1\}$ given by $f(x) = x^2$ is a different function from $f : \mathbb{R} \rightarrow \mathbb{R}$ given by $f(x) = x^2$. For the special case of the identity function, we sometimes write the domain of the function as a subscript. That is, for $\text{id} : A \rightarrow A$ we'd write id_A so it doesn't get confused with $\text{id} : B \rightarrow B$, which we'd write id_B .

Definition 1.1.13 — Inverse Function. Let $f : A \rightarrow B$ be a function. If there exists a function $g : B \rightarrow A$ such that

$$f \circ g = \text{id}_B \quad \text{and} \quad g \circ f = \text{id}_A,$$

we say f is *invertible* and we call g the *inverse* of f . If f is invertible, we notate its inverse by f^{-1} .

Inverses can be tricky some times. For example, consider $f(x) = x^2$ and $g(x) = \sqrt{x}$. Here $g \circ f(x) = \sqrt{x^2} = |x|$ and $f \circ g(x) = \sqrt{x^2} = x$. What's the deal? Well, it's all about domains.

⁶ It seems a little backward to write $f : A \rightarrow B$, $g : B \rightarrow C$ and then write $g \circ f$ instead of $f \circ g$. You can thank Euler for that. He decided to write functions with their input on the right instead of the left. If we wrote functions backwards, like $((x)f)g$ for “ g of f of x ,” they we could just *follow the arrows* and life would be simpler.

$f : \mathbb{R} \rightarrow [0, \infty)$ and $g : [0, \infty) \rightarrow [0, \infty)$. So, the domain of $g \circ f$ is \mathbb{R} and the domain of $f \circ g$ is $[0, \infty)$. The domains are different, and indeed f is not invertible. However, g is invertible, and $g^{-1} = f|_{[0, \infty)}$. If we only input non-negative numbers into f , then f exactly undoes what g did. This subtle domain trickery can cause us a lot of headaches if we're not used to thinking carefully, and many of our favorite functions that we're used to calling "inverse functions" are actually only inverses when paired with specific domains.

1.2 Proof

Mathematics has the highest standard of proof of any field. In the Platonic ideal of mathematics, we start from some basic assumptions, called *axioms*, that we have all agreed upon. Then from those axioms, using the rules of logic, we deduce *theorems*. Every single mathematical statement we make can be traced back from theorem to theorem and eventually to our initial axioms.

This is contrary to other disciplines, like physics. In physics, based on observation, we construct *laws*. Laws in physics are like axioms in mathematics, but they have an important difference—they can be disproven by observation. A mathematical axiom can never be disproven. One can certainly argue that an axiom is not *useful* or not *interesting*, but you cannot say it's *wrong*⁷. Of course, as human practitioners, we may misuse logic and be wrong ourselves, but that is no fault of the axioms.

But now, let's deviate from philosophical perfection and visit reality. In reality, *mathematics is a human pursuit to understand relationships between ideas and their consequences*. The key there is that *humans* do mathematics to *understand* relationships. If a theorem in math can ultimately be reduced to logical statements about axioms, but the argument is 100000 steps long, it doesn't help a human understand why something is true. Instead, a shorter argument that skips over some steps is more useful to us. And, indeed, most of our mathematics to date skips over some steps⁸.

We call a correct mathematical argument a *proof*. A proof starts from a set of assumptions, and following the rules of logic, arrives at a conclusion. Strictly speaking, a proof doesn't need to make sense or show motivation, applications, or examples. It just has to be a sequence of correct logical steps. However, for us, as humans studying mathematics, we prove things for two reasons: to understand why things are true and to avoid making mistakes.

Reconciling these two goals can be very hard for a novice mathematician. If you include *all* the steps, it won't help with understanding, but if you don't include enough steps, the argument may not be convincing and might contain mistakes. Even professionals struggle to balance these competing goals, and how you balance those goals depends on your audience—if

⁷ There are multiple ways to axiomatize geometry. In Euclidean geometry every pair of lines either coincides, intersects in exactly one place, or does not intersect. In spherical geometry, every pair of lines either coincides or intersects in exactly two places. Euclidean geometry is useful when your space looks flat. Spherical geometry is useful when your space is the surface of a sphere (like the Earth). Is one of these more *right* than the other? They're certainly contradictory.

⁸ There are some projects to prove all of mathematics directly from the axioms using computer assistance. They've made progress, but there are still theorems in calculus that have not been reduced to the axioms. We believe that they *could be* reduced to the axioms, but no one has taken the time to do so.

you're trying to convince your math professor of something your proof will need to have more detail than if you were trying to convince your friend (mathematicians are very skeptical!).

Enough talk, let's go through a 2000-year-old example of a proof.

Theorem 1.2.1 There is no rational number p/q such that $(p/q)^2 = 2$.

Proof. If p/q is a rational number, it can be expressed in lowest terms. Suppose p/q is in lowest terms and $(p/q)^2 = 2$. Then $p^2 = 2q^2$ and so p^2 is even. Since p^2 is even, it must be that p is even, and so by definition, $p = 2m$ for some integer m . Now,

$$\frac{p^2}{q^2} = \frac{(2m)^2}{q^2} = \frac{4m^2}{q^2} = 2,$$

with the last equality following by assumption. Multiplying both sides by q^2 and dividing by 2 we arrive at the equation

$$2m^2 = q^2,$$

and so q^2 is even which means q is even. By definition, this means $q = 2n$ for some integer n . But now,

$$\frac{p}{q} = \frac{2m}{2n}$$

is not in lowest terms! This is a contradiction and so it cannot be that $(p/q)^2 = 2$. ■

This is nearly identical to the argument the ancient Greeks gave. It's elegant, beautiful, and convincing. But, if we look closer, it does skip some steps. For example, it relies on the fact that there is such a thing as *lowest terms*. This is something that would need to be proven—a priori, the conclusion of the proof could be that the assumption that p/q could be in lowest terms is false.

You will not, overnight, become a master at understanding what steps you can leave out and what steps you must show. However, with feedback, you'll get better. For a detailed guide on writing good proofs, please see Appendix A.

Chapter 2

Vectors

A *vector* is a quantity which is characterized by a *magnitude* and a *direction*. Many quantities are best described by vectors rather than numbers. For example, when driving a car, it may be sufficient to know your speed, which can be described by a single number, but the motion of an airplane must be described by a vector quantity—velocity—which takes into account its direction as well as its speed.

Ordinary numerical quantities are called *scalars* when we want to emphasize that they are not vectors.

Whereas numbers allow us to specify relationships between single quantities (put in twice as much flour as sugar), vectors will allow us to specify relationships between geometric objects in space¹. If we have two points, $P = (1, 1)$ and $Q = (3, 2)$, we specify the *displacement* from P to Q as a vector.



We denote the displacement vector from P to Q by \overrightarrow{PQ} . The magnitude of \overrightarrow{PQ} is given by the Pythagorean theorem to be $\sqrt{5}$ and its direction is specified by the directed line segment from P to Q .

¹ Though in this book we will treat vectors as intertwined with Euclidean space, they are much more general. For instance, someone's internet browsing habits could be described by a vector—the topics they find most interesting might be the “direction” and the amount of time they browse might be the “magnitude.”

2.1 Vector Notation

There are many ways to represent vector quantities in writing. If we have two points, P and Q , we write \overrightarrow{PQ} to represent the vector from P to Q . Absent of points, bold-faced letters or a letter with an arrow over it are the most common typographical representations of vectors. For example, \vec{a} or \mathbf{a} may both be used to represent the vector quantity named “ a .” In this book we will use \vec{a} to represent a vector. The notation $\|\vec{a}\|$ represents the magnitude of the vector \vec{a} , which is sometimes called the *norm* of \vec{a} .

Graphically, vectors are represented as directed line segments (a line segment with an arrow at one end). The endpoints of the segment are called the *initial point* (the base) and the *terminal point* (the tip) of the vector.



Let $A = (1, 1)$, $B = (3, 2)$, $X = (1, 0)$, and $Y = (3, 1)$ and consider the vectors $\vec{a} = \overrightarrow{AB}$ and $\vec{x} = \overrightarrow{XY}$. Are these the same or different vectors? If we drew them as directed line segments, the drawings would be distinct. However, both \vec{a} and \vec{x} have equivalent magnitudes and directions. Thus, \vec{a} and \vec{x} are *equivalent*, and we would be justified writing $\vec{a} = \vec{x}$.



Alternatively, we could consider the *rooted vector* \vec{a} rooted at the point A . In this terminology, \vec{a} rooted at A is *different* than \vec{a} rooted at X . This idea of rooted vectors will occasionally be useful, but our primary study will be unrooted vectors.

Vectors and Points

The distinction between vectors and points is sometimes nebulous because they are so closely related to each other. A *point* in Euclidean space specifies an absolute position whereas a vector specifies a magnitude and direction. However, given a point P , one associates P with the vector $\vec{p} = \overrightarrow{OP}$, where O is the origin. Similarly, we associate the vector \vec{v} with the point V so that $\overrightarrow{OV} = \vec{v}$. Thus, we have a way to unambiguously go back and forth between vectors and points². As such, we will treat vectors and points as interchangeable.

² Mathematically, we say there is an *isomorphism* between vectors and points.

Exercises for 2.1

- 1 Find the displacement vector from the point $P(7, 2, 9)$ to the point $Q(-2, 1, 4)$ in the form $a\hat{x} + b\hat{y} + c\hat{z}$.

Solution: $-9\hat{x} - 1\hat{y} - 5\hat{z}$

- 2^{OS} For the following exercises, consider points $P(-1, 3)$, $Q(1, 5)$, and $R(-3, 7)$. Determine the requested vectors and express each of them i. in component form and ii. by using the standard unit vectors.

- a) \overrightarrow{PQ}
- b) \overrightarrow{PR}
- c) \overrightarrow{QP}
- d) \overrightarrow{RP}

2.2 Vector Arithmetic

Vectors provide a natural way to give directions. For example, suppose \hat{x} points one mile eastwards and \hat{y} points one mile northwards. Now, if you were standing at the origin and wanted to move to a location 3 miles east and 2 miles north, you might say: “Walk 3 times the length of \hat{x} in the \hat{x} direction and 2 times the length of \hat{y} in the \hat{y} direction.” Mathematically, we express this as

$$3\hat{x} + 2\hat{y}.$$

Of course, we’ve incidentally described a new vector. Namely, let P be the point at 3-east and 2-north. Then


$$\overrightarrow{OP} = 3\hat{x} + 2\hat{y}.$$

If the vector \vec{r} points north but has a length of 10 miles, we have a similar formula:


$$\overrightarrow{OP} = 3\hat{x} + \frac{1}{5}\vec{r},$$

and we have the relationship $\vec{r} = 10\hat{y}$. Our notation here is very suggestive. Indeed, if we could make sense of what $\alpha\vec{v}$ is for any scalar α and vector \vec{v} , and we could make sense of what $\vec{v} + \vec{w}$ means for any vectors \vec{v} and \vec{w} , we would be able to do algebra with vectors. We might even say we have *an algebra of vectors*.

Intuitively, for a vector \vec{v} and a scalar $\alpha > 0$, the vector $\vec{w} = \alpha\vec{v}$ should point in the same direction as \vec{v} but have magnitude scaled up by α . That is, $\|\vec{w}\| = \alpha\|\vec{v}\|$. Similarly, $-\vec{v}$ should be the vector of the same length as \vec{v} but pointing in the exact opposite direction.



For two vectors \vec{u} and \vec{v} , the sum $\vec{w} = \vec{u} + \vec{v}$ should be the displacement vector created by first displacing along \vec{u} and then displacing along \vec{v} .



Now, there is one snag. What should $\vec{v} + (-\vec{v})$ be? Well, first we displace along \vec{v} and then we displace in the exact opposite direction by the same amount. So, we have gone nowhere. This corresponds to a displacement with zero magnitude. But, what direction did we displace? Here we make a philosophical stand.

Definition 2.2.1 — Zero Vector. The *zero vector*, notated as $\vec{0}$, is the vector with no magnitude.

We will be pragmatic about the direction of the zero vector and say, *the zero vector does not have a well-defined direction*³. That means sometimes we consider the zero vector to point in every direction and sometimes we consider it to point in no directions. It depends on our mood—but we must never talk about *the* direction of the zero vector, since it's not defined.

We need the zero vector if we are to make precise mathematical sense of vector arithmetic. Further along this line of thinking, we can define precisely how vector arithmetic should behave. Specifically, if \vec{u} , \vec{v} , \vec{w} are vectors and α and β are scalars, the following conditions should be satisfied:

$$(\vec{u} + \vec{v}) + \vec{w} = \vec{u} + (\vec{v} + \vec{w}) \quad (\text{Associativity})$$

$$\vec{u} + \vec{v} = \vec{v} + \vec{u} \quad (\text{Commutativity})$$

$$\alpha(\vec{u} + \vec{v}) = \alpha\vec{u} + \alpha\vec{v} \quad (\text{Distributivity})$$

and

$$(\alpha\beta)\vec{v} = \alpha(\beta\vec{v}) \quad (\text{Associativity II})$$

$$(\alpha + \beta)\vec{v} = \alpha\vec{v} + \beta\vec{v} \quad (\text{Distributivity II})$$

Indeed, if we intuitively think about vectors in flat (Euclidean) space, all of these properties are satisfied⁴. From now on, these properties of vector operations will be considered the *laws (or axioms) of vector arithmetic*.

We'll be talking about these vector operations (scalar multiplication and vector addition) a lot. So much so that the concept is worth naming.

Definition 2.2.2 — Linear Combination. A *linear combination* of the vectors $\vec{v}_1, \dots, \vec{v}_n$ is any vector expressible as

$$\alpha_1 \vec{v}_1 + \dots + \alpha_n \vec{v}_n,$$

where $\alpha_1, \alpha_2, \dots, \alpha_n$ are scalars.

We've given laws for linear combinations of vectors, but what about for magnitudes of vectors? We'd like the magnitude (or norm) of a vector to obey the following laws.

$$\|\vec{v}\| \geq 0 \quad (\text{Non-negativity})$$

$$\|\vec{v}\| = 0 \text{ only when } \vec{v} = \vec{0} \quad (\text{Definiteness})$$

$$\|\alpha\vec{v}\| = |\alpha|\|\vec{v}\| \quad (\text{Homogeneity})$$

$$\|\vec{v} + \vec{w}\| \leq \|\vec{v}\| + \|\vec{w}\| \quad (\text{Triangle Inequality})$$

for all \vec{v} , \vec{w} , and scalars α . Any function on vectors satisfying those four properties is called a *norm*, and our usual notion of length in three-dimensional space indeed obeys those properties⁵.

³ In the mathematically precise definition of vector, the idea of “magnitude” and “direction” are dropped. Instead, a set of vectors is defined to be a set over which you can reasonably define addition and scalar multiplication.

⁴ If we deviate from flat space, some of these rules are no longer respected. Consider moving 100 miles north then 100 miles east on a sphere. Is this the same as moving 100 miles east and then 100 miles north?

⁵ The Euclidean norm comes from the Pythagorean theorem $a^2 + b^2 = c^2$. However, by changing the exponent, we have a whole family of norms coming from the equations $|a|^p + |b|^p = |c|^p$.

Homogeneity is a particularly special property of a norm. It allows us to easily create *unit vectors*.

■ **Definition 2.2.3 — Unit Vector.** A *unit vector* is a vector \vec{u} satisfying $\|\vec{u}\| = 1$.

Unit vectors are handy because if \vec{u} is a unit vector, then $k\vec{u}$ has length $|k|$. Further, we can always turn a vector into a unit vector.

■ **Example 2.1** The vector $\vec{v}/\|\vec{v}\|$ is always a unit vector in the direction of \vec{v} . Computing,


$$\left\| \frac{\vec{v}}{\|\vec{v}\|} \right\| = \left| \frac{1}{\|\vec{v}\|} \right| \|\vec{v}\| = \frac{1}{\|\vec{v}\|} \|\vec{v}\| = 1.$$

■

2.3 Coordinates

Recall that a rectangular coordinate system in the plane is specified by choosing an origin O and then choosing two perpendicular axes meeting at the origin. These axes are chosen in some order so that we know which axis (usually the x -axis) comes first and which (usually the y -axis) second. Note that there are many different coordinate systems which could be used although we often draw pictures as if there were only one.

In physics, one often has to think carefully about the coordinate system because choosing one appropriately may greatly simplify the analysis. Note that axes for coordinate systems are usually drawn with *right-hand orientation*, where the right angle from the positive x -axis to the positive y -axis is in the counter-clockwise direction. However, it would be equally valid to use the *left-hand orientation* in which that angle is in the clockwise direction. One can easily switch the orientation of a coordinate system by reversing one of the axes⁶.



For any coordinate system, there are special vectors associated with it. For the plane, the vector pointing one unit along the positive x -axis is called \hat{x} and the vector pointing one unit along the positive y -axis is called \hat{y} . The vectors \hat{x} and \hat{y} are called the *standard basis vectors* for \mathbb{R}^2 .

⁶ The concept of orientation is quite fascinating and it arises in mathematics, physics, chemistry, and even biology in many interesting ways. Note that almost all of us base our intuitive concept of orientation on our inborn notion of “right” versus “left”.

Notice that every point (or vector) in the plane can be represented as a linear combination of $\hat{\mathbf{x}}$ and $\hat{\mathbf{y}}$, and the vector $\alpha\hat{\mathbf{x}} + \beta\hat{\mathbf{y}}$ is the vector \overrightarrow{OP} where $P = (\alpha, \beta)$. Now, to state an intuitive fact: if \vec{w} is a vector in the plane, *there is only one way to write a vector as a linear combination of $\hat{\mathbf{x}}$ and $\hat{\mathbf{y}}$* . This means, if $\vec{w} = \alpha\hat{\mathbf{x}} + \beta\hat{\mathbf{y}}$, the pair (α, β) captures all information⁷ about \vec{w} .

For a vector $\vec{w} = \alpha\hat{\mathbf{x}} + \beta\hat{\mathbf{y}}$, we call the pair (α, β) the *components* of the vector \vec{w} . There are many equivalent notations used to represent components.

(α, β)	parenthesis
$\langle \alpha, \beta \rangle$	angle brackets
$\begin{bmatrix} \alpha & \beta \end{bmatrix}$	square brackets in a row (a row matrix)
$\begin{bmatrix} \alpha \\ \beta \end{bmatrix}$	square brackets in a column (a column matrix)

Given what we now know about representing vectors and their equivalency with points, we can now dissect the notation \mathbb{R}^2 . On the one hand, \mathbb{R}^2 is the set of vectors in two-dimensional Euclidean space. On the other hand $\mathbb{R}^2 = \mathbb{R} \times \mathbb{R}$ is the set of all pairs of real numbers. Via the use of coordinates, we know these concepts represent the same thing! Further, since vectors in \mathbb{R}^2 are equivalent to their representation in coordinates, we will often write

$$\vec{v} = (\alpha, \beta)$$

as a shorthand for $\vec{v} = \alpha\hat{\mathbf{x}} + \beta\hat{\mathbf{y}}$.

Breaking vectors into components, and in particular, viewing vectors as linear combinations of the standard basis vectors, allows us to solve problems that were difficult before. For instance, suppose we have vectors \vec{v} and \vec{w} . How can we compute $\|\vec{v} + \vec{w}\|$? With components, it's easy.

■ **Example 2.2** Suppose $\vec{v} = \alpha_1\hat{\mathbf{x}} + \beta_1\hat{\mathbf{y}}$ and $\vec{w} = \alpha_2\hat{\mathbf{x}} + \beta_2\hat{\mathbf{y}}$. By the laws of vector arithmetic we have

$$\vec{v} + \vec{w} = (\alpha_1\hat{\mathbf{x}} + \beta_1\hat{\mathbf{y}}) + (\alpha_2\hat{\mathbf{x}} + \beta_2\hat{\mathbf{y}}) = (\alpha_1 + \alpha_2)\hat{\mathbf{x}} + (\beta_1 + \beta_2)\hat{\mathbf{y}}.$$

Now, since $\hat{\mathbf{x}}$ and $\hat{\mathbf{y}}$ are orthogonal to each other, the Pythagorean theorem gives

$$\|\vec{v} + \vec{w}\| = \sqrt{(\alpha_1 + \alpha_2)^2 + (\beta_1 + \beta_2)^2}.$$

■

Writing things in terms of the standard basis allowed us to make easy work of computing $\|\vec{v} + \vec{w}\|$ in Example 2.2. We can use the laws of vector arithmetic to produce rules for working with components.

The rules are likely familiar:

$$\begin{bmatrix} a \\ b \end{bmatrix} + \begin{bmatrix} \alpha \\ \beta \end{bmatrix} = \begin{bmatrix} a + \alpha \\ b + \beta \end{bmatrix} \quad \text{and} \quad \alpha \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} \alpha a \\ \alpha b \end{bmatrix}.$$


⁷ Maybe you already knew this because the point (α, β) is described by the pair of numbers (α, β) , duh! But consider, what would we do if we didn't know about coordinates at all? One approach is to *define* coordinates in terms of vectors, which is really what we're doing.

Exercise 2.1 Prove the rules for adding the component representation of vectors and multiplying the component representation of vectors directly from the laws of vector arithmetic.

Armed with these rules, we will be able to tackle sophisticated vector problems.

Three-dimensional Coordinates

In three-dimensional space, the story is very similar. Again, we imagine three perpendicular axes, the x , y , and z axes. To draw consistent pictures, we have a notion of a right-handed three-dimensional coordinate system given by the *right-hand rule*.



We now have three standard basis vectors, \hat{x} , \hat{y} , and \hat{z} , each pointing one unit in the positive direction of their respective axes. Any vector in three-dimensional space can be represented in exactly one way as a linear combination $\alpha\hat{x} + \beta\hat{y} + \gamma\hat{z}$. Thus, vectors in three-dimensional space, notated \mathbb{R}^3 , are synonymous with triplets (α, β, γ) of real numbers. With some clever geometry, we deduce

$$\|\alpha\hat{x} + \beta\hat{y} + \gamma\hat{z}\| = \sqrt{\alpha^2 + \beta^2 + \gamma^2}.$$

Historically, three-dimensional space has been studied a lot and there are several notations for the standard basis vectors still in use.

The following is a non-exhaustive list.

\hat{x}	\hat{y}	\hat{z}
\hat{i}	\hat{j}	\hat{k}
\mathbf{i}	\mathbf{j}	\mathbf{k}
\vec{e}_1	\vec{e}_2	\vec{e}_3

⁸ Image credit: Acdx, from Wikipedia https://en.wikipedia.org/wiki/Cross_product

Keep these notations in the back of your mind. You might see them in other classes.

Higher dimensions

One can't progress very far in the study of science and mathematics without encountering a need for higher dimensional "vectors." For example, physicists have known since Einstein that the physical universe is best thought of as a four-dimensional entity called spacetime in which time plays a role close to that of the three spatial coordinates. Since we don't have any way to deal with \mathbb{R}^n intuitively, we must proceed by analogy with two and three dimensions. The easiest way to proceed is to generalize the idea of a standard basis. From there, we can represent vectors in \mathbb{R}^n as n -tuples of real numbers. We then define

$$\|(x_1, x_2, \dots, x_n)\| = \sqrt{x_1^2 + x_2^2 + \dots + x_n^2}.$$

We've now unified our theory of vectors across all integer dimensions $n > 0$. The case $n = 1$ yields "geometry" on a line, the cases $n = 2$ and $n = 3$ geometry in the plane and in space, and the case $n = 4$ yields the geometry of "4-vectors" which are used in the special theory of relativity. Larger values of n are used in a variety of contexts, some of which we shall encounter later.

Exercises for 2.3

- 1 Find $\|\vec{a}\|$, $5\vec{a} - 2\vec{b}$, and $-3\vec{b}$ for each of the following vector pairs.

a) $\vec{a} = 2\hat{x} + 3\hat{y}$, $\vec{b} = 4\hat{x} - 9\hat{y}$

b) $\vec{a} = (1, 2, -1)$, $\vec{b} = (2, -1, 0)$

Solution:

a) $\|\vec{a}\| = \sqrt{13}$, $5\vec{a} - 2\vec{b} = 2\hat{x} + 33\hat{y}$,
 $-3\vec{b} = -12\hat{x} + 27\hat{y}$.

b) $\|\vec{a}\| = \sqrt{6}$, $5\vec{a} - 2\vec{b} = (1, 12, -5)$,
 $-3\vec{b} = (-6, 3, 0)$.

- 2 Let $P = (7, 2, 9)$ and $Q = (-2, 1, 4)$. Find \vec{PQ} as a linear combination of \hat{x} , \hat{y} , and \hat{z} .

Solution: $\vec{PQ} = -9\hat{x} - \hat{y} - 5\hat{z}$.

- 3 Find unit vectors with the same direction as the vectors a) $(-4, -2)$, b) $3\hat{x} + 5\hat{y}$, c) $(1, 3, -2)$.

Solution:

a) $\frac{1}{\sqrt{5}}(-2, -1)$

b) $\frac{1}{\sqrt{34}}(3, 5)$

c) $\frac{1}{\sqrt{14}}(1, 3, -2)$

- 4 Show by direct calculation that the rule $\vec{AB} + \vec{BC} + \vec{CA} = \vec{0}$ holds for the three points $A(2, 1, 0)$, $B(-4, 1, 3)$, and $C(0, 12, 0)$. Can you prove the general rule for any three points in space?

Solution: In the specific case, $\vec{AB} = (-6, 0, 3)$, $\vec{BC} = (4, 11, -3)$, $\vec{CA} = (2, -11, 0)$, so that

$$\vec{AB} + \vec{BC} + \vec{CA} = \begin{bmatrix} -6 + 4 + 2 \\ 0 + 11 - 11 \\ 3 - 3 + 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

More generally, suppose that we have the points $A(a_1, a_2, a_3)$, $B(b_1, b_2, b_3)$, and $C(c_1, c_2, c_3)$. Then $\vec{AB} = (b_1 - a_1, b_2 - a_2, b_3 - a_3)$, $\vec{BC} = (c_1 - b_1, c_2 - b_2, c_3 - b_3)$, and $\vec{CA} = (a_1 - c_1, a_2 - c_2, a_3 - c_3)$. Thus,

$$\vec{AB} + \vec{BC} + \vec{CA} = \begin{bmatrix} (b_1 - a_1) + (c_1 - b_1) + (a_1 - c_1) \\ (b_2 - a_2) + (c_2 - b_2) + (a_2 - c_2) \\ (b_3 - a_3) + (c_3 - b_3) + (a_3 - c_3) \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

- 5 If an airplane flies with apparent velocity \vec{v}_a relative to air, and the wind velocity is denoted \vec{w} , then the plane's true velocity relative to the ground, is $\vec{v}_g = \vec{v}_a + \vec{w}$. Draw the diagram to assure yourself of this.

- A farmer wishes to fly his crop duster at 80 km/h north over his fields. If the weather vane atop the barn shows easterly winds at 10 km/h, what should his apparent velocity, \vec{v}_a be?
- What if the wind were northeasterly? Southeasterly?

Solution:

- $10\sqrt{65}$ km/h, $\tan^{-1}\left(\frac{1}{8}\right)^\circ$ east of north.
- for a northeasterly wind:
 $2\sqrt{\left(40 + \frac{5\sqrt{2}}{2}\right)^2 + \frac{25}{2}}$ at $\tan^{-1}\left(\frac{5\sqrt{2}}{80+5\sqrt{2}}\right)^\circ$
 east of north
 for a southeasterly wind:
 $2\sqrt{\left(40 - \frac{5\sqrt{2}}{2}\right)^2 + \frac{25}{2}}$ at $\tan^{-1}\left(\frac{5\sqrt{2}}{80-5\sqrt{2}}\right)^\circ$
 east of north

- 6 Suppose a right handed coordinate system has been set up in space. What happens to the orientation of the coordinate system if you make the following changes?

- Change the direction of one axis.
- Change the direction of two axes.
- Change the direction of all three axes.
- Interchange the x and y -axes.

Solution:

- the orientation changes.
- the orientation stays the same.
- the orientation changes.
- the orientation changes.

- 7 In body-centered crystals, a large atom (assumed spherical) is surrounded by eight smaller atoms. If the structure is placed in a "box" that just contains the large atom, the smaller atoms each occupy a corner. If the central atom has radius R , what is the

greatest atomic radii the smaller atoms may have?

Solution: $(2 - \sqrt{3})R$

- 8 Prove that the diagonals of a parallelogram bisect each other. (Hint: show that the position vectors from the origin to the midpoints are equal).

Solution: Suppose that the parallelogram is defined by two vectors \vec{a} and \vec{b} . Then the first diagonal (which originates from the origin) is given by $\vec{d}_1 = \vec{a} + \vec{b}$ and the second diagonal (which originates from the tip of \vec{b}) is given by $\vec{d}_2 = \vec{a} - \vec{b}$. Then, the position of the midpoint of the first diagonal is given by $\vec{d}_1/2 = 1/2(\vec{a} + \vec{b})$ and the position of the midpoint of the second diagonal is given by $\vec{b} + \vec{d}_2/2 = \vec{b} + 1/2(\vec{a} - \vec{b}) = 1/2(\vec{a} + \vec{b})$. Thus, the position vectors of the midpoints are equal, and so the diagonals must bisect each other.


2.4 Dot Products & Projections

Dot Product

Let \vec{a} and \vec{b} be vectors. We assume they are placed so their tails coincide. Let θ denote the *smaller* of the two angles between them, so $0 \leq \theta \leq \pi$. The *dot product* of \vec{a} and \vec{b} is defined to be

$$\vec{a} \cdot \vec{b} = \|\vec{a}\| \|\vec{b}\| \cos \theta.$$

We will call this the *geometric definition of the dot product*. The dot product is also sometimes called the *scalar product* because the result is a scalar. Note that $\vec{a} \cdot \vec{b} = 0$ when either \vec{a} or \vec{b} is zero or, more interestingly, if their directions are perpendicular.



Algebraically, we can define the dot product in terms of components:

$$\begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix} \cdot \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix} = a_1 b_1 + a_2 b_2 + \cdots + a_n b_n.$$

We will call this the *algebraic definition of the dot product*⁹.

By switching between algebraic and geometric definitions, we can use the dot product to find quantities that are otherwise difficult to find.

■ **Example 2.3** Find the angle between the vectors $\vec{v} = (1, 2, 3)$ and $\vec{w} = (1, 1, -2)$.

From the algebraic definition of the dot product, we know

$$\vec{v} \cdot \vec{w} = 1(1) + 2(1) + 3(-2) = -3.$$

From the geometric definition, we know

$$\vec{v} \cdot \vec{w} = \|\vec{v}\| \|\vec{w}\| \cos \theta = \sqrt{14} \sqrt{6} \cos \theta = 2\sqrt{21} \cos \theta.$$

Equating the two definitions of $\vec{v} \cdot \vec{w}$, we see

$$\cos \theta = \frac{-3}{2\sqrt{21}}$$

and so $\theta = \arccos\left(\frac{-3}{2\sqrt{21}}\right)$. ■

⁹ Philosophically, every object should have only one definition from which equivalent characterizations can be deduced as theorems. If you're bothered, pick your favorite definition to be the "true" definition and consider the other definition a theorem.

Recall that for vectors \vec{a} and \vec{b} , the relationship $\vec{a} \cdot \vec{b} = 0$ can hold for two reasons: (i) either $\vec{a} = \vec{0}$, $\vec{b} = \vec{0}$, or both or (ii) \vec{a} and \vec{b} meet at 90° . Thus, the dot product can be used to tell if two vectors are perpendicular. There is some strangeness with the zero vector here, but it turns out this strangeness simplifies our lives mathematically.

Definition 2.4.1 — Orthogonal. The vectors \vec{u} and \vec{v} are *orthogonal* if $\vec{u} \cdot \vec{v} = 0$.

The definition of orthogonal encapsulates both the idea of two vectors being perpendicular and the idea of one of them being $\vec{0}$.

Before we continue, let's pin down the idea of one vector pointing in the *direction* of another. There are many ways we could define this idea, but we'll go with this one.

Definition 2.4.2 The vector \vec{u} points in the *direction* of the vector \vec{v} if $k\vec{u} = \vec{v}$ for some scalar k .

A simple example is that $2\hat{x}$ points in the direction of \hat{x} since $\frac{1}{2}(2\hat{x}) = \hat{x}$. However, nothing in the definition says the scalar needs to be positive, so $-\hat{x}$ also points in the direction \hat{x} .


Projection

Another common vector operation is *projection*. Projection measures how much a vector points in the direction of another. This quantity is encoded as a vector. We make this definition mathematically precise as follows.

Definition 2.4.3 — Projection. For a vector \vec{u} and a non-zero vector \vec{v} , the *projection* of \vec{u} onto \vec{v} is written as $\text{proj}_{\vec{v}} \vec{u}$ and is a vector in the direction of \vec{v} with the property that $\vec{u} - \text{proj}_{\vec{v}} \vec{u}$ is orthogonal to \vec{v} .

The vector $\vec{u} - \text{proj}_{\vec{v}} \vec{u}$ is called the *perpendicular component* of the projection of \vec{u} onto \vec{v} and is notated $\text{perp}_{\vec{v}} \vec{u}$.

We can visualize projections with the following diagram.



From the picture, it appears that \vec{u} , $\text{proj}_{\vec{v}} \vec{u}$, and $\text{perp}_{\vec{v}} \vec{u}$ form a right triangle. Of course, we shouldn't trust the picture. We should verify this mathematically.

Theorem 2.4.1 If \vec{u} and \vec{v} are non-zero vectors, then \vec{v} , $\text{proj}_{\vec{v}} \vec{u}$, and $\text{perp}_{\vec{v}} \vec{u}$ form a (possibly degenerate) right triangle.

Proof. We need to verify that the sides $\text{proj}_{\vec{v}} \vec{u}$ and $\text{perp}_{\vec{v}} \vec{u}$ meet at a right angle and that the hypotenuse \vec{u} meets the sides. That is, $\text{perp}_{\vec{v}} \vec{u} + \text{proj}_{\vec{v}} \vec{u} = \vec{u}$.

By the definition of projection, $\text{perp}_{\vec{v}} \vec{u} = \vec{u} - \text{proj}_{\vec{v}} \vec{u}$ is orthogonal to \vec{v} . Since $\text{proj}_{\vec{v}} \vec{u}$ points in the direction of \vec{v} , we have $\text{proj}_{\vec{v}} \vec{u} = k\vec{v}$ and so $\text{perp}_{\vec{v}} \vec{u}$ is orthogonal to $\text{proj}_{\vec{v}} \vec{u}$.

Finally, consider

$$\text{perp}_{\vec{v}} \vec{u} + \text{proj}_{\vec{v}} \vec{u} = (\vec{u} - \text{proj}_{\vec{v}} \vec{u}) + \text{proj}_{\vec{v}} \vec{u} = \vec{u},$$

so indeed the vectors form a right triangle. ■

Now that we've proved \vec{u} , $\text{proj}_{\vec{v}} \vec{u}$, and $\text{perp}_{\vec{v}} \vec{u}$ form a right triangle, we are free to use trigonometry to compute projections. If θ is the angle between \vec{u} and \vec{v} and $0 \leq \theta \leq \pi/2$, we know $\|\text{proj}_{\vec{v}} \vec{u}\| = \|\vec{u}\| \cos \theta$. This means

$$\text{proj}_{\vec{v}} \vec{u} = k\vec{v} = \|\vec{u}\| \cos \theta \frac{\vec{v}}{\|\vec{v}\|}$$

(Recall that $\vec{v}/\|\vec{v}\|$ is a unit vector in the direction of \vec{v}). But $\cos \theta$ appears in the formula for the dot product. Solving for $\cos \theta$ in the dot product formula, we see $\cos \theta = \frac{\vec{u} \cdot \vec{v}}{\|\vec{u}\| \|\vec{v}\|}$. Thus,

$$\text{proj}_{\vec{v}} \vec{u} = \|\vec{u}\| \cos \theta \frac{\vec{v}}{\|\vec{v}\|} = \frac{\vec{u} \cdot \vec{v}}{\|\vec{u}\| \|\vec{v}\|} \left(\frac{\vec{v}}{\|\vec{v}\|} \right) = \frac{\vec{u} \cdot \vec{v}}{\|\vec{v}\|^2} \vec{v}.$$

Upon close inspection, we see $\|\vec{v}\|^2 = \vec{v} \cdot \vec{v}$ (since $\cos 0 = 1$) and so we finally arrive at the formula

$$\text{proj}_{\vec{v}} \vec{u} = \frac{\vec{u} \cdot \vec{v}}{\vec{v} \cdot \vec{v}} \vec{v}$$

Incredibly, if we use the algebraic definition of the dot product, we can compute a projection without computing cosine of anything!

Exercises for 2.4

- 1 In each case determine if the given pair of vectors is orthogonal.
 - a) $\vec{a} = 3\hat{x} + 4\hat{y}$, $\vec{b} = -4\hat{x} + 3\hat{y}$
 - b) $\vec{a} = (4, -1, 2)$, $\vec{b} = (3, 0, -6)$
 - c) $\vec{a} = 3\hat{x} - 2\hat{y}$, $\vec{b} = -2\hat{x} - 4\hat{z}$

Solution:

 - a) $\vec{a} \cdot \vec{b} = -12 + 12 = 0$, so \vec{a} and \vec{b} are orthogonal.
 - b) $\vec{a} \cdot \vec{b} = 12 + 0 - 12 = 0$, so \vec{a} and \vec{b} are orthogonal.
 - c) $\vec{a} \cdot \vec{b} = -6 + 8 = 2$, so \vec{a} and \vec{b} are not orthogonal.
- 2 Assume $\vec{u} = (a, b)$ is a non-zero plane vector. Show that $\vec{v} = (-b, a)$ is orthogonal to \vec{u} . By examining all possible signs for a and b , convince yourself that the 90 degree angle between \vec{u} and \vec{v} is in the counter-clockwise direction.

Solution: $\vec{u} \cdot \vec{v} = -ab + ba = 0$, so \vec{u} and \vec{v} are orthogonal.
- 3 The methane molecule, CH_4 , has four Hydrogen atoms at the vertices of a regular tetrahedron and a Carbon atom at its center. Choose as the vertices of this tetrahedron the points $(0, 0, 0)$, $(1, 1, 0)$, $(1, 0, 1)$, and $(0, 1, 1)$.
 - a) Find the angle between two edges of the tetrahedron.

- b) Find the bond angle between two Carbon–Hydrogen bonds.

Solution:

a) 60°

b) $\cos^{-1}\left(\frac{1}{3}\right)^\circ$

- 4 An inclined plane makes an angle of 30° degrees with the horizontal. Use vectors and the dot product to find the projection of the gravitational acceleration vector $-g\hat{y}$ along a unit vector pointing along the inclined plane.

Solution: $-\frac{g\sqrt{3}}{4}\hat{x} - \frac{g}{4}\hat{y}$

- 5 Derive the formula

$$\|\vec{a} + \vec{b}\|^2 = \|\vec{a}\|^2 + \|\vec{b}\|^2 + 2\|\vec{a}\|\|\vec{b}\|\cos\theta.$$

How is it related to the Law of Cosines? A picture might help.

Solution:

$$\begin{aligned}\|\vec{a} + \vec{b}\|^2 &= (\vec{a} + \vec{b}) \cdot (\vec{a} + \vec{b}) \\ &= \vec{a} \cdot \vec{a} + \vec{a} \cdot \vec{b} + \vec{b} \cdot \vec{a} + \vec{b} \cdot \vec{b} \\ &= \|\vec{a}\|^2 + \|\vec{b}\|^2 + 2\vec{a} \cdot \vec{b} \\ &= \|\vec{a}\|^2 + \|\vec{b}\|^2 + 2\|\vec{a}\|\|\vec{b}\|\cos\theta\end{aligned}$$

- 6 Use the dot product to determine if the points $P(3, 1, 2)$, $Q(-1, 0, 2)$, and $R(11, 3, 2)$ are collinear.

- 7^{OS} For the following exercises, the vectors \vec{u} and \vec{v} are given. Find the projection $\vec{w} = \text{proj}_{\vec{u}} \vec{v}$ of vector \vec{v} onto vector \vec{u} . Express your answer in component form.

a) $\vec{u} = 5\hat{x} + 2\hat{y}$, $\vec{v} = 2\hat{x} + 3\hat{y}$.

b) $\vec{u} = (4, 7)$, $\vec{v} = (3, 5)$

c) $\vec{u} = 3\hat{x} + 2\hat{z}$, $\vec{v} = 2\hat{y} + 4\hat{z}$

d) $\vec{u} = (4, 4, 0)$, $\vec{v} = (0, 4, 1)$

- 8^{OS} Consider the vectors $\vec{u} = 4\hat{x} + 3\hat{y}$ and $\vec{v} = 3\hat{x} + 2\hat{y}$

- a) Find the components of vector $\vec{w} = \text{proj}_{\vec{u}} \vec{v}$ that represents the projection of \vec{v} onto \vec{u} .

- b) Write the decomposition $\vec{v} = \vec{w} + \vec{q}$ of vector \vec{v} into the orthogonal components \vec{w} and \vec{q} , where \vec{w} is the projection of \vec{v} onto \vec{u} and \vec{q} is a vector orthogonal to the direction of \vec{u} .

- 9^{OS} Consider the vectors $\vec{u} = 2\hat{x} + 4\hat{y}$ and $\vec{v} = 4\hat{x} + 2\hat{z}$

- a) Find the components of vector $\vec{w} = \text{proj}_{\vec{u}} \vec{v}$ that represents the projection of \vec{v} onto \vec{u} .

- b) Write the decomposition $\vec{v} = \vec{w} + \vec{q}$ of vector \vec{v} into the orthogonal components \vec{w} and \vec{q} , where \vec{w} is the projection of \vec{v} onto \vec{u} and \vec{q} is a vector orthogonal to the direction of \vec{u} .

2.5 The Cross Product

For vectors \vec{a} and \vec{b} , the dot product $\vec{a} \cdot \vec{b}$ measures how close \vec{a} and \vec{b} are to being orthogonal. In contrast, the *cross product* of \vec{a} and \vec{b} , written $\vec{a} \times \vec{b}$, will measure the *area* of the parallelogram whose sides are given by \vec{a} and \vec{b} .

Let's explore this idea. Since the cross product is a *product*, we will demand it follow reasonable distribution laws¹⁰:

$$\vec{a} \times (\vec{b} + \vec{c}) = \vec{a} \times \vec{b} + \vec{a} \times \vec{c}$$


$$(\vec{a} + \vec{b}) \times \vec{c} = \vec{a} \times \vec{c} + \vec{b} \times \vec{c}$$

$$(\alpha \vec{a}) \times \vec{b} = \alpha(\vec{a} \times \vec{b})$$

$$\vec{a} \times (\alpha \vec{b}) = \alpha(\vec{a} \times \vec{b})$$

for vectors \vec{a} , \vec{b} , \vec{c} and scalars α .

Now, suppose $\vec{a} \times \vec{b}$ indeed encapsulates the area of the parallelogram with sides \vec{a} and \vec{b} . If we slide the tip of \vec{b} parallel to the vector \vec{a} , we should not change the area. Thus the cross product of \vec{a} and \vec{b} should be the same as that of \vec{a} and $\vec{b} + \alpha \vec{a}$.



Using this invariance along with our distributive rules, we now see

$$\vec{a} \times \vec{b} = \vec{a} \times (\vec{b} + \alpha \vec{a}) = \vec{a} \times \vec{b} + \alpha(\vec{a} \times \vec{a}),$$

and so $\vec{a} \times \vec{a} = 0$. We can apply this newly-found fact to the vector $\vec{a} + \vec{b}$ to deduce

$$\begin{aligned} 0 &= (\vec{a} + \vec{b}) \times (\vec{a} + \vec{b}) = \vec{a} \times \vec{a} + \vec{a} \times \vec{b} + \vec{b} \times \vec{a} + \vec{b} \times \vec{b} \\ &= 0 + \vec{a} \times \vec{b} + \vec{b} \times \vec{a} + 0 \\ &= \vec{a} \times \vec{b} + \vec{b} \times \vec{a}, \end{aligned}$$

and so

$$\vec{a} \times \vec{b} = -\vec{b} \times \vec{a}.$$

Products with this property are called *anti-commutative*. Now for an incredible fact of the universe: the result of the cross product of two vectors in \mathbb{R}^3 can be represented by another vector in \mathbb{R}^3 whose magnitude corresponds to the area of the parallelogram with sides \vec{a} and \vec{b} .¹¹ Using trigonometry, we deduce

$$\|\vec{a} \times \vec{b}\| = \|\vec{a}\| \|\vec{b}\| \sin \theta$$

¹⁰ The technical term for satisfying these laws is *bilinearity*.

¹¹ This is *only* true in \mathbb{R}^3 . In \mathbb{R}^4 a product that produces area-like quantities does exist, but the output cannot be described by a vector. In higher dimensions, the cross product is called the *wedge product*.

where $0 \leq \theta \leq \pi$ is smaller of the two angles between \vec{a} and \vec{b} . What remains to be seen is what direction $\vec{a} \times \vec{b}$ points in. For this, we use the standard basis for \mathbb{R}^3 as a launching point. Recall \hat{x} , \hat{y} , and \hat{z} are all unit vectors and all orthogonal to each other. Thus, any of them cross each other must result in a unit vector. By convention,

$$\begin{aligned}\hat{x} \times \hat{y} &= \hat{z}, \\ \hat{y} \times \hat{z} &= \hat{x} \\ \hat{z} \times \hat{x} &= \hat{y}.\end{aligned}$$

Let $\vec{a} = a_x \hat{x} + a_y \hat{y} + a_z \hat{z}$ and $\vec{b} = b_x \hat{x} + b_y \hat{y} + b_z \hat{z}$. Using the distributive laws of the cross product we see,

$$\begin{aligned}\vec{a} \times \vec{b} &= (a_x \hat{x} + a_y \hat{y} + a_z \hat{z}) \times (b_x \hat{x} + b_y \hat{y} + b_z \hat{z}) \\ &= a_x b_x \hat{x} \times \hat{x} + a_x b_y \hat{x} \times \hat{y} + a_x b_z \hat{x} \times \hat{z} \\ &\quad + a_y b_x \hat{y} \times \hat{x} + a_y b_y \hat{y} \times \hat{y} + a_y b_z \hat{y} \times \hat{z} \\ &\quad + a_z b_x \hat{z} \times \hat{x} + a_z b_y \hat{z} \times \hat{y} + a_z b_z \hat{z} \times \hat{z} \\ &= \vec{0} + a_x b_y \hat{z} - a_x b_z \hat{y} \\ &\quad - a_y b_x \hat{z} + \vec{0} + a_y b_z \hat{x} \\ &\quad + a_z b_x \hat{y} - a_z b_y \hat{x} + \vec{0}\end{aligned}$$

so

$$\vec{a} \times \vec{b} = (a_y b_z - a_z b_y) \hat{x} - (a_x b_z - a_z b_x) \hat{y} + (a_x b_y - a_y b_x) \hat{z}.$$

Exercise 2.2 Verify that $\|\vec{a} \times \vec{b}\| = \|\vec{a}\| \|\vec{b}\| \sin \theta$. (Hint: you can use $\vec{a} \cdot \vec{b} = \|\vec{a}\| \|\vec{b}\| \cos \theta$ to solve for θ and then proceed using components.)


Now that we know what the cross product is and how to compute it, let's explore some of its incredible properties. First,

$$(\vec{a} \times \vec{b}) \cdot \vec{a} = (a_y b_z - a_z b_y) a_x - (a_x b_z - a_z b_x) a_y + (a_x b_y - a_y b_x) a_z = 0$$

and

$$(\vec{a} \times \vec{b}) \cdot \vec{b} = (a_y b_z - a_z b_y) b_x - (a_x b_z - a_z b_x) b_y + (a_x b_y - a_y b_x) b_z = 0.$$

Thus, $\vec{a} \times \vec{b}$ is orthogonal to both \vec{a} and \vec{b} . Just based on this property, since the length of $\vec{a} \times \vec{b}$ is fixed, $\vec{a} \times \vec{b}$ can be one of two vectors in space. If we investigate further, we'll see that $\vec{a} \times \vec{b}$ is the vector that satisfies the *right-hand rule*.



A vector that encodes area, points orthogonally to others, and obeys the right-hand rule is handy indeed, and the cross product will be a useful tool for solving many problems.

Exercises for 2.5

1 Find $\vec{a} \times \vec{b}$ for the following pairs:

- a) $\vec{a} = (4, -2, 0)$, $\vec{b} = (2, 1, -1)$
- b) $\vec{a} = (3, 3, 0)$, $\vec{b} = (4, -3, 2)$
- c) $\vec{a} = 2\hat{x} + 3\hat{y} + 4\hat{z}$, $\vec{b} = \hat{x} - 3\hat{y} + 4\hat{z}$.

Solution:

- a) $(2, 4, 8)$
- b) $(15, 6, -21)$
- c) $24\hat{x} - 4\hat{y} - 9\hat{z}$

2 Use the cross product to find the areas of the following figures.

- a) The parallelogram with vertices $(0, 0, 0)$, $(1, 1, 0)$, $(1, 2, 1)$ and $(0, 1, 1)$. (Perhaps you should first check that this is a parallelogram.)
- b) The triangle with vertices $(1, 0, 0)$, $(0, 1, 0)$, and $(0, 0, 1)$.

Solution:

- a) $\sqrt{3}$
- b) $\frac{\sqrt{3}}{2}$

3 Prove that the cross product is not associative by calculating $\vec{a} \times (\vec{b} \times \vec{c})$ and $(\vec{a} \times \vec{b}) \times \vec{c}$ for $\vec{a} = \hat{x}$, $\vec{b} = \hat{c} = \hat{y}$.

Solution: $\hat{x} \times (\hat{y} \times \hat{y}) = \hat{x} \times \vec{0} = \vec{0}$ while $(\hat{x} \times \hat{y}) \times \hat{y} = \hat{z} \times \hat{y} = -\hat{x}$

4 Verify the formula.

$$\|\vec{a} \times \vec{b}\|^2 = \|\vec{a}\|^2 \|\vec{b}\|^2 - (\vec{a} \cdot \vec{b})^2.$$

Hint: Use the definitions in terms of the sine and cosine of the included angle θ .

Solution:

$$\begin{aligned} \|\vec{a} \times \vec{b}\|^2 &= (\|\vec{a}\| \|\vec{b}\| \sin \theta)^2 \\ &= \|\vec{a}\|^2 \|\vec{b}\|^2 \sin^2 \theta \\ &= \|\vec{a}\|^2 \|\vec{b}\|^2 (1 - \cos^2 \theta) \\ &= \|\vec{a}\|^2 \|\vec{b}\|^2 - \|\vec{a}\|^2 \|\vec{b}\|^2 \cos^2 \theta \\ &= \|\vec{a}\|^2 \|\vec{b}\|^2 - (\|\vec{a}\| \|\vec{b}\| \cos \theta)^2 \\ &= \|\vec{a}\|^2 \|\vec{b}\|^2 - (\vec{a} \cdot \vec{b})^2. \end{aligned}$$

¹² Image credit: Acdx, from Wikipedia https://en.wikipedia.org/wiki/Cross_product

2.6 Lines and Planes

With a handle on vectors, we can now use them to describe some common geometric objects: lines and planes.


Lines

Consider for a moment the line ℓ through the points P and Q . If $P, Q \in \mathbb{R}^2$, we could describe this line in $y = mx + b$ form (provided it isn't a vertical line), but if $P, Q \in \mathbb{R}^3$, it's much harder to describe ℓ with an equation. Using vectors provides an easier way.

Let $\vec{d} = \overrightarrow{PQ}$ and consider the set of points (or vectors)

$$\vec{x} = t\vec{d} + P$$

for $t \in \mathbb{R}$. Geometrically, this is the set of all points we get by starting at P and displacing by some multiple of the direction \vec{d} . This is a line!



Note that sometimes when we draw pictures of vectors, drawing them as line segments is illuminating. Sometimes, however, drawing them as line segments can make it hard to see what's going on, and it is better to draw each vector by marking only its terminal point.

Call the line described above ℓ . In set-builder notation, we would write

$$\ell = \{\vec{x} : \vec{x} = t\vec{d} + P \text{ for some } t \in \mathbb{R}\}.$$

Notice that in set-builder notation, we write “for some $t \in \mathbb{R}$.” Make sure you understand why replacing “for some $t \in \mathbb{R}$ ” with “for all $t \in \mathbb{R}$ ” would be incorrect.

Writing lines with set-builder notation all the time can be overkill, so we will allow ourselves to describe lines in a shorthand called *vector form*¹³.

Definition 2.6.1 — Vector form of a Line. A line ℓ is described in *vector form* if there are two vectors $\vec{d} \neq \vec{0}$ and \vec{p} so that

$$\vec{x} = t\vec{d} + \vec{p}$$

satisfies $\vec{x} \in \ell$ for all $t \in \mathbb{R}$. In this case we call \vec{d} the *direction* of ℓ and the equation

¹³ $y = mx + b$ form of a line is also shorthand. The line ℓ described by the equation $y = mx + b$ is actually the set $\{(x, y) \in \mathbb{R}^2 : y = mx + b\}$.

■ $\vec{x} = t\vec{d} + \vec{p}$ the *vector equation* or *vector form* of ℓ .

Note that if $\vec{x} = t\vec{d} + \vec{p}$ is the vector equation of a line ℓ , by setting $t = 0$ we necessarily have $\vec{p} \in \ell$.

The direction of a line is easily obtained by finding the displacement vector between two points on the line. Thus, given a line in another form, computing its vector form is straightforward.

■ **Example 2.4** Find vector form of the line $\ell \subseteq \mathbb{R}^2$ with equation $y = 2x + 3$.

First, we find two points on the line. By guess-and-check we see $P = (0, 3)$ and $Q = (1, 5)$ are on ℓ . Thus, a direction vector for ℓ is given by

$$\vec{d} = (1, 5) - (0, 3) = (1, 2).$$

We may now write the vector equation of ℓ as

$$\vec{x} = t\vec{d} + P$$

or, in components,

$$\begin{bmatrix} x \\ y \end{bmatrix} = t \begin{bmatrix} 1 \\ 2 \end{bmatrix} + \begin{bmatrix} 0 \\ 3 \end{bmatrix}.$$

■

The downside of writing lines in vector form is that there are multiple direction vectors and multiple points for every line. Thus, merely by looking at the vector equation for two lines, it can be hard to tell if they're equal.

For example,

$$\begin{bmatrix} x \\ y \end{bmatrix} = t \begin{bmatrix} 1 \\ 2 \end{bmatrix} + \begin{bmatrix} 0 \\ 3 \end{bmatrix}, \quad \begin{bmatrix} x \\ y \end{bmatrix} = t \begin{bmatrix} 2 \\ 4 \end{bmatrix} + \begin{bmatrix} 0 \\ 3 \end{bmatrix}, \quad \text{and} \quad \begin{bmatrix} x \\ y \end{bmatrix} = t \begin{bmatrix} 1 \\ 2 \end{bmatrix} + \begin{bmatrix} 1 \\ 5 \end{bmatrix}$$

all represent the same line. In the second equation, the direction is parallel but scaled, and in the third equation, a different point on the line was chosen.

In vector form, the variable t is called the *parameter variable*. It is an instance of a *dummy variable*; that is, it is mostly there as a placeholder. Remember, vector form is shorthand for a set described in set-builder notation.

Let $\vec{d}_1, \vec{d}_2 \neq \vec{0}$ and \vec{p}_1, \vec{p}_2 be vectors and define the lines

$$\ell_1 = \{\vec{x} : \vec{x} = t\vec{d}_1 + \vec{p}_1 \text{ for some } t \in \mathbb{R}\}$$

$$\ell_2 = \{\vec{x} : \vec{x} = t\vec{d}_2 + \vec{p}_2 \text{ for some } t \in \mathbb{R}\}.$$

These lines have vector equations $\vec{x} = t\vec{d}_1 + \vec{p}_1$ and $\vec{x} = t\vec{d}_2 + \vec{p}_2$. However, declaring that $\ell_1 = \ell_2$ if and only if $t\vec{d}_1 + \vec{p}_1 = t\vec{d}_2 + \vec{p}_2$ does *not* make sense. Instead $\ell_1 = \ell_2$ if $\ell_1 \subseteq \ell_2$ and $\ell_2 \subseteq \ell_1$. If $\vec{x} \in \ell_1$ then $\vec{x} = t\vec{d}_1 + \vec{p}_1$ for some $t \in \mathbb{R}$. If $\vec{x} \in \ell_2$ then $\vec{x} = t\vec{d}_2 + \vec{p}_2$ for some possibly different $t \in \mathbb{R}$. This can get confusing really quickly. The easiest solution is to use different parameter variables if we want to compare lines in vector form.

■ **Example 2.5** Determine if the lines ℓ_1 and ℓ_2 , represented in vector form by the equations

$$\vec{x} = t \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 2 \\ 1 \end{bmatrix} \quad \text{and} \quad \vec{x} = t \begin{bmatrix} 2 \\ 2 \end{bmatrix} + \begin{bmatrix} 4 \\ 3 \end{bmatrix}$$

are the same line. To determine this, we need to figure out if $\vec{x} \in \ell_1$ implies $\vec{x} \in \ell_2$ and if $\vec{x} \in \ell_2$ implies $\vec{x} \in \ell_1$.

If $\vec{x} \in \ell_1$, then $\vec{x} = t \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 2 \\ 1 \end{bmatrix}$ for some $t \in \mathbb{R}$. If $\vec{x} \in \ell_2$, then $\vec{x} = s \begin{bmatrix} 2 \\ 2 \end{bmatrix} + \begin{bmatrix} 4 \\ 3 \end{bmatrix}$ for some $s \in \mathbb{R}$. Thus if


$$t \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 2 \\ 1 \end{bmatrix} = \vec{x} = s \begin{bmatrix} 2 \\ 2 \end{bmatrix} + \begin{bmatrix} 4 \\ 3 \end{bmatrix}$$

always has a solution, $\ell_1 = \ell_2$. Moving everything to one side we see

$$\begin{aligned} \vec{0} &= \begin{bmatrix} 4 \\ 3 \end{bmatrix} - \begin{bmatrix} 2 \\ 1 \end{bmatrix} + s \begin{bmatrix} 2 \\ 2 \end{bmatrix} - t \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 2 \\ 2 \end{bmatrix} + s \begin{bmatrix} 2 \\ 2 \end{bmatrix} - t \begin{bmatrix} 1 \\ 1 \end{bmatrix} \\ &= (s+1) \begin{bmatrix} 2 \\ 2 \end{bmatrix} - \frac{t}{2} \begin{bmatrix} 2 \\ 2 \end{bmatrix} \\ &= (s+1 - \frac{t}{2}) \begin{bmatrix} 2 \\ 2 \end{bmatrix}. \end{aligned}$$

This has a solution whenever $0 = s+1 - t/2$. Since for every $t \in \mathbb{R}$ we can find an $s \in \mathbb{R}$ and for every $s \in \mathbb{R}$ we can find a $t \in \mathbb{R}$ satisfying this equation, we know $\ell_1 = \ell_2$. ■

The geometry of lines in space is a bit more complicated than that of lines in the plane. Lines in the plane either intersect or are parallel. In space, we have to be a bit more careful about what we mean by “parallel lines,” since lines with entirely different directions can still fail to intersect¹⁴.



¹⁴ Recall that in Euclidean geometry two lines are defined to be parallel if they coincide or never intersect.

■ **Example 2.6** Consider the lines described by

$$\vec{x} = t(1, 3, -2) + (1, 2, 1)$$

$$\vec{x} = t(-2, -6, 4) + (3, 1, 0).$$

They have parallel directions since $(-2, -6, 4) = -2(1, 3, -2)$. Hence, in this case, we say the lines are *parallel*. (How can we be sure the lines are not the same?) ■

■ **Example 2.7** Consider the lines described by

$$\vec{x} = t(1, 3, -2) + (1, 2, 1)$$

$$\vec{x} = t(0, 2, 3) + (0, 3, 9).$$

They are not parallel because neither of the direction vectors is a multiple of the other. They may or may not intersect. (If they don't, we say the lines are *skew*.) How can we find out? Mirroring our earlier approach, we can set their equations equal and see if we can solve for the point of intersection *after ensuring we give their parametric variables different names*. We'll keep one parametric variable named t and name the other one s . Thus, we want

$$\vec{x} = t(1, 3, -2) + (1, 2, 1) = s(0, 2, 3) + (0, 3, 9),$$

which after collecting terms yields

$$(t + 1, 3t + 2, -2t + 1) = (0, 2s + 3, 3s + 9).$$

Picking out the components yields three equations

$$t + 1 = 0$$

$$3t + 2 = 2s + 3$$

$$-2t + 1 = 3s + 9$$

in 2 unknowns s and t . This is an *overdetermined* system, and it may or may not have a consistent solution. The first two equations yield $t = -1$ and $s = -2$. Putting these values in the last equation yields $(-2)(-1) + 1 = 3(-2) + 9$, which is indeed true. Hence, the equations are consistent, and the lines intersect. To find the point of intersection, put $t = -1$ in the equation for the first line (or $s = -2$ in that for the second) to obtain $(0, -1, 3)$. ■

Planes


Any two distinct points define a line. To define a plane, we need three points. But there's a caveat: the three points cannot be on the same line, otherwise they'd define a line and not a plane. Let $A, B, C \in \mathbb{R}^3$ be three points that are not collinear and let \mathcal{P} be the plane that passes through A , B , and C .

Just like lines, planes have direction vectors. For \mathcal{P} , both $\vec{d}_1 = \overrightarrow{AB}$ and $\vec{d}_2 = \overrightarrow{AC}$ are direction vectors for \mathcal{P} . Of course, \vec{d}_1 , \vec{d}_2 and their multiples are not the only direction vectors for \mathcal{P} . There are infinitely many more, including $\vec{d}_1 + \vec{d}_2$, and $\vec{d}_1 - 7\vec{d}_2$, and so on. However, since a plane is a *two-dimensional* object, we only need two different direction vectors to describe it.

Again like lines, planes have a vector form. \mathcal{P} can be written in vector form as

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = t\vec{d}_1 + s\vec{d}_2 + A.$$

Vector form of \mathcal{P} is not unique. Any two different directions in \mathcal{P} suffice for defining \mathcal{P} in vector form.




Definition 2.6.2 — Vector form of a plane. The plane \mathcal{P} is described in *vector form* if there are three vectors \vec{d}_1 , \vec{d}_2 , and \vec{p} where $\vec{d}_1, \vec{d}_2 \neq \vec{0}$ point in different directions and

$$\vec{x} = t\vec{d}_1 + s\vec{d}_2 + \vec{p}$$

satisfies $\vec{x} \in \mathcal{P}$ for all scalars $t, s \in \mathbb{R}$. The vectors \vec{d}_1 and \vec{d}_2 are called *direction vectors* for the plane \mathcal{P} .

Since we will commonly be working in \mathbb{R}^3 there is another way to define a plane. Given any vector $\vec{n} \in \mathbb{R}^3$, we can consider the set $\mathcal{Q} \subseteq \mathbb{R}^3$ of vectors orthogonal to \vec{n} . If $\vec{n} = \vec{0}$, then $\mathcal{Q} = \mathbb{R}^3$. Otherwise, \mathcal{Q} is a plane through the origin. In this case, \vec{n} is called the *normal vector* of the plane \mathcal{Q} .



Definition 2.6.3 — Normal form of a plane. The plane \mathcal{P} is described in *normal form* if for some \vec{n} and \vec{p} , the equation

$$\vec{n} \cdot (\vec{x} - \vec{p}) = 0$$

if and only if $\vec{x} \in \mathcal{P}$. Equivalently, \mathcal{P} is described in normal form if for some \vec{n} and scalar $\alpha \in \mathbb{R}$ the equation

$$\vec{n} \cdot \vec{x} = \alpha$$

is satisfied if and only if $\vec{x} \in \mathcal{P}$. In either case, the vector \vec{n} is called a *normal vector* for \mathcal{P} .

Normal form of a plane only exists in \mathbb{R}^3 , but it is often useful¹⁵. The equivalence of the two ways to write a normal form of a plane is straight forward.

$$\vec{n} \cdot (\vec{x} - \vec{p}) = 0$$

if and only if

$$\vec{n} \cdot \vec{x} = \vec{n} \cdot \vec{p} = \alpha.$$

Since \vec{n} and \vec{p} are fixed, α is a constant. Expanding normal form in terms of components we see

$$\vec{n} \cdot (\vec{x} - \vec{p}) = \vec{n} \cdot \vec{x} - \alpha = n_x x + n_y y + n_z z - \alpha = 0$$

and so

$$n_x x + n_y y + n_z z = \alpha \tag{2.1}$$

is another way to write a plane. Equation (2.1) is sometimes called *scalar form* of a plane. For us, it will not be important to distinguish between scalar and normal form.

It should be noted that like vector form of a plane, normal form of a plane is not unique. For example, the plane described by $\vec{n} \cdot (\vec{x} - \vec{p}) = 0$ is the same as the plane $(2\vec{n}) \cdot (\vec{x} - \vec{p}) = 0$.

¹⁵ Just like $y = mx + b$ form of a line only exists in \mathbb{R}^2 .

■ **Example 2.8** Find vector form and normal form of the plane \mathcal{P} passing through the point $A = (1, 0, 0)$, $B = (0, 1, 0)$ and $C = (0, 0, 1)$.

To find vector form of \mathcal{P} , we need a point on the plane and two direction vectors. We have three points on the plane, so we can obtain two direction vectors by subtracting these points in different ways. Let

$$\vec{d}_1 = \vec{AB} = \begin{bmatrix} -1 \\ 1 \\ 0 \end{bmatrix} \quad \vec{d}_2 = \vec{AC} = \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}.$$

Using the point A , we may now write vector form of \mathcal{P} as

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = t \begin{bmatrix} -1 \\ 1 \\ 0 \end{bmatrix} + s \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}.$$

To write normal form we need to find a normal vector to \mathcal{P} . By symmetry, we can see that $\vec{n} = (1, 1, 1)$ is a normal vector to \mathcal{P} . If we weren't so insightful, we could also compute $\vec{d}_1 \times \vec{d}_2 = (1, 1, 1)$ to find a normal vector. Now, we may express \mathcal{P} in normal form as

$$\begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \cdot \left(\begin{bmatrix} x \\ y \\ z \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \right) = 0$$

or equivalently,

$$x + y + z = 1.$$

■

■ **Example 2.9** Find the line $\mathcal{P}_1 \cap \mathcal{P}_2$ where \mathcal{P}_1 is the plane given by the equation

$$x + y + z = 2$$

and \mathcal{P}_2 is the plane given by the equation

$$2x - y + z = 0.$$

Let $\ell = \mathcal{P}_1 \cap \mathcal{P}_2$. Since $\ell \subseteq \mathcal{P}_1$ and $\ell \subseteq \mathcal{P}_2$, every direction vector for ℓ is also a direction vector for \mathcal{P}_1 and \mathcal{P}_2 .

Let $\vec{n}_1 = (1, 1, 1)$ be a normal vector for \mathcal{P}_1 and $\vec{n}_2 = (2, -1, 1)$ be a normal vector for \mathcal{P}_2 . If \vec{d} is a direction vector for ℓ , then $\vec{n}_1 \cdot \vec{d} = 0$ and $\vec{n}_2 \cdot \vec{d} = 0$. Thus,

$$\vec{d} = \vec{n}_1 \times \vec{n}_2 = \begin{bmatrix} 2 \\ 1 \\ -3 \end{bmatrix}$$

is a direction vector for ℓ . By guess and check we find that $\vec{p} = (0, 1, 1)$ satisfies $\vec{p} \in \mathcal{P}_1$ and $\vec{p} \in \mathcal{P}_2$ and so $\vec{p} \in \ell$. Thus, we may write ℓ in vector form as

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = t \begin{bmatrix} 2 \\ 1 \\ -3 \end{bmatrix} + \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix}.$$

■

Exercises for 2.6

Chapter 3

Parameterization

Parameterization is a mouthful, but the fundamental idea of a parameterization is to describe one object in terms of another. For example, consider the line ℓ described by the equation $y = 2x$. By its nature, ℓ is a set. Using set-builder notation, we could write

$$\ell = \left\{ \begin{bmatrix} x \\ y \end{bmatrix} : y = 2x \right\}.$$

But, we could also write ℓ in vector form as

$$\begin{bmatrix} x \\ y \end{bmatrix} = t \begin{bmatrix} 1 \\ 2 \end{bmatrix}.$$

Writing ℓ in vector form shows a pairing between scalars $t \in \mathbb{R}$ and points on ℓ . In many ways, ℓ is the same as \mathbb{R} —it's just sitting in two-dimensional space instead of being on its own.

Taking a more technical viewpoint, we may consider ℓ to be the range of a vector-valued function. Define $\vec{p}(t) = t \begin{bmatrix} 1 \\ 2 \end{bmatrix}$. Then,

$$\ell = \text{range}(\vec{p}) = \{ \vec{x} : \vec{x} = \vec{p}(t) \text{ for some } t \in \mathbb{R} \}.$$

Now we have something special. The function $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^2$ has domain \mathbb{R} and outputs every point on the line ℓ exactly once. In other words, we've described ℓ in terms of \mathbb{R} and \vec{p} . We could make a further assertion that *anything that you could learn by studying ℓ , you could learn by studying \mathbb{R} and \vec{p} .*

However, there are other ways to create functions that describe ℓ . For example, consider $\vec{q} : \mathbb{R} \rightarrow \mathbb{R}^2$ where $\vec{q}(t) = 2t\vec{d}$. Again, $\ell = \text{range}(\vec{q})$ and so everything we could possibly learn about ℓ , we could learn by studying \mathbb{R} and \vec{q} . We call both \vec{p} and \vec{q} *parameterizations of ℓ by \mathbb{R} .*

Definition 3.0.1 — Parameterization. A *parameterization* of an object X by an object Y is a continuous function $p : Y \rightarrow X$ with the added conditions that p is one-to-one^a and $\text{range}(p) = X$. In this case p is called a *parameterization* and Y is called the *parameter*.

^a Sometimes we will drop the requirement that a parameterization be one-to-one, but for now we'll be strict about it.

This definition is fairly abstract, which will come in handy later. For now, we will think of X as being some curve in \mathbb{R}^n and Y as being an interval of real numbers.

■ **Example 3.1 — A Circle.** Let $\mathcal{C} \subseteq \mathbb{R}^2$ be the unit circle centered at the origin. We can parameterize \mathcal{C} by angles in $[0, 2\pi)$. Consider the function $\vec{p} : [0, 2\pi) \rightarrow \mathcal{C}$ defined by

$$\vec{p}(\theta) = \begin{bmatrix} \cos \theta \\ \sin \theta \end{bmatrix}.$$

Here, \vec{p} traces out \mathcal{C} starting at the point $(1, 0)$ and moving counter clockwise as the parameter θ increases. ■

■ **Example 3.2 — A Circle Again.** Let $\mathcal{C} \subseteq \mathbb{R}^2$ be the unit circle centered at the origin. We will parameterize \mathcal{C} by the interval $[0, 1)$. Here we might imagine that our parameter $t \in [0, 1)$ represents a point that is t -percentage around the circle.

Recall $\vec{p}(\theta) = \begin{bmatrix} \cos \theta \\ \sin \theta \end{bmatrix}$, which parameterizes \mathcal{C} based on angles.

Now, consider the function $w(t) = 2\pi t$. w inputs numbers in $[0, 1)$ and outputs angles in $[0, 2\pi)$. We should now be able to use w to parameterize \mathcal{C} in the desired way. After all, if we convert $[0, 1)$ to $[0, 2\pi)$ to \mathcal{C} , we win!

Let the parameterization $\vec{q} : [0, 1) \rightarrow \mathcal{C}$ be defined as $\vec{q} = \vec{p} \circ w$. Explicitly,

$$\vec{q}(t) = \vec{p} \circ w(t) = \vec{p}(2\pi t) = \begin{bmatrix} \sin 2\pi t \\ \cos 2\pi t \end{bmatrix}.$$

■

Exercise 3.1 Parameterize the unit circle $\mathcal{C} \subseteq \mathbb{R}^2$ by the interval $[1/2, 1)$.

Exercise 3.2 Let ℓ be the line segment connecting $(0, 0)$ and $(1, 1)$. Explain why $\vec{p} : [-1, 1] \rightarrow \ell$ given by $p(t) = (t^2, t^2)$ is *not* a parameterization.

3.1 Speed and Velocity of a Parameterization

In our day-to-day life, almost without thinking, we make a comparison between real numbers and time. Time has a forwards and backwards, which we equate to the real number's increasing and decreasing. We might even say we parameterize *time* by the real numbers. Thus, if $\vec{p} : [a, b] \rightarrow \mathcal{S}$ is a parameterization of the curve \mathcal{S} by the interval $[a, b]$, we could think of \vec{p} as describing the motion of a particle—at time $t \in [a, b]$ the particle is at $\vec{p}(t)$.

Interpreting parameterizations in this way, the *speed* of a parameterization should be the rate of change of distance with respect to time and the *velocity* of a parameterization should be the rate of change of displacement with respect to time.

Suppose $\vec{p} : [a, b] \rightarrow \mathcal{S}$ is a parameterization of \mathcal{S} and $t \in [a, b]$ represents time. The *displacement* of \vec{p} from time t to time $t + \Delta t$ is $\vec{p}(t + \Delta t) - \vec{p}(t)$ and the change in *distance* is $\|\vec{p}(t + \Delta t) - \vec{p}(t)\|$. Thus, if Δt is small, the velocity at time t can be approximated by

$$\text{velocity } \vec{p}(t) \approx \frac{\vec{p}(t + \Delta t) - \vec{p}(t)}{\Delta t}$$

and the speed¹ by

$$\text{speed } \vec{p}(t) \approx \frac{\|\vec{p}(t + \Delta t) - \vec{p}(t)\|}{|\Delta t|}.$$

Taking limits, we arrive at exact rates of change, which leads us to the following definitions.

Definition 3.1.1 — Speed. Let $\vec{p} : [a, b] \rightarrow \mathcal{S}$ be a parameterization of \mathcal{S} . The *speed* of \vec{p} at the time $t \in [a, b]$ is

$$\text{speed } \vec{p}(t) = \lim_{\Delta t \rightarrow 0} \frac{\|\vec{p}(t + \Delta t) - \vec{p}(t)\|}{|\Delta t|}.$$

Definition 3.1.2 — Velocity. Let $\vec{p} : [a, b] \rightarrow \mathcal{S}$ be a parameterization of \mathcal{S} . The *velocity* of \vec{p} at the time $t \in [a, b]$ is

$$\text{velocity } \vec{p}(t) = \lim_{\Delta t \rightarrow 0} \frac{\vec{p}(t + \Delta t) - \vec{p}(t)}{\Delta t}.$$

Both the definition of speed and the definition of velocity look a lot like the definition of the derivative. In fact, if \vec{p} were a scalar valued function, the velocity of \vec{p} would be exactly the derivative of \vec{p} . For this reason, we will define a notation similar to that of the derivative you're familiar with. From now on, the following notations mean the same thing:

$$\text{velocity } \vec{p}(t) = \vec{p}'(t) = \frac{d}{dt} \vec{p}(t) = \frac{d\vec{p}}{dt}(t).$$

Let's try to use our new definition. Let $\vec{r}(t) = \begin{bmatrix} \cos t \\ \sin t \end{bmatrix}$. Now,

$$\begin{aligned} \text{velocity } \vec{r}(t) &= \lim_{\Delta t \rightarrow 0} \frac{\begin{bmatrix} \cos(t + \Delta t) \\ \sin(t + \Delta t) \end{bmatrix} - \begin{bmatrix} \cos t \\ \sin t \end{bmatrix}}{\Delta t} \\ &= \lim_{\Delta t \rightarrow 0} \begin{bmatrix} \frac{\cos(t + \Delta t) - \cos t}{\Delta t} \\ \frac{\sin(t + \Delta t) - \sin t}{\Delta t} \end{bmatrix}. \end{aligned}$$

At this point, we should pause. We don't know how to take limits of vectors. Fortunately the rule is simple enough—to take a limit of a vector, take the limit of each of its components².

¹ Recall that speed is always positive; if a particle is moving with speed 2 and we then ran the particle back in time, it would still move at speed 2, so speed is not distance/ Δt , it is distance/ $|\Delta t|$.

² As intuitive as it sounds, this rule actually has a proof which relies on the definition of limit and the continuity of $\|\cdot\|$.

Thus we see

$$\begin{aligned}
 \text{velocity } \vec{r}(t) &= \lim_{\Delta t \rightarrow 0} \left[\frac{\cos(t + \Delta t) - \cos t}{\frac{\sin(t + \Delta t) - \sin t}{\Delta t}} \right] \\
 &= \left[\frac{\lim_{\Delta t \rightarrow 0} \frac{\cos(t + \Delta t) - \cos t}{\frac{\sin(t + \Delta t) - \sin t}{\Delta t}}}{\lim_{\Delta t \rightarrow 0} \frac{\sin(t + \Delta t) - \sin t}{\Delta t}} \right] \\
 &= \begin{bmatrix} \cos'(t) \\ \sin'(t) \end{bmatrix} = \begin{bmatrix} -\sin t \\ \cos t \end{bmatrix}.
 \end{aligned}$$

Our use of the notation $\vec{r}'(t)$ for velocity $\vec{r}(t)$ seems further justified.

Speed also appears to be a derivative. From physics, we know that speed is the magnitude of velocity. We can prove it mathematically.

Theorem 3.1.1 For a parameterization $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^n$ where velocity $\vec{p}'(t)$ exists, we have

$$\text{speed } \vec{p}(t) = \|\text{velocity } \vec{p}(t)\| = \|\vec{p}'(t)\|.$$

Proof. The proof relies on the continuity of $\|\cdot\|$. Since $\|\cdot\|$ is continuous, we may freely move limits in and out³. Thus

$$\begin{aligned}
 \text{speed } \vec{p}(t) &= \lim_{\Delta t \rightarrow 0} \frac{\|\vec{p}(t + \Delta t) - \vec{p}(t)\|}{|\Delta t|} \\
 &= \lim_{\Delta t \rightarrow 0} \left\| \frac{\vec{p}(t + \Delta t) - \vec{p}(t)}{\Delta t} \right\| \\
 &= \left\| \lim_{\Delta t \rightarrow 0} \frac{\vec{p}(t + \Delta t) - \vec{p}(t)}{\Delta t} \right\| = \|\text{velocity } \vec{p}(t)\|.
 \end{aligned}$$


■

Arc-length

Let $S \subseteq \mathbb{R}^n$ be a curve parameterized by $\vec{p} : [a, b] \rightarrow \mathbb{R}^n$. The *arc-length* of S should be the length of S if you somehow untwisted S into a straight line without stretching anything. One of the big ideas of calculus is that we can handle curvy things by chopping them up into little pieces, computing for each piece, and then adding them back together. We use the same principle to define arc-length.

In essence, we will divide our curve S into many tiny line segments, add up the lengths of those line segments and take a limit as our line segments get tinier. A parameterization provides us with a way to do this. Since parameterizations are continuous, if we chop the domain of the parameterization into tiny pieces, we will have chopped the range into tiny pieces.

³ This limit rule actually says if f is continuous and $\lim_{x \rightarrow a} g(x)$ exists, then $\lim_{x \rightarrow a} f(g(x)) = f(\lim_{x \rightarrow a} g(x))$.



To find the arc-length of a curve, we will approximate each tiny piece with a straight line segment connecting the endpoints. We then add up the lengths of all tiny segments and take a limit as our segment's length goes to zero.

Definition 3.1.3 — Arc-length. Let $S \subseteq \mathbb{R}^n$ be a curve parameterized by $\vec{p} : [a, b] \rightarrow S$. The *arc-length* of S is

$$\text{arclen } S = \lim_{\Delta t \rightarrow 0^+} \sum_{i=1}^{\frac{b-a}{\Delta t}} \|\vec{p}(a + (i-1)\Delta t) - \vec{p}(a + i\Delta t)\|.$$

There's something unsatisfying about this definition, though. We used a parameterization of S to compute the arc-length of S . But $S \subseteq \mathbb{R}^n$ is a curve regardless of whether or not it has a parameterization, and if you use a different parameterization, you should get the same arc length for S . If you're worried about this, good! You're thinking carefully! We won't show it here, but in fact no matter what parameterization you use for a curve, this definition will always produce the same arc length.

There's another reason we might be unhappy with this definition. Limits of sums are hard to compute! However, the sum involved in arc-length looks very close to a Riemann sum. If we can rewrite it exactly as a Riemann sum, we can replace it with an integral. With some superficial manipulation we see

$$\begin{aligned} \text{arclen } S &= \lim_{\Delta t \rightarrow 0^+} \sum_{i=1}^{\frac{b-a}{\Delta t}} \|\vec{p}(a + (i-1)\Delta t) - \vec{p}(a + i\Delta t)\| \\ &= \lim_{\Delta t \rightarrow 0^+} \sum_{i=1}^{\frac{b-a}{\Delta t}} \frac{\|\vec{p}(a + (i-1)\Delta t) - \vec{p}(a + i\Delta t)\|}{\Delta t} \Delta t \\ &= \int_a^b \text{speed } \vec{p}(t) dt. \end{aligned}$$

For the last equality, we noticed that $\lim_{\Delta t \rightarrow 0^+} \frac{\|\vec{p}(t) - \vec{p}(t + \Delta t)\|}{\Delta t} = \text{speed } \vec{p}(t)$, which involved switching a limit and an infinite sum. In order to do this rigorously, we need a mathematical

proof that it is logically valid. Such a proof is the subject of a course in *real analysis*, and won't be covered here, but it's always good to keep track of what you've actually proved and what you've been told is true⁴.

Speed is easy to calculate, and we have a better handle on calculating integrals than we do limits of sums, so now we have a chance of calculating arc-length.

■ **Example 3.3** We shall find the length of the parabola with equation $y = x^2$ on the interval $-1 \leq x \leq 1$. A parametric representation of the parabola is $\vec{p}(t) = (t, t^2)$ where $-1 \leq t \leq 1$. Now,

$$\frac{d\vec{p}}{dt} = \begin{bmatrix} 1 \\ 2t \end{bmatrix},$$

so speed $\vec{p}(t) = \|\vec{p}'(t)\| = \sqrt{1 + 4t^2}$. Hence,

$$L = \int_{-1}^1 \sqrt{1 + 4t^2} dt = \sqrt{5} + \frac{1}{4} \ln \frac{\sqrt{5} + 2}{\sqrt{5} - 2}.$$

■


As you can see from Example 3.3, the integrals involved in computing arc length can be difficult. In fact, most of them don't have an elementary form, which means in the real world we often use a computer to approximate arc length directly from the Riemann sum rather than calculate it exactly.

Exercises for 3.1


⁴ In order to prove that swapping the limit and sum is valid, we actually need extra assumptions on \vec{p} . If we make \vec{p} *differentiable* rather than merely *continuous*, we can prove that the swap is valid.

3.2 Arc-length Parameterization

Recall a parameterization is a relation between two objects. If a curve $\mathcal{S} \subseteq \mathbb{R}^n$ is parameterized by \mathbb{R} , it means there is a continuous, one-to-one function $\vec{p} : \mathbb{R} \rightarrow \mathcal{S}$. This function can be thought of as a map from \mathbb{R} to \mathcal{S} . Any interval $[a, b] \subseteq \mathbb{R}$ corresponds to a segment $\vec{p}([a, b]) \subseteq \mathcal{S}$.



Alternatively, we may think of $\vec{p} : \mathbb{R} \rightarrow \mathcal{S}$ as a function that picks up the real line, stretches, twists, and warps it, and sticks it into \mathbb{R}^n in the shape of \mathcal{S} . In this sense, not all parameterizations are created equally. Some significantly stretch and warp and others barely do at all.



The least stretchy type of parameterization is called an *arc-length parameterization*.

Before we define arc-length parameterization, let's introduce some notation. If $\mathcal{S} \subseteq \mathbb{R}^n$ is a curve parameterized by $\vec{p} : \mathbb{R} \rightarrow \mathcal{S}$, then

$$\text{arclen } \vec{p} \Big|_a^b = \text{arc length of } \mathcal{S} \text{ between } \vec{p}(a) \text{ and } \vec{p}(b).$$

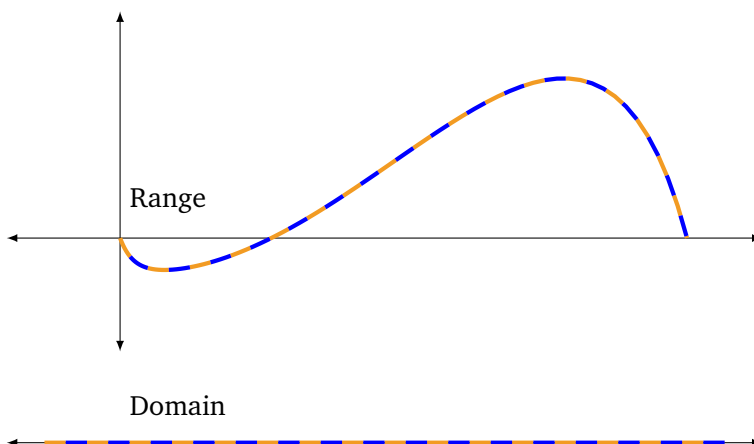
We might read $\text{arclen } \vec{p} \Big|_a^b$ as the “arc length of the curve traced by $\vec{p}(t)$ from $t = a$ to $t = b$.” Using notion for the *image* of a set, we can also write

$$\text{arclen } \vec{p} \Big|_a^b = \text{arclen } \vec{p}([a, b]).$$

Definition 3.2.1 Let $S \subseteq \mathbb{R}^n$ be a curve and $\vec{p} : \mathbb{R} \rightarrow S$ be a parameterization. The parameterization \vec{p} is called an *arc-length parameterization* if for $b \geq a$,

$$\text{arclen } \vec{p} \Big|_a^b = b - a.$$

In plain language, if \vec{p} is an arc-length parameterization, then the distance traveled in parameter space is the same as the distance traveled along the curve.



■ **Example 3.4** Find an arc-length parameterization of the line ℓ which passes through the origin and has a direction vector $\vec{d} = \hat{x} + 2\hat{y} + \hat{z}$.

Let's name our parameterization \vec{p} and see if we can build up a formula for \vec{p} . Since ℓ passes through the origin, let's have $\vec{p}(0) = \vec{0}$. Now, $\vec{p}(1)$ has to be a point on ℓ that is distance 1 from $\vec{0}$. There are two such points, so we'll arbitrarily declare $\vec{p}(1) = \frac{1}{\sqrt{6}}(1, 2, 1)$. The point $\vec{p}(2)$ must be distance 1 from $\vec{p}(1)$ and distance 2 from $\vec{p}(0)$. Since parameterizations aren't allowed to double-back, we only have one choice. Namely, $\vec{p}(2) = \frac{2}{\sqrt{6}}(1, 2, 1)$. Noticing the pattern, let's guess

$$\vec{p}(t) = \frac{t}{\sqrt{6}} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}.$$

Now we'll verify that \vec{p} is an arc-length parameterization. Since ℓ is a straight line this is easy:

$$\text{arclen } \vec{p} \Big|_a^b = \|\vec{p}(b) - \vec{p}(a)\| = \left\| \frac{b-a}{\sqrt{6}} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \right\| = |b-a|,$$

and so \vec{p} is an arc-length parameterization. ■

Arc-length parameterizations can also be characterized by their speed. For an arc-length parameterization, the distance traveled in parameter space must be equal to the distance traveled along the curve. Therefore, an arc-length parameterization must also move at speed 1.

Theorem 3.2.1 Let $\mathcal{S} \subseteq \mathbb{R}^n$ be a curve and $\vec{p} : \mathbb{R} \rightarrow \mathcal{S}$ a parameterization. The parameterization \vec{p} is an arc-length parameterization if and only if speed $\vec{p} = 1$.

Proof. Suppose \vec{p} is a parameterization of \mathcal{S} and speed $\vec{p} = 1$. Then

$$\text{arclen } \vec{p} \Big|_a^b = \int_a^b \text{speed } \vec{p}(t) dt = \int_a^b 1 dt = b - a.$$

Now, suppose that \vec{p} satisfies $\text{arclen } \vec{p} \Big|_a^b = b - a$. We then know if $\Delta t > 0$ is small

$$\frac{\|\vec{p}(t + \Delta t) - \vec{p}(t)\|}{\Delta t} \approx \frac{\text{arclen } \vec{p} \Big|_t^{t+\Delta t}}{\Delta t} = \frac{(t + \Delta t) - t}{\Delta t} = 1.$$

After taking a limit as $\Delta t \rightarrow 0$, “ \approx ” will turn into “ $=$ ”.

To make this argument completely rigorous, we will have to do a little bit of mathematical gymnastics. Since the shortest distance between two points is a straight line, what we really know is

$$\|\vec{p}(t + \Delta t) - \vec{p}(t)\| \leq \text{arclen } \vec{p} \Big|_t^{t+\Delta t} = \Delta t$$

for all Δt . Now, taking a limit as $\Delta t \rightarrow 0$, we deduce speed $\vec{p}(t) \leq 1$. But,

$$b - a = \text{arclen } \vec{p} \Big|_a^b = \int_a^b \text{speed } \vec{p}(t) dt \leq \int_a^b 1 dt = b - a,$$

and so if speed \vec{p} is a continuous function, speed $\vec{p}(t) = 1$ for all t . Even if speed is not continuous, we can argue that speed \vec{p} is *essentially* 1.⁵ ■

Now we have a new way of thinking about arc-length parameterizations and a new way of creating them. Instead of working with arc length directly, we can attempt to manipulate *speed*.

■ **Example 3.5** Find an arc-length parameterization of \mathcal{C} , the circle of radius 2 centered at the origin.

We already know a parameterization of \mathcal{C} . Namely,

$$\vec{r}(t) = \begin{bmatrix} 2 \cos t \\ 2 \sin t \end{bmatrix}.$$

However, speed $\vec{r}(t) = \|\vec{r}'(t)\| = \sqrt{(-2 \sin t)^2 + (2 \cos t)^2} = 2$. If we could somehow *slow down* time, we could slow down the speed to be 1, giving an arc-length parameterization.

⁵ Here, the word *essentially* is a technical term coming from real analysis.

Let $w(t) = t/2$. The function w stretches time by a factor of 2. Define

$$\vec{p}(t) = \vec{r} \circ w(t) = \begin{bmatrix} 2 \cos t/2 \\ 2 \sin t/2 \end{bmatrix}.$$

Now,

$$\begin{aligned} \text{speed } \vec{p}(t) &= \|\vec{p}'(t)\| = \|(\vec{r} \circ w)'(t)\| = \|(\vec{r}' \circ w(t))w'(t)\| \\ &= |w'(t)| \|\vec{r}' \circ w(t)\| = \frac{1}{2} 2 = 1. \end{aligned}$$

In this computation we made judicious use of the chain rule, however we could have computed directly from our formula for \vec{p} . Now, since $\text{speed } \vec{p}(t) = 1$, the function \vec{p} is an arc-length parameterization. ■

Exercise 3.3 Let \mathcal{C} and \vec{r} be as in Example 3.5. The function $\vec{q}(t) = \frac{1}{2}\vec{r}(t)$ has speed 1 and so is an arc-length parameterization of something. Explain why \vec{q} is *not* an arc-length parameterization of \mathcal{C} .

In Example 3.5, we adjusted the speed of a parameterization by warping time. Suppose $S \subseteq \mathbb{R}^n$ is a curve parameterized by $\vec{p} : \mathbb{R} \rightarrow S$ and let $w : \mathbb{R} \rightarrow \mathbb{R}$ be a parameterization of \mathbb{R} . In other words, w stretches or squishes (or flips) \mathbb{R} by varying amounts. Now consider

$$\vec{r} = \vec{p} \circ w.$$

The function \vec{r} has domain \mathbb{R} and range S . Further, since \vec{p} and w are both one-to-one and continuous, we know \vec{r} is one-to-one and continuous. Thus, \vec{r} is a parameterization of S . And, by using the chain rule,

$$\vec{r}'(t) = (\vec{p} \circ w)'(t) = w'(t) [\vec{p}' \circ w(t)],$$

so


$$\text{speed } \vec{r} \text{ at time } t = |w'(t)| [\text{speed } \vec{p} \text{ at time } w(t)].$$

It is important to note that we must compose \vec{p} and w in order to get a parameterization of S . See Exercise 3.3 for an example of why you cannot multiply \vec{p} and w .

With the idea of stretching time in the back of our minds, let's work through a hypothetical example.

Suppose $S \subseteq \mathbb{R}^n$ is a curve parameterized by $\vec{p} : \mathbb{R} \rightarrow S$ and we've computed the following table of values.

t	$\left. \text{arclen } \vec{p} \right _0^t$
0	0
1	3
2	3.5
3	6
4	7



We'd like to find a time-stretching function $w : \mathbb{R} \rightarrow \mathbb{R}$ so that $\vec{r} = \vec{p} \circ w$ is an arc-length parameterization of S . In other words, we need

$$\text{arclen } \vec{r} \Big|_0^t = \text{arclen}(\vec{p} \circ w) \Big|_0^t = t.$$

In particular, $\text{arclen } \vec{r} \Big|_0^0 = 0$ (this we get for free) and $\text{arclen } \vec{r} \Big|_0^3 = 3$. Looking at the table, we see that the arc length from $\vec{p}(0)$ to $\vec{p}(1)$ is 3, and so we need $w(3) = 1$. This way $\vec{r}(3) = \vec{p}(w(3)) = \vec{p}(1)$. Continuing in this manner, we get the following table for w .

t	$w(t)$
0	0
3	1
3.5	2
6	3
7	4

The function w is just the inverse of the function $\text{arclen } \vec{p} \Big|_0^t$! This also makes sense from a purely algebraic perspective. Consider

$$x = \text{arclen } \vec{r} \Big|_0^x = \text{arclen}(\vec{p} \circ w) \Big|_0^x = \text{arclen}(\vec{p}) \Big|_0^{w(x)}.$$

Replacing x with $w^{-1}(t)$, we see

$$w^{-1}(t) = \text{arclen } \vec{p} \Big|_0^{w \circ w^{-1}(t)} = \text{arclen } \vec{p} \Big|_0^t.$$

This means the inverse of w is $\text{arclen } \vec{p} \Big|_0^t$ and so the inverse of $\text{arclen } \vec{p} \Big|_0^t$ must be w .⁶

We now have a concrete way to find an arc-length parameterization.

■ **Example 3.6** Let \mathcal{R} be the ray parameterized by $\vec{p} : [0, \infty) \rightarrow \mathbb{R}^2$ where $\vec{p}(t) = (t^2, 2t^2)$. Find an arc-length parameterization of \mathcal{R} .

⁶ Recall that for an invertible function f , we have $(f^{-1})^{-1} = f$.

We'll start by finding the arc-length function for \vec{p} .

$$\begin{aligned} a(t) = \text{arclen } \vec{p} \Big|_0^t &= \int_0^t \|\vec{p}'(x)\| dx \\ &= \int_0^t \sqrt{(2x)^2 + (4x)^2} dx = t^2 \sqrt{5}. \end{aligned}$$

Now define $w(t) = a^{-1}(t)$. Since $t \geq 0$, the inverse of a is well defined and is given by

$$w(t) = a^{-1}(t) = \sqrt{t/\sqrt{5}}.$$

Now, define $\vec{r}(t) = \vec{p} \circ w(t) = (t/\sqrt{5}, 2t/\sqrt{5})$. The parameterization \vec{r} is an arc-length parameterization!

Of course we didn't need to go through all this work in this case. If all we wanted was to find an arc-length parameterization of \mathcal{R} we could create one directly using geometry, since we know how to parameterize lines and rays. ■

Explicit Arc-length Parameterizations

The idea of arc-length parameterization is very important. However, for most curves, we're hopeless in finding a formula for the arc-length parameterization. We can for a hand full of

⁷ If you don't believe me, go ahead and try to find the inverse of the function $f(x) = xe^x$.

curves, like a circle, a line, a helix, but even something as simple as an ellipse cannot be arc-length parameterized with elementary functions.

Why is it so hard? Well, integrals are hard in general. Most formulas cannot be integrated in closed form. Integrals involving square roots are even harder to evaluate. And, even if you manage to integrate to find the arc-length function, you still have to invert that function⁷. So, if finding formulas for arc-length parameterizations is so hard, why do we bother with them at all? The answer is that the *idea* of an arc-length parameterization is incredibly useful. Its mere existence will aid our thinking. And, in many of the problems we will be solving, the arc-length parameterization will somehow get canceled out and we won't ever need to find a formula for it.

Exercises for 3.2

3.3 Acceleration and Curvature

In the Newtonian mechanics of one-dimensional motion, acceleration is the second derivative of position with respect to time. In \mathbb{R}^n , we define it in the same way.

Definition 3.3.1 — Acceleration. Let $\vec{p} : \mathbb{R} \rightarrow S$ be a parameterization of S . The *acceleration* of \vec{p} is

$$\text{accel } \vec{p}(t) = (\text{velocity } \vec{p})'(t) = \vec{p}''(t).$$


Just like velocity, acceleration is a *vector*. Let's consider two examples. Define

$$\vec{l}(t) = \begin{bmatrix} t^2 \\ t^2 \end{bmatrix} \quad \text{and} \quad \vec{c}(t) = \begin{bmatrix} \cos t \\ \sin t \end{bmatrix}.$$

Here \vec{l} parameterizes a ray and \vec{c} a circle. Further, \vec{l} traces along the ray faster and faster, whereas \vec{c} has constant speed as it traces the circle. We compute

$$\text{accel } \vec{l}(t) = \begin{bmatrix} 2 \\ 2 \end{bmatrix} \quad \text{and} \quad \text{accel } \vec{c}(t) = \begin{bmatrix} -\cos t \\ -\sin t \end{bmatrix},$$

and see that the acceleration of \vec{l} is constant whereas the acceleration of \vec{c} is not. Further, $\|\text{accel } \vec{l}(t)\| = 2\sqrt{2}$ and $\|\text{accel } \vec{c}(t)\| = 1$, and so the magnitude of the acceleration of both \vec{l} and \vec{c} is constant.



In the past, you might have distinguished linear acceleration (running faster and faster along a straight line) from centripetal acceleration (the acceleration you experience by moving at a constant speed around a circle). With vectors, these two types of acceleration are unified.

In the previous example, \vec{l} had purely linear acceleration and the acceleration vector pointed tangent to the path \vec{l} traces. In contrast, \vec{c} has an acceleration vector that is orthogonal

to the curve it traces (its acceleration is purely centripetal). What happens if we mix the two types of acceleration?

Consider

$$\vec{r}(t) = \begin{bmatrix} \cos t^2 \\ \sin t^2 \end{bmatrix}.$$

Computing,

$$\text{accel } \vec{r}(t) = \begin{bmatrix} -2(\sin t^2 + 2t^2 \cos t^2) \\ 2(\cos t^2 - 2t^2 \sin t^2) \end{bmatrix}.$$

There is no clear relationship between the curve \vec{r} traces and the acceleration vector for \vec{r} . To see this relationship, we need to decompose $\text{accel } \vec{r}$ into its tangential and normal components.

Definition 3.3.2 — Tangential and Normal Acceleration. Let $\vec{p} : \mathbb{R} \rightarrow S$ be a parameterization of S . Then $\text{accel } \vec{p}$ can be written as

$$\text{accel } \vec{p}(t) = \vec{a}_T(t) + \vec{a}_N(t)$$

where $\vec{a}_T(t)$ is tangent to S at the point $\vec{p}(t)$ and $\vec{a}_N(t)$ is orthogonal to S at the point $\vec{p}(t)$. In this case, $\vec{a}_T(t)$ is called the *tangential component of the acceleration* and $\vec{a}_N(t)$ is called the *normal component of the acceleration* of \vec{p} .

■ **Example 3.7** Let $\vec{r}(t) = (\cos t^2, \sin t^2)$. Find the tangential and normal components of the acceleration of \vec{r} .

Earlier we computed

$$\text{accel } \vec{r}(t) = \begin{bmatrix} -2(\sin t^2 + 2t^2 \cos t^2) \\ 2(\cos t^2 - 2t^2 \sin t^2) \end{bmatrix}.$$

We can use projections to split $\text{accel } \vec{r}$ into its tangential and normal components.

Recall $\vec{r}'(t)$ is tangent to the curve \vec{r} traces at the point $\vec{r}(t)$. Thus,

$$\vec{a}_T = \text{proj}_{\vec{r}'(t)} \text{accel } \vec{r}(t) = \begin{bmatrix} -2 \sin t^2 \\ 2 \cos t^2 \end{bmatrix},$$

and

$$\vec{a}_N = \text{accel } \vec{r}(t) - \vec{a}_T(t) = \begin{bmatrix} -4t^2 \cos t^2 \\ -4t^2 \sin t^2 \end{bmatrix}.$$

■

Changing gears, suppose that $\vec{q} : \mathbb{R} \rightarrow S$ is an arc-length parameterization of S . Since the speed of \vec{q} is constant, it is intuitive that the tangential component of the acceleration of \vec{q} is zero. Equipped with our knowledge of vectors, it won't be so hard to prove our intuition. But, it will be helpful to establish the product rule for dot products.

Exercise 3.4 Let $\vec{a}(t) = (a_x(t), a_y(t), a_z(t))$ and $\vec{b}(t) = (b_x(t), b_y(t), b_z(t))$ be parameter-

izations. Establish the *product rule for dot products*. That is, show that

$$\left[\vec{a}(t) \cdot \vec{b}(t) \right]' = \vec{a}'(t) \cdot \vec{b}(t) + \vec{a}(t) \cdot \vec{b}'(t).$$

Theorem 3.3.1 If $\vec{p} : \mathbb{R} \rightarrow \mathcal{S}$ is an arc-length parameterization of \mathcal{S} then $\text{accel} \vec{p}$ is always orthogonal to \mathcal{S} . Equivalently,

$$\vec{p}''(t) \cdot \vec{p}'(t) = 0.$$

Proof. Since \vec{p} is an arc-length parameterization, we know

$$\sqrt{\vec{p}'(t) \cdot \vec{p}'(t)} = \|\vec{p}'(t)\| = 1.$$

Squaring both sides, we get the relationship

$$\vec{p}'(t) \cdot \vec{p}'(t) = 1. \quad (3.1)$$

Now, we may take the derivative of both sides of Equation (3.1) and apply the product rule for dot products to find

$$0 = [\vec{p}' \cdot \vec{p}']'(t) = \vec{p}''(t) \cdot \vec{p}'(t) + \vec{p}'(t) \cdot \vec{p}''(t) = 2\vec{p}''(t) \cdot \vec{p}'(t),$$

and so $\vec{p}''(t)$ and $\vec{p}'(t)$ are orthogonal. ■

If you examine the proof of Theorem 3.3.1 closely, you'll notice that we didn't actually need the speed of our parameterization to be 1. The proof works just as well if the speed is some other constant.

Curvature


For any curve \mathcal{S} , there are infinitely many choices of parameterizations, but in some sense, there is only one arc-length parameterization. An arc-length parameterization of \mathcal{S} is uniquely determined by a direction (forwards or backwards along \mathcal{S}) and a starting position. Thus, we might think of an arc-length parameterization as *intrinsic* to a curve.

The *curvature* of a curve is a measure of how sharply a curve bends or twists. Curvature is a property of a curve—you don't need a parameterization to define curvature—but it is a lot easier to define with reference to an arc-length parameterization.

Definition 3.3.3 — Curvature. Let $\mathcal{S} \subseteq \mathbb{R}^n$ be a curve and let $\vec{p} : \mathbb{R} \rightarrow \mathcal{S}$ be an arc-length parameterization of \mathcal{S} . The *curvature* of \mathcal{S} at the point $\vec{p}(t)$ is

$$\|\text{accel} \vec{p}(t)\| = \|\vec{p}''(t)\|.$$

This definition of curvature can be made intuitive. If $\vec{p} : \mathbb{R} \rightarrow \mathcal{S}$ is an arc-length parameterization, all velocity vectors are unit length. Therefore, all acceleration of \vec{p} must come from the velocity vectors changing direction (and not changing length). If a curve has a sharp bend (high curvature), the velocity vectors will rapidly change direction. If a curve is generally flat (low curvature), the velocity vectors hardly change direction at all.



Curvature can be hard to calculate exactly, but it isn't so hard to eyeball.

■ **Example 3.8** Estimate the curvature the parabola $y = x^2$ at the point $(0, 0)$.

Let $\vec{p}(t) = (t, t^2)$ be a parameterization of our parabola and let \vec{q} be an arc-length parameterization of our parabola satisfying $\vec{p}(0) = \vec{q}(0) = (0, 0)$. We can check that \vec{p} is *not* an arc-length parameterization. However, \vec{p} and \vec{q} have a lot in common.

For example, if $\vec{p}(t) = \vec{q}(s)$, then $\vec{p}'(t)$ is parallel to $\vec{q}'(s)$. Further $\vec{q}'(s)$ is a unit vector, so

$$\vec{q}'(s) = \frac{\vec{p}'(t)}{\|\vec{p}'(t)\|},$$

provided that \vec{p} and \vec{q} trace the parabola in the same direction.

We will estimate the curvature at $(0, 0)$ by first estimating \vec{q}' at two close-by points and then using those estimates to estimate $\vec{q}''(0)$. Consider $\vec{p}(0) = (0, 0)$ and $\vec{p}(0.1) = (0.1, 0.01)$. The approximate arc-length along the parabola from $\vec{p}(0)$ to $\vec{p}(0.1)$ is

$$\text{arc-length from } \vec{p}(0) \text{ to } \vec{p}(0.1) \approx \|\vec{p}(0.1) - \vec{p}(0)\| = \sqrt{0.0101} \approx 0.1005.$$

Thus, $\vec{p}(0.1) \approx \vec{q}(0.1005)$. Using the fact that \vec{q}' is always a unit vector, we compute

$$\vec{q}'(0) = \frac{\vec{p}'(0)}{\|\vec{p}'(0)\|} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \text{and} \quad \vec{q}'(0.1005) \approx \frac{\vec{p}'(0.1)}{\|\vec{p}'(0.1)\|} \approx \begin{bmatrix} 0.9806 \\ 0.1961 \end{bmatrix}.$$

Finally, we use another numerical approximation to get


$$\vec{q}''(0) \approx \frac{\vec{q}'(0.1005) - \vec{q}'(0)}{0.1005} \approx \begin{bmatrix} -0.1932 \\ 1.9514 \end{bmatrix},$$

and so the curvature of the parabola at $(0, 0)$ is approximately $\|\vec{q}''(0)\| \approx 1.96$, which is quite close to the true value of 2. ■


Exercises for 3.3

3.4 Line Integrals

Suppose that you're walking around in a dome-shaped circus tent. The circus tent has a radius of 3 and if you're standing at the coordinates (x, y) , the height of the circus tent is given by $f(x, y) = 9 - x^2 - y^2$.



Now, suppose you want to measure the area above the ground and below the circus tent as you walk from $(-3, 0)$ to $(3, 0)$ in a straight line (perhaps you're trying to figure out how many square feet of streamers you need to form a wall in the circus tent). We can compute this quantity in a familiar way.



Let \mathcal{X} be the straight-line path from $(-3, 0)$ to $(3, 0)$. Since the y -coordinate along this path is always 0, the height of the tent along this path is given by $f(x, 0) = 9 - x^2$, and the area question simplifies to “what is the area above the x -axis and below the graph $z = 9 - x^2$ from $x = -3$ to $x = 3$?” This is a familiar calculus question and is answered with an integral:

$$\text{area} = \int_{-3}^3 (9 - x^2) dx = 36.$$

Similarly, we could ask how much area is below the tent and above the straight line segment \mathcal{Y} from $(0, -3)$ to $(0, 3)$. This is answered with an analogous integral:

$$\text{area} = \int_{-3}^3 9 - y^2 \, dy = 36.$$

Okay, so how about walking diagonally along the straight line segment \mathcal{D} from $(-3/\sqrt{2}, -3/\sqrt{2})$ to $(3/\sqrt{2}, 3/\sqrt{2})$? Since the tent is radially symmetric, we know this area should be 36, but it is unclear how to compute it with an integral—both coordinates are changing, so the trick of substituting one variable for a constant won't work. To rephrase this question as an integral, we need the concept of a *line integral*.


Definition 3.4.1 — Line Integral. Let $\mathcal{C} \subseteq \mathbb{R}^n$ be a curve and let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function. The *line integral* of f along the curve \mathcal{C} is the signed area above \mathcal{C} and below the surface defined by the graph of f . The line integral of f along \mathcal{C} is written

$$\int_{\mathcal{C}} f \quad \text{or} \quad \int_{\mathcal{C}} f \, ds.$$

Using the notation of a line integral, the area below the tent and above the line segment \mathcal{X} is $\int_{\mathcal{X}} f = 36$, the area below the tent and above \mathcal{Y} is $\int_{\mathcal{Y}} f = 36$, and the area below the tent and above \mathcal{D} is $\int_{\mathcal{D}} f$. However, we seemed to have side-stepped the issue. We've defined a notation, but how do you evaluate $\int_{\mathcal{D}} f$? For this, we need to think about our path and the Riemann sums that integrals come from.


■ **Example 3.9** Let $f(x, y) = 9 - x^2 - y^2$ and let \mathcal{D} be the straight line segment from $(-3/\sqrt{2}, -3/\sqrt{2})$ to $(3/\sqrt{2}, 3/\sqrt{2})$. Compute $\int_{\mathcal{D}} f$.

Let \mathcal{R} be the region below the graph of f and above the curve \mathcal{D} .



We will find the area of \mathcal{R} by dividing \mathcal{R} up into tiny rectangles, summing the area of these tiny rectangles, and taking a limit as the width of each rectangle goes to zero. But how should we create these rectangles? With a parameterization, of course!

\mathcal{D} is parameterized by the function $\vec{r}(t) = \begin{bmatrix} t \\ t \end{bmatrix}$ where $-\frac{3}{\sqrt{2}} \leq t \leq \frac{3}{\sqrt{2}}$. Let's find the area of the rectangle below the graph of f and above the line segment from $\vec{r}(t)$ to $\vec{r}(t + \Delta t)$.



At position $\vec{r}(t)$, the height of f is $f \circ \vec{r}(t) = 9 - t^2 - t^2 = 9 - 2t^2$. And, if we move from $\vec{r}(t)$ to $\vec{r}(t + \Delta t)$, we've moved by

$$\|\vec{r}(t + \Delta t) - \vec{r}(t)\| = \left\| \begin{bmatrix} \Delta t \\ \Delta t \end{bmatrix} \right\| = \sqrt{2}\Delta t.$$

Therefore, the area of this rectangle is $(9 - 2t^2)\sqrt{2}\Delta t$. Summing the area of these rectangles gives us

$$\text{total area} \approx \sum_{t=-\frac{3}{\sqrt{2}}, -\frac{3}{\sqrt{2}}+\Delta t, \dots, \frac{3}{\sqrt{2}}} (9 - 2t^2)\sqrt{2}\Delta t,$$

which looks strikingly like a Riemann sum. After squinting for a while and taking a limit, we see

$$\text{total area} = \lim_{\Delta t \rightarrow 0} \sum_{t=-\frac{3}{\sqrt{2}}, -\frac{3}{\sqrt{2}}+\Delta t, \dots, \frac{3}{\sqrt{2}}} (9 - 2t^2)\sqrt{2}\Delta t = \int_{-3/\sqrt{2}}^{3/\sqrt{2}} (9 - 2t^2)\sqrt{2} dt = 36.$$

■

In Example 3.9, we used a parameterization of \mathcal{D} to compute $\int_{\mathcal{D}} f$, but there are infinitely many ways to parameterize \mathcal{D} . Of course $\int_{\mathcal{D}} f$ does not depend on any particular parameterization of \mathcal{D} —it depends only on \mathcal{D} and f .

General Line Integrals


Now that we have an example to ground us, let's tackle line integrals in greater generality. Suppose that $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and \mathcal{C} is some bounded curve parameterized by a function $\vec{r} : [0, 1] \rightarrow \mathbb{R}^n$. We will attempt to compute $\int_{\mathcal{C}} f$ in the same way we did in Example 3.9.

First, we find the area of each rectangle of concern. At the point $\vec{r}(t) \in \mathcal{C}$, the height of f is $f \circ \vec{r}(t)$, and, as we traverse from $\vec{r}(t)$ to $\vec{r}(t + \Delta t)$, we've traveled a distance of $\|\vec{r}(t + \Delta t) - \vec{r}(t)\|$. Thus,

$$\text{area of rectangle} = (f \circ \vec{r}(t)) \|\vec{r}(t + \Delta t) - \vec{r}(t)\|,$$

and so

$$\text{total area} \approx \sum_{t=0, \Delta t, 2\Delta t, \dots, 1} (f \circ \vec{r}(t)) \|\vec{r}(t + \Delta t) - \vec{r}(t)\|. \quad (3.2)$$



Equation (3.2) doesn't yet look like a Riemann sum that we could convert into an integral. It's not multiplied by Δt ! But, we can fix that by multiplying by $\frac{\Delta t}{\Delta t}$. Now we have

$$\text{total area} \approx \sum_{t=0, \Delta t, 2\Delta t, \dots, 1} (f \circ \vec{r}(t)) \frac{\|\vec{r}(t + \Delta t) - \vec{r}(t)\|}{\Delta t} \Delta t.$$

We're all set to take a limit. But wait, $\lim_{\Delta t \rightarrow 0} \frac{\|\vec{r}(t + \Delta t) - \vec{r}(t)\|}{\Delta t} = \|\vec{r}'(t)\|$, and so⁸

$$\text{total area} = \lim_{\Delta t \rightarrow 0} \sum_{t=0, \Delta t, 2\Delta t, \dots, 1} (f \circ \vec{r}(t)) \frac{\|\vec{r}(t + \Delta t) - \vec{r}(t)\|}{\Delta t} \Delta t = \int_0^1 (f \circ \vec{r}(t)) \|\vec{r}'(t)\| dt.$$

We can encapsulate all this hard work into a theorem.

Theorem 3.4.1 Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and let \mathcal{C} be a bounded curve with finite arc-length in \mathbb{R}^n parameterized by the differentiable function $\vec{r} : [a, b] \rightarrow \mathcal{C}$. Then

$$\int_{\mathcal{C}} f = \int_a^b (f \circ \vec{r}(t)) \|\vec{r}'(t)\| dt.$$

⁸ We exchanged a limit and an infinite sum when we replaced $\|\vec{r}(t + \Delta t) - \vec{r}(t)\|/\Delta t$ with $\|\vec{r}'(t)\|$. This isn't always allowed, but if \vec{r} is differentiable and \mathcal{C} has finite arc-length, it's okay.

Now that we have Theorem 3.4.1, we can interpret it another way. In the integrand, $f \circ \vec{r}(t)$ can still be thought of as the height of f at the point $\vec{r}(t)$. The quantity $\|\vec{r}'(t)\|$ is the *speed* of \vec{r} as it traces through the point $\vec{r}(t)$, all so $\|\vec{r}'(t)\|\Delta t$ is of the form *speed* \times *time*. Thus, we may interpret $\|\vec{r}'(t)\|\Delta t$ as the approximate amount of distance that \vec{r} traverses between the times t and $t + \Delta t$. Therefore, $(f \circ \vec{r}(t))\|\vec{r}'(t)\|\Delta t$ could be thought of as an approximate area of a rectangle. Again, adding $(f \circ \vec{r}(t))\|\vec{r}'(t)\|\Delta t$ up over all relevant t values and taking a limit gives the same integral as in Theorem 3.4.1.

■ **Example 3.10** Let $f(x, y) = 9 - x^2 - y^2$ and let \mathcal{D} be the straight line segment from $(-3/\sqrt{2}, -3/\sqrt{2})$ to $(3/\sqrt{2}, 3/\sqrt{2})$. Compute $\int_{\mathcal{D}} f$.

This is the same example as before, but now we have a different perspective. We can parameterize \mathcal{D} by the function $\vec{r} : [-3/\sqrt{2}, 3/\sqrt{2}] \rightarrow \mathcal{D}$ given by the formula $\vec{r}(t) = \begin{bmatrix} t \\ t \end{bmatrix}$.

Computing, $\|\vec{r}'(t)\| = \sqrt{2}$ regardless of t , and so by Theorem 3.4.1 we have

$$\int_{\mathcal{D}} f = \int_{-3/\sqrt{2}}^{3/\sqrt{2}} f \circ \vec{r}(t) \sqrt{2} dt = \int_{-3/\sqrt{2}}^{3/\sqrt{2}} (9 - 2t^2) \sqrt{2} dt = 36.$$

■

The Vector Line Integral and Work

A common physics problem is to compute the total amount of work done moving between two places while subject to changing forces. This, of course, is a line integral in disguise.

Suppose that $\vec{F} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ is a force field. That is, if you stand at the point $\vec{x} \in \mathbb{R}^2$, you're subjected to the force $\vec{F}(\vec{x})$. Now, if you move from point A to point B along the curve \mathcal{C} , you will be subjected to different forces at different points along your journey.

XXX Figure

Let $\vec{r} : [0, 1] \rightarrow \mathcal{C}$ parameterize \mathcal{C} , and recall that the work done by a constant force \vec{A} to during a displacement by \vec{D} is $\vec{A} \cdot \vec{D}$ (which could be negative). Our force, \vec{F} , is changing, but if we displace by a tiny amount, we can think of \vec{F} as constant.

Let $\Delta\vec{r}(t) = \vec{r}(t + \Delta t) - \vec{r}(t)$ be a tiny displacement. We can approximate the work done by \vec{F} from time t to $t + \Delta t$ by $(\vec{F} \circ \vec{r}(t)) \cdot \Delta\vec{r}(t)$. Therefore,

$$\text{Work along } \mathcal{C} \approx \sum_{\text{relevant } t} (\vec{F} \circ \vec{r}(t)) \cdot \Delta\vec{r}(t).$$

That doesn't quite look like a Riemann sum yet, so we'll again multiply by $\frac{\Delta t}{\Delta t}$ and notice that $\lim_{\Delta t \rightarrow 0} \frac{\Delta\vec{r}(t)}{\Delta t} = \vec{r}'(t)$. Therefore,

$$\text{Work along } \mathcal{C} = \lim_{\Delta t \rightarrow 0} \sum_{\text{relevant } t} (\vec{F} \circ \vec{r}(t)) \cdot \vec{r}'(t) \Delta t = \int_0^1 (\vec{F} \circ \vec{r}(t)) \cdot \vec{r}'(t) dt.$$

Whereas the line integral from Theorem 3.4.1 is often called the *scalar* line integral, the line integral we just derived is called the *vector* line integral. The key idea in both cases is to take a path, chop it up into little pieces via a parameterization, and then take a limit to recognize the result as an integral you're familiar with.

Line Integral Notations

Different disciplines and fields have developed a handful of notations for line integrals. The notation $\int_C f$ is standard in high-level math. Physicists often prefer $\int_C f \, ds$, where the added ds stands for “with respect to arc-length.” The notation $\oint_C f$ is sometimes used if C is a closed curve (one that starts and ends in the same position). Physicists also have a special notation for vector line integrals. The expression

$$\int_C \vec{F} \cdot d\vec{r}$$

is often used to notate the vector line integral of \vec{F} over C . Here $d\vec{r}$ isn’t related to the name of any particular parameterization, but just stands for “an infinitesimal displacement along the curve C .”

Whatever notation you use, the important thing is to know the line between when notation should be taken literally, when it is a mnemonic, and what the real meaning behind it is.

3.5 Multi-dimensional Parameterizations

Parameterizations aren’t just for curves—higher dimensional objects can also be parameterized. For example, let $\mathcal{P} \subseteq \mathbb{R}^3$ be the plane with equation $x + y + z = 0$. We can write this plane in vector form as

$$\vec{x} = t \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} + s \begin{bmatrix} 0 \\ 1 \\ -1 \end{bmatrix}.$$

Written in vector form, we see that \mathcal{P} is the range of a function $\vec{p} : \mathbb{R}^2 \rightarrow \mathcal{P}$ defined by

$$\vec{p}(t, s) = t \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} + s \begin{bmatrix} 0 \\ 1 \\ -1 \end{bmatrix}.$$

We could further verify that \vec{p} is one-to-one and continuous, making \vec{p} a parameterization of \mathcal{P} .

Just like parameterizations of curves, for any given multi-dimensional object, there are infinitely many parameterizations. For instance,

$$\vec{q}(t, s) = t^3 \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} + (1-s) \begin{bmatrix} 0 \\ 1 \\ -1 \end{bmatrix} \quad \text{and} \quad \vec{h}(t, s) = t \begin{bmatrix} 3 \\ -3 \\ 0 \end{bmatrix} + s \begin{bmatrix} 1 \\ 1 \\ -2 \end{bmatrix}$$

are also both parameterizations of \mathcal{P} .

Converting Between Parameterizations

Suppose $\vec{p} : \mathbb{R} \rightarrow C$ and $\vec{q} : \mathbb{R} \rightarrow C$ are both parameterizations of the same curve C and that $\vec{p}(0) = \vec{q}(0)$. Based on our previous explorations, we know that we can convert from \vec{p} to \vec{q} by

stretching, warping, or flipping “time” in some way. That is, we can find a function $w : \mathbb{R} \rightarrow \mathbb{R}$ so that $\vec{p} \circ w = \vec{q}$.

Such time-warping functions, which we should really call *domain-warping* functions, are themselves parameterizations of \mathbb{R} . We can also convert between parameterizations of multi-dimensional objects by conjuring up the appropriate domain-warping function.

Theorem 3.5.1 If $\vec{p} : \mathbb{R}^n \rightarrow \mathcal{S}$ and $\vec{q} : \mathbb{R}^n \rightarrow \mathcal{S}$ are both parameterizations of an n -dimensional object \mathcal{S} , then there exists a parameterization $\vec{w} : \mathbb{R}^n \rightarrow \mathbb{R}^n$ so that

$$\vec{p} \circ \vec{w} = \vec{q}.$$

Proof. First, recall that the composition of two continuous functions is continuous. Further, if a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is continuous and invertible, then f^{-1} is also continuous.

Now, since $\vec{p} : \mathbb{R}^n \rightarrow \mathcal{S}$ is one-to-one, $\vec{p}^{-1}|_{\mathcal{S}}$ exists and so $\vec{w} : \mathbb{R}^n \rightarrow \mathbb{R}^n$ given by $\vec{w} = \vec{p}^{-1} \circ \vec{q}$ is well defined. Further, \vec{w} is a composition of parameterizations and therefore a parameterization. It now directly follows that $\vec{p} \circ \vec{w} = \vec{p} \circ (\vec{p}^{-1} \circ \vec{q}) = \vec{q}$. ■

■ **Example 3.11** Let \mathcal{P} be a plane parameterized by $\vec{p} : \mathbb{R}^2 \rightarrow \mathcal{P}$ and $\vec{q} : \mathbb{R}^2 \rightarrow \mathcal{P}$. Given that

$$\vec{p}(t, s) = t \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} + s \begin{bmatrix} 0 \\ 1 \\ -1 \end{bmatrix} \quad \text{and} \quad \vec{q}(t, s) = t^3 \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} + s \begin{bmatrix} 1 \\ 1 \\ -2 \end{bmatrix},$$

find a function $\vec{w} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ such that $\vec{p} \circ \vec{w} = \vec{q}$.

We will find \vec{w} by analyzing how it needs to adjust each of its parameters. Suppose $\vec{w}(t, s) = (a, b)$. Then, if $\vec{p} \circ \vec{w} = \vec{q}$, we must have

$$\vec{p} \circ \vec{w}(t, s) = \vec{p}(a, b) = a \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} + b \begin{bmatrix} 0 \\ 1 \\ -1 \end{bmatrix} = t^3 \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} + s \begin{bmatrix} 1 \\ 1 \\ -2 \end{bmatrix} = \vec{q}(t, s).$$

If we can find a formula for a and b in terms of t and s , we can write down a formula for \vec{w} .

From the vector equation above, we get the following system of equations.

$$\begin{aligned} a &= t^3 + s \\ -a + b &= -t^3 + s \\ -b &= -2s \end{aligned}$$

We conclude that $a = t^3 + s$ and $b = 2s$. Therefore,

$$\vec{w}(t, s) = (t^3 + s, 2s).$$

■

Isometric Parameterizations

When parameterizing curves, arc-length parameterizations stand out as special. If you have an arc-length parameterization of a curve, line integrals become easy, speed is always one, and you can imagine the curve as the real line picked up and bent, but *not stretched*, and placed into space. The higher-dimensional analog of an arc-length parameterization is an *isometric parameterization*.

Definition 3.5.1 — Isometric Parameterization. Let $\vec{p} : \mathbb{R}^n \rightarrow S$ be a parameterization of S . The parameterization \vec{p} is called an *isometric parameterization* if the speed of \vec{p} with respect to each parameter is one and \vec{p} preserves area.

Suppose S is a surface (that is, a two dimensional object) and that $\vec{p} : \mathbb{R}^2 \rightarrow S$ is an isometric parameterization. We can then view S as the plane \mathbb{R}^2 picked up, bent, but *not stretched*, and placed into space. Any calculations you might want to do on the surface S can be done equally well by using \vec{p} to convert the calculation to one in \mathbb{R}^2 .

■ **Example 3.12** Show that $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ given by $\vec{p}(t, s) = \frac{t}{\sqrt{2}} \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \frac{s}{\sqrt{2}} \begin{bmatrix} 1 \\ -1 \end{bmatrix}$ is an isometric parameterization of \mathbb{R}^2 .

To show that \vec{p} is an isometric parameterization of \mathbb{R}^2 , we need to show that the speed with respect to each parameter is one and that it preserves area. We'll verify the speed first. Let t_0 and s_0 be fixed and consider

$$\begin{aligned} \text{speed with respect to } t &= \lim_{h \rightarrow 0} \frac{\left\| \left(\frac{t_0+h}{\sqrt{2}} \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \frac{s_0}{\sqrt{2}} \begin{bmatrix} 1 \\ -1 \end{bmatrix} \right) - \left(\frac{t_0}{\sqrt{2}} \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \frac{s_0}{\sqrt{2}} \begin{bmatrix} 1 \\ -1 \end{bmatrix} \right) \right\|}{|h|} \\ &= \frac{\left| \frac{h}{\sqrt{2}} \right|}{|h|} \left\| \begin{bmatrix} 1 \\ 1 \end{bmatrix} \right\| = 1. \end{aligned}$$

Computing the speed with respect to s is similar and, indeed, the speed with respect to s is 1.

Now, to verify that \vec{p} preserves area, we only need to consider the image of rectangles under \vec{p} .⁹ Recall that $[t, t + \Delta t] \times [s, s + \Delta s] \subseteq \mathbb{R}^2$ is a rectangle with lower left corner at (t, s) and sides of length Δt and Δs . We only need to verify that $\vec{p}([t, t + \Delta t] \times [s, s + \Delta s])$ has area $\Delta t \Delta s$.

Graphing, we see that the corners of $\vec{p}([t, t + \Delta t] \times [s, s + \Delta s])$ are at
XXX Finish ■

■ **Example 3.13** Find an isometric parameterization the cylinder with radius one and height one centered at the origin in \mathbb{R}^3 and oriented with height along the z -axis.

XXX Finish ■

Isometric parameterizations are really neat, and as we'll see in the future, if we have an isometric parameterization, many computations become easier. Alas, though for most curves

⁹ We can decompose an area into the limit of the sum of a bunch of rectangles, so if we prove the area of rectangles are preserved, we've proved all areas are preserved.

there exists an arc-length parameterization, very few multi-dimensional objects have isometric parameterizations. If an object has *intrinsic* curvature, like a sphere, there could never be an isometric parameterization. Asking that both speed and area be preserved is just too much.

However, there are several other special types of parameterizations for multi-dimensional objects.

Definition 3.5.2 — Conformal Mapping. A parameterization $\vec{p} : \mathbb{R}^n \rightarrow S$ is called a *conformal map*^a if it preserves angles.

^a *Map* is just another word for function.

Definition 3.5.3 — Volume Preserving Transformation. A parameterization $\vec{p} : \mathbb{R}^n \rightarrow S$ is called a *volume-preserving transformation*^a if \vec{p} preserves volume.

^a *Transformation* is just another word for function

Every isometric parameterization is both a conformal mapping and a volume-preserving transformation, but there can be conformal mappings which aren't volume-preserving or isometric and volume-preserving transformations which aren't conformal or isometric. We won't study these types of parameterizations in detail, but they show up in the domains of *complex analysis* and the study of *general relativity*.


Exercises for 3.5

3.6 Coordinate Systems

Coordinate systems are an important special case of multi-dimensional parameterizations. In fact, any parameterization $\vec{p} : \mathbb{R}^n \rightarrow \mathbb{R}^n$ gives rise to a coordinate system. Rather than starting off with a theoretical discussion of coordinate systems, let's look at an important coordinate system for \mathbb{R}^2 .

Polar Coordinates

Polar coordinates describe points in \mathbb{R}^2 by a *distance*, denoted r , and an *angle*, denoted θ . A point $\vec{x} \in \mathbb{R}^2$ specified by the polar-coordinates (r, θ) is the point at distance r from the origin at an angle of θ measured counter clockwise from the x -axis.



Using trigonometry, we can deduce the rectangular coordinates of a point in polar coordinates. In particular,

$$x = r \cos \theta \quad \text{and} \quad y = r \sin \theta. \quad (3.3)$$


We can also write the polar coordinates of a point $\vec{x} = (x, y)$ from its rectangular coordinates via

$$r = \sqrt{x^2 + y^2} \quad \text{and} \quad \theta = \arctan\left(\frac{y}{x}\right).$$

Typically when we use polar coordinates, we assume $r \in [0, \infty)$ and $\theta \in [0, 2\pi)$, however Equations (3.3) are valid for all $r, \theta \in \mathbb{R}$, and so we will let r and θ take non-traditional values when it suits us.

There are a couple of things to notice about polar coordinates. First, if you list the coordinates in the order (r, θ) , then polar coordinates form a right-handed coordinate system¹⁰.

¹⁰ We haven't defined what it means for a coordinate system whose "axes" are curved to be right-handed. For polar coordinates it means that if you take velocity vectors for the r direction and the θ direction, the θ velocity vector is oriented counter clockwise compared to the r velocity vector. A more appropriate name for this would be *locally* right handed.



Secondly, polar coordinates are not unique. If we allow r to take negative values, the point described by $(r, \theta) = (1, \pi)$ is the same point described by $(r, \theta) = (-1, 0)$. Most non-uniqueness issues are solved by insisting $0 \leq r < \infty$ and $0 \leq \theta < 2\pi$. However, even then $(r, \theta) = (0, \theta_0)$ describes the same point (the origin) no matter the value of θ_0 . For this reason, we call the origin in polar coordinates a *singularity*¹¹.

Shapes with radial symmetry can often be described more easily in polar coordinates than in rectangular. For instance, let \mathcal{C} be the circle of radius 7 centered at the origin. \mathcal{C} can be described as $\mathcal{C} = \{(r, \theta) : r = 7\}$ in polar coordinates. This doesn't have a dramatic advantage over the representation $\mathcal{C} = \{(x, y) : x^2 + y^2 = 7^2\}$ until you try to write \mathcal{C} as the graph of a function. \mathcal{C} cannot be described as the graph of a function in rectangular coordinates, but it is the graph of the function $r = 7$ in polar coordinates.

Circles and polar coordinates go together hand in hand. Even circles centered away from the origin have nice descriptions in polar coordinates.

■ **Example 3.14** Graph the function whose equation is given in polar coordinates by $r = \sin \theta$.
XXX Finish ■

Multiple Perspectives on Coordinate Systems

We now have two coordinate systems for \mathbb{R}^2 under our belt: rectangular and polar. It's time to get philosophical so that we avoid getting confused later on.

Recall that we've been using \mathbb{R}^2 to represent two-dimensional Euclidean space—the plane. We also use \mathbb{R}^2 to represent *pairs of real numbers*. If we were really pedantic, we would use a different symbol for these two concepts because they really are different things. However, once we draw x and y axes in the Euclidean plane, we see that pairs of real numbers are *equivalent* to points in the Euclidean plane. That is, given any point in the Euclidean plane, we can assign it exactly one pair of real numbers and any pair of real numbers corresponds to exactly one point in the Euclidean plane.

This justified our use of \mathbb{R}^2 for both pairs of real numbers and the Euclidean plane. However, our situation now is more complicated. Polar coordinates are *also* pairs of real numbers

¹¹ In mathematics, the word *singularity* is used to describe a point where “bad things” happen. Ironically, the singularity of polar coordinates is the place where there *isn't* a singular description in terms of polar coordinates.

and polar coordinates *also* describe points in the Euclidean plane. Further, consider the function $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ defined by

$$\vec{p}(r, \theta) = (r \cos \theta, r \sin \theta). \quad (3.4)$$

On the one hand, \vec{p} can be interpreted as a function that inputs polar coordinates and outputs rectangular coordinates. Under this interpretation, \vec{p} *doesn't do anything* to the Euclidean plane. On the other hand, we could interpret \vec{p} as a function that takes in rectangular coordinates and outputs rectangular coordinates. Remember, r and θ in Equation (3.4) are dummy variables. They only have meaning insofar as they allow us to make sense of the definition. But we could have just as easily defined \vec{p} by $\vec{p}(x, y) = (x \cos y, x \sin y)$. It would have made no difference. Under this interpretation \vec{p} twists and warps the Euclidean plane significantly.

When we start writing functions down in multiple coordinate systems, we need to be careful that we know how to interpret these functions. There are systems of notation what would allow us to deal with multiple coordinate systems in an unambiguous manner, but we will take a more pragmatic approach and have variable names carry extra meaning when talking about multiple coordinate systems.

For example, when we defined polar coordinates, we introduced the parameters of polar coordinates as r and θ . Therefore, if we write an equation using those variables, for example $r = \sin 2\theta$, that equation should be interpreted as a relationship in polar coordinates and given a point $\vec{x} \in \mathbb{R}^2$, if we write $\vec{x} = (r, \theta)$ we interpret the description of \vec{x} as given in polar coordinates. Similarly, if we write $\vec{x} = (x, y)$ we interpret the description of \vec{x} in rectangular coordinates. This can lead to unfortunate-looking equations like $(r, \theta) = \vec{x} = (x, y)$. This equation would be more clearly expressed as

$$(r, \theta)_{\text{polar}} = \vec{x} = (x, y)_{\text{rectangular}},$$

but we will be sloppy if the context allows. Other times we will just say in words, for example, “the point (a, b) interpreted in polar coordinates,” to make it clear what we mean.

There are also times when we want to interpret (r, θ) as a pair of real numbers instead of as describing a point in Euclidean space. In this situation, we may talk about the point (r, θ) in the $r\theta$ -plane. This is code for, “interpret (r, θ) as a pair of numbers that describe the rectangular coordinates of a point in the plane.” However, by distinguishing the $r\theta$ -plane from the usual xy -plane, we can draw helpful distinctions.

For example, recall \vec{p} from earlier defined by $\vec{p}(r, \theta) = (r \cos \theta, r \sin \theta)$. The function \vec{p} converts from polar coordinates to rectangular coordinates, but will also interpret it as mapping the $r\theta$ -plane to the xy -plane.

XXX Figure

While the graph of $r = 2\theta$ is a line in the $r\theta$ -plane, it is a spiral in the xy -plane.

XXX Figure

Despite this discussion, expect coordinate systems to occasionally be confusing. Still, their usefulness in applications outweighs the inconvenience of being confused. We will examine a couple more of the most popular non-rectangular coordinate systems.

Cylindrical Coordinates

Cylindrical coordinates is a coordinate system for \mathbb{R}^3 that arises as a straightforward extension of polar coordinates. Every point in \mathbb{R}^3 is described by three numbers denoted by r , θ , and z .

Recall that \mathbb{R}^3 can be described as $\mathbb{R}^2 \times \mathbb{R}$. That is, any point in \mathbb{R}^3 can be described as a point in \mathbb{R}^2 coupled with a “z-height.” Cylindrical coordinates are obtained by writing points in \mathbb{R}^2 in polar coordinates and then adding a z component.

XXX Figure

It's related to rectangular coordinates by the formulas

$$x = r \cos \theta \quad y = r \sin \theta \quad z = z$$

where (r, θ, z) represents a point in cylindrical coordinates and (x, y, z) is the same point in rectangular coordinates.

■ **Example 3.15** $r = a$ describes an (infinite) cylinder of radius a centered on the z -axis. If we let a vary, we obtain an infinite family of concentric cylinders. We can treat the case $a = 0$ (the z -axis) as a degenerate cylinder of radius 0.

XXX Figure

■ **Example 3.16** $\theta = \alpha$ describes a *half plane* making angle α with the positive xz -plane. In this half plane, r can assume any non-negative value and z can assume any value.

XXX Figure

■ **Example 3.17** $z = mr$ describes an (infinite) cone centered on the z -axis with vertex at the origin. For a fixed value of θ , we obtain a ray in this cone which starts at the origin and extends to infinity. This ray makes angle $\arctan m$ with the z -axis, and if we let θ vary, the ray rotates around the z -axis generating the cone. Note also that if $m > 0$, the angle with the z -axis is acute and the cone lies above the x, y -plane. If $m < 0$, the angle is obtuse, and the cone lies below the xy -plane. The case $m = 0$ yields the x, y -plane ($z = 0$) which may be considered a special ‘cone’.

Note that in rectangular coordinates, $z = mr$ becomes $z = m\sqrt{x^2 + y^2}$.

XXX Figure

■ **Example 3.18** $r^2 + z^2 = a^2$ describes a sphere of radius a centered at the origin. The easiest way to see this is to put $r^2 = x^2 + y^2$ whence the equation becomes $x^2 + y^2 + z^2 = a^2$. The top hemisphere of the sphere would be described by $z = \sqrt{a^2 - r^2}$ and the bottom hemisphere by $z = -\sqrt{a^2 - r^2}$.

Spherical Coordinates

Cylindrical coordinates are one way to generalize polar coordinates to \mathbb{R}^3 , but there is another way that is more useful for problems with spherical symmetry. *Spherical coordinates* are denoted with the variables ρ , ϕ , and θ .

A point P in space is defined by the coordinates (ρ, θ, ϕ) as follows. The coordinate ρ gives the distance $\|\vec{OP}\|$ of the point to the origin. It is always non-negative, and it should be distinguished from the cylindrical coordinate r which is the distance from the z -axis. The coordinate ϕ gives the *azimuthal angle*, which is the angle of declination between \vec{OP} and the positive z -axis. We will assume $\phi \in [0, \pi]$. Finally, the coordinate θ denotes the *longitudinal angle*, and is the same as the θ from cylindrical coordinates. Again, we assume $\theta \in [0, 2\pi)$.

XXX Figure of spherical coords

Note the reason we restrict ϕ to the interval $[0, \pi]$ (rather than $[0, 2\pi]$). Fix ρ and θ . If $\phi = 0$, the point is on the positive z -axis, and, as ϕ increases, the point swings down toward the negative z -axis. However, it stays in the half plane determined by that value of θ . For $\phi = \pi$, the point is on the negative z -axis, but if we allow ϕ to increase further, the point swings into the *opposite* half plane with longitudinal angle $\theta + \pi$. Such points can be obtained just as well by swinging down from the positive z -axis in the opposite half plane determined by $\theta + \pi$.

XXX Figure

The following relationships hold between spherical coordinates, cylindrical coordinates, and rectangular coordinates.

$$r = \rho \sin \phi$$

$$z = \rho \cos \phi$$

so

$$x = \rho \sin \phi \cos \theta$$

$$y = \rho \sin \phi \sin \theta$$

$$z = \rho \cos \phi$$

and

$$\rho = \sqrt{r^2 + z^2} = \sqrt{x^2 + y^2 + z^2}$$

$$\tan \phi = \frac{r}{z} \quad \text{if } z \neq 0.$$

It is important to note that unlike rectangular, polar, or cylindrical coordinates, mathematicians and physicists use two different conventions for spherical coordinates. We have introduced spherical coordinates in the order (ρ, θ, ϕ) . This paints spherical coordinates as an extension of polar coordinates in the xy -plane by the addition of a third coordinate, ϕ . However, ordered this way, spherical coordinates produce a *left-handed* coordinate system. Since physicists prefer right-handed coordinate systems, they tend to use the order (ρ, ϕ, θ) when describing points in spherical coordinates¹². We will primarily use the mathematical convention, but when you read problems presented in spherical coordinates in other contexts, be aware of what convention the author is using!

■ **Example 3.19** $\rho = a$ describes a sphere of radius a centered at the origin. ■

■ **Example 3.20** $\phi = \alpha$ describes a cone making angle α with the positive z -axis. If $\phi < \pi/2$ the cone lies above the xy -plane; if $\phi > \pi/2$, the cone lies below the xy -plane; and if $\phi = \pi/2$, the (degenerate) cone is the xy -plane. ■

■ **Example 3.21** $\theta = \beta$ describes a half plane starting from the z -axis as before. ■

¹² Mathematics and physicists differ in their conventions in several other regards. Often times a force that is negative to a physicist is positive to a mathematician and mathematicians tend to reverse the temperature scale when doing thermodynamics. All of the theorems come out the same, of course, but it takes some interpretation to understand theorems from another field.

■ **Example 3.22** $\rho = 2a \cos \phi$ describes a sphere of radius a centered at $(0, 0, a)$. You can see this by looking at the half plane determined by fixing θ . In that half plane, the locus is the *semi-circle* with the given radius and center. If we then let θ vary, the effect is to rotate the semi-circle about the z -axis and generate the sphere

XXX Figure ■

If we fix $\rho = a$, we obtain a sphere of radius a . Then (ϕ, θ) specify the position of a point on that sphere.

For $\theta = \text{constant}$, we obtain the semi-circle which is the intersection of the half plane for that θ with the sphere. That circle is called a *meridian of longitude*. This is exactly the concept of longitude used to measure position on the surface of the Earth, except that we use radians instead of degrees. Earth's longitude is usually measured in degrees east or west of the Greenwich Meridian. That corresponds in our case to the positive and negative directions from the 0-meridian.

XXX Figure

For $\phi = \text{constant}$, we obtain the circle which is the intersection of the cone for that ϕ with the sphere. Such circles are called *circles of latitude*. The coordinate ϕ is related to the notion of latitude on the surface of the Earth, except that the latter is an angle in degrees north or south of the *equatorial plane*. The spherical coordinate ϕ is sometimes called *co-latitude*, and we have $\phi = \pi/2 - \lambda$, where λ is the latitude measures from the equatorial plane (assuming both λ and ϕ are measured in radians). The unique point with $\phi = 0$ is called the *north pole*, that with $\phi = \pi$ is called the *south pole*, and at the poles θ is not well defined.

XXX Figure

Coordinate Systems from Parameterizations

We've looked at several important coordinate systems for \mathbb{R}^2 and \mathbb{R}^3 , but any parameterization of \mathbb{R}^2 or \mathbb{R}^3 gives rise to a coordinate system.

Consider $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ given by $\vec{p}(a, b) = (a + b, a - b)$.

Exercise 3.5 Show that $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ given by $\vec{p}(a, b) = (a + b, a - b)$ is a parameterization. That is, that it is a continuous one-to-one and onto function.

Geometrically, we can view \vec{p} as taking the plane, rotating it by 45° , and then stretching it uniformly in all directions by a factor of $\sqrt{2}$.

XXX Figure

Alternatively, we can imagine that \vec{p} describes how to take points described in “ ab -coordinates” and rewrite them in rectangular coordinates. Call this new coordinate system \mathcal{A} . We then get that \mathcal{A} coordinates relate to rectangular coordinates by

$$x = a + b \quad \text{and} \quad y = a - b.$$

Drawing the lines $a = 0$ and $b = 0$ we see that \mathcal{A} coordinates are very similar to rectangular coordinates but with rotated and scaled axes.

XXX Figure

There might be good reason to use \mathcal{A} coordinates, depending on your application. For example, suppose you were working on a construction project and you are using rectangular

coordinates (x, y) to represent the position north and east from the origin of your construction lot. If you're building a diagonal building, it might be nicer to describe locations in your building using \mathcal{A} coordinates.

Now, we might be tempted to propose that all coordinate systems come from parameterizations, and this is *almost* true. However, coordinate systems like polar, cylindrical, and spherical all have *singularities*. That is, there are points in space that are not uniquely describable in those coordinate systems. Thus, if we consider, for example, the function $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ which converts polar coordinates to rectangular coordinates, it isn't a true parameterization because it isn't one-to-one at the origin. We won't let this detail bother us—polar, cylindrical, and spherical coordinates are still useful and they *almost* come from parameterizations. Occasionally we will even slip up and proclaim that \mathbb{R}^2 is parameterized by polar coordinates or that \mathbb{R}^3 is parameterized by spherical coordinates¹³.

Exercises for 3.6

¹³ To a novice mathematician, it can be hard to discern the patterns in what determines where a professional mathematician allows herself to be sloppy and where she maintains excruciating logical precision. Rest assured, after enough mistakes, one learns where one needs to tread carefully and where one can be more lax.

Chapter 4

Multi-variable Functions

We want to develop the calculus necessary to discuss functions of many variables. We shall start with functions $f(x, y)$ of two independent variables and functions $f(x, y, z)$ of three independent variables. However, in general, we need to consider functions $f(x_1, x_2, \dots, x_n)$ of any number of independent variables. Recall, \mathbb{R}^n stands for the set of all n -tuples (x_1, x_2, \dots, x_n) with real entries x_i . We shall use the old fashioned term *locus* to denote the set of all points satisfying some equation or condition.

4.1 Graphing in Many Variables

We shall encounter equations involving two, three, or more variables. As you know, an equation of the form

$$f(x, y) = c$$

may be viewed as defining a curve in the plane. For example, $ax + by = c$ defines a line in \mathbb{R}^2 , while $x^2 + y^2 = r^2$ defines a circle of radius r centered at the origin. Similarly, an equation involving three variables

$$f(x, y, z) = c$$

may be thought of as defining a *surface* in space. We saw previously that the locus in \mathbb{R}^3 of a linear equation

$$ax + by + cz = d$$

(where not all a , b , and c are zero) is a plane. If we use more complicated equations, we get more complicated surfaces.

■ **Example 4.1** The equation

$$x^2 + y^2 + z^2 = r^2$$

may be rewritten $\|\vec{r}\| = \sqrt{x^2 + y^2 + z^2} = R$, so it asserts that the point with position vector \vec{r} is at distance r from the origin. Hence, the locus of all such points is a *sphere* of radius r centered at the origin. ■

■ **Example 4.2** Consider the locus of the equation


$$x^2 + 2x + y^2 - 4y + z^2 = 20.$$

This is also a sphere, but one not centered at the origin. To see this, *complete the squares* for the terms involving x and y .

$$\begin{aligned}x^2 + 2x + 1 + y^2 - 4y + 4 + z^2 &= 10 + 1 + 4 = 25 \\(x + 1)^2 + (y - 2)^2 + z^2 &= 5^2.\end{aligned}$$


This asserts that the point with position vector $\vec{r} = (x, y, z)$ is 5 units from the point $(-1, 2, 0)$, i.e., it lies on a sphere of radius 5 centered at $(-1, 2, 0)$. ■

■ **Example 4.3** Consider the locus of the equation $z = x^2 + y^2$ (which could also be written $x^2 + y^2 - z = 0$). To see what this looks like, we consider its intersection with various planes. Its intersection with the yz -plane is obtained by setting $x = 0$ to get $z = y^2$. This is a parabola in the yz -plane. Similarly, its intersection with the xz -plane is the parabola given by $z = x^2$. To fill in the picture, consider intersections with planes parallel to the xy -plane. Any such plane has the equation $z = h$, so the intersection has the equation $x^2 + y^2 = h = (\sqrt{h})^2$, which you should recognize as a circle of radius \sqrt{h} , if $h > 0$. Note that the circle is centered at $(0, 0, h)$ on the z -axis since it lies in the plane $z = h$. If $z = h = 0$, then the circle reduces to a single point, and for $z = h < 0$, there is no locus. The surface is “bowl” shaped and is called a *circular paraboloid*.




Graphing a surface in \mathbb{R}^3 by sketching its traces on various planes is a useful strategy. In order to be good at it, you need to know the basics of plane analytic geometry so you can recognize the resulting curves. In particular, you should be familiar with the elementary facts concerning *conic sections*, i.e., ellipses, hyperbolas, and parabolas. Edwards and Penney, 3rd Edition, Chapter 10 is a good reference for this material.

■ **Example 4.4** Consider the locus in space of $\frac{x^2}{4} + \frac{y^2}{9} = 1$. Its intersection with a plane $z = h$ parallel to the xy -plane is an ellipse centered on the z -axis and with semi-minor and semi-major axes 2 and 3, respectively. The surface is a *cylinder* perpendicular to the xy -plane with elliptical cross sections. Note that the locus *in space* is not just the ellipse in the xy -plane with the same equation.




XXX Figure

■ **Example 4.5** Consider the locus in space of the equation $z = \frac{1}{x^2 + y^2}$. Its intersection with the plane $z = h$ (for $h > 0$) is the circle with equation $x^2 + y^2 = 1/h = (\sqrt{1/h})^2$. The surface does not intersect the xy -plane itself ($z = 0$) nor any plane below the xy -plane. Its intersection with the xz -plane ($y = 0$) is the curve $z = 1/x^2$, which is asymptotic to the x -axis and to the positive z -axis. Similarly, for its intersection with the yz -plane, the surface flattens out and approaches the xy -plane as $r = \sqrt{x^2 + y^2} \rightarrow \infty$. It approaches the positive z -axis as $r \rightarrow 0$.



■ **Example 4.6** Consider the locus in space of the equation $yz = 1$. Its intersection with a plane parallel to the yz -plane ($x = d$) is a hyperbola asymptotic to the y and z axes. The surface is perpendicular to the yz -plane. Such a surface is also called a *cylinder*, although it doesn't close upon itself as the elliptical cylinder considered above.

■ **Example 4.7** Consider the locus of the equation $x^2 + z^2 = y^2 - 1$. For each plane parallel to the xz -plane ($y = c$), the intersection is a circle $x^2 + z^2 = c^2 - 1 = \left(\sqrt{c^2 - 1}\right)^2$ centered on the y -axis, at least of $c^2 > 1$. For $y = c = \pm 1$, the locus is a point, and for $-1 < y = c < 1$, the locus is empty. In addition, the intersection of the surface with the xy -plane ($z = 0$) is the hyperbola with equation $x^2 - y^2 = -1$, and similarly for its intersection with the yz -plane. The surface comes in two pieces which open up as “bowls” centered on the positive and negative y -axes. The surface is called a *hyperboloid of 2 sheets*.



■

Graphing Functions

For a scalar function f of one independent variable, the *graph of the function* is the set of all points in \mathbb{R}^2 of the form $(x, f(x))$ for x in the domain of the function. (The domain of a function is the set of values of the independent variable for which the function is defined). In other words, it is the locus of the equation $y = f(x)$. It is generally a curve in the plane.


We can define a similar notion for a scalar function f of two independent variables. The graph is the set of points in \mathbb{R}^3 of the form $(x, y, f(x, y))$ for (x, y) , a point in the domain of the function. In other words, it is the locus of the equation $z = f(x, y)$, and it is generally a surface in space. The graph of a function is often useful in understanding the function.

We have already encountered several examples of graphs of functions. For example, the locus of $z = x^2 + y^2$ is the graph of the function f defined by $f(x, y) = x^2 + y^2$. Similarly, the locus of $z = 1/(x^2 + y^2)$ is the graph of the function f defined by $f(x, y) = 1/(x^2 + y^2)$ for $(x, y) \neq (0, 0)$. Note that in the first case there does not need to be a restriction on the domain of the function, but in the second case $(0, 0)$ was omitted.

In some of the other examples, the locus of the equation cannot be considered the graph of a function. For example, the equation $x^2 + y^2 + z^2 = R^2$ cannot be solved uniquely for z in terms of (x, y) . Indeed, we have $z = \pm\sqrt{R^2 - x^2 - y^2}$, so that two possible functions are suggested. $z = f_1(x, y) = \sqrt{R^2 - x^2 - y^2}$ defines a function with graph the *top hemisphere*, while $z = f_2(x, y) = -\sqrt{R^2 - x^2 - y^2}$ yields the lower hemisphere. (Note that for either of

the functions, the relevant domain is the set of points on or inside the circle $x^2 + y^2 = R^2$. For points outside that circle, the expression inside the square root is negative, and since we are only talking about functions assuming real values, such points must be excluded).

■ **Example 4.8** Let $f(x, y) = xy$ for all (x, y) in \mathbb{R}^2 . The graph is the locus of the equation $z = xy$. We can sketch it by considering traces on various planes. Its intersection with a plane parallel to the xy -plane ($z = \text{constant}$) is a hyperbola asymptotic to lines parallel to the x and y axes. For $z > 0$, the hyperbola is in the first and third quadrants of the plane, but for $z < 0$, it is in the second and fourth quadrants. For $z = 0$, the equation is $xy = 0$ with locus consisting of the x -axis ($y = 0$) and the y -axis ($x = 0$). Thus, the graph intersects the xy -plane in two straight lines. The surface is generally shaped like an “infinite saddle.” It is called a *hyperbolic paraboloid*. It is clear where the term “hyperbolic” comes from. Can you see any parabolas? (Hint: Try planes perpendicular to the xy -plane with equations of the form $y = mx$).



■ **Example 4.9** Let $f(x, y) = x/y$ for $y \neq 0$. Thus, the domain of this function consists of all points (x, y) not on the x -axis ($y = 0$). The trace in the plane $y = c, c \neq 0$ is the line $z = (1/c)x$ with slope $1/c$. Similarly, the trace in the plane $z = c, c \neq 0$ is the line $y = (1/c)x$. Finally, the trace in the plane $x = c$, is the hyperbola $z = c/y$. Even with this information you will have some trouble visualizing the graph. However, the equation $z = x/y$ can be rewritten $yz = x$. By permuting the variables, you should see that the locus of $yz = x$ is similar to the saddle-shaped surface we just described but oriented differently in space. However, the saddle is not quite the graph of the function since it contains the z -axis ($y = x = 0$) but the graph of the function does not. In general, the graph of a function, since it consists of points of the form $(x, y, f(x, y))$, cannot contain points with the same values for x and y but different values for z . In other words, any line parallel to the z -axis can intersect such a graph at most once. ■

Sketching graphs of functions, or more generally loci of equations in x, y , and z , is not easy. One approach drawn from the study of topography is to interpret the equation $z = f(x, y)$ as giving the *elevation* of the surface, viewed as a hilly terrain, above a reference plane. (Negative elevation $f(x, y)$ is interpreted to mean that the surface dips below the reference plane.) For each possible elevation c , the intersection of the plane $z = c$ with the graph yields a curve

$f(x, y) = c$. This curve is called a *level curve*, and we draw a 2-dimensional map of the graph by sketching the level curves and labeling each by the appropriate elevation c . Of course, there are generally infinitely many level curves since there are infinitely many possible values of z , but we select some subset to help us understand the topography of the surface.

XXX Figure

■ **Example 4.10** The level curves of the surface $z = xy$ have equations $xy = c$ for various c . They form a family of hyperbolas, each with two branches. For $c > 0$, these hyperbolas fill the first and third quadrants, and for $c < 0$ they fill the second and fourth quadrants. For $c = 0$ the x and y axes together constitute the level “curve.” See the diagram.

You can see that the region around the origin $(0, 0)$ is like a “mountain pass” with the topography rising in the first and third quadrants and dropping off in the second and fourth quadrants. In general, a point where the graph behaves this way is called a *saddle point*. Saddle points indicate the added complexity which can arise when one goes from functions of one variable to functions of two or more variables. At such points, the function can be considered as having a maximum or a minimum depending on where you look. ■

Quadric Surfaces

One important class of surfaces are those defined by quadratic equations. These are analogs in three dimensions of conics in two dimensions. They are called *quadric surfaces*. We describe here *some* of the possibilities. You can verify the pictures by using the methods described above.

Consider first equations of the form

$$\pm \frac{x^2}{a^2} \pm \frac{y^2}{b^2} \pm \frac{z^2}{c^2} = 1$$

If all the signs are positive, the surface is called an *ellipsoid*. Planes perpendicular to one of the coordinate axes intersect it in ellipses (if they intersect at all). However, at the extremes these ellipses degenerate into the points $(\pm a, 0, 0)$, $(0, \pm b, 0)$, and $(0, 0, \pm c)$.

XXX Figure

If exactly one of the signs are negative, the surface is called a *hyperboloid of one sheet*. It is centered on one axis (the one associated with the negative coefficient in the equation) and it opens up in both positive and negative directions along that axis. Its intersection with planes perpendicular to that axis are ellipses. Its intersections with planes perpendicular to the other axes are hyperbolas.

If exactly two of the signs are negative, the surface is called a *hyperboloid of two sheets*. It is centered on one axis (associated with the positive coefficient). For example, suppose the equation is

$$-\frac{x^2}{a^2} + \frac{y^2}{b^2} - \frac{z^2}{c^2} = 1.$$

For $y < -b$ or $y > b$, the graph intersects a plane perpendicular to the y -axis in an ellipse. For $y = \pm b$, the intersection is the point $(0, \pm b, 0)$. (These two points are called vertices of the surface.) For $-b < y < b$, there is no intersection with a plane perpendicular to the y -axis.

The above surfaces are called *central quadrics*. Note that for the hyperboloids, with equations in standard form as above, the number of sheets is the same as the number of minus signs.

Consider next equations of the form

$$z = \pm \frac{x^2}{a^2} \pm \frac{y^2}{b^2}$$

(or similar equations obtained by permuting x, y and z).

If both signs are the same, the surface is called an *elliptic paraboloid*. If both signs are positive, it is centered on the positive z -axis and its intersections with planes perpendicular to the positive z -axis are a family of similar ellipses which increase in size as z increases. If both signs are negative, the situation is similar, but the surface lies below the x, y plane.

If the signs are different, the surface is called a *hyperbolic paraboloid*. Its intersection with planes perpendicular to the z -axis are hyperbolas asymptotic to the lines in those planes parallel to the lines $x/a = \pm y/b$. Its intersection with the xy -plane is just those two lines. The surface has a saddle point at the origin.

The locus of the equation $z = cxy$ when $c \neq 0$ is also a hyperbolic paraboloid, but rotated so it intersects the xy -plane in the x and y axes.

XXX Figure

Finally, we should note that many so called “degenerate conics” are loci of quadratic equations. For example, consider

$$\frac{x^2}{a^2} + \frac{y^2}{b^2} - \frac{z^2}{c^2} = 0$$

which may be solved to obtain

$$z = \pm c \sqrt{\frac{x^2}{a^2} + \frac{y^2}{b^2}}.$$

The locus is a double cone with elliptical cross sections and vertex at the origin.

Generalizations

In general, we will want to study functions of any number of independent variables. For example, we may define the graph of a scalar valued function f of three independent variables to be the set of all points in \mathbb{R}^4 of the form $(x, y, z, f(x, y, z))$. Such an object should be considered a three dimensional subset of \mathbb{R}^4 , and it is certainly not easy to visualize. It is more useful to consider the analogs of level curves for such functions. Namely, for each possible value c attained by the function, we may consider the locus in \mathbb{R}^3 of the equation $f(x, y, z) = c$. This is generally a surface called a *level surface* for the function.

■ **Example 4.11** For $f(x, y, z) = x^2 + y^2 + z^2$, the level surfaces are concentric spheres centered at the origin if $c > 0$. For $c = 0$ the level ‘surface’ is not really a surface at all; it just consists of the point at the origin. (What if $c < 0$?)

For $f(x, y, z) = x^2 + y^2 - z^2$, the level surfaces are either hyperboloids of one sheet if $c > 0$ or hyperboloids of two sheets if $c < 0$. (What if $c = 0$?) ■

For functions of four or more variables, geometric interpretations are even harder to come by. If $f(x_1, x_2, \dots, x_n)$ denotes a function of n variables, the locus in \mathbb{R}^n of the equation $f(x_1, x_2, \dots, x_n) = c$ is called a level set, but one doesn't ordinarily try to visualize it geometrically.

Instead of talking about many independent variables, it is useful to think instead of a single independent variable which is a *vector*, i.e., an element of \mathbb{R}^n for some n . In the case $n = 2, 3$, we usually write $\vec{r} = (x, y)$ or $\vec{r} = (x, y, z)$ so $f(x, y)$ or $f(x, y, z)$ would be written simply $f(\vec{r})$. If $n > 3$, then one often denotes the variables x_1, x_2, \dots, x_n and denotes the vector (i.e., element of \mathbb{R}^n) by $\vec{x} = (x_1, x_2, \dots, x_n)$. Then $f(x_1, x_2, \dots, x_n)$ becomes simply $f(\vec{x})$. The case of a function of a single real variable can be subsumed in this formalism by allowing the case $n = 1$. That is, we consider a scalar x to be just a vector of dimension 1, i.e., an element of \mathbb{R}^1 .

When we talked about kinematics, we considered *vector valued* functions $\vec{r}(t)$ of a single independent variable. Thus we see that it makes sense to consider in general functions of a vector variable which can also assume vector values. We indicate this by the notation $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$. That shall mean that the domain of the function f is a subset of \mathbb{R}^n while the set of values is a subset of \mathbb{R}^m . Thus, $n = m = 1$ would yield a scalar function of one variable, $n = 2, m = 1$ a scalar function of two variables, and $n = 1, m = 3$ a vector valued function of one scalar variable. We shall have occasion to consider several other special cases in detail.

There is one slightly non-standard aspect to the above notation. In ordinary usage in mathematics, " $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ " means that \mathbb{R}^n is the entire domain of the function f , whereas we are taking it to mean that the domain is some subset. We do this mostly to save writing since usually the domain will be almost all of \mathbb{R}^n or at least some significant chunk of it. What we want to make clear by the notation is the dimensionality of both the independent and dependent variables.

Exercises for 4.1

You are encouraged to make use of the available computer software (e.g., Maple, Mathematica, etc.) to help you picture the graphs in the following problems.

- State the largest possible domain for the function

- $f(x, y) = e^{x^2 - y^2}$
- $f(x, y) = \ln(y^2 - x^2 - 2)$
- $f(x, y) = \frac{x^2 - y^2}{x - y}$
- $f(x, y, z) = \frac{1}{xyz}$
- $f(x, y, z) = \frac{1}{\sqrt{z^2 - x^2 - y^2}}$

- Describe the graph of the function described by

- $f(x, y) = 5$
- $f(x, y) = 2x - y$
- $f(x, y) = 1 - x^2 - y^2$
- $f(x, y) = 4 - \sqrt{x^2 + y^2}$
- $f(x, y) = \sqrt{24 - 4x^2 - 6y^2}$

- Sketch selected level curves for the functions given by

- $f(x, y) = x + y$
- $f(x, y) = x^2 + 9y^2$
- $f(x, y) = x - y^2$
- $f(x, y) = x - y^3$
- $f(x, y) = x^2 + y^2 + 4x + 2y + 9$

- Describe selected level surfaces for the functions given by

- $f(x, y, z) = x^2 + y^2 - z$

- b) $f(x, y, z) = x^2 + y^2 + z^2 + 2x - 2y + 4z$
c) $f(x, y, z) = z^2 - x^2 - y^2$
5. Describe the quadric surfaces which are loci in \mathbb{R}^3 of the following equations.
- a) $x^2 + y^2 = 16$
b) $z^2 = 49x^2 + y^2$
c) $z = 25 - x^2 - y^2$
d) $x = 4y^2 - z^2$
e) $4x^2 + y^2 + 9z^2 = 36$
f) $x^2 + y^2 - 4z^2 = 4$
g) $9x^2 + 4y^2 - z^2 = 36$
h) $9x^2 - 4y^2 - z^2 = 36$
6. Describe the traces of the following functions in the given planes
- a) $z = xy$, in horizontal planes $z = c$
b) $z = x^2 + 9y^2$ in vertical planes $x = c$ or $y = c$
c) $z = x^2 + 9y^2$ in horizontal planes $z = c$
7. Describe the intersection of the cone $x^2 + y^2 = z^2$ with the plane $z = x + 1$.
8. Let $\vec{f}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$.
- a) Try to invent a definition of 'graph' for such a function. For what n would it be a subset of \mathbb{R}^n ?
b) Try to invent a definition of 'level set' for such a function. For what n would it be a subset of \mathbb{R}^n ?

4.2 Rates of Change and the Directional Derivative

For a function $f : \mathbb{R} \rightarrow \mathbb{R}$ of one variable, we have the idea of the *rate of change* of f . We might ask, “rate of change with respect to what?” For a function of one variable, this question has an obvious answer—the rate of change with respect to the only thing that varies! If $g : \mathbb{R}^n \rightarrow \mathbb{R}$ is a multi-variable function, the question of “what is the rate of change of g ?” has a less obvious answer. After all, there are multiple variables and infinitely many directions in the domain.

Consider the function $g : \mathbb{R}^2 \rightarrow \mathbb{R}$ where $g(x, y) = (x - 2)^2 + y^2$. Let \vec{u} be a unit vector. Now, starting at $(0, 0)$, we might ask approximately how much g changes when we move α units away from $(0, 0)$ in the direction \vec{u} .

XXX Figure

Since \vec{u} is a unit vector, the displacement in the domain is $\alpha\vec{u}$, and so the exact change in g from $\vec{0}$ to $\alpha\vec{u}$ is $g(\alpha\vec{u}) - g(\vec{0})$. The approximate *rate of change* is then $\frac{g(\alpha\vec{u}) - g(\vec{0})}{\alpha}$. Taking a limit as $\alpha \rightarrow 0$ will give us an instantaneous rate of change—one worth name.

Definition 4.2.1 — Rate of change with respect to distance. For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, the *rate of change with respect to distance* of f in the direction \vec{v} at the point \vec{a} is

$$\lim_{h \rightarrow 0^+} \frac{f(\vec{a} + h\vec{v}) - f(\vec{a})}{\|h\vec{v}\|}.$$

We must divide by $\|h\vec{v}\|$ in the definition of rate of change with respect to distance because $\|h\vec{v}\|$ is exactly how far we’ve displaced. Of course, if we were forward-thinking enough to pick \vec{v} to be a unit vector, then $\|h\vec{v}\| = |h|$.

We can think of the rate of change with respect to distance in another way. Suppose that $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and that $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^n$ is the arc-length parameterization of the with direction vector \vec{v} and where $\vec{p}(0) = \vec{a}$. Then $f \circ \vec{p} : \mathbb{R} \rightarrow \mathbb{R}$ is a single-variable function and the rate of change of \vec{f} at \vec{a} in the direction \vec{v} is just $(f \circ \vec{p})'(0)$. Think for a moment about why this is.

The rate of change of a function with respect to distance is, in some sense, the most natural notion of “rate of change” for a multivariable function. (If you’re handed a function, what would you measure the rate of change against if not distance?) However, it turns out not to be the most useful notion of rate of change. For that, we introduce the *directional derivative*.

Definition 4.2.2 — Directional Derivative. For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, the *directional derivative* of f at the point \vec{a} in the direction \vec{v} is

$$D_{\vec{a}}f(\vec{v}) = \lim_{h \rightarrow 0} \frac{f(\vec{a} + h\vec{v}) - f(\vec{a})}{h}.$$

The directional derivative $D_{\vec{a}}f(\vec{v})$ corresponds to the rate of change of f at the point \vec{a} if you moved in the direction of \vec{v} with speed $\|\vec{v}\|$. This seems like a less intuitive notion than the rate of change with respect to distance, but its virtues will soon become clear¹.

■ **Example 4.12** Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be given by $f(x, y) = (x - 2)^2 - y^2$. Compute $D_{\vec{0}}f(1, 2)$.

¹ For starters, $D_{\vec{a}}f$ is *linear*. Namely, if f is differentiable, then $D_{\vec{a}}f(\vec{u} + \vec{v}) = D_{\vec{a}}f(\vec{u}) + D_{\vec{a}}f(\vec{v})$. The same cannot be said for the rate of change with respect to distance.

Since we have no further theory to lean on, we must use the definition of the directional derivative directly.

$$\begin{aligned} D_{\vec{0}}f(1, 2) &= \lim_{h \rightarrow 0} \frac{f(\vec{0} + h(1, 2)) - f(\vec{0})}{h} = \lim_{h \rightarrow 0} \frac{((h-2)^2 - (2h)^2) - 4}{h} \\ &= \lim_{h \rightarrow 0} \frac{(-3h^2 - 4h + 4) - 4}{h} = -4. \end{aligned}$$

■

The directional derivative can be used to approximate a function. Recall that if $g : \mathbb{R} \rightarrow \mathbb{R}$ is a differentiable function of one variable, then

$$g(a + h) \approx g(a) + hg'(a).$$

Similarly, for $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ we have²

$$f(\vec{a} + h\vec{v}) \approx f(\vec{a}) + hD_{\vec{a}}f(\vec{v}). \quad (4.1)$$

The similarity in these two formulas is one reason why directional derivatives are more useful than rates of change with respect to distance³.

Exercise 4.1 Create the analogous equation to Equation (4.1) but approximate $f(\vec{a} + h\vec{v})$ using the rate of change with respect to distance instead of the directional derivative.

Partial Derivatives

There are some particularly common directional derivatives. Namely, those in the directions of the standard basis. Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ and let $\vec{a} = (a_x, a_y)$. Now,

$$D_{\vec{a}}f(\vec{x}) = \lim_{h \rightarrow 0} \frac{f(\vec{a} + h\vec{x}) - f(\vec{a})}{h} = \lim_{h \rightarrow 0} \frac{f(a_x + h, a_y) - f(a_x, a_y)}{h}.$$

This limit looks quite similar to a one-dimensional derivative. However, f does not have a one-dimensional domain. Accordingly we call $D_{\vec{a}}f(\vec{x})$ a *partial derivative* and we have a special notation for it.

Definition 4.2.3 — Partial Derivative. For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, let x_i denote the i th input variable. The *partial derivative* of f at $\vec{a} = (a_1, \dots, a_n)$ with respect to x_i is notated $\frac{\partial f}{\partial x_i}(\vec{a})$ and is defined to be

$$\frac{\partial f}{\partial x_i}(\vec{a}) = \lim_{h \rightarrow 0} \frac{f(a_1, \dots, a_{i-1}, a_i + h, a_{i+1}, \dots, a_n) - f(\vec{a})}{h}.$$

² Assuming f has directional derivatives, of course.

³ If we think about vectors in \mathbb{R}^1 , the number 1 is actually a vector. Therefore, for $g : \mathbb{R} \rightarrow \mathbb{R}$, we have $g'(1) = D_a g(1)$. Try writing out the limit expression. It's the same as the usual definition of derivative!

Sometimes we'll write $\frac{\partial f}{\partial x_i}$, omitting the point where the partial derivative is taken, if we want to save space, or if we want to view a partial derivative as a function of where it is taken. For functions whose domains are \mathbb{R}^2 or \mathbb{R}^3 , we often use x, y, z instead of x_1, x_2, x_3 to denote the first, second, or third variable. It is not important how you write your variable as long as it is clear to you and your reader which variable is changing.

Tangent Vectors to Surfaces

When we had a curve S parameterized by $\vec{p} : \mathbb{R} \rightarrow S$, we could find tangent vectors to S by computing the velocity vectors of \vec{p} . In particular, $\vec{p}'(t_0)$ is a tangent vector to S at the point $\vec{p}(t_0)$. We will find tangent vectors to surfaces in a similar way—by parameterizing a path on the surface.

Suppose $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ is a function and let S be the surface given by $z = f(x, y)$. Formally,

$$S = \left\{ \begin{bmatrix} x \\ y \\ z \end{bmatrix} : z = f(x, y) \right\}.$$

Since S is the graph of a function whose domain is \mathbb{R}^2 we have a natural parameterization of S coming from f . That is,

$$\vec{p} : \mathbb{R}^2 \rightarrow S \quad \text{where} \quad \vec{p}(x, y) = \begin{bmatrix} x \\ y \\ f(x, y) \end{bmatrix}$$

is a parameterization of S .

Let $\vec{a} = (a_x, a_y, f(a_x, a_y)) \in S$ and consider the function $\vec{r}_x : \mathbb{R} \rightarrow \mathbb{R}^2$ given by $\vec{r}_x(t) = (a_x + t, a_y)$. The function \vec{r}_x starts at the point (a_x, a_y) in the xy -plane and moves parallel to the x -axis with speed one. Since we have parameterized S by the xy -plane, by composing \vec{p} and \vec{r}_x , we will get a parameterization of a curve on S . That is, $\vec{c} = \vec{p} \circ \vec{r}_x$ parameterizes the curve on the surface S that when viewed from above looks like the line $y = a_y$.

XXX Figure

Further, $\vec{c}(0) = (a_x, a_y, f(a_x, a_y)) = \vec{a}$. Thus, if we want a tangent vector to S at \vec{a} , all we need to do is compute $\vec{c}'(0)$.

Computing,

$$\vec{c}'(0) = (\vec{p} \circ \vec{r}_x)'(0) = \left(\begin{bmatrix} a_x + t \\ a_y \\ f(a_x + t, a_y) \end{bmatrix}' \text{ at } t = 0 \right) = \begin{bmatrix} 1 \\ 0 \\ \frac{\partial f}{\partial x}(a_x, a_y) \end{bmatrix}.$$

Of course, this was just one choice of a curve along S . We could have just as easily used $\vec{p} \circ \vec{r}_y$ where $\vec{r}_y(t) = (a_x, a_y + t)$ as a parameterization of a curve. We could have parameterized paths that weren't parallel to the x or y axes. For instance $\vec{p} \circ \vec{r}_d$ where $\vec{r}_d(t) = (a_x + t d_x, a_y + t d_y)$ for non-zero constants d_x, d_y .

■ **Example 4.13** Find at least three tangent vectors at the point $\vec{a} = (1, 2, 12)$ to the surface \mathcal{S} defined by $z = f(x, y)$ where $f(x, y) = (x - 3)^2 + y^3$.

Let $\vec{r}_a(t) = (1 + t, 2)$, $\vec{r}_b(t) = (1, 2 + t)$, and $\vec{r}_c(t) = (1 + t, 2 + t)$. Further, let $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ be defined by $\vec{p}(x, y) = (x, y, f(x, y))$. Since \vec{p} is a parameterization of \mathcal{S} , $\vec{p} \circ \vec{r}_a$, $\vec{p} \circ \vec{r}_b$, and $\vec{p} \circ \vec{r}_c$ all parameterize paths in \mathcal{S} . Further,

$$\vec{p} \circ \vec{r}_a(0) = \vec{p} \circ \vec{r}_b(0) = \vec{p} \circ \vec{r}_c(0) = \vec{a},$$

so the velocity vectors at $t = 0$ of these parameterizations will give tangent vectors to \mathcal{S} at \vec{a} .

Computing, we see

$$(\vec{p} \circ \vec{r}_a)'(0) = \begin{bmatrix} 1 \\ 0 \\ \frac{\partial f}{\partial x}(1, 2) \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ -4 \end{bmatrix},$$

$$(\vec{p} \circ \vec{r}_b)'(0) = \begin{bmatrix} 0 \\ 1 \\ \frac{\partial f}{\partial y}(1, 2) \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 12 \end{bmatrix},$$

and

$$(\vec{p} \circ \vec{r}_c)'(0) = \begin{bmatrix} 1 \\ 1 \\ D_{(1,2)}f(1, 1) \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 8 \end{bmatrix}$$

are all tangent vectors to \mathcal{S} at \vec{a} . ■

Exercises for 4.2

4.3 Tangent Planes and Differentiability

Suppose $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ is a function and consider the surface $S \subseteq \mathbb{R}^3$ given by the equation $z = f(x, y)$. If f is a “nice” function, like a polynomial, the surface S will be smooth. Visually, at each point on S we could imagine a tangent plane—the analog of a tangent line to a curve.

XXX Figure

More formally, fix a point $\vec{a} \in S$ and consider the set \mathcal{P} of all tangent vectors to S at \vec{a} . If S is a smooth surface, these tangent vectors will lie in a single plane. In other words, \mathcal{P} is a plane, and in this case we call it the *tangent plane* to S at \vec{a} . (Fantastically, if \mathcal{P} is a plane, it can be fully described by the point \vec{a} and *two* non-parallel tangent vectors. We’ll make use of this idea later.)

XXX Figure

Tangent planes to surfaces, like tangent lines to curves, have the property that they are a very good approximation to a surface at their point of tangency. We will leverage this idea to define what it means to be differentiable for a multivariable function.

Directional Derivatives and Planes

A plane \mathcal{P} is a *set* and is distinct from any equations or parameterizations. However, if $\mathcal{P} \subseteq \mathbb{R}^3$ is not a vertical plane (i.e., one for which \hat{z} is a direction vector), then it can be represented uniquely by an equation of the form $z = \alpha x + \beta y + c$. In other words, a non-vertical plane in \mathbb{R}^3 is the graph of a function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ which takes the form $f(x, y) = \alpha x + \beta y + c$. Functions of this form are known as *affine functions*, and they often show up in the context of *linear approximations*⁴.

Definition 4.3.1 — Affine Function. A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is an *affine function* if it can be expressed as

$$f(x_1, x_2, \dots, x_n) = \alpha_1 x_1 + \alpha_2 x_2 + \dots + \alpha_n x_n + c$$

for some $\alpha_1, \dots, \alpha_n, c \in \mathbb{R}$. Equivalently, $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is an affine function if it can be expressed as

$$f(\vec{x}) = \vec{a} \cdot \vec{x} + c$$

for some $\vec{a} \in \mathbb{R}^n$ and $c \in \mathbb{R}$.

Directional derivatives of affine functions take a particularly nice form. Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be an affine function. That is, $f(x, y) = \alpha x + \beta y + c$ for some $\alpha, \beta, c \in \mathbb{R}$. Let $\vec{a} = (a_x, a_y)$ and $\vec{v} = (v_x, v_y)$. Now,

⁴ In many contexts in calculus, the term “linear approximation” is used where “affine approximation” would be more accurate. Sometimes, though, it’s easiest not to fight the tide of established culture.

$$\begin{aligned}
D_{\vec{a}}f(\vec{v}) &= \lim_{h \rightarrow 0} \frac{f(\vec{a} + h\vec{v}) - f(\vec{a})}{h} = \lim_{h \rightarrow 0} \frac{f(a_x + hv_x, a_y + hv_y) - f(a_x, a_y)}{h} \\
&= \lim_{h \rightarrow 0} \frac{\alpha(a_x + hv_x) + \beta(a_y + hv_y) + c - (\alpha a_x + \beta a_y + c)}{h} \\
&= \lim_{h \rightarrow 0} \frac{\alpha hv_x + \beta hv_y}{h} = \alpha v_x + \beta v_y \\
&= \begin{bmatrix} \alpha \\ \beta \end{bmatrix} \cdot \vec{v}.
\end{aligned}$$

We have written a formula for the directional derivative of f in an arbitrary direction! There's more.

$$\alpha = \begin{bmatrix} \alpha \\ \beta \end{bmatrix} \cdot \hat{\mathbf{x}} = \frac{\partial f}{\partial x}(\vec{a}) \quad \text{and} \quad \beta = \begin{bmatrix} \alpha \\ \beta \end{bmatrix} \cdot \hat{\mathbf{y}} = \frac{\partial f}{\partial y}(\vec{a})$$

and so

$$D_{\vec{a}}f(\vec{v}) = \begin{bmatrix} \frac{\partial f}{\partial x}(\vec{a}) \\ \frac{\partial f}{\partial y}(\vec{a}) \end{bmatrix} \cdot \vec{v}.$$

The vector $\begin{bmatrix} \frac{\partial f}{\partial x}(\vec{a}) \\ \frac{\partial f}{\partial y}(\vec{a}) \end{bmatrix}$ comes up often. It is called the *gradient*.

Definition 4.3.2 — Gradient. For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, the *gradient* of f at the point a is written $\nabla f(\vec{a})$ and is defined to be the vector

$$\nabla f(\vec{a}) = \left(\frac{\partial f}{\partial x_1}(\vec{a}), \dots, \frac{\partial f}{\partial x_n}(\vec{a}) \right).$$

Using the notation of *gradient*, we conclude (using the same $f(x, y) = \alpha x + \beta y + c$ from earlier) that

$$D_{\vec{a}}f(\vec{v}) = \nabla f(\vec{a}) \cdot \vec{v}. \quad (4.2)$$

All directional derivatives of affine functions can be expressed as a dot product with the gradient. So far, we have only proven this for affine function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$, but it is true for affine functions $g : \mathbb{R}^n \rightarrow \mathbb{R}$ as well. This fact is important enough to write down as a theorem.

Theorem 4.3.1 If $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is an affine function, then $D_{\vec{a}}f(\vec{v}) = \nabla f(\vec{a}) \cdot \vec{v}$ for all $\vec{a}, \vec{v} \in \mathbb{R}^n$.

Note that at the moment, the gradient is just a bag of derivatives—it doesn't yet have meaning for us beyond a convenient notational trick.

So far the rule in Equation (4.2) only applies to affine functions. It would be nice if it worked for other functions. So, in true mathematical fashion, we will define a class of functions based on whether or not Equation (4.2) applies⁵.

⁵ For technical reasons we won't do exactly this. If we considered the set of functions for which Equation (4.2) holds at every point, we'd never have a reasonable chain rule for multivariable functions. Instead, we'll use the slightly stronger notation of *linear approximations*.

Linear and First-Order Approximations

Recall from single variable calculus, if $f : \mathbb{R} \rightarrow \mathbb{R}$ was a differentiable function, a *linear approximation* of f at the point $a \in \mathbb{R}$ was a function $L : \mathbb{R} \rightarrow \mathbb{R}$ whose graph is the tangent line to $y = f(x)$ at the point $(a, f(a))$. A linear approximation is equivalently called a *first-order approximation*⁶.

XXX Figure

Further, L always takes the form $L(x) = \alpha(x - a) + c$. In fact, if f' exists, we can be more specific:

$$L(x) = f'(a)(x - a) + f(a).$$

The approximation L is quite good around the point a , so we have

$$f(a + \Delta x) \approx L(a + \Delta x) = f'(a)\Delta x + f(a)$$

if Δx is small. It's a little unclear though what “ \approx ” exactly means. In your single-variable calculus class, you may or may not have made this precise. Visually, the line given by $y = L(x) = f'(a)(x - a) + f(a)$ is distinguished from all other lines passing through the point $(a, f(a))$ because it is *tangent*. Analytically, the linear approximation $L(x) = f'(a)(x - a) + f(a)$ is distinguished from all other linear approximations to f at a because it satisfies the property

$$\lim_{\Delta x \rightarrow 0} \frac{f(a + \Delta x) - L(a + \Delta x)}{\Delta x} = 0. \quad (4.3)$$

Exercise 4.2 Let $L(x) = \alpha(x - a) + c$ and let $f : \mathbb{R} \rightarrow \mathbb{R}$ be a differentiable function. Show that if L satisfies Equation (4.3) then it must be the case that $c = f(a)$ and $\alpha = f'(a)$.

Equation (4.3) encapsulates such a strong idea that it could, in fact, be used (and sometimes is used) to define the derivative of a single-variable function. We will use the multi-dimensional analog of Equation (4.3) to define differentiability.

Definition 4.3.3 — Differentiable. The function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is *differentiable* at $\vec{a} = (a_1, \dots, a_n)$ if there exists an affine function $p(x_1, \dots, x_n) = c + \sum \alpha_i(x_i - a_i) = c + \vec{a} \cdot (\vec{x} - \vec{a})$ so that

$$f(\vec{a}) = p(\vec{a}) \quad \text{and} \quad \lim_{\vec{u} \rightarrow 0} \frac{f(\vec{a} + \vec{u}) - p(\vec{a} + \vec{u})}{\|\vec{u}\|} = 0.$$

The definition of differentiability can be interpreted in two equivalent ways. One is that a function is differentiable at a point \vec{a} if there exists a “good” linear approximation to the function at \vec{a} (where “good” is defined to mean that it satisfies the limit). Alternatively, one could think about the graph of a function as a surface. A function f is then differentiable at \vec{a} if there exists a tangent plane to the surface given by the graph of f at the point $(\vec{a}, f(\vec{a}))$. Being able to switch back between these two perspectives will aid your intuition.

However, there's a big problem with this definition of differentiability. It involves a multivariable limit⁷. Multivariable limits are much more subtle than single-variable limits, and they are covered in the next section. But before we do that, we will further explore the consequences of differentiability.

⁶ This phrasing comes from language used in Taylor approximations.

⁷ If you already know how to evaluate multi-dimensional limits, this isn't much of a problem.

Consequences of Differentiability

Recall, directional derivatives of affine functions can be expressed as dot products with the gradient vector. This is true in general for differentiable functions.

Theorem 4.3.2 If $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is differentiable at the point \vec{a} , then

$$D_{\vec{a}}f(\vec{v}) = \nabla f(\vec{a}) \cdot \vec{v}$$

for all $\vec{v} \in \mathbb{R}^n$.

Proof. Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function that is differentiable at the point \vec{a} , and let $L : \mathbb{R}^n \rightarrow \mathbb{R}$ be the unique affine function such that $f(\vec{a}) = L(\vec{a})$ and

$$\lim_{\vec{u} \rightarrow \vec{0}} \frac{f(\vec{a} + \vec{u}) - L(\vec{a} + \vec{u})}{\|\vec{u}\|} = 0. \quad (4.4)$$

By Theorem 4.3.1, $D_{\vec{a}}L(\vec{v}) = \nabla L(\vec{a}) \cdot \vec{v}$ for all \vec{v} .

Now, let us recall a fact about limits. If h, g are functions and $\lim_{x \rightarrow a} h(x)$ exists and $\lim_{x \rightarrow a} (h(x) + g(x))$ exists, then we must have $\lim_{x \rightarrow a} g(x)$ exists and

$$\lim_{x \rightarrow a} (h(x) + g(x)) = \lim_{x \rightarrow a} h(x) + \lim_{x \rightarrow a} g(x).$$

Putting it all together, Equation (4.4) implies for any \vec{v} ,

$$0 = 0\|\vec{v}\| = \|\vec{v}\| \lim_{h \rightarrow 0} \frac{f(\vec{a} + h\vec{v}) - L(\vec{a} + h\vec{v})}{\|h\vec{v}\|} = \lim_{h \rightarrow 0} \frac{f(\vec{a} + h\vec{v}) - L(\vec{a} + h\vec{v})}{h}$$

exists. In particular, this shows that $D_{\vec{a}}(f - L)(\vec{v}) = 0$. Applying our fact about limits, since $D_{\vec{a}}L(\vec{v})$ exists and $D_{\vec{a}}(f - L)(\vec{v})$ exists, we must have that $D_{\vec{a}}f(\vec{v})$ exists and that

$$D_{\vec{a}}f(\vec{v}) - D_{\vec{a}}L(\vec{v}) = 0.$$

In other words,

$$D_{\vec{a}}f(\vec{v}) = D_{\vec{a}}L(\vec{v}) = \nabla L(\vec{a}) \cdot \vec{v}. \quad (4.5)$$

To finish off the proof, notice that

$$\frac{\partial L}{\partial x}(\vec{a}) = D_{\vec{a}}L(\hat{\mathbf{x}}) = D_{\vec{a}}f(\hat{\mathbf{x}}) = \frac{\partial f}{\partial x}(\vec{a}).$$

Repeating this argument with $\hat{\mathbf{y}}, \hat{\mathbf{z}}$, etc. shows that the partial derivatives of f at \vec{a} must equal the partial derivatives of L at \vec{a} . Therefore,

$$\nabla L(\vec{a}) = \nabla f(\vec{a}).$$

Substituting $\nabla L(\vec{a})$ with $\nabla f(\vec{a})$ in Equation (4.5) completes the proof. ■

Theorem 4.3.2 shows that if a function is differentiable, directional derivatives can be computed with dot products instead of limits!

■ **Example 4.14** Find the directional derivative of the function $f : \mathbb{R}^3 \rightarrow \mathbb{R}$ given by $f(x, y, z) = x^2 - z^3$ at the point $\vec{a} = (1, 1, 1)$ in the direction $\vec{v} = (1, 2, 3)$.

Since f is a polynomial it is differentiable. Computing,

$$\nabla f(1, 1, 1) = \begin{bmatrix} 2 \\ 0 \\ -3 \end{bmatrix}$$

and so

$$D_{\vec{a}}f(\vec{v}) = \nabla f(\vec{a}) \cdot \vec{v} = -7.$$

■

Geometric Interpretation of the Gradient

Suppose $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is a differentiable function and \vec{u} is a unit vector. We know that

$$D_{\vec{a}}f(\vec{u}) = \nabla f(\vec{a}) \cdot \vec{u} = \|\nabla f(\vec{a})\| \|\vec{u}\| \cos \theta = \|\nabla f(\vec{a})\| \cos \theta,$$

where θ is the angle between \vec{u} and $\nabla f(\vec{a})$. Since $\cos \theta \leq 1$, we see that $D_{\vec{a}}f(\vec{u})$ is maximized when $\theta = 0$. In other words, *the largest directional derivative of f occurs in the direction of the gradient of f* . Thus, the gradient may be interpreted as pointing in the direction of greatest change.

XXX Figure

Continuing this further, if $D_{\vec{a}}f(\vec{u}) = 0$ then one of two things must occur. Either $\nabla f(\vec{a}) = \vec{0}$ or $\theta = 90^\circ$. Therefore, if $\nabla f(\vec{a}) \neq \vec{0}$, the direction in which the function doesn't change at all is orthogonal to the gradient.

We have another word for paths on which a function doesn't change—level curves. If we plot level curves for a function, the gradient of the function must always be orthogonal to the level curves.

XXX Figure

Exercises for 4.3

4.4 Multidimensional Limits and Continuity

Most users of mathematics don't worry about things that might go wrong with the functions they use to represent physical quantities. They tend to assume that functions are differentiable when derivatives are called for (except possibly for a finite set of isolated points), and they assume all functions which need to be integrated are continuous so the integrals will exist. For much of the period during which Calculus was developed (during the 17th and 18th centuries), mathematicians also did not bother themselves with such matters. Unfortunately, during the 19th century, mathematicians discovered that general functions could behave in unexpected and subtle ways, so they began to devote much more time to careful formulation of definitions and careful proofs in analysis. This is an aspect of mathematics which is covered in courses in real analysis, so we won't devote much time to it in this course. (You may have noticed that we didn't worry about the existence of derivatives in our discussion of velocity and acceleration). However, for functions of several variables, lack of rigor can be more troublesome than in the one variable case, so we briefly devote some attention to such questions. In this section, we shall discuss the concepts of *limit* and *continuity* for functions $f : \mathbb{R}^2 \rightarrow \mathbb{R}$. The big step, it turns out, is going from one independent variable to two. Once you understand that, going to three or more independent variables introduces few additional difficulties.

Let $\vec{r}_0 = (x_0, y_0)$ be (the position vector of) a point in the domain of the function f . We want to define the concept to be expressed symbolically

$$\lim_{\vec{r} \rightarrow \vec{r}_0} f(\vec{r}) = L \quad \text{or} \quad \lim_{(x,y) \rightarrow (x_0,y_0)} f(x,y) = L.$$

We start with two examples which illustrate the concept and some differences from the single variable case.

■ **Example 4.15** Let $f(x, y) = x^2 + 2y^2$, and consider the nature of the graph of f near the point $(1, 2)$. As we saw in the previous section, the graph is an elliptic paraboloid, the locus of $z = x^2 + 2y^2$. In particular, the surface is quite smooth, and if (x, y) is a point in the domain *close to* $(1, 2)$, then $f(x, y)$ will be very close to the value of the function there, $f(1, 2) = 1^2 + 2(2^2) = 9$. Thus, it makes sense to assert that

$$\lim_{(x,y) \rightarrow (1,2)} x^2 + 2y^2 = 9.$$

XXX Figure

■

In Example 4.15, the limit was determined simply by evaluating the function at the desired point. You may remember that in the single variable case, you cannot always do that. For example, putting $x = 0$ in $\sin x/x$ yields the meaningless expression $0/0$, but $\lim_{x \rightarrow 0} \sin x/x$ is known to be 1. Usually, it requires some ingenuity to find such examples in the single variable case, but the next example shows that fairly simple formulas can lead to unexpected difficulties for functions of two or more variables.

■ **Example 4.16** Let

$$f(x, y) = \frac{x^2 - y^2}{x^2 + y^2} \quad \text{for } (x, y) \neq (0, 0).$$

What does the graph of this function look like in the vicinity of the point $(0, 0)$? (Since, $(0, 0)$ is not in the domain of the function, it does not make sense to talk about $f(0, 0)$, but we can

still seek a “limit.”) The easiest way to answer this question is to switch to polar coordinates. Using $x = r \cos \theta$, $y = r \sin \theta$, we find

$$f(\vec{r}) = f(x, y) = \frac{r^2 \cos^2 \theta - r^2 \sin^2 \theta}{r^2 \cos^2 \theta + r^2 \sin^2 \theta} = \cos^2 \theta - \sin^2 \theta = \cos 2\theta.$$

Thus, $f(\vec{r}) = f(x, y)$ is independent of the polar coordinate r and depends only on θ . As $r = \|\vec{r}\| \rightarrow 0$ with θ fixed, $f(\vec{r})$ is constant, and equal to $\cos 2\theta$, so, if it ‘approaches’ a limit, that limit would have to be $\cos 2\theta$. Unfortunately, $\cos 2\theta$ varies between -1 and 1 , so it does not make sense to say $f(\vec{r})$ has a limit as $\vec{r} \rightarrow \vec{0}$. You can get some idea of what the graph looks like by studying the level curves which are pictured in the diagram. For each value of θ , the function is constant, so the level curves consist of rays emanating from the origin, as indicated. On any such ray, the graph is at some constant height z with z taking on *every value* between -1 and $+1$. ■

In general, the statement

$$\lim_{\vec{r} \rightarrow \vec{r}_0} f(\vec{r}) = L$$

will be taken to mean that $f(\vec{r})$ is *close to* L whenever \vec{r} is *close to* \vec{r}_0 . To make this completely rigorous, we will give an “ ε - δ definition” of limit.

Definition 4.4.1 — Limit. Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function. For a point \vec{a} , we say *the limit of f as \vec{x} approaches \vec{a} is $L \in \mathbb{R}$* , written

$$\lim_{\vec{x} \rightarrow \vec{a}} f(\vec{x}) = L,$$

if for all numbers $\varepsilon > 0$ there exists a number $\delta > 0$ so that whenever

$$0 < \|\vec{x} - \vec{a}\| < \delta \quad \text{we have} \quad |f(\vec{x}) - L| < \varepsilon.$$

In this statement, $0 < \|\vec{x} - \vec{a}\| < \delta$ asserts that the distance from \vec{x} to \vec{a} is less than δ . Since δ is thought of as small, the inequality makes precise the meaning of “ \vec{x} is close to \vec{a} , but not equal to it.” Similarly, $|f(\vec{x}) - L| < \varepsilon$ catches the meaning of “ $f(\vec{x})$ is close to L .” Note that we never consider the case $\vec{x} = \vec{a}$, so the value of $f(\vec{r}_0)$ is not relevant in checking the limit as $\vec{x} \rightarrow \vec{a}$. (It is not even necessary that $f(\vec{r})$ be well defined at $\vec{x} = \vec{a}$.)

Limits for functions of several variables behave formally much the same as limits for functions of one variable. Thus, you may calculate the limit of a sum by taking the sum of the limits, and similarly for products and quotients (except that for quotients the limit of the denominator should not be zero). The understanding you gained of these matters in the single variable case should be an adequate guide to what to expect for several variables. If you never really understood all this before, we won’t enlighten you much here. You will have to wait for a course in real analysis for real understanding.

Continuity

In Example 4.15, the limit was determined simply by evaluating the function at the point. This is certainly not always possible because the value of the function may be irrelevant or there may be no meaningful way to attach a value. Functions for which it is always possible to find

the limit this way are called *continuous*. (This is the same notion as for functions of a single scalar variable). More precisely, we say that f is continuous at a point \vec{r}_0 if the point is in its domain (i.e., $f(\vec{r}_0)$ is defined) and

$$\lim_{\vec{r} \rightarrow \vec{r}_0} f(\vec{r}) = f(\vec{r}_0).$$

Points at which this fails are called *discontinuities* or sometimes *singularities*. (The latter term is also sometimes reserved for less serious kinds of mathematical pathology.) It sometimes happens, that a function f has a well defined limit L at a point \vec{r}_0 which does not happen to be in the domain of the function, i.e., $f(\vec{r}_0)$ is not defined. (In the single variable case, $\sin x/x$ at $x = 0$ is a good example). Then we can extend the domain of the function to include the point \vec{r}_0 by defining $f(\vec{r}_0) = L$. Thus the original function had a discontinuity, but it can be eliminated simply by extending the definition of the function. In this case, the discontinuity is called *removable*. As Example 4.16 shows, there are functions with discontinuities which cannot be defined away no matter what you try.

A function without discontinuities is called continuous. Continuous functions have graphs which look reasonably smooth. They don't have big holes or sudden jumps, but as we shall see later, they can still look pretty bizarre. Usually, just knowing that a function is continuous won't be enough to make it a good candidate to represent a physical quantity. We shall also want to be able to take derivatives and do the usual things one does in differential calculus, but as you might expect, this is somewhat more involved than it is in the single variable case.

Exercises for 4.4

4.5 Polynomial Approximation

When we try to model the real world with mathematics, we often encounter very complicated functions. The more complicated a function, the more difficult it is to analyze, so a common mathematical strategy is to substitute complicated functions for simple ones.

Among the simplest mathematical functions are *polynomials*, and in the multivariable context, this means polynomials of multiple variables. Before we dive into the details of polynomial approximations, let us introduce some terminology.

Definition 4.5.1 — Polynomial. Let x_1, x_2, \dots, x_n denote variables. A *degree- m term* is a product of the form

$$x_1^{\alpha_1} x_2^{\alpha_2} \cdots x_n^{\alpha_n}$$

where $\alpha_i \geq 0$ and $\sum_{i=1}^n \alpha_i = m$.

A *degree- m polynomial* is a linear combination of terms of degree at most m . A *degree-0 polynomial* is constant.

In the definition of polynomial, we always interpret x_i^0 as 1. You're intimately familiar with polynomials of a single variable. For instance, $x^3 - 2x$ is a degree-3 polynomial. However, when it comes to the multivariable context, $x^2y - 2x$ and $xyz - 2x$ are also degree-3 polynomials.

Affine functions are all degree-1 polynomials, and the definition of $f : \mathbb{R}^n \rightarrow \mathbb{R}$ being differentiable depended on being able to find an affine function that approximated f very well. There is a similar concept for higher degree polynomial functions.

Definition 4.5.2 — Taylor Approximation. Given a continuous function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, a degree m polynomial p , and a point $\vec{a} \in \mathbb{R}^n$, we call p the *degree- m Taylor approximation of f at \vec{a}* if

$$\lim_{\vec{r} \rightarrow \vec{0}} \frac{f(\vec{a} + \vec{r}) - p(\vec{a} + \vec{r})}{\|\vec{r}\|^m} = 0.$$

Note that if $\|\vec{r}\| < 1$, then $\|\vec{r}\| > \|\vec{r}\|^2 > \cdots > \|\vec{r}\|^m$. Thus,

$$\frac{f(\vec{a} + \vec{r}) - p(\vec{a} + \vec{r})}{\|\vec{r}\|^m} < \cdots < \frac{f(\vec{a} + \vec{r}) - p(\vec{a} + \vec{r})}{\|\vec{r}\|},$$

and so having a degree- m Taylor approximation is a more stringent demand than merely being differentiable.

■ **Example 4.17** Let $f(x, y) = (x - 2)^2 + y^2$. Find the degree 1, 2, and 3 Taylor approximations to f at $(0, 0)$.

On the one hand, since f is itself a polynomial, we know it is differentiable. Therefore it has an affine approximation coming from the definition of differentiability which coincides with a degree-1 Taylor approximation. Further f is a degree-2 polynomial and so it is its own degree-2 and degree-3 Taylor approximation.

On the other hand, let's imagine we weren't so clever. Let $p_1(x, y) = ax + by + c$ be a degree-1 polynomial. If p is a degree-1 Taylor approximation to f at $\vec{0}$, then

$$\lim_{\vec{r} \rightarrow \vec{0}} \frac{f(\vec{0} + \vec{r}) - p_1(\vec{0} + \vec{r})}{\|\vec{r}\|} = 0.$$

Let r_x and r_y stand for the x and y components of \vec{r} . Expanding, we see

$$\begin{aligned}\lim_{\vec{r} \rightarrow \vec{0}} \frac{f(\vec{0} + \vec{r}) - p_1(\vec{0} + \vec{r})}{\|\vec{r}\|} &= \lim_{\vec{r} \rightarrow \vec{0}} \frac{(r_x - 2)^2 + r_y^2 - (ar_x + br_y + c)}{\sqrt{r_x^2 + r_y^2}} \\ &= \lim_{\vec{r} \rightarrow \vec{0}} \frac{r_x^2 - 4r_x + 4 + r_y^2 - ar_x - br_y - c}{\sqrt{r_x^2 + r_y^2}} = 0\end{aligned}$$

if and only if $a = -4$, $b = 0$, and $c = 4$. Thus $p_1(x, y) = -4x + 4$ is the unique degree-1 Taylor approximation to f .

We can use a similar process to find the degree-2 Taylor approximation to f . Let $p_3(x, y) = ax^2 + by^2 + cxy + dx + ey + h$ be a degree-2 polynomial. Then

$$\begin{aligned}\lim_{\vec{r} \rightarrow \vec{0}} \frac{f(\vec{0} + \vec{r}) - p_3(\vec{0} + \vec{r})}{\|\vec{r}\|^2} &= \lim_{\vec{r} \rightarrow \vec{0}} \frac{(r_x - 2)^2 + r_y^2 - (ar_x^2 + br_y^2 + cr_xr_y + dr_x + er_y + h)}{r_x^2 + r_y^2} \\ &= \lim_{\vec{r} \rightarrow \vec{0}} \frac{r_x^2 - 4r_x + 4 + r_y^2 - (ar_x^2 + br_y^2 + cr_xr_y + dr_x + er_y + h)}{r_x^2 + r_y^2} = 0\end{aligned}$$

if and only if $a = 1$, $b = 1$, $c = 0$, $d = -4$, $e = 0$, and $h = 4$. Therefore the degree-2 Taylor approximation of f is $p_3(x, y) = x^2 - 4x + 4 + y^2 = (x - 2)^2 + y^2$, which is just f itself.

We could continue on to use the limit definition to find a degree-3 Taylor approximation to f , or we could notice that $f(x, y) - p_2(x, y) = 0$ for all x and y . Since $\lim_{t \rightarrow 0} 0/t = 0$ and p_2 satisfies the definition of a degree-3 polynomial (none of its terms exceed degree 3), p_2 must be a degree-3 Taylor approximation of f . Further, if we used the limit definition, we would conclude that p_2 is actually the *unique* degree-3 Taylor approximation of f . ■

There's a faster way to find degree- m Taylor approximations than taking a limit each time. We can rely on higher-order partial derivatives.

Higher-order Partial Derivatives

Recall that if $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is a differentiable function, then $\frac{\partial f}{\partial x_i}(\vec{a})$ is a well-defined number for all $\vec{a} \in \mathbb{R}^n$. Thus, we may view $\frac{\partial f}{\partial x_i}(\vec{a})$ as a function of \vec{a} . From this perspective,

$$\frac{\partial f}{\partial x_i} : \mathbb{R}^n \rightarrow \mathbb{R}$$

is just another multivariable function. If this function is differentiable, we may consider

$$\frac{\partial}{\partial x_j} \frac{\partial f}{\partial x_i}(\vec{a}),$$

and so on. A function is called an *mth order partial derivative* if this process has been repeated m times. The nesting fraction notation can get pretty awkward, pretty quickly, so, modeling

after the notation from single-variable calculus, we write

$$\frac{\partial^2 f}{\partial x_j \partial x_i}(\vec{a}) = \frac{\partial}{\partial x_j} \left(\frac{\partial f}{\partial x_i}(\vec{a}) \right).$$

The expression $\frac{\partial^2 f}{\partial x_i \partial x_j}(\vec{a})$ is an *iterated partial derivative*. If $x_i \neq x_j$, it is called a *mixed partial derivative*. If $x_i = x_j$, we further simplify our notation and write

$$\frac{\partial^2 f}{\partial x_i^2}(\vec{a}).$$

■ **Example 4.18** Compute $\frac{\partial^2 f}{\partial x \partial y}(x, y)$ where $f(x, y) = e^{x-2xy}$.

XXX Finish ■

Strictly speaking, $\frac{\partial^2 f}{\partial x_j \partial x_i}(\vec{a})$ means take a derivative with respect to x_i first and then with respect to x_j , and a priori, it could be that $\frac{\partial^2 f}{\partial x_j \partial x_i}(\vec{a}) \neq \frac{\partial^2 f}{\partial x_i \partial x_j}(\vec{a})$. However, if our function has a degree-2 Taylor approximation, these two quantities are always equal.

Theorem 4.5.1 — Clairaut's Theorem. If $f : \mathbb{R}^n \rightarrow \mathbb{R}$ has a degree-2 Taylor approximation at $\vec{a} \in \mathbb{R}^n$, then

$$\frac{\partial^2 f}{\partial x_j \partial x_i}(\vec{a}) = \frac{\partial^2 f}{\partial x_i \partial x_j}(\vec{a})$$

for all choices of variables x_i, x_j .

Clairaut's Theorem will apply to most functions we encounter, but some fairly reasonable looking functions can fail.

■ **Example 4.19** Consider $f(x, y) = \begin{cases} \frac{xy(x^2-y^2)}{x^2+y^2} & \text{if } (x, y) \neq (0, 0) \\ 0 & \text{if } (x, y) = (0, 0) \end{cases}$.

XXX Example where mixed partials aren't equal ■

Higher order partial derivatives are interesting in their own right, but the following theorem makes them useful for finding Taylor approximations.

Theorem 4.5.2 Suppose $f : \mathbb{R}^n \rightarrow \mathbb{R}$ has a degree- m Taylor approximation $p : \mathbb{R}^n \rightarrow \mathbb{R}$ at $\vec{a} \in \mathbb{R}^n$. Then, every i th-order partial derivative of f at the point \vec{a} agrees with every i th-order partial derivative of p at the point \vec{a} for $i \leq m$. Symbolically,

$$\frac{\partial^i f}{\partial x_{j_1} \cdots \partial x_{j_i}}(\vec{a}) = \frac{\partial^i p}{\partial x_{j_1} \cdots \partial x_{j_i}}(\vec{a})$$

for $i \leq m$.

Theorem 4.5.2 is not difficult to prove, but it is a notational nightmare, so we won't prove it here. Theorem 4.5.2 also simplifies the task of finding a Taylor approximation (provided we know one already exists).

■ **Example 4.20** XXX Finish

Uses of Polynomial Approximations

As we have seen previously if $f : \mathbb{R}^n \rightarrow \mathbb{R}$ has a degree-1 Taylor approximation at \vec{a} , that approximation captures all information about directional derivatives of f at \vec{a} (we used this fact to prove $D_{\vec{a}}f(\vec{v}) = \nabla f(\vec{a}) \cdot \vec{v}$). Similarly, a degree- m Taylor approximation to f at \vec{a} captures all m -th order derivative information. This statement can be made precise with a theorem.

Theorem 4.5.3 Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function and let $p : \mathbb{R}^n \rightarrow \mathbb{R}$ be the degree- m Taylor approximation to f at \vec{a} . If $\vec{r} : \mathbb{R} \rightarrow \mathbb{R}^n$ is an m -times differentiable parameterization of a curve with $\vec{p}(0) = \vec{a}$, then

$$\begin{aligned}(f \circ \vec{r})'(0) &= (p \circ \vec{r})'(0) \\ (f \circ \vec{r})''(0) &= (p \circ \vec{r})''(0) \\ &\vdots \\ (f \circ \vec{r})^{(m)}(0) &= (p \circ \vec{r})^{(m)}(0).\end{aligned}$$

That is, the first m derivatives of $f \circ \vec{r}$ and $p \circ \vec{r}$ agree at 0.

The implications of Theorem 4.5.3 are far reaching. You can use polynomial approximations to study curvature, concavity, local minimums and maximums, etc.—and we will.

Approximations of Parameterizations

Thus far we have only concerned ourselves with functions $f : \mathbb{R}^n \rightarrow \mathbb{R}$ where the co-domain is \mathbb{R} . However, there are a host of interesting functions of the form $\vec{g} : \mathbb{R}^n \rightarrow \mathbb{R}^m$. In particular, a parameterization of a curve takes the form $\vec{r} : \mathbb{R} \rightarrow \mathbb{R}^m$ and a parameterization of a surface takes the form $\vec{s} : \mathbb{R}^2 \rightarrow \mathbb{R}^m$.

We can define first-order and higher-order approximations to functions $\vec{g} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ by using components.

Definition 4.5.3 Let $\vec{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ have component functions $f_1, f_2, \dots, f_m : \mathbb{R}^n \rightarrow \mathbb{R}$. The order n approximation to f at the point $\vec{a} \in \mathbb{R}^n$ is the function

$$\vec{g}(\vec{x}) = \begin{bmatrix} g_1(\vec{x}) \\ g_2(\vec{x}) \\ \vdots \\ g_m(\vec{x}) \end{bmatrix}$$

where g_i is an n th order polynomial approximation to f_i at \vec{c} for each $i = 1, \dots, m$.

■ **Example 4.21** Let $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ be the parameterization of \mathbb{R}^2 coming from polar coordinates. That is $\vec{p}(r, \theta) = (r \cos \theta, r \sin \theta)$. Find the first-order approximation to \vec{p} at $(2, 0)$.

Let $x(r, \theta) = r \cos \theta$ and $y(r, \theta) = r \sin \theta$ be the component functions of \vec{p} . We need to find first-order approximations to x and y at the point $(2, 0)$. Using gradients we can quickly compute these linear approximations.

$$L_x(r, \theta) = \nabla x(2, 0) \cdot (r - 2, \theta - 0) + x(2, 0) = r$$

and

$$L_y(r, \theta) = \nabla y(2, 0) \cdot (r - 2, \theta - 0) + y(2, 0) = 2\theta.$$

Therefore, the first-order approximation to \vec{p} at the point $(2, 0)$ is

$$\vec{L}_{\vec{p}}(r, \theta) = \begin{bmatrix} r \\ 2\theta \end{bmatrix}.$$

■

In a similar way, we can define what it means for a vector-valued function $\vec{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ to be differentiable.

Definition 4.5.4 — Differentiable (Vector-valued Function). Let $\vec{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ have component functions $f_1, f_2, \dots, f_m : \mathbb{R}^n \rightarrow \mathbb{R}$. The vector-valued function \vec{f} is *differentiable* if f_i is differentiable for $i = 1, \dots, m$.

4.6 Optimization

Optimization is a fancy word for finding minimums and maximums. If $f : \mathbb{R} \rightarrow \mathbb{R}$ is a function of one variable, the procedure should be familiar: find critical points of f and check whether they are minimums, maximums, or neither; then, check the boundary. For a multivariable function, the procedure is similar.

Let's consider an example. Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be defined by $f(x, y) = 4 - x^2 - y^2$. This example is nice because we already know where the maximum occurs!

XXX Figure

A global maximum of 4 for f occurs at $(0, 0)$. Now, suppose $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^2$, given by

$$\vec{p}(t) = t\vec{d} + \vec{0}$$

is the parameterization of a line passing through $(0, 0)$. Since 4 is a global maximum of f achieved at $(0, 0)$ and \vec{p} passes through $(0, 0)$, the function $f \circ \vec{p}$ must attain a global maximum of 4 exactly when $\vec{p}(t) = \vec{0}$.

XXX Figure

Since $f \circ \vec{p} : \mathbb{R} \rightarrow \mathbb{R}$ is a function of one variable, all of its maximums must occur at either a critical point or a boundary point of the domain. Since the domain is unbounded and $(f \circ \vec{p})'$ always exists, we know a maximum of $f \circ \vec{p}$ must occur when

$$(f \circ \vec{p})'(t) = 0. \tag{4.6}$$

If both f and \vec{p} were single-variable functions, we could apply the chain rule to Equation (4.6). Unfortunately they are not. So, before we continue, let's see if we can figure out a multivariable chain rule.

The Multivariable Chain Rule

In mathematics, when you don't know the answer, a technique⁸ is to replace your objects of study with something simpler. In differential calculus, the simplest thing you could hope for is something flat (a line, a plane, etc.). These are precisely characterized by affine functions.

Suppose $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ is an affine function and $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^2$ is a differentiable parameterization of a curve. Since f is an affine function, it can be written as $f(x, y) = \alpha_x x + \alpha_y y + c$ or alternatively as $f(\vec{x}) = \vec{\alpha} \cdot \vec{x} + c$. Let $p_x : \mathbb{R} \rightarrow \mathbb{R}$ and $p_y : \mathbb{R} \rightarrow \mathbb{R}$ be the components of \vec{p} . Now we have

$$f \circ \vec{p}(t) = f(\vec{p}(t)) = f(p_x(t), p_y(t)) = \alpha_x p_x(t) + \alpha_y p_y(t) + c.$$

This is a function of one variable so we can take a derivative as usual:

$$(f \circ \vec{p})'(t) = \alpha_x p'_x(t) + \alpha_y p'_y(t) = \vec{\alpha} \cdot \vec{p}'(t).$$

But, since f is an affine function, $\nabla f(\vec{x}) = \vec{\alpha}$ for all \vec{x} , and so we in fact have

$$(f \circ \vec{p})'(t) = \nabla f(\vec{x}) \cdot \vec{p}'(t). \quad (4.7)$$

For an affine function, ∇f is constant, so we don't need to specify what \vec{x} is for Equation (4.7) to make sense. However, without knowing what \vec{x} should be, it's difficult to see the *meaning* behind Equation (4.7).

Let's consider a more general case. Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be a function that is differentiable at the point \vec{a} and let $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^2$ be the parameterization of a curve so that $\vec{p}(t_0) = \vec{a}$. Since f is differentiable at \vec{a} , we know there is an affine function $L : \mathbb{R}^2 \rightarrow \mathbb{R}$ that approximates f very well near the point \vec{a} . This means,

$$f \circ \vec{p}(t) \approx L \circ \vec{p}(t) \quad \text{when} \quad t \approx t_0.$$

From Equation (4.7), we know $(L \circ \vec{p})'(t) = \nabla L(\vec{x}) \cdot \vec{p}'(t)$ for any \vec{x} . Further, $\nabla L(\vec{a}) = \nabla f(\vec{a})$ and $\vec{a} = \vec{p}(t_0)$. Putting this all together,

$$\begin{aligned} (L \circ \vec{p})'(t_0) &= \nabla L(\vec{a}) \cdot \vec{p}'(t_0) \\ &= \nabla f(\vec{p}(t_0)) \cdot \vec{p}'(t_0) = D_{\vec{p}(t_0)} f(\vec{p}'(t_0)). \end{aligned}$$

That equation is a mouthful, but the right hand side involves only f and \vec{p} . If we could show that $(L \circ \vec{p})'(t_0) = (f \circ \vec{p})'(t_0)$, we would have a formula for the multivariable chain rule. Fortunately, we have the following theorem.

Theorem 4.6.1 Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function that is differentiable at \vec{a} and let $L_{\vec{a}} : \mathbb{R}^n \rightarrow \mathbb{R}$ be the affine function coming from the definition of differentiability. Let $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^n$ satisfy $\vec{p}(t_0) = \vec{a}$. Then,

$$(f \circ \vec{p})'(t_0) = (L_{\vec{a}} \circ \vec{p})'(t_0).$$

⁸ In math, some things are called tricks and some techniques. The difference is: a technique is used more than once.

We will not prove Theorem 4.6.1 in the general context. However, if \vec{p} parameterizes a line, the proof is straightforward. Suppose \vec{p} parameterizes a line, and for simplicity, assume $t_0 = 0$. Then $\vec{p}(t) = \vec{a} + t\vec{v}$ for some \vec{v} . Now,

$$(f \circ \vec{p})'(0) = D_{\vec{a}}f(\vec{v}) \quad \text{and} \quad (L_{\vec{a}} \circ \vec{p})'(0) = D_{\vec{a}}L_{\vec{a}}(\vec{v}),$$

and we have previously established that $D_{\vec{a}}f(\vec{v}) = D_{\vec{a}}L_{\vec{a}}(\vec{v})$. Proving Theorem 4.6.1 in the general case amounts to showing that you can replace \vec{p} by its affine approximation, which parameterizes a line.

We may now write down the multivariable chain rule.

Theorem 4.6.2 — Multivariable Chain Rule. If $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is differentiable at \vec{a} and $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^n$ is differentiable and satisfies $\vec{p}(t_0) = \vec{a}$, then

$$(f \circ \vec{p})'(t_0) = D_{\vec{p}(t_0)}f(\vec{p}'(t_0)).$$

Equivalently,

$$(f \circ \vec{p})'(t_0) = \nabla f(\vec{p}(t_0)) \cdot \vec{p}'(t_0).$$

The Multivariable First-Derivative Test

We can now develop the multivariable first derivative test for local extrema. Recall that if $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and $\vec{p} : \mathbb{R} \rightarrow \mathbb{R}^n$ is a parameterization such that f attains a local maximum at $\vec{p}(0)$, then

$$(f \circ \vec{p})'(0) = 0.$$

Using the multivariable chain rule we see

$$(f \circ \vec{p})'(0) = D_{\vec{p}(0)}f(\vec{p}'(0)) = 0.$$

This is true for any parameterization that passes through a local extreme. Since we can pass through the local extreme from any direction, we have the following theorem.

Theorem 4.6.3 — Multivariable First Derivative Test. If $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is differentiable at $\vec{a} \in \mathbb{R}^n$ and f attains a local maximum or local minimum at \vec{a} then

$$D_{\vec{a}}f(\vec{v}) = 0$$

for all vectors \vec{v} . Equivalently,

$$\nabla f(\vec{a}) = \vec{0}.$$

Note that Theorem 4.6.3 does not say that if $\nabla f(\vec{a}) = \vec{0}$ then f must have a local max or min at \vec{a} . It only says that if you have a local max or min at \vec{a} , then you have $\nabla f(\vec{a}) = \vec{0}$.

This is just like the single-variable first derivative test. If $g : \mathbb{R} \rightarrow \mathbb{R}$ and $g'(a) = 0$, it need not be the case that g has a local max or min at a .

XXX Figure

Definition 4.6.1 — Critical Point. For a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, the point $\vec{a} \in \mathbb{R}^n$ is called a *critical point* if f is not differentiable at \vec{a} or if $\nabla f(\vec{a}) = \vec{0}$.

■ **Example 4.22** Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be given by $f(x, y) = 4 - x^2 - y^2$. Find all critical points of f .

Computing, $\nabla f(x, y) = (2x, -2y)$. Since f is differentiable everywhere, the only critical points of f must occur when $\nabla f(x, y) = (0, 0)$. Thus the only critical points are at $(x, y) = (0, 0)$. ■

■ **Example 4.23** XXX Finish (Find Local Max) ■

■ **Example 4.24** XXX Finish Find (Local saddle point) ■

The Multivariable Second-Derivative Test

Inspired by the single-variable second-derivative test, we will create a multivariable second-derivative test. First, a reminder about the single-variable case.

Theorem 4.6.4 — Single-variable Second-Derivative Test. Suppose $f : \mathbb{R} \rightarrow \mathbb{R}$ is a twice differentiable function and $f'(a) = 0$. Then, if $f''(a) > 0$, it must be that $f(a)$ is a local minimum and if $f''(a) < 0$, it must be that $f(a)$ is a local maximum. If $f''(a) = 0$, nothing can be concluded about whether $f(a)$ is a local maximum, local minimum, or neither.

Now, imagine $f : \mathbb{R}^n \rightarrow \mathbb{R}$ has a local maximum at \vec{a} . Let \vec{u} be a unit vector and let ℓ be the line with direction \vec{u} passing through \vec{a} . Then, ℓ has a parameterization $\vec{p}(t) = \vec{a} + t\vec{u}$. The function $f \circ \vec{p}$ can be thought of as the restriction of f to ℓ . Since f has a local max at \vec{a} , we know $f \circ \vec{p}$ has a local max at 0 (since $\vec{p}(0) = \vec{a}$), and we would expect $f \circ \vec{p}$ to be a concave-down function.

XXX Figure

Thus, we expect $(f \circ \vec{p})''(0) < 0$. We call $(f \circ \vec{p})''(0)$ a *second order directional derivative in the direction \vec{u}* . This inspires the multivariable second derivative test.

Theorem 4.6.5 — Multivariable Second-Derivative Test. Suppose $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is differentiable at \vec{a} and $D_{\vec{a}}f(\vec{v}) = 0$ for all \vec{v} . Let $\vec{p}_{\vec{v}}(t) = \vec{a} + t\vec{v}$. If

$$(f \circ \vec{p}_{\vec{v}})''(0) < 0$$

for all \vec{v} , then f has a local maximum at \vec{a} and if

$$(f \circ \vec{p}_{\vec{v}})''(0) > 0$$

for all \vec{v} , then f has a local minimum at \vec{a} . Further, if there are \vec{v}_1 and \vec{v}_2 such that

$$(f \circ \vec{p}_{\vec{v}_1})''(0) < 0 < (f \circ \vec{p}_{\vec{v}_2})''(0),$$

then f has neither a local minimum or a local maximum at \vec{a} . In any other case, $f(\vec{a})$ may have a local minimum, a local maximum, or neither.

The multivariable second-derivative test is more complicated because more things can happen in higher dimensions. Consider the following examples.

■ **Example 4.25** $f(x, y) = x^2 + y^2$

XXX Finish

■

■ **Example 4.26** $f(x, y) = x^2 - y^2$

XXX Finish

■

■ **Example 4.27** $f(x, y) = x^3 + y^3$

XXX Finish

■

For complicated functions, using the multivariable second-derivative test involves a lot of computation. Fortunately, we can use Theorem 4.5.3 to replace our function with a degree-2 Taylor approximation at a critical point. Since all second-order directional derivatives of our function and our Taylor approximation will agree, we can instead use the Taylor approximation to determine whether we have a local minimum, local maximum, or neither.

■ **Example 4.28** $f(x, y) = \frac{1}{1+x^2+y^2}$

XXX Finish

■

Lagrange Multipliers

We've just explored how to find local minimums and maximums on an unconstrained domain, but what if we wish to restrict the domain.

■ **Example 4.29** Let $f(x, y) = (x - 2)^2 + (y - 3)^2$ and let \mathcal{D} be the disk of radius one centered at $\vec{0}$. We wish to find the minimum of $f|_{\mathcal{D}}$. That is, we wish to find the minimum of f restricted to the disk \mathcal{D} .

Computing, the only critical point of f occurs at $(2, 3)$ which is outside of \mathcal{D} . We can conclude that $f|_{\mathcal{D}}$ does not have a minimum on the interior of \mathcal{D} , but $f|_{\mathcal{D}}$ still may have a minimum on the boundary of \mathcal{D} .

The boundary of \mathcal{D} is the unit circle which can be parameterized by $\vec{r}(t) = (\cos t, \sin t)$. Thus, we need to find the minimum of

$$\begin{aligned} f \circ \vec{r}(t) &= (\cos t - 2)^2 + (\sin t - 3)^2 = \cos^2 t + \sin^2 t - 6 \sin t - 4 \cos t + 13 \\ &= -6 \sin t - 4 \cos t + 14. \end{aligned}$$

Since this is a single-variable differentiable function, we need only find the places where $(f \circ \vec{r})'(t) = 0$. In other words, we need to solve

$$4 \sin t - 6 \cos t = 0.$$

This simplifies to $\tan t = 3/2$ and so the equation is satisfied when $t = t_1 = \arctan \frac{3}{2}$ or $t = t_2 = \pi + \arctan \frac{3}{2}$.

Evaluating, $f \circ \vec{r}(t_1) = 14 - 2\sqrt{13}$ and $f \circ \vec{r}(t_2) = 14 + 2\sqrt{13}$, and so $f|_{\mathcal{D}}$ attains a minimum of $14 - 2\sqrt{13}$ at $(\cos t_1, \sin t_1) = (\frac{2}{\sqrt{13}}, \frac{3}{\sqrt{13}})$. ■

Example 4.29 involved complicated, but doable, trigonometry. But, we can make our lives much easier by thinking geometrically for a while.

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and suppose \mathcal{R} is a region whose boundary is parameterized by $\vec{r} : \mathbb{R} \rightarrow \mathbb{R}^n$. To find the extrema of $f|_{\mathcal{R}}$, we first check for critical points in \mathcal{R} , and then we need to check

for minimums and maximums along the boundary of \mathcal{R} . Since \vec{r} parameterizes the boundary of \mathcal{R} , we know

$$0 = (f \circ \vec{r})'(t) = D_{\vec{r}(t)}f(\vec{r}'(t)) = \nabla f(\vec{r}(t)) \cdot \vec{r}'(t)$$

whenever $\vec{r}(t)$ is at a local minimum or local maximum along the boundary. Now, observe that $\vec{r}'(t)$ is one of many vectors tangent to the boundary of \mathcal{R} at the point $\vec{r}(t)$. Let $\vec{T}_{\vec{r}(t)}$ be an arbitrary tangent vector to the boundary of \mathcal{R} . We now have that

$$0 = D_{\vec{r}(t)}f(\vec{T}_{\vec{r}(t)}) = \nabla f(\vec{r}(t)) \cdot \vec{T}_{\vec{r}(t)}$$

whenever $\vec{r}(t)$ is at a local minimum or local maximum along the boundary.

XXX Figure

We're going to generalize slightly more. Suppose $\vec{T} : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is a function such that if \vec{a} is on the boundary of \mathcal{R} , we have $\vec{T}(\vec{a})$ is a tangent vector to the boundary of \mathcal{R} . We won't care about how $\vec{T}(\vec{a})$ is defined if \vec{a} is not on the boundary of \mathcal{R} .

XXX Figure of such a function

Now, if $f|_{\mathcal{R}}$ has a local extreme at some point \vec{a} on the boundary of \mathcal{R} , we know

$$0 = \nabla f(\vec{a}) \cdot \vec{T}(\vec{a}).$$

This gives us another way to approach the problem of finding extrema of $f|_{\mathcal{R}}$.

■ **Example 4.30** Let $f(x, y) = (x - 2)^2 + (y - 3)^2$ and let \mathcal{D} be the disk of radius one centered at $\vec{0}$. Find the minimum of $f|_{\mathcal{D}}$.

As before, we know that $f|_{\mathcal{D}}$ attains its minimum on the boundary of \mathcal{D} which can be parameterized by $\vec{r}(t) = (\cos t, \sin t)$. Now, consider the function $\vec{T}(x, y) = (-y, x)$. \vec{T} has the property that $\vec{T}(\vec{r}(t)) = \vec{r}'(t)$. Thus, if

$$0 = (f \circ \vec{r})'(t)$$

we must have

$$0 = \nabla f(\vec{a}) \cdot \vec{T}(\vec{a})$$

for some $\vec{a} = \vec{r}(t)$ on the boundary of \mathcal{D} . Let $\vec{a} = (x, y)$. Solving, we see

$$\begin{aligned} 0 = \nabla f(x, y) \cdot \vec{T}(x, y) &= \begin{bmatrix} 2x - 4 \\ 2y - 6 \end{bmatrix} \cdot \begin{bmatrix} -y \\ x \end{bmatrix} \\ &= -2xy + 4y + 2xy - 6x = -6x + 4y \end{aligned}$$

and so $y = \frac{3}{2}x$. Finally, since $\vec{a} = (x, y)$ is on the unit circle we must have $x^2 + y^2 = 1$ and so $(x, y) = (\frac{2}{\sqrt{13}}, \frac{3}{\sqrt{13}})$ or $(x, y) = (-\frac{2}{\sqrt{13}}, -\frac{3}{\sqrt{13}})$. Checking each of these points, we see $f|_{\mathcal{D}}$ attains a minimum at $(x, y) = (\frac{2}{\sqrt{13}}, \frac{3}{\sqrt{13}})$. ■

Now, suppose we wish to find the extrema of $f|_{\mathcal{C}}$ and \mathcal{C} can be expressed as a level curve of some function $h : \mathbb{R}^n \rightarrow \mathbb{R}$. Without loss of generality, we may assume that $\mathcal{C} = \{\vec{x} : h(\vec{x}) = 0\}$. (Make sure you understand why we can always assume the level curve takes the form $h(\vec{x}) = 0$.) Again, we have the same geometric intuition as before—we are looking for a place where $\nabla f \cdot \vec{T} = 0$ where \vec{T} is a tangent vector to \mathcal{C} .

Since \mathcal{C} is a level set of h , we know every tangent vector to \mathcal{C} must be orthogonal to ∇h . Thus we have $\nabla h \cdot \vec{T} = 0$ and $\nabla f \cdot \vec{T} = 0$, at a local extreme. Since this is true for any \vec{T} that is a tangent vector at a local extreme, we must have that both ∇h and ∇f are normal vectors for the tangent line/tangent plane to \mathcal{C} at the local extreme. Thus, ∇f and ∇h are parallel and so either $\nabla f = 0$ or $\nabla h = 0$ or $\nabla f = \lambda \nabla h$ for some $\lambda \neq 0$. This observation leads to the method of *Lagrange Multipliers*.

■ **Example 4.31** Let $f(x, y) = (x - 2)^2 + (y - 3)^2$ and let \mathcal{D} be the disk of radius one centered at $\vec{0}$. Find the minimum of $f|_{\mathcal{D}}$.

As before, we know that $f|_{\mathcal{D}}$ attains its minimum on the boundary of \mathcal{D} . Let $h(x, y) = x^2 + y^2 - 1$. Then the boundary of \mathcal{D} is precisely the level curve $h(x, y) = 0$. We therefore seek to find all solutions to the system

$$\begin{cases} \nabla f(x, y) = \lambda \nabla h(x, y) \\ h(x, y) = 0 \end{cases}.$$

Computing ∇f and ∇h our system becomes

$$\begin{cases} 2x - 4 = 2\lambda x \\ 2y - 6 = 2\lambda y \\ x^2 + y^2 = 1 \end{cases}.$$

We don't actually care about what λ is, so we can quickly eliminate it from the equations. Assuming that $x \neq 0$ and $y \neq 0$ we obtain

$$\frac{2x - 4}{2x} = \lambda = \frac{2y - 6}{2y}$$

and so $y = \frac{3}{2}x$. Since we also have $x^2 + y^2 = 1$ we have candidate local extremes at $(x, y) = (\frac{2}{\sqrt{13}}, \frac{3}{\sqrt{13}})$ or $(x, y) = (-\frac{2}{\sqrt{13}}, -\frac{3}{\sqrt{13}})$. Now we must handle the case where $x = 0$ or $y = 0$. These cases give us candidate points of $(x, y) = (0, \pm 1)$ or $(x, y) = (\pm 1, 0)$. Testing these six points we see that a local minimum occurs at $(x, y) = (\frac{2}{\sqrt{13}}, \frac{3}{\sqrt{13}})$. ■

Exercises for 4.6


Chapter 5

Integrals

5.1 Multivariable Integrals

Abstractly, the processes of integration is: (i) chop a region into easy to approximate pieces, (ii) approximate each tiny piece, (iii) add together the approximations, and (iv) take a limit as the size of each piece tends to zero¹. Thus far, the domain of this procedure—the thing we chop into tiny pieces—has been a line or curve. We will now consider integrating over multi-dimensional domains.

We will start by finding the area of a region in the plane. Let $\mathcal{R} \subseteq \mathbb{R}^2$ be the region below the line $y = 1$ and above the curve $y = x^2$. We wish to find the area of \mathcal{R} .




Using the usual calculus strategy, we will chop \mathcal{R} up into little rectangles of width Δx and height Δy . Then

$$\text{area of } \mathcal{R} \approx \sum_{\text{tiny rectangles}} \text{area of tiny rectangle}$$

and

$$\text{area of } \mathcal{R} = \lim_{\Delta x, \Delta y \rightarrow 0} \sum_{\text{tiny rectangles}} \text{area of tiny rectangle}.$$

¹ If your approximations get better and better as the size of each piece shrinks, the whole process works. This isn't the case for all objects, but conditions like differentiability and piecewise smooth are often sufficient. It's annoying to specify exactly what is needed over and over again, so sometimes we just say an object is "nice."



Using integral notation, we would write

$$\text{area of } \mathcal{R} = \int_{\mathcal{R}} dA.$$

Here, dA represents a “tiny area,” the subscript \mathcal{R} represents the region of integration, and the integral sign means we’re adding things up. In this case, we’re finding area, so $dA = 1dA$ is exactly what we’re adding up. In other situations, we’ll be adding up more complicated quantities.

This is all well and good, but how do we actually *find* the area? To do this, we’ll need to convert $\int_{\mathcal{R}} dA$ into a more traditional-looking integral—one that we know how to evaluate.

Let’s write down our sum more carefully. We need to sum over all tiny rectangles that fit inside \mathcal{R} . To do so, we can take a systematic approach: sum all the rectangles in a column first and then sum up all the columns. The lower left corner of all rectangles in a single column share a common x -coordinate. Consider the column with lower left corner at (x_0, y_0) . Counting, we see there are approximately $(1 - x_0^2)/\Delta y$ rectangles in this column. Further, there are approximately $2/\Delta x$ columns. Therefore,

$$\text{area of } \mathcal{R} \approx \sum_{i=1}^{2/\Delta x} \sum_{j=1}^{(1-(i\Delta x)^2)/\Delta y} \Delta y \Delta x.$$

That sum is really hard to parse, so we’ll write it another way.

$$\text{area of } \mathcal{R} \approx \sum_{x_0=-1, -1+\Delta x, \dots, 1} \sum_{y_0=x_0^2, x_0^2+\Delta y, \dots, 1} \Delta y \Delta x.$$

This is still hard to read, but it’s looking more like an integral. The inner sum is adding up things from $y_0 = x_0^2$ to $y_0 = 1$ and the outer sum is adding up things from $x_0 = -1$ to $x_0 = 1$. After taking the limit as $\Delta x, \Delta y \rightarrow 0$, we may rewrite this as an integral, giving

$$\text{area of } \mathcal{R} = \int_{\mathcal{R}} dA = \int_{x=-1}^{x=1} \int_{y=x^2}^{y=1} dy dx. \quad (5.1)$$

Now, we should take a moment to make sure we understand what we’ve just written. The right side of Equation (5.1) is an *iterated integral*. That is,

$$\int_{x=-1}^{x=1} \int_{y=x^2}^{y=1} dy dx = \int_{x=-1}^{x=1} \left(\int_{y=x^2}^{y=1} dy \right) dx$$

and so the integral with respect to y must be done *before* the integral with respect to x . To be clear, dy and dx are *not* being multiplied. However, Δy and Δx *were* being multiplied in our sum expression. What happened? The answer is some slight of hand. We can write a more complete list of steps:²

$$\lim_{\Delta x, \Delta y \rightarrow 0} \sum_{x_i} \sum_{y_i} \Delta y \Delta x = \lim_{\Delta x, \Delta y \rightarrow 0} \sum_{x_i} \left(\sum_{y_i} \Delta y \right) \Delta x = \int_{x=-1}^{x=1} \left(\int_{y=x^2}^{y=1} dy \right) dx.$$

Now we can evaluate this iterated integral to conclude

$$\text{area of } \mathcal{R} = \int_{\mathcal{R}} dA = \int_{x=-1}^{x=1} \left(\int_{y=x^2}^{y=1} dy \right) dx = \frac{4}{3}.$$

But, there was another way we could have divided up our original sum. We could have summed all tiny rectangles in a single row first and then added up all the row sums. Using this approach, we see

$$\lim_{\Delta x, \Delta y \rightarrow 0} \sum_{y_i} \sum_{x_i} \Delta x \Delta y = \lim_{\Delta x, \Delta y \rightarrow 0} \sum_{y_i} \left(\sum_{x_i} \Delta x \right) \Delta y = \int_{y=0}^{y=1} \left(\int_{x=-\sqrt{y}}^{x=\sqrt{y}} dx \right) dy.$$

Computing this integral, we again get $4/3$, as expected.

Note that when we swapped the order of the sums (and hence the order of the integrals) the bounds changed. This is worth considering carefully.

When we integrate with respect to y and then x , it means that we first imagine x is fixed and we let y vary between bounds (which may depend on the fixed x). When we integrate with respect to x first, the situation is reversed. Geometrically, it's easy to see why our bounds change.

XXX Figure

However, there's also an algebraic argument for why our bounds change. Suppose $(x, y) \in \mathcal{R}$. Then $-1 \leq x \leq 1$ and $x^2 \leq y \leq 1$. These were the bounds we used adding up along columns first. However, we also know that $0 \leq y \leq 1$ and $-\sqrt{x} \leq y \leq \sqrt{x}$, which gives the bounds when adding up along rows first. These two system of inequalities don't contradict each other; they contain the exact same information. The difference is in which coordinates you can verify without the others.


Given $-1 \leq x \leq 1$ and $x^2 \leq y \leq 1$, you can check the x -coordinate of the point (x_0, y_0) without needing the y -coordinate. To check the y -coordinate, you need to compute x_0^2 and so you need the x -coordinate. On the other hand, if you were given $0 \leq y \leq 1$ and $-\sqrt{x} \leq y \leq \sqrt{x}$, you could check the y -coordinate without needing to know what the x -coordinate was, but you would need the y -coordinate to check the x coordinate.

When computing an iterated integral, we expect to end up with a number. Thus, the bounds of the outer-most integral cannot depend on any other variables. The bounds of the second outer-most integral can only depend on variables "further out." Keeping this in mind will give us a way to quickly judge if we've written down correct bounds for iterated integrals.

Now we turn to integrating other functions over regions.

² Even still, we skipped steps. In particular, we split up $\lim_{\Delta x, \Delta y \rightarrow 0}$ into two limits $\lim_{\Delta x \rightarrow 0} \lim_{\Delta y \rightarrow 0}$. Then, we exchanged a limit and a sum. These steps require justification, but we won't trouble ourselves with that here.

■ **Example 5.1** Let $f(x, y) = x^2 + y^2$, and let \mathcal{D} be the region between the parabola $x = y^2$ on the left and the line $x = y + 2$ on the right. These curves intersect when $y^2 = y + 2$. In other words, the curves intersect when $y = 2$ or $y = -1$. Thus, \mathcal{D} also lies between $y = -1$ below and $y = 2$ above.




Thus

$$\begin{aligned}
 \int_{\mathcal{D}} f \, dA &= \int_{y=-1}^{y=2} \left(\int_{x=y^2}^{x=y+2} x^2 + y^2 \, dx \right) dy \\
 &= \int_{-1}^2 \left(\frac{x^3}{3} + xy^2 \Big|_{x=y^2}^{x=y+2} \right) dy \\
 &= \int_{-1}^2 \left(\frac{(y+2)^3}{3} + (y+2)y^2 - \frac{y^6}{3} - y^4 \right) dy \\
 &= \left(\frac{(y+2)^4}{12} + \frac{y^4}{4} + \frac{2y^3}{3} - \frac{y^7}{21} - \frac{y^5}{5} \right) \Big|_{-1}^2 \\
 &= \frac{256}{12} + \frac{16}{4} + \frac{16}{3} - \frac{128}{21} - \frac{32}{5} - \frac{1}{12} - \frac{1}{4} + \frac{2}{3} - \frac{1}{21} - \frac{1}{5} \\
 &= \frac{639}{35} \approx 18.26.
 \end{aligned}$$

Note that the region is also bounded vertically by graphs, so in principle the integral could be evaluated in the other order. However, there is a serious hindrance to trying this. The top graph is that of $y = h_{\text{top}}(x) = \sqrt{x}$, but the bottom graph is described by two different formulas depending on what x is. It is a parabola to the left of the point $(1, -1)$ and a line to the right of that point, so

$$h_{\text{bot}}(x) = \begin{cases} -\sqrt{x} & \text{if } 0 \leq x \leq 1 \\ x - 2 & \text{if } 1 < x \leq 4 \end{cases}.$$



(The x -values at the relevant points are determined from the corresponding y -values which were calculated above.) That means to actually compute the integral we must decompose the region \mathcal{D} into two subregions meeting along the line $x = 1$ and treat each one separately. Doing so,

$$\int_{\mathcal{D}} f \, dA = \int_{x=0}^{x=1} \left(\int_{y=-\sqrt{x}}^{y=\sqrt{x}} x^2 + y^2 \, dy \right) dx + \int_{x=1}^{x=4} \left(\int_{y=x-2}^{y=\sqrt{x}} x^2 + y^2 \, dy \right) dx.$$

You should work out the two iterated integrals on the right just to check that their sum gives the same answer as above. ■

Integral Notation

In single-variable calculus, you used the notation $\frac{d}{dx}f(x) = \frac{df}{dx}(x)$ to represent the derivative of the function f and $\int_a^b f(x)dx$ to represent the integral from a to b of the function f . This notation was invented by Gottfried Leibniz in the 1600s and was inspired by geometric thinking. Though the infinitely large and infinitely small need to be handled with mathematical care, the idea that “ dx ” represents an infinitesimally small change in the variable x gives rise to fruitful intuitions.

We will extend on this idea. Consider $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and a region $\mathcal{R} \subseteq \mathbb{R}^n$. We will write

$$\int_{\mathcal{R}} f \, dV$$

to notate *the integral of f over the region \mathcal{R} with respect to volume*. If $f : \mathbb{R}^2 \rightarrow \mathbb{R}$, we may write dA in place of dV , since a two-dimensional volume is traditionally called an area. If the context is clear, we might even omit the dV entirely, opting to write

$$\int_{\mathcal{R}} f.$$

This notation can replace single-variable calculus notation. For example

$$\int_a^b f(x) \, dx = \int_{[a,b]} f.$$

As always, the purpose of notation is to facilitate our thinking, not to be a formal grammar³, and so you should use whatever notation most clearly conveys your thoughts. For us, $\int_{\mathcal{R}} f$ will often suffice.

When we write iterated integrals, we will omit parenthesis. For example, we write

$$\int \left(\int f(x, y) dx \right) dy \quad \text{as} \quad \int \int f(x, y) dx dy,$$

and to keep ourselves from getting the bounds of integration mixed up, we will label the bounds of integration for iterated integrals with the variable they bound. For example, we write

$$\int_{y=a}^{y=b} \int_{x=c}^{x=d} f(x, y) dx dy \quad \text{instead of} \quad \int_a^b \int_c^d f(x, y) dx dy.$$

This will prevent us from becoming confused when we swap the order of integration.

Finally, it is worth being aware of other notation that you might encounter. For example, some textbooks write

$$\iint f \, dA \quad \text{and} \quad \iiint f \, dV$$

for integrals with respect to area and volume. The motivation being that the number of integral signs should match the number of dimensions you're integrating over.

Iterated Integrals


We've already seen how an integral over a region can be turned into an iterated integral in multiple ways. We will turn this into a technique for evaluating some difficult iterated integrals. Namely, if we are given an iterated integral, we can try to convert it into an integral over a region, and then convert it back to a different iterated integral.

■ **Example 5.2** Consider the iterated integral

$$\int_{x=0}^{x=1} \int_{y=x}^{y=1} \frac{\sin y}{y} dy dx.$$

This is the iterated integral obtained from the integral of the function $f(x, y) = \sin y / y$ over the triangular region \mathcal{D} contained between the vertical lines $x = 0, x = 1$, the line $y = x$ below, and the line $y = 1$ above.


³ Sometimes notation's dual purpose is to be a formal grammar. For instance, most programming languages could be thought of as very stringent systems of notation.



The indefinite integral (anti-derivative)

$$\int \frac{\sin y}{y} dy$$

cannot be expressed in terms of known elementary functions. (Try to integrate it or look in an integral table if you don't believe that.) Hence, the iterated integral cannot be evaluated by anti-derivatives. However, the triangular region may be described just as well by bounding it horizontally by graphs: it lies between $y = 0$ and $y = 1$ and for each y between $x = 0$ and $x = y$.



Thus, the double integral can be evaluated from the iterated integral

$$\begin{aligned} \int_{y=0}^{y=1} \int_{x=0}^{x=y} \frac{\sin y}{y} dx dy &= \int_0^1 \frac{\sin y}{y} x \Big|_{x=0}^{x=y} dy \\ &= \int_0^1 \frac{\sin y}{y} y dy = \int_0^1 \sin y dy \\ &= -\cos y \Big|_0^1 = 1 - \cos 1 \approx 0.46. \end{aligned}$$

Note that in order to set up the iterated integral in the other order, we had to draw a diagram and work directly from that. There are no algebraic rules which will allow you to switch orders without using a diagram. ■

Swapping the order of integration is a neat trick, and if an iterated integral comes from the integral of some function over a region, then we can always swap the order of integration⁴. However, if we're given an iterated integral, a priori, we don't know that we can swap the order of integration.

■ **Example 5.3** Let $f(x, y) = \frac{x^2 - y^2}{(x^2 + y^2)^2} = -\frac{\partial}{\partial x} \left(\frac{\partial}{\partial y} \arctan(y/x) \right)$. Now,

$$\int_{x=0}^{x=1} \int_{y=1}^{y=1} f(x, y) dy dx = \frac{\pi}{4}$$

and

$$\int_{y=0}^{y=1} \int_{x=1}^{x=1} f(x, y) dx dy = -\frac{\pi}{4}.$$

The order of integration matters for f and so the iterated integrals $\iint f dx dy$ and $\iint f dy dx$ cannot have come from the integration of f over a region—an integral of f over a region cannot be reasonably defined. ■

⁴ Depending on how we define things, this may be a circular argument. In our definition of an integral over a region, we had to sum up a bunch of tiny areas in an unspecified order. We might *define* this to make sense only if every possible order we could sum results in the same number.

For us, it will be uncommon to come across iterated integrals where the order of integration cannot be changed. However, the following theorem from analysis gives us explicit conditions for when it is okay.

Theorem 5.1.1 — Fubini's Theorem. Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ and suppose that

$$\int_{x=a}^{x=b} \int_{y=c}^{y=d} |f(x, y)| \, dy \, dx < \infty.$$

Then,

$$\int_{x=a}^{x=b} \int_{y=c}^{y=d} f(x, y) \, dy \, dx = \int_{y=c}^{y=d} \int_{x=b}^{x=a} f(x, y) \, dx \, dy.$$

Exercises for 5.1

5.2 The Volume Form

Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ and $\mathcal{D} \subseteq \mathbb{R}^2$ be the unit disk centered at the origin. If we wanted to compute $\int_{\mathcal{D}} f \, dA$, we could split it up into an iterated integral like so:

$$\int_{\mathcal{D}} f \, dA = \int_{x=-1}^{x=1} \int_{y=-\sqrt{1-x^2}}^{y=\sqrt{1-x^2}} f(x, y) \, dy \, dx.$$

However, the thought of integrating square roots makes me cringe. Especially, when \mathcal{D} has a simple description in polar coordinates. If we could use polar coordinates, our iterated integral might look more like $\int_{\theta=0}^{\theta=2\pi} \int_{r=0}^{r=1} f$. Much less intimidating!

Let's see if we can figure out how to integrate in polar coordinates. The principle should be the same. If we divide \mathcal{D} into tiny *sectors*, then

$$\int_{\mathcal{D}} f \, dA = \lim_{\text{sector size} \rightarrow 0} \sum_{\text{sectors} \in \mathcal{D}} (\text{sector area}) f(\text{sector location}).$$

When we previously used iterated integrals to compute integrals over regions, we chopped our region into little rectangles of width Δx and height Δy . This meant that the area of each sector was $\Delta x \Delta y$, and so we had the formula

$$\int_{\mathcal{D}} f \, dA = \lim_{\text{sector size} \rightarrow 0} \sum_{\text{sectors} \in \mathcal{D}} (\Delta x \Delta y) f(x, y).$$

For polar coordinates, the natural way to chop up \mathcal{D} is not into rectangles, but into *annular sectors*.

XXX Figure comparing rectangular chopping and annular sector chopping.

When we compute the exact area of the annular sector between the curves $\theta = \theta_0$, $\theta = \theta_0 + \Delta\theta$, $r = r_0$, and $r = r_0 + \Delta r$, we get

$$\text{annular sector area} = r_0 \Delta r \Delta\theta + \frac{1}{2} \Delta\theta (\Delta r)^2.$$

The first thing to note is that the quantity r_0 shows up in the area formula. This makes sense, because if you fix two θ values, the further “out” (i.e., the larger r_0 is) you are, the larger the area. This didn't happen in rectangular coordinates—all the rectangular sectors had exactly the same area.

Now, just as before, we can write a sum expression for our integral.

$$\int_{\mathcal{D}} f \, dA = \lim_{\Delta r, \Delta\theta \rightarrow 0} \sum_{r=0, \Delta r, 2\Delta r, \dots, 1} \sum_{\theta=0, \Delta\theta, 2\Delta\theta, \dots, 2\pi} \left(r \Delta r \Delta\theta + \frac{1}{2} \Delta\theta (\Delta r)^2 \right) f(r, \theta) \quad (5.2)$$

However, it isn't clear how to rewrite Equation (5.2) as an iterated integral for easy evaluation.

Order Analysis

We can rewrite Equation (5.2) into one that looks more like a Riemann sum by doing some *order analysis*. Order analysis allows us to determine with precision which parts of an expression become negligible and which parts are non-negligible⁵. Often times order analysis will show that we can ignore or erase large parts of a formula and still get the correct answer in the end.

Consider the right hand side of Equation (5.2).

$$\lim_{\Delta r, \Delta \theta \rightarrow 0} \sum_{r=0, \Delta r, 2\Delta r, \dots, 1} \sum_{\theta=0, \Delta \theta, 2\Delta \theta, \dots, 2\pi} \left(r \Delta r \Delta \theta + \frac{1}{2} \Delta \theta (\Delta r)^2 \right) f(r, \theta)$$

In this expression, both Δr and $\Delta \theta$ are tending towards 0 and so

$$r \Delta r \Delta \theta + \frac{1}{2} \Delta \theta (\Delta r)^2 \rightarrow 0.$$

But, we don't expect the entire limit to go to zero because as $\Delta \theta$ and Δr get smaller, each of the two sums get more and more terms. The outer sum will have roughly $1/\Delta r$ terms and the inner sum will have roughly $1/\Delta \theta$ terms⁶. Now, $\sum_{i=1}^n K = nK$ if K is a constant. The inside of our sums are not constant, but if they were bounded, we could replace them with constants and apply squeeze-theorem arguments to them. For order analysis, since we are only interested in whether a term *must* go to zero, we bypass the detailed squeeze theorem arguments by just imagining all bounded terms inside our sum are constant and seeing what happens. To proceed, we first split up our sum to figure out which parts are negligible.

$$\begin{aligned} & \sum_{r=0, \Delta r, 2\Delta r, \dots, 1} \sum_{\theta=0, \Delta \theta, 2\Delta \theta, \dots, 2\pi} \left(r \Delta r \Delta \theta + \frac{1}{2} \Delta \theta (\Delta r)^2 \right) f(r, \theta) \\ &= \sum_{\theta} \sum_r r \Delta r \Delta \theta f(r, \theta) + \sum_{\theta} \sum_r \frac{1}{2} (\Delta r)^2 \Delta \theta f(r, \theta) \end{aligned}$$

Doing order analysis on the first part of the sum, we see

$$\begin{aligned} \lim_{\Delta r, \Delta \theta \rightarrow 0} \sum_{\theta} \sum_r r \Delta r \Delta \theta f(r, \theta) &\sim \lim_{\Delta r, \Delta \theta \rightarrow 0} \frac{1}{\Delta r} \left(\frac{1}{\Delta \theta} \left(r \Delta r \Delta \theta f(r, \theta) \right) \right) \\ &= r f(r, \theta), \end{aligned}$$

where \sim means *order equivalent*⁷. Since this term is not order-equivalent to zero, it is non-negligible, and so we must keep it.

Analyzing the other component of the sum tells a different story:

$$\begin{aligned} \lim_{\Delta r, \Delta \theta \rightarrow 0} \sum_{\theta} \sum_r \frac{1}{2} (\Delta r)^2 \Delta \theta f(r, \theta) &\sim \lim_{\Delta r, \Delta \theta \rightarrow 0} \frac{1}{\Delta r} \left(\frac{1}{\Delta \theta} \left(\frac{1}{2} (\Delta r)^2 \Delta \theta f(r, \theta) \right) \right) \\ &= \lim_{\Delta r, \Delta \theta \rightarrow 0} \frac{1}{2} \Delta r f(r, \theta) = 0, \end{aligned}$$

⁵ *Negligible* is a fancy way of saying “this can be safely ignored.”

⁶ We say roughly because the number of terms involved is always an integer and the quantities $1/\Delta r$ and $1/\Delta \theta$ may not be integers.

⁷ There is a formal theory of order equivalence and a slew of specialized notations. We don't need those formalities right now.

and so this part of the sum is negligible. We may therefore conclude,

$$\begin{aligned} \lim_{\Delta r, \Delta \theta \rightarrow 0} \sum_{r=0, \Delta r, 2\Delta r, \dots, 1} \sum_{\theta=0, \Delta \theta, 2\Delta \theta, \dots, 2\pi} \left(r \Delta r \Delta \theta + \frac{1}{2} \Delta \theta (\Delta r)^2 \right) f(r, \theta) \\ = \lim_{\Delta r, \Delta \theta \rightarrow 0} \sum_{r=0, \Delta r, 2\Delta r, \dots, 1} \sum_{\theta=0, \Delta \theta, 2\Delta \theta, \dots, 2\pi} r \Delta r \Delta \theta f(r, \theta), \end{aligned}$$

which is looking a lot more like a Riemann sum that can be converted into an iterated integral. In fact, if f is a reasonable function (for example, one where Fubini's theorem would hold),

$$\lim_{\Delta r, \Delta \theta \rightarrow 0} \sum_{r=0, \Delta r, 2\Delta r, \dots, 1} \sum_{\theta=0, \Delta \theta, 2\Delta \theta, \dots, 2\pi} r \Delta r \Delta \theta f(r, \theta) = \int_{\theta=0}^{\theta=2\pi} \int_{r=0}^{r=1} f(r, \theta) r dr d\theta.$$

Order analysis can be very powerful, but we should take note about the assumptions we made. First off, for the order analysis that we did, we assume f was *bounded* on the region \mathcal{D} . We could still do order analysis on an unbounded function, but it would take more care⁸. Further, it'd be nice if there were a quicker way to come up with an iterated integral in various coordinate systems, rather than lugging around Riemann sums all the time—and there is.

The Volume Form

The *volume form* gives a formal way to rewrite integrals as iterated integrals with respect to different coordinate systems. What it really does is give a way to relate areas or volumes described in a particular coordinate system to “true” areas and volumes in Euclidean space—and if your mind is tingling with thoughts of parameterizations, you're on the right track.

Let's think back to line integrals. Suppose $\vec{p} : [a, b] \rightarrow \mathbb{R}^n$ is an arc-length parameterization of the curve S . Then, if $f : \mathbb{R}^n \rightarrow \mathbb{R}$,

$$\int_S f = \int_a^b f \circ \vec{p}(t) dt.$$

However, if $\vec{q} : [c, d] \rightarrow \mathbb{R}^n$ is a non-arc-length parameterization of S , we have to use a different formula. Namely,

$$\int_S f = \int_c^d f \circ \vec{q}(t) \|\vec{q}'(t)\| dt.$$

In this expression, $\|\vec{q}'(t)\|$ compensates for the fact that as the parameter t marches on, the arc-length traced out by \vec{q} is non-uniform. The volume form does the same thing but for parameterizations of multi-dimensional regions.

Definition 5.2.1 — Volume Form. Suppose \mathcal{F} is a coordinate system for \mathbb{R}^2 with a relationship to rectangular coordinates given by $(x, y) = \vec{f}(a, b)$. The *pre-volume form* associated with \mathcal{F} at the point (a, b) is written $\Delta V(a, b)$ and is defined to be the area of $\vec{f}([a, a + \Delta a] \times [b, b + \Delta b])$.

The *volume form* associated with \mathcal{F} at the point (a, b) is the infinitesimal $dV(a, b) =$

⁸ Order analysis on $\sum_{i=1}^n i$ shows that it is order n^2 , not order n . This is because the part inside the sum, i , is not bounded.

$V(a, b)da db$ where $V : \mathbb{R}^2 \rightarrow \mathbb{R}$ is the unique function satisfying

$$\lim_{\Delta a, \Delta b \rightarrow 0} \frac{\Delta V(a, b)}{V(a, b)\Delta a \Delta b} = 1.$$

We only wrote down the definition of the volume form for two-dimensional coordinate systems, but the definition for three and higher dimensions is similar. To define the volume form in an n -dimensional coordinate system, you take the image of n -dimensional rectangles instead of two-dimensional ones and you replace $\Delta a \Delta b$ with $\Delta x_1 \Delta x_2 \Delta x_3 \cdots \Delta x_n$.

There are a few peculiarities in the definition of the volume form. First, the word “infinitesimal” shows up in the definition. As of yet, this has not been a formal term for us. Further, the volume form has “ $da db$ ” built into it. These features hark back to the origins of the volume form. The term *volume form* comes from the field of differential geometry, and in differential geometry, da has a unique meaning all by itself and is called a *differential form*. Differential forms can be multiplied with each other and added together and the result still has meaning. For us, da by itself has no precise meaning and only appears as a notational book-keeping device in integrals and derivatives. Of course, the intuitive meaning of da is that it is an infinitesimal quantity. At this point, we won’t dive into the details of how and why differential geometry and differential forms work, but we will use the language of differential geometry.

■ **Example 5.4** Compute the volume form for polar coordinates.

We know that $p : [0, \infty) \times [0, 2\pi) \rightarrow \mathbb{R}^2$ given by $\vec{p}(r, \theta) = (r \cos \theta, r \sin \theta)$ converts from polar to rectangular coordinates. Further, we already computed the pre-volume form for polar coordinates. Namely, $\Delta V(r, \theta) = r \Delta r \Delta \theta + \frac{1}{2}(\Delta r)^2 \Delta \theta$.

XXX Figure

To compute the volume form for polar coordinates, we need to find a function V so that

$$\begin{aligned} \lim_{\Delta r, \Delta \theta \rightarrow 0} \frac{r \Delta r \Delta \theta + \frac{1}{2}(\Delta r)^2 \Delta \theta}{V(r, \theta) \Delta r \Delta \theta} &= \lim_{\Delta r, \Delta \theta \rightarrow 0} \frac{r}{V(r, \theta)} + \frac{\Delta r}{2V(r, \theta)} \\ &= \frac{r}{V(r, \theta)} = 1. \end{aligned}$$

There’s a clear choice. Namely, $V(r, \theta) = r$. Therefore, the volume form for polar coordinates is $r dr d\theta$. ■

Using the Volume Form

Once we have gone through the work of finding the volume form for a particular coordinate system, writing down iterated integrals in that coordinate system becomes easy—you can just replace “ dV ” with the volume form and integrate away.

Intuitively the volume form plays the same role that $\|\vec{q}'(t)\|dt$ did in line integrals. It *compensates* for any stretching or warping that an alternate coordinate system may do to area.

■ **Example 5.5** Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be defined by $f(x, y) = x^2 + y^2$ and let \mathcal{D} be the unit disk centered at the origin. Compute $\int_{\mathcal{D}} f dV$.

We’ve already computed the volume form in polar coordinates. It is $r dr d\theta$. Further, f is easy to describe in polar coordinates, since $x^2 + y^2 = r^2$. Lastly, in polar coordinates, \mathcal{D} is the

region where $r \in [0, 1]$ and $\theta \in [0, 2\pi)$. Therefore,


$$\int_{\mathcal{D}} f \, dV = \int_{\theta=0}^{\theta=2\pi} \int_{r=0}^{r=1} (r^2)r \, dr \, d\theta = \int_{\theta=0}^{\theta=2\pi} \int_{r=0}^{r=1} r^3 \, dr \, d\theta.$$

■

■ **Example 5.6** Consider the two-dimensional skewed coordinate system \mathcal{S} . The \mathcal{S} coordinate system uses the variables a and b and relates to rectangular coordinates by the equations $x = a + b$ and $y = a - b$. Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be given by the formula $f(a, b) = a + 6b$ and let $\mathcal{R} \subseteq \mathbb{R}^2$ be the region defined in \mathcal{S} coordinates by the inequalities $1 \leq a \leq 2$ and $4 \leq b \leq 6$. Compute $\int_{\mathcal{R}} f \, dV$.

At present, we have two options. We could rewrite everything in rectangular coordinates or we could figure out the volume form for \mathcal{S} coordinates and set our iterated integral up with respect to \mathcal{S} coordinates. Let's do the second way—it sounds more fun.

Finding the pre-volume form is often the hardest part. It helps to write down change-of-coordinate functions and draw a picture. First off, let $\vec{s} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ be the function that takes points written in \mathcal{S} coordinates and rewrites them in rectangular coordinates. From the relationships given above, we see $\vec{s}(a, b) = (a + b, a - b)$. Now we can draw $[a_0, a_0 + \Delta a] \times [b_0, b_0 + \Delta b]$ in the ab -coordinate plane and draw $\vec{s}([a_0, a_0 + \Delta a] \times [b_0, b_0 + \Delta b])$ in the xy -coordinate plane.



Conveniently enough, the result is still a rectangle in the xy -coordinate plane, so its area is easy to compute (after finding the lengths of each side of the rectangle): $2\Delta a\Delta b$.

Now, we're looking for a function $V : \mathbb{R}^2 \rightarrow \mathbb{R}$ so that

$$\lim_{\Delta a, \Delta b \rightarrow 0} \frac{2\Delta a\Delta b}{V(a, b)\Delta a\Delta b} = 1,$$

and such a function can be directly guessed. It's given by $V(a, b) = 2$. Thus, the volume form for the \mathcal{S} coordinate system is $2da\,db$.

Since everything else is already written in terms of the \mathcal{S} coordinate system, we can just write down the iterated integral.

$$\begin{aligned}\int_{\mathcal{R}} f \, dV &= \int_{b=4}^{b=6} \int_{a=1}^{a=2} f(a, b) 2 \, da \, db \\ &= \int_{b=4}^{b=6} \int_{a=1}^{a=2} (2a + 12b) \, da \, db = 126.\end{aligned}$$

■

Volume forms are really convenient, but you may be asking: didn't you just stress that in an expression like $\iint da \, db$, the symbols da and db are not being multiplied and are just notation? When we use volume forms aren't we treating $da \, db$ like it has meaning and then working backwards to write down an integral? This is an astute observation. Yes, when we use volume forms, we're assuming that everything will work out as purported. We won't prove the details (if you're interested, look in a book on differential geometry), but it is worth noting that volume forms only exist for *smooth* parameterizations. That means, every component function of the parameterization is differentiable in every coordinate. Most coordinate systems that we will come across have this property.

When trying to write down an iterated integral in non-rectangular coordinates, computing the volume form can be the hardest part. We've been computing volume forms directly from the definition, but there are several ways that avoid the trouble of finding the pre-volume form. They include the algebra of differential forms and the Jacobian. We won't discuss the details of these computational tools here, but it is worth listing the volume forms for some common coordinate systems.

Coordinate System	Volume Form
Rectangular Coordinates in \mathbb{R}^2	$dx \, dy$
Polar Coordinates	$r \, dr \, d\theta$
Rectangular Coordinates in \mathbb{R}^3	$dx \, dy \, dz$
Cylindrical Coordinates	$r \, dr \, d\theta \, dz$
Spherical Coordinates	$r^2 \sin \phi \, dr \, d\phi \, d\theta$

The Chain Rule for Volume Forms

Suppose $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ is a parameterization. We can think of \vec{p} as a function that takes points in one coordinate system and outputs points in rectangular coordinates. Call the coordinate system that \vec{p} inputs the \mathcal{P} coordinate system with variables α and β . Now, there is some volume form associated with the \mathcal{P} coordinate system and it takes the form $V_{\vec{p}}(\alpha, \beta) \, d\alpha \, d\beta$. We interpret $V_{\vec{p}}(\alpha, \beta)$ as the amount that \vec{p} stretches or shrinks volume by at the point (α, β) (in \mathcal{P} coordinates).

Now, suppose $\vec{q} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ is another parameterization. Similarly to \vec{p} , associate \vec{q} with the \mathcal{Q} coordinate system with variables γ and δ . Let $V_{\vec{q}}(\gamma, \delta) \, d\gamma \, d\delta$ be the volume form for \mathcal{Q} coordinates.

Now, \vec{p} and \vec{q} relate to coordinate systems, but they are really just parameterizations with domain \mathbb{R}^2 and range \mathbb{R}^2 . So, $\vec{s} = \vec{q} \circ \vec{p}$ is another parameterization. Associate it with the

\mathcal{S} coordinate system with the variables a and b . Can we write the volume form for the \mathcal{S} coordinate system in terms of the volume form for \mathcal{P} and \mathcal{Q} coordinates?

Yes, we can! The process is fairly straightforward. We know $V_{\vec{p}}(\alpha, \beta)$ is how much \vec{p} changes area by at the point (α, β) and we know $V_{\vec{q}}(\gamma, \delta)$ is how \vec{q} changes area at the point (γ, δ) . Further, considering the composition $\vec{s} = \vec{q} \circ \vec{p}$, we could interpret \vec{p} as inputting \vec{P} coordinates and then outputting \vec{Q} coordinates, which then get taken by \vec{q} and finally converted to rectangular coordinates. Thus, at the point (α, β) , area first gets changed by $V_{\vec{p}}(\alpha, \beta)$ and then at the point $\vec{p}(\alpha, \beta)$, area gets changed by $V_{\vec{q}} \circ \vec{p}(\alpha, \beta)$. Change in area is multiplicative, so

$$V_{\vec{s}}(\alpha, \beta) = V_{\vec{p}}(\alpha, \beta) \left(V_{\vec{q}} \circ \vec{p}(\alpha, \beta) \right). \quad (5.3)$$

Does that look a little like the chain rule?

Let's try this again in one dimension. Let $f : \mathbb{R} \rightarrow \mathbb{R}$ and $g : \mathbb{R} \rightarrow \mathbb{R}$ both be parameterizations. In one dimension the limit involved in finding a volume form turns out to be equivalent to taking a derivative (go ahead, compute the one-dimension volume form of a few of your favorite functions!). Therefore,

$$V_f(a) = f'(a) \quad \text{and} \quad V_g(b) = g'(b).$$

Let $h = g \circ f$. Expanding with Equation (5.3), we get

$$h'(a) = V_h(a) = V_f(a) \left(V_g \circ f(a) \right) = f'(a) \left(g' \circ f(a) \right),$$

⁹ Hold the phone! There's something strange going on. The derivative of a function can take negative values but by definition a volume form must take only positive values. So, on the surface $w' = V_w$ can only hold for special parameterizations $w : \mathbb{R} \rightarrow \mathbb{R}$. This is true, and this is all fixed in the land of differential forms where volumes and areas are allowed to be *signed*.

which is exactly what the chain rule says! This is no coincidence. Equation (5.3) really is the chain rule in disguise⁹.

Let's use the chain rule for volume forms to solve a problem.

■ **Example 5.7** XXX Finish. Maybe an isometric rotation and a stretching of some sort? ■

Exercises for 5.2

5.3 Surface Integrals

We will now turn our attention to integrals over non-flat regions. Suppose we wanted to find the average temperature of the surface of a perfectly spherical earth. Conceptually, we can compute this with a Riemann sum just like before. We chop the surface of the earth up into little regions and find the temperature in each region. We then compute

$$\text{average temp} \approx \frac{1}{\text{total surface area}} \sum_{\text{sectors}} (\text{temp at sector})(\text{area of sector}),$$

with equality after we take a limit as the sector size tends towards zero.

■ **Example 5.8** Let $S \subseteq \mathbb{R}^3$ be the unit sphere centered at the origin, and consider the parameterization $\vec{s}: [0, 2\pi] \times [0, \pi] \rightarrow S$ of S given by

$$\vec{s}(\theta, \phi) = \begin{bmatrix} \sin \phi \cos \theta \\ \sin \phi \sin \theta \\ \cos \phi \end{bmatrix}.$$

We wish to compute the average value of $T: S \rightarrow \mathbb{R}$ over the surface S .

The parameterization \vec{s} gives us a natural way to chop S up into tiny pieces. In particular, if $\Delta\theta$ and $\Delta\phi$ are tiny then $\vec{s}([\theta_0, \theta_0 + \Delta\theta] \times [\phi_0, \phi_0 + \Delta\phi]) \subseteq S$ will be a tiny sector of S .

XXX Figure

The exact area of this sector turns out to be $(\cos \phi_0 - \cos(\phi_0 + \Delta\phi))\Delta\theta$. The total surface area of S is 4π and so we have

$$\text{average of } T = \frac{1}{4\pi} \lim_{\Delta\theta, \Delta\phi \rightarrow 0} \sum_{\theta_i} \sum_{\phi_i} T \circ \vec{s}(\theta_i, \phi_i) (\cos \phi_i - \cos(\phi_i + \Delta\phi)) \Delta\theta,$$

where the sums are taken over $\theta_i = 0, \Delta\theta, 2\Delta\theta, \dots, 2\pi$ and $\phi_i = 0, \Delta\phi, 2\Delta\phi, \dots, \pi$.

Evaluating that Riemann sum looks challenging, but we can use the idea of volume forms to rewrite it as an iterated integral. Notice

$$\lim_{\Delta\theta, \Delta\phi \rightarrow 0} \frac{(\cos \phi - \cos(\phi + \Delta\phi))\Delta\theta}{\Delta\phi \Delta\theta} = - \lim_{\Delta\theta, \Delta\phi \rightarrow 0} \frac{\cos(\phi + \Delta\phi) - \cos \phi}{\Delta\phi} = \sin \phi,$$

since the expression in the middle is just the derivative of cosine in disguise. Now, we can view $T \circ \vec{s}$ as a function from $\mathbb{R}^2 \rightarrow \mathbb{R}$ being integrated in a different coordinate system. And, we just computed the volume form of that different coordinate system to be $\sin \phi d\phi d\theta$. Therefore,

$$\text{average of } T = \frac{1}{4\pi} \int_{\theta=0}^{\theta=2\pi} \int_{\phi=0}^{\phi=\pi} T \circ \vec{s}(\theta, \phi) \sin \phi d\phi d\theta.$$

Given a formula for T , we now have a much better chance of being able to evaluate its average!

■

The reasoning used in Example 5.8 closely follows the reasoning used for integrating functions over flat regions, and we would be justified in extending our integral notation. If $S \subseteq \mathbb{R}^3$ is a surface and $f : S \rightarrow \mathbb{R}$ is a function, then

$$\int_S f = \int_S f \, dA$$

is defined to be result of chopping up S into tiny sectors, summing up the area of each sector multiplied by the value of f on that sector, and then taking the limit as the size of each sector goes to zero. Just like before, since we don't explicitly define the way that S must be chopped up into sectors and we don't define the order in which we should sum the sectors, the expression $\int_S f$ is only defined if the way we chop and the order we sum doesn't matter. Again, if f and S are nice (for example, S has a tangent plane at every point and f satisfies the conditions of Fubini's theorem), then $\int_S f$ is defined. Most functions and surfaces we will encounter will be nice.

Another thing to notice about Example 5.8 is that doing the computation was somewhat difficult. For starters, how could we find the sector area if we weren't given an explicit formula for it? This is as difficult, if not more difficult, than finding the pre-volume form¹⁰. Like we did with volume forms, we will leverage our intuition and leave it to other mathematicians to verify the details.

Approximating Surfaces

Suppose $S \subset \mathbb{R}^3$ and that $\vec{s} \in S$. Let $\mathcal{P}_{\vec{s}}$ be the tangent plane to S at the point \vec{s} . There are many ways we could write down a formula for $\mathcal{P}_{\vec{s}}$, but if we were given a parameterization for S , there would be a particularly direct way to describe $\mathcal{P}_{\vec{s}}$ in vector form.

Definition 5.3.1 — Canonical Representation of the Tangent Plane. Given a surface $S \subseteq \mathbb{R}^3$ parameterized by $\vec{p} : \mathbb{R}^2 \rightarrow S$, the *canonical representation* of the tangent plane to S relative to the parameterization \vec{p} at the point $\vec{p}(a, b) \in S$ is given by the first-order (affine) approximation to \vec{p} at (a, b) . In vector form, this representation is

$$\vec{x} = t \frac{\partial \vec{p}}{\partial a}(a, b) + s \frac{\partial \vec{p}}{\partial b}(a, b) + \vec{p}(a, b).$$

There are two ways to unpack the canonical representation of a tangent plane. Let $S \subseteq \mathbb{R}^3$ be parameterized by $\vec{p} : \mathbb{R}^2 \rightarrow S$ as before and let $\vec{s} \in S$. Further, suppose that $\vec{p}(a, b) = \vec{s}$. Now, to write down a formula for the tangent plane $\mathcal{P}_{\vec{s}}$, all we need is to find two vectors, \vec{d}_1 and \vec{d}_2 , that are tangent to S at the point \vec{s} . Then, we can describe $\mathcal{P}_{\vec{s}}$ in vector form by

$$\vec{x} = t\vec{d}_1 + s\vec{d}_2 + \vec{s}.$$

Further, the curve traced out by $\vec{p}(t, b)$ as t varies and b is held constant is contained in S and passes through \vec{s} exactly when $t = a$. Therefore, we can find a tangent vector to S at \vec{s} by

¹⁰ At this point, we don't actually have a definition of the "surface area" of a curved surface, so how could we compute the area of a sector?

finding a velocity vector for this curve.

$$\left. \frac{d}{dt} \vec{p}(t, b) \right|_{t=a} = \frac{\partial \vec{p}}{\partial a}(a, b).$$

Similarly, we can find a second tangent vector by considering the curve traced out by $\vec{p}(a, t)$ as t varies and a is held constant. This gives us a tangent vector of

$$\left. \frac{d}{dt} \vec{p}(a, t) \right|_{t=b} = \frac{\partial \vec{p}}{\partial b}(a, b).$$

XXX Figure

For all reasonable parameterizations, $\frac{\partial \vec{p}}{\partial a}(a, b)$ and $\frac{\partial \vec{p}}{\partial b}(a, b)$ will point in different directions, and so we may represent $\mathcal{P}_{\vec{p}}$ in vector form as

$$\begin{aligned} \vec{x} &= t \frac{\partial \vec{p}}{\partial a}(a, b) + s \frac{\partial \vec{p}}{\partial b}(a, b) + \vec{p} \\ &= t \frac{\partial \vec{p}}{\partial a}(a, b) + s \frac{\partial \vec{p}}{\partial b}(a, b) + \vec{p}(a, b). \end{aligned}$$

This is the *canonical* way to write down a tangent plane if you are given a parameterization.

There's another way to view the canonical representation of the tangent plane with respect to \vec{p} , and that is as a first-order approximation. Recall that to find the first-order approximation of a function $\vec{p} : \mathbb{R}^2 \rightarrow \mathbb{R}^3$, you first find the first-order approximations of each of \vec{p} 's component functions. These component approximations become the component functions of the first order approximation to \vec{p} . When all is said and done, we arrive at the following formula for $\vec{L}_{\vec{p}}$, the first-order approximation of \vec{p} at the point (a, b) .

$$\vec{L}_{\vec{p}}(t, s) = t \frac{\partial \vec{p}}{\partial a}(a, b) + s \frac{\partial \vec{p}}{\partial b}(a, b) + \vec{p}(a, b)$$

Let $\vec{L}_{\vec{p}} : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ be the canonical representation of the tangent plane to \mathcal{S} , relative to \vec{p} , at the point $\vec{p}(a, b)$. Since $\vec{L}_{\vec{p}}$ is a first-order approximation to \vec{p} , we know that if (t, s) is close to (a, b) , then

$$\vec{L}_{\vec{p}}(t, s) \approx \vec{p}(t, s).$$

It's now not so far fetched to to believe

$$\text{area of } \vec{L}_{\vec{p}}([a, a + \Delta a] \times [b, b + \Delta b]) \approx \text{area of } \vec{p}([a, a + \Delta a] \times [b, b + \Delta b]).$$

That is, the area of the tiny sector in \mathcal{S} close to the point $\vec{p}(a, b)$ is approximately equal to the area of the corresponding tiny sector on the tangent plane to \mathcal{S} at the point $\vec{p}(a, b)$.

XXX Figure

In fact, if \vec{p} is differentiable, we have

$$\lim_{\Delta a, \Delta b \rightarrow 0} \frac{\text{area of } \vec{L}_{\vec{p}}([a, a + \Delta a] \times [b, b + \Delta b])}{\Delta a \Delta b} = \lim_{\Delta a, \Delta b \rightarrow 0} \frac{\text{area of } \vec{p}([a, a + \Delta a] \times [b, b + \Delta b])}{\Delta a \Delta b},$$

which is exactly what we need to compute a volume form!

Exercises for 5.3

Chapter 6

Vector Fields

A *scalar field* is a function that assigns each point in space a scalar value. For example, given latitude and longitude (x, y) , the function that outputs the elevation of the surface of the earth at position (x, y) is a scalar field. However, we might ask for a function that when given (x, y) returns a vector (maybe the velocity of the wind at that point on the earth). Functions that assign a vector to each point in space are called *vector fields*.

Vector fields are key in physics because they model vector quantities that are non-constant with respect to space. For example, force is a vector! The force due to of gravity around a point mass is always directed towards the origin—that means, as you move in space, the direction of the force changes. Suppose there is a point mass at the origin. We might model the force¹ due to this point mass acting on a particle at position (x, y, z) with the vector field

$$\vec{F}(x, y, z) = \frac{1}{(x^2 + y^2 + z^2)^{3/2}} \begin{bmatrix} -x \\ -y \\ -z \end{bmatrix}.$$

Indeed, the vector $\vec{F}(x, y, z)$ always points towards the origin, and $\|\vec{F}(x, y, z)\| = \frac{1}{x^2 + y^2 + z^2}$ is $1/d^2$ where d is the distance from the origin.

Notation

Notationally, $\vec{F} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ indicates that \vec{F} inputs points or vectors in \mathbb{R}^n and outputs points or vectors in \mathbb{R}^m . Oftentimes in physics $n = m$, but doesn't need to.

When defining a vector field, sometimes it's convenient to use components and sometimes it's easier to work with vectors. For example, the force field $\vec{F} : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ above was defined in terms of components as

$$\vec{F}(x, y, z) = \frac{1}{(x^2 + y^2 + z^2)^{3/2}} \begin{bmatrix} -x \\ -y \\ -z \end{bmatrix},$$

but we could have equivalently defined it using vectors by the equation

$$\vec{F}(\vec{v}) = -\frac{\vec{v}}{\|\vec{v}\|^3}.$$

¹ Despite what is peddled in Star Wars, this is a real *force field*.

6.1 Graphing Vector Fields

Traditionally, a function $f : \mathbb{R} \rightarrow \mathbb{R}$ is graphed on a coordinate plane by picking one axis to be the domain, one axis to be the range, and then putting a point at each pair $(x, f(x))$. For a function $g : \mathbb{R}^2 \rightarrow \mathbb{R}$, we do something similar, except, this time, the domain is a plane. But, no matter, we arrange a third axis orthogonal to a plane and graph points $(x, y, f(x, y))$ in this three-dimensional space.

With vector fields, we cannot graph in the same way. A vector field $\vec{F} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ inputs vectors in a plane and outputs vectors in a plane. To traditionally “graph” this function, we would need to plot points in $\mathbb{R}^2 \times \mathbb{R}^2 = \mathbb{R}^4$. Four-dimensional plots are hard to make! Fortunately, the goal of graphing is rarely to make a faithful representation of a function—it is to make a visual representation of a function to aid our thinking. With this perspective, representing four-dimensional graphs won’t be hard.

If we think of $\vec{F} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ as attaching a vector to each point in \mathbb{R}^2 , we might graph \vec{F} by drawing vectors at each point in \mathbb{R}^2 .

XXX Figure

This type of plot is called a *quiver plot*, and it will be the main way we visualize vector fields.

■ **Example 6.1** Graph the vector field $\vec{F} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ defined by $\vec{F}(x, y) = \begin{bmatrix} x \\ y \end{bmatrix}$.

XXX Finish

When making a quiver plot, there are two things to consider: (i) how densely do you draw vectors, and (ii) how do you scale the vectors? The answer to these two questions depends on the situation—remember, the goal of plotting a vector field is to visualize what is happening.

Consider the vector field $\vec{A} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ defined by $\vec{A}(x, y) = \begin{bmatrix} -y \\ x \end{bmatrix}$. If we plot vectors every 0.5 units and plot the vectors at the same scale as the domain, we get the following, hard to follow, figure.

XXX Figure

However, if we plot vectors every 0.5 units but scale the vectors so they are 1/4 the size they’d conventionally be, we get a much more understandable figure.

XXX Figure

We might even decide that the magnitude of each vector is not so important and totally ignore it in our plot.

XXX Figure with unit vectors

Exercises for 6.1

6.2 The Gradient

We already have familiarity with one class of vector fields—gradients. Given a differentiable function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, we can always consider the vector field $\nabla f : \mathbb{R}^n \rightarrow \mathbb{R}^n$. And, ∇f has some nice geometric properties. For instance, ∇f always points in the direction of greatest change for f .

XXX Figure surface with vector field at base

The connections between f and ∇f run deep. Suppose for a moment that $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ is affine. That is, $f(x, y) = \alpha x + \beta y + \gamma$. Now, $\nabla f(x, y) = \begin{bmatrix} \alpha \\ \beta \end{bmatrix}$ is constant. If $\vec{a}, \vec{b} \in \mathbb{R}^2$, we can consider the change in height of f when moving from \vec{a} to \vec{b} . This is,

$$f(\vec{b}) - f(\vec{a}) = \alpha b_x + \beta b_y - \alpha a_x - \beta a_y = \nabla f(\vec{a}) \cdot (\vec{b} - \vec{a}).$$

Of course this is no surprise—we first encountered gradients when finding affine approximations of functions. Baked into the very core of the gradient is the fact that for a differentiable function $h : \mathbb{R}^n \rightarrow \mathbb{R}$,

$$h(\vec{b}) - h(\vec{a}) = \text{change in } h \text{ from } \vec{a} \text{ to } \vec{b} \approx \nabla h(\vec{a}) \cdot (\vec{b} - \vec{a}),$$

so long as \vec{a} and \vec{b} are close together. Now, let's apply some calculus reasoning.

Let $h : \mathbb{R}^n \rightarrow \mathbb{R}$ be a differentiable function and let $\vec{a}, \vec{b} \in \mathbb{R}^n$ be points that *aren't* close to each other. Let $\vec{a} = \vec{x}_1, \vec{x}_2, \dots, \vec{x}_m = \vec{b}$ be a sequence of points so that $\|\vec{x}_{k+1} - \vec{x}_k\|$ is small. We may think of \vec{x}_i as a sequence of tiny steps used to get from \vec{a} to \vec{b} . Now, we have

$$h(\vec{b}) - h(\vec{a}) = \text{change in } h \text{ from } \vec{a} \text{ to } \vec{b} \approx \sum \nabla h(\vec{x}_k) \cdot (\vec{x}_{k+1} - \vec{x}_k).$$

If we define $\Delta \vec{x}_k = \vec{x}_{k+1} - \vec{x}_k$, we get an eerily familiar formula:

$$h(\vec{b}) - h(\vec{a}) \approx \sum \nabla h(\vec{x}_k) \cdot \Delta \vec{x}_k.$$

On the left we have the scalar quantity $h(\vec{b}) - h(\vec{a})$ and on the right we have a Riemann sum approximation of the vector line integral of ∇h along some path starting at \vec{a} and ending at \vec{b} . After taking a limit, this approximation becomes exact.

Theorem 6.2.1 Let $h : \mathbb{R}^n \rightarrow \mathbb{R}$ be a differentiable function and let $\mathcal{C} \subset \mathbb{R}^n$ be a parameterized, bounded curve which starts at $\vec{a} \in \mathbb{R}^n$ and ends at $\vec{b} \in \mathbb{R}^n$. Then,

$$h(\vec{b}) - h(\vec{a}) = \int_{\mathcal{C}} \nabla h \cdot d\vec{r},$$

where $\int_{\mathcal{C}} \nabla h \cdot d\vec{r}$ is the vector line integral of ∇h along \mathcal{C} .

Theorem 6.2.1 should remind you of the fundamental theorem of calculus, which states for a differentiable function $f : \mathbb{R} \rightarrow \mathbb{R}$,

$$f(b) - f(a) = \int_a^b f'.$$

In one dimension, there aren't many ways to get between two points. In contrast, there are infinitely many ways to get between two points in higher dimensions. Theorem 6.2.1 says that all these ways are equivalent, and that ∇h takes the place of h' .

Orientation

As we delve into the realm of vector fields, we need to take note of something new—orientation. In a scalar line integral of f along the curve \mathcal{C} , we added up the “height” of f along the curve \mathcal{C} . It didn't matter how we traverse \mathcal{C} , and so it doesn't matter how we parameterize \mathcal{C} when computing $\int_{\mathcal{C}} f$. The area of the graph of f above \mathcal{C} depends only on f and \mathcal{C} . However, things are different for vector line integrals.

If \mathcal{C} connects the points \vec{a} and \vec{b} and we wish to compute the amount of work done by the force field \vec{h} when moving from \vec{a} to \vec{b} along \mathcal{C} , not all parameterizations of \mathcal{C} will do. We must parameterize \mathcal{C} such that we start at \vec{a} and end at \vec{b} . Thus, the vector integral $\int_{\mathcal{C}} \vec{h} \cdot d\vec{r}$ depends not only on \vec{h} and \mathcal{C} , but also on the direction or *orientation* of \mathcal{C} .

For curves, there are only two orientations, forwards and backwards, and reversing the orientation of a curve in a vector line integral just multiplies the result by -1 .

■ **Example 6.2** XXX simple example with orientation flipped ■

The concept of orientation is familiar from single-variable calculus. After all,

$$\int_a^b f = - \int_b^a f.$$

From now on, if it matters, we will always specify a curve with its orientation. If the curve is closed, we will specify its orientation as *clockwise* or *counterclockwise*.

Conservative Vector Fields

Theorem 6.2.1 shows that gradients and vector line integrals interact nicely. So nicely, in fact, that vector fields coming from gradients have a special name.

Definition 6.2.1 — Conservative Vector Field & Potential Function. A vector field $\vec{h} : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is called *conservative* if there is some function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ so that $\nabla f = \vec{h}$. Such an f is called a *potential function* for \vec{h} .

We can think of conservative vector fields as those having anti-derivatives, called *potential functions*². Once we have a potential function f for a vector field \vec{h} , Theorem 6.2.1 shows that we can compute vector line integrals of \vec{h} by evaluating f on the endpoints of our path. An immediate consequence is that the integral around a closed path in a conservative vector field is zero.

■ **Example 6.3** XXX Example — work on a hill via complicated path and sub for and easy path. ■

² In math, if $\nabla f = \vec{h}$, we say f is a potential function for \vec{h} . In physics you say that $-f$ is a potential function for \vec{h} —this aligns more closely with a physicist's notion of *potential energy*.

Since conservative vector fields make vector line integrals so easy, it's useful to determine whether a given vector field is conservative. The only sure-fire way to determine a vector field is conservative is to find a potential function. However, there are several heuristics for showing that a vector field is not conservative.

Having established that conservative vector fields are useful, it would be nice to be able to decide if a given vector field is conservative. The only sure-fire way to do so is to produce a potential function. However, there methods we can use to quickly determine if a particular vector field *cannot* be conservative.

The *screening test* is one such method. Suppose that $\vec{F} = \begin{bmatrix} F_x \\ F_y \end{bmatrix}$ and that \vec{F} is conservative. That means $\vec{F} = \nabla f$ for some potential function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$, and so

$$F_x = \frac{\partial f}{\partial x} \quad \text{and} \quad F_y = \frac{\partial f}{\partial y}.$$

If \vec{F} is a *nice* vector field (that is, it comes from a sufficiently differentiable f), then

$$\frac{\partial F_x}{\partial y} = \frac{\partial^2 f}{\partial y \partial x} = \frac{\partial^2 f}{\partial x \partial y} = \frac{\partial F_y}{\partial x}.$$

This means that if $\frac{\partial F_x}{\partial y} \neq \frac{\partial F_y}{\partial x}$, we can conclude that \vec{F} is *not* conservative.

■ **Example 6.4** Is the vector field $\vec{F}(x, y) = \begin{bmatrix} x+1 \\ y+x^2 \end{bmatrix}$ conservative?

Define $F_x(x, y) = x+1$ and $F_y(x, y) = y+x^2$ so that $\vec{F} = \begin{bmatrix} F_x \\ F_y \end{bmatrix}$. Now,

$$\frac{\partial F_x}{\partial y} = 0 \neq 2x = \frac{\partial F_y}{\partial x},$$

and so \vec{F} is not conservative. ■

It should be noted that the screening test can only prove vector fields are not conservative. It cannot be used to show that they are!

■ **Example 6.5** Is the vector field $\vec{F}(x, y) = \begin{bmatrix} \frac{-y}{x^2+y^2} \\ \frac{x}{x^2+y^2} \end{bmatrix}$ conservative?

Defining $F_x(x, y) = \frac{-y}{x^2+y^2}$ and $F_y(x, y) = \frac{x}{x^2+y^2}$ so that $\vec{F} = [F_x, F_y]$, we compute

$$\frac{\partial F_x}{\partial y} = \frac{y^2 - x^2}{(x^2 + y^2)^2} = \frac{\partial F_y}{\partial x},$$

and so the screening test cannot be used to conclude anything about whether or not \vec{F} is conservative.

However, letting \mathcal{C} be the circle of radius 1 centered at the origin and oriented counter clockwise, we compute $\int_{\mathcal{C}} \vec{F} \cdot d\vec{r} = 2\pi$. If \vec{F} were a conservative vector field, the vector line integral around any closed path would be 0. Thus, we conclude that \vec{F} is not conservative. ■

The screening test works analogously for higher-dimensional vector fields. One just needs to test all combinations of mixed partials—if any combination fails, the vector field is not conservative.

Exercises for 6.2

6.3 Flux and Divergence

Suppose that the vector field $\vec{F}(x, y, z) = (4, 0, 0)$ describes the velocity of water in a river of uniform density. You cast a 1×1 square net in the river with normal vector $(1, 0, 0)$ and ask the following question: how much water per unit time is flowing through the net?

XXX Figure

In this situation, it is easy to compute. Let us assume the density of the water is 1. In this case, since the direction of flow is orthogonal to the net (because the normal vector is parallel to the direction of flow), we know

$$\text{mass/unit time} = (\text{net area})(\text{water speed}) = (1)(4) = 4.$$

If we angle the net slightly, so that it has a normal vector $(1, 1, 0)$, computing the amount of water that goes through the net in a time unit is still not so bad.

$$\text{mass/unit time} = (\text{perceived net area})(\text{water speed}).$$

Here, the *perceived net area* is the area that the net appears to the water. Since the net is angled at 45° relative to the flow of the water, the net *appears* smaller, from the perspective of the water.

XXX Figure showing “apparent” height

To the water, the net appears to have width 1 and height $\sqrt{2}/2$, giving the net a perceived area of $\sqrt{2}/2$. Now we can compute

$$\text{mass/unit time} = (\text{perceived net area})(\text{water speed}) = \frac{\sqrt{2}}{2}(4) = 2\sqrt{2}.$$

Equivalently, we could have thought about water from the perspective of the net. We know

$$\text{mass/unit time} = (\text{net area})(\text{water speed normal to net}),$$

and computing the velocity of water normal to the net can be done with a dot product. In particular,

$$\text{water speed normal to net} = \text{proj}_{\vec{n}} \vec{v} = \begin{bmatrix} 2 \\ 2 \\ 0 \end{bmatrix},$$

where $\vec{n} = (1, 1, 0)$ is the normal vector for the net and $\vec{v} = (4, 0, 0)$ is the velocity vector of the water. Now

$$\begin{aligned} \text{mass/unit time} &= (\text{net area})(\text{water speed normal to net}) \\ &= (1) \left\| \begin{bmatrix} 2 \\ 2 \\ 0 \end{bmatrix} \right\| = 2\sqrt{2}, \end{aligned}$$

which is the same quantity we computed before.

Measuring how much a vector field “flows” through a surface S is called the *flux*³ of the vector field through S , and comes up in applications ranging from fluid mechanics to electromagnetism and relativity gravity.

³ The term *flux* comes from *fluxus*, which means “flow” in Latin.

Computing Flux

Given constant vector fields and flat surfaces, flux is not difficult to compute. However, the vector fields we encounter in nature are rarely constant and the surfaces are rarely flat. But, the ideas of calculus are here to help us!

By now, we are familiar with the idea of chopping things up into little pieces, approximating each piece, adding them together again, and taking a limit. Finding the flux of a vector field through an arbitrary surface follows this pattern.

Let $\vec{F} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ be a smooth vector field and let S be a surface⁴. Let ΔA represent a tiny patch of the surface S . We can then say

$$\text{Flux}_S \vec{F} = \sum_{\Delta A} \text{Flux}_{\Delta A} \vec{F},$$

where $\text{Flux}_A \vec{B}$ means the flux of the vector field \vec{B} through the surface A . This looks a lot like a setup for a surface integral! Let's make the connection exact.

Let $\vec{p} : \mathbb{R}^2 \rightarrow S$ be a parameterization of the surface S , and let $\Delta A_{(x_0, y_0)} = \vec{p}([x_0, x_0 + \Delta x] \times [y_0, y_0 + \Delta y])$ be a tiny sector of S with “lower-left corner” at $\vec{p}(x_0, y_0)$. As long as Δx and Δy are small, $\Delta A_{(x_0, y_0)}$ will be small.

Now, let's approximate. Since $\Delta A_{(x_0, y_0)}$ is small, we could approximate \vec{F} near $\Delta A_{(x_0, y_0)}$ by the constant vector field $\vec{F}_{(x_0, y_0)}(x, y) = \vec{F} \circ \vec{p}(x_0, y_0)$. Further, if Δx and Δy are small enough, we can approximate $\Delta A_{(x_0, y_0)}$ by a tiny parallelogram coming from the canonical parameterization of the tangent plane to S at $\vec{p}(x_0, y_0)$.

Recall, given the parameterization \vec{p} , the canonical parameterization of the tangent plane to S at $\vec{p}(x_0, y_0)$ is

$$\vec{P}(t, s) = t \frac{\partial \vec{p}}{\partial x}(x_0, y_0) + s \frac{\partial \vec{p}}{\partial y}(x_0, y_0) + \vec{p}(x_0, y_0).$$

Since \vec{P} is a canonical parameterization, if Δx and Δy are small, then

$$\Delta A_{(x_0, y_0)} = \vec{p}([x_0, x_0 + \Delta x] \times [y_0, y_0 + \Delta y]) \approx \vec{P}([x_0, x_0 + \Delta x] \times [y_0, y_0 + \Delta y]).$$

Further, $\vec{P}([x_0, x_0 + \Delta x] \times [y_0, y_0 + \Delta y])$ is a parallelogram with normal vector

$$\vec{n} = \frac{\partial \vec{P}}{\partial x}(x_0, y_0) \times \frac{\partial \vec{P}}{\partial y}(x_0, y_0)$$

and area $\|\vec{n}\| \Delta x \Delta y$.

We have developed several approximations. Before we put them all together, let's give them simple names. Fix x_0 and y_0 , and let $\Delta A = \Delta A_{x_0, y_0}$; let $Q(x, y) = \vec{F} \circ \vec{p}(x_0, y_0)$ be a constant vector field approximating \vec{F} ; and, let $\Delta R = \vec{P}([x_0, x_0 + \Delta x] \times [y_0, y_0 + \Delta y])$ be a parallelogram approximating ΔA . We then have, so long as Δx and Δy are small

$$\begin{aligned} \Delta A &\approx \Delta R \\ \vec{F} &\approx \vec{Q}. \end{aligned}$$

⁴ An orientable surface, to be precise.

Putting this all together,

$$\text{Flux}_{\Delta A} \vec{F} \approx \text{Flux}_{\Delta A} \vec{Q} \approx \text{Flux}_{\Delta R} \vec{Q}$$

Now, $\frac{\vec{n}}{\|\vec{n}\|}$ is a unit normal vector to ΔR , and so

$$\text{Flux}_{\Delta R} \vec{Q} = \underbrace{\|\vec{n}\| \Delta x \Delta y}_{\text{area of } \Delta R} \frac{\vec{n}}{\|\vec{n}\|} \cdot \vec{Q} = \vec{n} \cdot \vec{Q} \Delta x \Delta y.$$

Substituting in for \vec{n} and \vec{Q} , we have

$$\text{Flux}_{\Delta A} \vec{F} \approx \left(\frac{\partial \vec{p}}{\partial x}(x_0, y_0) \times \frac{\partial \vec{p}}{\partial y}(x_0, y_0) \right) \cdot \left(\vec{F} \circ \vec{p}(x_0, y_0) \right) \Delta x \Delta y. \quad (6.1)$$

We worked pretty hard to get Equation (6.1), but let's remind ourselves that the idea was simple. We approximated a changing vector field by a constant one and a curvy shape by a flat one. The rest of the formula comes from exploiting the relationship between the cross product and the area of a parallelogram and the computation of the flux of a constant vector field through a flat surface.

Finally, since

$$\text{Flux}_S \vec{F} = \sum_{\Delta A} \text{Flux}_{\Delta A} \vec{F}$$

and we have approximated $\text{Flux}_{\Delta A} \vec{F}$, we may take a limit to obtain

$$\begin{aligned} \text{Flux}_S \vec{F} &= \sum_{\Delta A} \text{Flux}_{\Delta A} \vec{F} \\ &= \lim_{\Delta x, \Delta y \rightarrow 0} \sum_{\Delta R} \text{Flux}_{\Delta R} \vec{F} \\ &= \int_S \left(\frac{\partial \vec{p}}{\partial x}(x, y) \times \frac{\partial \vec{p}}{\partial y}(x, y) \right) \cdot \left(\vec{F} \circ \vec{p}(x, y) \right) dx dy. \end{aligned}$$

Of course, $\int_S \cdot dx dy$ is actually a double integral over the domain of the parameterization \vec{p} of S .

■ **Example 6.6** XXX flux computation example ■

Flux and Surface Orientation

When computing the flux, a surface's orientation is significant. For example, the flux of a vector field $\vec{F} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ through the xy -plane oriented upward (i.e., upward-pointing normal vector) is negative the flux of \vec{F} through the xy -plane oriented downward.

The relationship between an upward or downward orientation for the xy -plane is pretty clear, as is the relationship between a sphere oriented outward or inward, but for more complicated surfaces, orientation an orientation can be harder to specify, or might not even exist!

XXX Figure upward/downward normal for plane, inward/outward normal for sphere.

Definition 6.3.1 — Consistent Orientation. A surface S is endowed with a *consistent orientation* if each point in S is assigned a unit normal vector and if every unit normal vector can be *continuously transported* to every other unit normal vector.

The rigorous definition of *consistent orientation* requires a rigorous definition of *continuous transport*, which we won't try to pin down⁵. However, the idea is intuitive enough. Let S be a surface equipped with an orientation—a choice of a unit normal vector at each point. Let \vec{n}_1 be the unit normal vector to S at s_1 and \vec{n}_2 be the unit normal vector to S at s_2 . The vector \vec{n}_1 can be continuously transported to \vec{n}_2 if it can be “slid” from s_1 to s_2 while staying normal to the surface, and, after it has “slid” it is identical to \vec{n}_2 .

XXX Figure sliding vectors on cylinder

If a surface admits a consistent orientation, we call it *orientable*. We will exclusively deal with orientable surfaces, and, for an orientable surface, to specify an orientation, all you need to do is specify a normal vector at a single point.

■ **Example 6.7** Let \mathcal{H} be the hemisphere defined by $z = \sqrt{1 - x^2 - y^2}$ with a normal vector $\vec{n} = (0, 0, -1)$ at the point $(0, 0, 1)$. Find a unit normal vector to \mathcal{H} at $(\sqrt{2}/2, \sqrt{2}/2, 0)$ consistent with this choice.

XXX Finish ■

It is mathematically interesting to consider the fact that there exist *non-orientable* surfaces. One example is the Möbius strip. A Möbius strip can be created by gluing the ends of a strip of paper together with a half twist, and it cannot be oriented!

XXX Figure

If a surface cannot be oriented, one cannot define the flux “through” that surface.

Flux in Other Dimensions

We've explored flux in three dimensions, but flux can be defined in any dimension for a “surface” in that dimension, where “surface” means something for which a consistent choice of normal vectors can be defined. This means *volumes* in four-dimensional space and curves in two-dimensional space can be considered surfaces. In other dimensions, one does not have the cross product, so an appeal to basic principles may be needed to compute flux.

■ **Example 6.8** Let \mathcal{B} be the boundary of $[0, 1] \times [0, 1]$ oriented outward and let $\vec{F}(x, y) = (x^2, -y^2)$. Compute the flux of \vec{F} through \mathcal{B} .

XXX Finish ■

Divergence

Let $\vec{F} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ be a vector field and suppose S is a sphere oriented outwards. As we move S around, the flux through S might change from positive to negative or negative to positive. If

⁵ For a surface $S \subset \mathbb{R}^3$, you can consider the *unit normal bundle* $\mathcal{N} \subset S \times \mathbb{R}^3$. That is, \mathcal{N} is the set of pairs (s, \vec{n}) where $s \in S$ and \vec{n} is a unit normal vector to S at the point s . Since $S \subset \mathbb{R}^3$, we may view $\mathcal{N} \subset \mathbb{R}^6$. Now, we say that (s_1, \vec{n}_1) can be continuously transported to (s_2, \vec{n}_2) if there is a continuous (connected) path in \mathcal{N} connecting (s_1, \vec{n}_1) and (s_2, \vec{n}_2) .

we change the size of S , again we might see a change in flux, and if we changed both the size and the position of S , we might be able to pinpoint places in \vec{F} where vectors are generally pointing away from each other (diverging) and where they are generally pointing towards each other (converging).

The *divergence* of a vector field precisely measures where a vector field is diverging or converging.

Definition 6.3.2 — Divergence. Let $\vec{F} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ be a vector field and let $R_{\Delta x, \Delta y, \Delta z}(x, y, z)$ be the surface of the box with lower left corner at (x, y, z) , side-lengths $\Delta x, \Delta y, \Delta z$, and oriented outwards^a. The *divergence* of \vec{F} at the point (x, y, z) , notated as $\nabla \cdot \vec{F}(x, y, z)$ is

$$\nabla \cdot \vec{F}(x, y, z) = \lim_{\Delta x, \Delta y, \Delta z \rightarrow 0} \frac{\text{flux of } \vec{F} \text{ through } R_{\Delta x, \Delta y, \Delta z}(x, y, z)}{\text{area of } R_{\Delta x, \Delta y, \Delta z}(x, y, z)}.$$

^a Precisely, $R_{\Delta x, \Delta y, \Delta z}(x, y, z)$ is the surface of $[x, x + \Delta x] \times [y, y + \Delta y] \times [z, z + \Delta z]$ oriented outwards.

In plain language, the divergence of a vector field can be thought of as the flux per unit volume of that vector field. Divergence in dimensions other than three is defined similarly.

■ **Example 6.9** Find the divergence of $\vec{F}(x, y, z) = (x, y, z)$ at the origin.

XXX Finish ■

■ **Example 6.10** Estimate regions where the divergence is positive, negative, or zero for the following vector field plots.

XXX Finish ■

The notation $\nabla \cdot \vec{F}$ comes from physics and is a handy mnemonic for remembering quick way to compute divergence. In order to see why this notation makes sense, we need to work in some generality.

To keep the number of variables down, we will explore divergence in \mathbb{R}^2 for the moment. Let $\vec{F} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ be a vector field defined by $\vec{F}(x, y) = \begin{bmatrix} f_x(x, y) \\ f_y(x, y) \end{bmatrix}$ and consider $R_{\Delta x, \Delta y}(x_0, y_0)$. $R_{\Delta x, \Delta y}(x_0, y_0)$ has four sides, which we will label L, R, T, B .

XXX Figure

We will compute the approximate flux through these sides separately. Consider $L = \{x_0\} \times [y_0, y_0 + \Delta y]$ oriented to the left. If Δy is small, \vec{F} (along L) can be reasonably approximated by the constant vector $\vec{F}(x_0, y_0)$. Therefore,

$$\text{Flux}_L \vec{F} \approx \Delta y \begin{bmatrix} -1 \\ 0 \end{bmatrix} \cdot \vec{F}(x_0, y_0) = -\Delta y f_x(x_0, y_0).$$

Similarly, near $R = \{x_0 + \Delta x\} \times [y_0, y_0 + \Delta y]$, \vec{F} can be approximated by $\vec{F}(x_0 + \Delta x, y_0)$. Since R is oriented to the right, we have

$$\text{Flux}_R \vec{F} \approx \Delta y \begin{bmatrix} 1 \\ 0 \end{bmatrix} \cdot \vec{F}(x_0 + \Delta x, y_0) = \Delta y f_x(x_0 + \Delta x, y_0).$$

Continuing the process, we compute

$$\text{Flux}_B \vec{F} \approx -\Delta x \begin{bmatrix} 0 \\ 1 \end{bmatrix} \cdot \vec{F}(x_0, y_0) = -\Delta x f_y(x_0, y_0),$$

and

$$\text{Flux}_T \vec{F} \approx \Delta x \begin{bmatrix} 0 \\ 1 \end{bmatrix} \cdot \vec{F}(x_0, y_0 + \Delta y) = \Delta x f_y(x_0, y_0 + \Delta y).$$

We're now ready to approximate the divergence of \vec{F} at (x_0, y_0) .

$$\begin{aligned} \nabla \cdot \vec{F} &\approx \frac{\text{Flux}_R \vec{F} + \text{Flux}_L \vec{F} + \text{Flux}_T \vec{F} + \text{Flux}_B \vec{F}}{\Delta x \Delta y} \\ &= \frac{\Delta y (f_x(x_0 + \Delta x, y_0) - f_x(x_0, y_0)) + \Delta x (f_y(x_0, y_0 + \Delta y) - f_y(x_0, y_0))}{\Delta x \Delta y} \\ &= \frac{f_x(x_0 + \Delta x, y_0) - f_x(x_0, y_0)}{\Delta x} + \frac{f_y(x_0, y_0 + \Delta y) - f_y(x_0, y_0)}{\Delta y}, \end{aligned}$$

with equality following after we take a limit as $\Delta x, \Delta y \rightarrow 0$. But some of these limits look familiar!

$$\lim_{\Delta x \rightarrow 0} \frac{f_x(x_0 + \Delta x, y_0) - f_x(x_0, y_0)}{\Delta x} = \frac{\partial f_x}{\partial x}(x_0, y_0)$$

and

$$\lim_{\Delta y \rightarrow 0} \frac{f_y(x_0, y_0 + \Delta y) - f_y(x_0, y_0)}{\Delta y} = \frac{\partial f_y}{\partial y}(x_0, y_0).$$

We conclude⁶

$$\nabla \cdot \vec{F}(x_0, y_0) = \frac{\partial f_x}{\partial x}(x_0, y_0) + \frac{\partial f_y}{\partial y}(x_0, y_0).$$

A very similar process shows that if $\vec{G} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ with component functions g_x, g_y, g_z , then

$$\nabla \cdot \vec{G}(x_0, y_0, z_0) = \frac{\partial g_x}{\partial x}(x_0, y_0, z_0) + \frac{\partial g_y}{\partial y}(x_0, y_0, z_0) + \frac{\partial g_z}{\partial z}(x_0, y_0, z_0).$$

In light of these formulas for divergence, the pneumatic $\nabla \cdot \vec{G}$ makes some sense. If we suspend our belief that every mathematical sentence we write should have meaning, we can think of $\nabla = (\frac{\partial}{\partial x}, \frac{\partial}{\partial y}, \frac{\partial}{\partial z})$ and $\vec{G} = (g_x, g_y, g_z)$, then

$$\nabla \cdot \vec{G} = (\frac{\partial}{\partial x}, \frac{\partial}{\partial y}, \frac{\partial}{\partial z}) \cdot (g_x, g_y, g_z) = \frac{\partial g_x}{\partial x} + \frac{\partial g_y}{\partial y} + \frac{\partial g_z}{\partial z}$$

is not bad as a memory aid⁷.

⁶ Of course, we had some sleight of hand in our derivation. In order to guarantee all the limits involved exist and don't depend on particular choices (like, why did we approximate \vec{F} along L by $\vec{F}(x_0, y_0)$ instead of $\vec{F}(x_0, y_0 + \frac{\Delta y}{2})$?) we need to assume that \vec{F} is differentiable. With that assumption, one can prove that the answer is the same regardless of choices for first-order approximations.

⁷ There is a theory of operators where $\frac{\partial}{\partial w}$ has meaning by itself. However, even in that framework, "multiplication" typically means scalar multiplication and not function application. There's a bit of twisting one needs to do to make something like $\nabla \cdot \vec{G}$ more than a pneumatic.

The Divergence Theorem

Imagine a pond with various holes in its basin. Water flows in from some of the holes and out from others. If we consider $\vec{V} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ to be the velocity of water in the pond, the divergence of \vec{V} would be positive at holes where water is flowing in and negative at holes where it is flowing out. For this reason, we call points in a vector field where the divergence is positive *sources* and points where the divergence is negative *sinks*.

If we enclose a region of the pond in a net and ask, “what is the total flux through the net?” it is reasonable to expect that the flux through the net is equal to the amount of water coming from sources minus the amount coming from sinks. The divergence theorem makes this reasoning precise.

In order to properly state the divergence theorem, we need some notation.

Definition 6.3.3 — Boundary. For a region $R \subseteq \mathbb{R}^n$, the *boundary* of R is notated ∂R and always comes with an outward orientation.

Precisely defining what it means to be a boundary of a region is tricky business. Fortunately, your intuition will take you quite far. For example, the boundary of a disk is a circle, the boundary of a filled in cube is the surface of that cube and the boundary of an interval consists of its endpoints.

XXX Figure

It’s important to note that boundaries can be empty. For example, the boundary of a single point is empty. The boundary of a circle is also empty (it has no “ends”). In fact, for any nice region R , it holds that $\partial(\partial R) = \emptyset$.

We can now precisely state the *divergence theorem* (also known as *Gauss’s theorem* or *Ostrogradsky’s theorem*).

Theorem 6.3.1 — Divergence Theorem. Let \vec{F} be a vector field and R be a region. Then

$$\text{Flux}_{\partial R} \vec{F} = \int_R \nabla \cdot \vec{F}.$$

The divergence theorem holds in all dimensions, but it is sometimes stated in \mathbb{R}^3 strictly in terms of integrals as

$$\iint_{\partial R} \vec{F} \cdot \hat{n} dV_{\partial R} = \iiint_R \nabla \cdot \vec{F} dV_R$$

where \hat{n} is a unit normal vector to ∂R , and $dV_{\partial R}$ and dV_R are the volume forms for ∂R and R , respectively.

■ **Example 6.11** Verify the divergence theorem for the vector field $\vec{F} = \dots$ on the region $R = \dots$
 XXX Finish ■

In words, the divergence theorem states that the flux through a boundary equals the sum of the divergence (sources and sinks) on its interior, and, though proving the divergence theorem can be challenging, the essential idea is simple and beautiful.

Let \vec{F} be a vector field, and consider two adjacent squares, A and B , both oriented outwards.

XXX Figure

Let $R = A \cup B$ be the rectangular region formed by joining A and B . Now,

$$\text{Flux}_R \vec{F} = \text{Flux}_A \vec{F} + \text{Flux}_B \vec{F}$$

because the contributions to flux on the common edge shared by A and B are opposite and cancel. Now consider a more complicated region X . If we chop X into tiny squares each called ΔX , we see that

$$\text{Flux}_{\partial X} \vec{F} = \sum_{\partial \Delta X} \text{Flux}_{\partial \Delta X} \vec{F},$$

because, again, flux through any interior edges cancels out. But, this means the only flux that actually contributes is the flux along the boundary.

XXX Figure

Of course, we were considering $\text{Flux}_{\partial X} \vec{F}$ to begin with, so only the boundary should matter. However, as the area of ΔX gets very tiny,

$$\text{Flux}_{\partial \Delta X} \vec{F} \approx (\text{area of } \Delta X)(\nabla \cdot \vec{F} \text{ at } \Delta X).$$

After taking a limit, we get the divergence theorem.

6.4 Circulation and Curl

Given a vector field and a surface, it is natural to ask what the flux of the vector field through the surface is. Given a vector field and a closed curve (a curve that loops back to itself), it is natural to ask what the *circulation* of the vector field through the curve is. That is, given a loop L , how much would a vector field “push” something around L .

Circulation is closely related to the concept of work, except that it only applies to closed curves.

Definition 6.4.1 — Circulation. Given an oriented, closed C and a vector field \vec{F} , the *circulation* of \vec{F} around C is the amount of work done by \vec{F} while traversing C . As a vector line integral,

$$\text{Circ}_C \vec{F} = \int_C \vec{F} \cdot d\vec{r}.$$

To visualize the circulation of \vec{F} around a curve C , imagine that C is a roller coaster track, weightless in space. You then place a car on this track and let it get pushed by \vec{F} . If it makes a complete loop and continues in the same direction that C is oriented, $\text{Circ}_C \vec{F}$ is positive. If it makes a loop and continues in the opposite direction to how C is oriented, $\text{Circ}_C \vec{F}$ is negative.

■ **Example 6.12** Find the circulation of ...

XXX Finish

Just like the divergence of a vector field is the amount of flux per unit area, we have an idea of “circulation per unit area,” which is called *curl*. However, there is a problem. At a point, there is fundamentally only *one* tiny volume centered at that point⁸, however, at a point, there are *infinitely* many tiny curves circulating that point—at a point \vec{p} , for every choice of vector \vec{n} ,

⁸ When we defined divergence, we used a cube, but we could have used a sphere.

we could consider a tiny curve going around \vec{p} and contained in the plane with normal vector \vec{n} . Thus, curl cannot be represented by a single scalar quantity.

Definition 6.4.2 — Curl. For a unit vector \vec{v} , let $R_{\Delta a, \Delta b}^{\vec{v}}(x, y, z)$ be the perimeter of a rectangle with lower left corner at (x, y, z) , side lengths Δa and Δb , contained in a plane with normal vector \vec{v} , and oriented counter-clockwise.

Let $\vec{F} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ be a vector field. The \vec{v} component of the curl of \vec{F} at the point (x, y, z) is

$$\text{curl}_{\vec{v}} \vec{F}(x, y, z) = \lim_{\Delta a, \Delta b \rightarrow 0} \frac{\text{circulation of } \vec{F} \text{ around } R_{\Delta a, \Delta b}^{\vec{v}}(x, y, z)}{\text{area of } R_{\Delta a, \Delta b}^{\vec{v}}(x, y, z)}.$$

The curl of \vec{F} at the point (x, y, z) , notated $\nabla \times \vec{F}(x, y, z)$, is the vector so that

$$(\nabla \times \vec{F}(x, y, z)) \cdot \vec{v} = \text{curl}_{\vec{v}} \vec{F}(x, y, z).$$

It's important to notice that for a vector field \vec{F} and a unit vector \vec{v} , the *component* of the curl of \vec{F} in the direction \vec{v} is a scalar, but *the* curl of \vec{F} is encapsulated as a vector.

A priori, $\nabla \times \vec{F}$ might be hard to compute, but its definition makes it workable. If we assume $\nabla \times \vec{F}(x, y, z) = (a, b, c)$, then

$$\begin{aligned} a &= \nabla \times \vec{F}(x, y, z) \cdot \hat{\mathbf{x}} = \text{curl}_{\hat{\mathbf{x}}} \vec{F}(x, y, z) \\ b &= \nabla \times \vec{F}(x, y, z) \cdot \hat{\mathbf{y}} = \text{curl}_{\hat{\mathbf{y}}} \vec{F}(x, y, z) \\ c &= \nabla \times \vec{F}(x, y, z) \cdot \hat{\mathbf{z}} = \text{curl}_{\hat{\mathbf{z}}} \vec{F}(x, y, z), \end{aligned}$$

and $R_{\Delta a, \Delta b}^{\hat{\mathbf{x}}}(x, y, z)$, $R_{\Delta a, \Delta b}^{\hat{\mathbf{y}}}(x, y, z)$, and $R_{\Delta a, \Delta b}^{\hat{\mathbf{z}}}(x, y, z)$ are workable rectangles⁹.

■ Example 6.13 XXX Finish

Again, like divergence, the $\nabla \times \vec{F}$ notation serves as a mnemonic for computing curl. The derivation of a formula for the components of curl is very similar to the derivation for the components of divergence, so we won't go through it here. When all is said and done, if $\vec{F}(x, y, z) = (f_x(x, y, z), f_y(x, y, z), f_z(x, y, z))$, then

$$\nabla \times \vec{F} = \begin{bmatrix} \frac{\partial f_z}{\partial y} - \frac{\partial f_y}{\partial z} \\ \frac{\partial f_x}{\partial z} - \frac{\partial f_z}{\partial x} \\ \frac{\partial f_y}{\partial x} - \frac{\partial f_x}{\partial y} \end{bmatrix}.$$

Closed Vector Fields

We already have the tools to deduce some theorems about curl and circulation. In particular, if $\vec{F} = \nabla f$ is a conservative vector field, then the circulation around any closed loop should be zero. Therefore, every component of the curl of \vec{F} is zero.

⁹ Pun intended.

Theorem 6.4.1 If $\vec{F} = \nabla f$ is a conservative vector field, then $\nabla \times \vec{F} = \vec{0}$.

The converse of this statement isn't necessarily true.

Definition 6.4.3 — Closed Vector Field. A vector field $\vec{F} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ is called *closed* if $\nabla \times \vec{F} = \vec{0}$.

If we write out in components what it means for a vector field to be closed, we see that \vec{F} is closed exactly when it passes the *screening test* from Section 6.2. We already know that passing the screening test is necessary but not sufficient for showing a vector field is conservative.

■ **Example 6.14** Show that ... is closed by not conservative

XXX Finish ■

However, there are times when closed does imply conservative. In particular, if a region could be deformed to a ball, a vector field being closed on that region implies that it's conservative.

Theorem 6.4.2 — Poincaré Lemma. Let R be a region that can be deformed into a ball. Then, for a vector field \vec{F} , if $\nabla \times \vec{F} = \vec{0}$ on the region R , it must be that \vec{F} is conservative.

■ **Example 6.15** XXX show that the region from the previous example cannot be deformed into a ball. ■

xx

Curl in Other Dimensions

The concept of circulation is defined for vector fields of any dimension. Therefore, we can define the analog of curl in any dimension. However, in other dimensions, curl looks a little different than in three dimensions.

For example, if $\vec{F} : \mathbb{R} \rightarrow \mathbb{R}^2$, then the curl of \vec{F} takes the form of a scalar. Why? Well, if we specify the normal vector of a plane in \mathbb{R}^2 , there aren't that many choices. We can either specify the normal vector to be "up," giving the plane a counter-clockwise orientation, or we can specify it to be "down," giving it a clockwise orientation. Therefore, there is only one *component* to the curl of a vector field in \mathbb{R}^2 .

■ **Example 6.16** XXX Example of curl of vector field in \mathbb{R}^2 ■

If, $G : \mathbb{R}^4 \rightarrow \mathbb{R}^4$, things are even stranger. To specify the orientation of a plane in \mathbb{R}^4 , you need a *bi-vector*. After some fiddling, you can conclude that to represent the curl of \vec{G} , you need *six* components¹⁰.

In general, to represent curl in \mathbb{R}^n , you need a $n(n-1)/2$ components. Of course, there are better ways to talk about curl in \mathbb{R}^n than vectors (for example, differential forms).

Stokes's Theorem

It may have struck you that circulation through adjacent regions cancels in a similar way to how flux canceled.

¹⁰ Surprise, it's not eight!

Let A and B be the perimeter of two adjacent squares oriented counter-clockwise, and let R be the perimeter of the rectangle containing A and B oriented counter-clockwise.

XXX Figure

We see that because the shared edge between A and B is oriented in opposite directions, the contribution to circulation along that edge cancels. Thus, for a vector field \vec{F} ,

$$\text{Circ}_R \vec{F} = \text{Circ}_A \vec{F} + \text{Circ}_B \vec{F}.$$

We will now play a nearly-identical trick as in the derivation of the divergence theorem. However, we need to sort out the issue of *orientation* again.

For a volume V in \mathbb{R}^3 or an area A in \mathbb{R}^2 , the orientation of ∂V and ∂A is “outward.” However, what if we have a surface S in \mathbb{R}^3 . If S is a surface, then ∂S is a curve, and there is no “outward” for a curve in \mathbb{R}^3 . There is however a forward and a backwards. We will define the orientation of ∂S according to the right-hand rule.

Without getting overly technical, we say that ∂S is *oriented consistently with S according to the right hand rule* if the following condition holds. Let $s \in \partial S$ and let \vec{t} be the tangent vector to ∂S aligned with ∂S ’s orientation; let \vec{n} be a normal vector to S aligned with its orientation; then, $\vec{t} \times \vec{n}$ points “away” from the boundary.

XXX Figure

Now we are ready explore the analog of the divergence theorem, Stokes’s theorem. Consider a region X . If we chop X into tiny squares each called ΔX , we see that

$$\text{Circ}_{\partial X} \vec{F} = \sum_{\partial \Delta X} \text{Circ}_{\partial \Delta X} \vec{F},$$

because circulation along any interior edges cancels out.

XXX Figure

However, as the area of ΔX gets very tiny,

$$\text{Circ}_{\partial \Delta X} \vec{F} \approx (\text{area of } \Delta X)(\nabla \times \vec{F} \text{ at } \Delta X) \cdot \vec{n},$$

where \vec{n} is a unit normal vector to ΔX . But,

$$(\text{area of } \Delta X)(\nabla \times \vec{F} \text{ at } \Delta X) \cdot \vec{n} = \text{Flux}_{\Delta X}(\nabla \times \vec{F} \text{ at } \Delta X).$$

After taking a limit, the approximation becomes exact and we get Stokes’s theorem relating flux, circulation, and curl.

Theorem 6.4.3 — Stokes’s Theorem. Let $S \subset \mathbb{R}^3$ be an oriented surface and let $\vec{F} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$. Then,

$$\text{Circ}_{\partial R} \vec{F} = \text{Flux}_R(\nabla \times \vec{F}).$$

We can rephrase Stokes’s theorem in terms of integrals. That is, for an oriented surface S and vector field $\vec{F} : \mathbb{R}^3 \rightarrow \mathbb{R}^3$,

$$\int_{\partial S} \vec{F} \cdot d\vec{r} = \iint_S (\nabla \times \vec{F}) \cdot \hat{n} dV_S,$$

where \hat{n} is the unit normal vector to S with appropriate orientation and dV_S is the volume form for S . Phrased this way, there is a striking similarity between Stokes’s theorem and the divergence theorem.

- **Example 6.17** Use Stoke's theorem to calculate ...
XXX Finish



Chapter 7

Differential Forms

Multivariable calculus has a long history with many refinements over the years. Every once in a while, a mathematician will notice something new—a connection between various concepts in mathematics. You may have noticed that much of what we have done was not particular to any specific dimension, but the problems we solved were usually in \mathbb{R}^2 and \mathbb{R}^3 . The theoretical framework for work in \mathbb{R}^4 (like is often required for relativity) is set, but how would we actually compute a surface integral in \mathbb{R}^4 ? We would need to find the volume form/surface element without the use of the cross product.

Differential forms are one solution to this issue. They recast many of the fundamental ideas, gradients, curl, divergence, line integrals, surface integrals, etc., in a unified framework that works regardless of dimension and *regardless of coordinate system*. Because of their abstract foundation, they are not typically encountered in a first course on multivariable calculus, but that has more to do with culture than anything else.

7.1 Abstract Foundation for Differential Forms

Differential forms have close ties with calculus and linear algebra. For that reason, we will use some terms from linear algebra, including the abstract idea of a *vector*.

Definition 7.1.1 — Vector Space. A set V , coupled with an addition and a scalar multiplication is called a *vector space* if the following properties hold for all $\vec{u}, \vec{v}, \vec{w} \in V$ and all scalars α and β :

$\vec{u} + \vec{v} \in V$ and $\alpha\vec{u} \in V$	(Closure)
$(\vec{u} + \vec{v}) + \vec{w} = \vec{u} + (\vec{v} + \vec{w})$	(Associativity)
$\vec{u} + \vec{v} - \vec{v} = \vec{u} + \vec{0} = \vec{u}$	(Identity)
$\vec{u} + \vec{v} = \vec{v} + \vec{u}$	(Commutativity)
$\alpha(\vec{u} + \vec{v}) = \alpha\vec{u} + \alpha\vec{v}$	(Distributivity)
$(\alpha\beta)\vec{v} = \alpha(\beta\vec{v})$	(Scalar Associativity)
$(\alpha + \beta)\vec{v} = \alpha\vec{v} + \beta\vec{v}$	(Scalar Distributivity)

These properties should look familiar from Section 2.2, and indeed \mathbb{R}^n is a real vector

space. However there are many seemingly non-geometric sets that are also vector spaces.

■ **Example 7.1** Let $V = \{\text{functions from } \mathbb{R} \text{ to } \mathbb{R}\}$ equipped with pointwise addition and scalar multiplication. Then V is a vector space.

We won't prove every property in detail, but if we pick $f, g \in V$, we need the quantities $f + g$ and αf to make sense and to also be in V . This becomes almost trivially true, thanks to the definition of addition of functions.

Let $h = f + g$. By definition $h(x) = f(x) + g(x)$, and so h is a function from \mathbb{R} to \mathbb{R} . Therefore $h \in V$. Further $q = \alpha h$ is defined by $q(x) = \alpha h(x)$, and so $q \in V$. Thus, V satisfies the closure property. The rest of the properties follow via similar arguments. ■

Recall that if $\vec{v}_1, \dots, \vec{v}_n \in V$ are vectors, then $\vec{w} = \alpha_1 \vec{v}_1 + \dots + \alpha_n \vec{v}_n$ is a *linear combination* of the vectors $\vec{v}_1, \dots, \vec{v}_n$. Another way to think of a vector space is as a set on which linear combinations make sense.

Now that we have abstract “vectors” (elements of a vector space), let's get abstract lines and planes. These objects can be defined using *linear functions* which are equivalently called *linear transformations*¹.

Definition 7.1.2 — Linear. Let V and W be vector spaces and let $T : V \rightarrow W$ be a function. T is called *linear* if for all $\vec{u}, \vec{v} \in V$ and all scalars α ,

$$T(\vec{u} + \vec{v}) = T(\vec{u}) + T(\vec{v}) \quad \text{and} \quad T(\alpha \vec{v}) = \alpha T(\vec{v}).$$

If a transformation is linear, it is typical to omit parenthesis when passing in a single vector. That is, instead of writing $T(\vec{v})$, we write $T\vec{v}$, but we would always write $T(\vec{u} + \vec{v})$ with parenthesis.

We have already encountered many linear functions. For example, consider $f : \mathbb{R}^3 \rightarrow \mathbb{R}$ where $f(\vec{v}) = \hat{\mathbf{x}} \cdot \vec{v}$. Based on the properties of the dot product,

$$f(\vec{u} + \vec{v}) = \hat{\mathbf{x}} \cdot (\vec{u} + \vec{v}) = \hat{\mathbf{x}} \cdot \vec{u} + \hat{\mathbf{x}} \cdot \vec{v} = f(\vec{u}) + f(\vec{v})$$

and

$$f(\alpha \vec{v}) = \hat{\mathbf{x}} \cdot (\alpha \vec{v}) = \alpha \hat{\mathbf{x}} \cdot \vec{v} = \alpha f(\vec{v}).$$

Another linear transformation that we have encountered is projection. For a fixed vector \vec{u} , let $f : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ be defined by $f(\vec{v}) = \text{proj}_{\vec{u}} \vec{v}$.

Exercise 7.1 Show that for a fixed \vec{u} , the function $f(\vec{v}) = \text{proj}_{\vec{u}} \vec{v}$ is linear.

Linear functions help us define lines, planes, etc. via their *kernel*.

Definition 7.1.3 — Kernel. For a function $T : V \rightarrow W$, the *kernel* (or *null space* if T is a linear function) of T is defined to be

$$\text{kernel}(T) = \{\vec{v} \in V : T(\vec{v}) = \vec{0}\}.$$

Thinking at this level of abstraction can be challenging at first, but we've really been working with these ideas for a long time.

¹ The words *function*, *transformation*, and *mapping* are synonymous in math. We like to mix it up some times.

For example, let $\vec{n} = (1, 2, 3)$ and let $P : \mathbb{R}^3 \rightarrow \mathbb{R}$ be defined by $P(\vec{v}) = \vec{n} \cdot \vec{v}$. We've already seen that P is linear. What is $\text{kernel}(P)$? Well,

$$P \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \end{bmatrix} = x + 2y + 3z = 0$$

exactly when (x, y, z) lies on the plane through the origin with normal vector \vec{n} .

Lines can be written as kernels as well.

Exercise 7.2 Let $\vec{d} = (1, 1, 1)$. Find a linear transformation L so that $\text{kernel}(L)$ is the line through $\vec{0}$ with direction vector \vec{d} . (It may help to think of a line as the intersection of two planes.)

Unfortunately, linear functions cannot describe *all* “flat” objects. The kernel of a linear function always contains $\vec{0}$, and so lines and planes that do not pass through the origin cannot be described as kernels of linear functions. However, they can be described as kernels of the related *affine functions*.

Definition 7.1.4 — Affine Function. Let V and W be vector spaces. A function $A : V \rightarrow W$ is called *affine* if there exists some $\vec{w} \in W$ and some linear transformation $T : V \rightarrow W$ so that

$$A(\vec{v}) = T(\vec{v}) + \vec{w}$$

for all \vec{v} .

You may recall the previously, in Section 4.3, we only defined what it means for a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ to be affine. We now have a fully general definition!

Exercise 7.3 Write the plane with normal vector $\vec{n} = (1, 2, 3)$ and which passes through the point $\vec{p} = (1, 1, 1)$ as the kernel of an affine function.

Covectors

Now that we have formulated standard geometric objects, we will develop a set of tools to *measure* those objects. The first property we will attempt to measure is displacement. From displacement, we will be able to measure area as “displacement in two directions” and volume as “displacement in three directions.”

Displacement is measured with objects called *covectors*. In other contexts, these objects are also known as *linear functionals* or *dual vectors*. We will typically write covectors in a bold font to distinguish them from vectors.

Definition 7.1.5 — Covector. Given a vector space V , a *covector* on the space V is a linear function $\alpha : V \rightarrow \mathbb{R}$. The set of all covectors on V is denoted V^* .

This definition is very abstract, but again, we've encountered covectors before! Let $\alpha : \mathbb{R}^3 \rightarrow \mathbb{R}$ be defined by $\alpha(\vec{v}) = \hat{x} \cdot \vec{v}$. We already know that α is linear and that it outputs real numbers. Therefore it is a covector. However, α also has geometric meaning. It $\alpha(\vec{v})$ outputs the x -coordinate of \vec{v} . We could similarly think of covectors that when given a vector output its y or z coordinates.

Any function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ defined by $f(\vec{v}) = \vec{r} \cdot \vec{v}$ for some fixed \vec{r} is a covector. Further, we can interpret f to have geometric meaning. The quantity $f(\vec{v})$ is how much \vec{v} points in the direction \vec{r} . In other words, $f(\vec{v})$ is the “ \vec{r} component” of \vec{v} (is anyone else thinking about projections?).

A theorem from linear algebra shows us that this type of covector is all there is (in \mathbb{R}^n , at least).

Theorem 7.1.1 Let $\alpha : \mathbb{R}^n \rightarrow \mathbb{R}$ be a covector for \mathbb{R}^n . Then, there exists a $\vec{a} \in \mathbb{R}^n$ such that

$$\alpha(\vec{v}) = \vec{a} \cdot \vec{v}$$

for all $\vec{v} \in \mathbb{R}^n$.

Now we know all covectors for \mathbb{R}^n have a geometric meaning as measuring some component of an input vector. The following theorem says that the term “vector” in *covector* is justified.

Theorem 7.1.2 If V is a vector space, then V^* , the set of all covectors for V is also a vector space.

Proof. XXX Finish ■

We won’t be doing much with the abstract vector space V^* , but it is nice to know that we can take linear combinations of covectors and still get a covector².

It is, however, worth our time to consider $(\mathbb{R}^n)^*$. Recall that in \mathbb{R}^3 we have special vectors \hat{x} , \hat{y} , and \hat{z} , called the standard basis for \mathbb{R}^3 , and all other vectors can be written as a linear combination of these. Higher dimensions also have the idea of a standard basis, and there is a similar concept for sets of covectors.

Definition 7.1.6 — Standard Basis. The *standard basis* for \mathbb{R}^n is notated by $\vec{e}_1, \vec{e}_2, \dots, \vec{e}_n$ and satisfies that every vector in \mathbb{R}^n can be written as a linear combination of $\vec{e}_1, \vec{e}_2, \dots, \vec{e}_n$ and that

$$\vec{e}_i \cdot \vec{e}_j = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}.$$

The standard basis for $(\mathbb{R}^n)^*$ is notated e^1, e^2, \dots, e^n and satisfies that every covector for \mathbb{R}^n can be written as a linear combination of e^1, e^2, \dots, e^n and that

$$e^i(\vec{e}_j) = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}.$$

Following the convention from physics, we index covectors using superscripts and regular vectors using subscripts. In plain terms, the standard basis for \mathbb{R}^n is a collection of unit vectors that point along the coordinate axes³, and the standard basis for $(\mathbb{R}^n)^*$ picks off the standard coordinates of input vectors.

² If we wanted to investigate V^* , we might start by saying, since V^* is a vector space, there is a set V^{**} of covectors for V^* . Another linear algebra theorem says $V \subseteq V^{**}$, if you interpret “subset” in the right sense.

³ Technically, the standard basis *defines* the coordinate axes.

It's now easy to appreciate that we can take linear combinations of covectors. For example, let $\alpha = e^1 + 2e^3$. We can compute

$$\alpha \begin{bmatrix} a \\ b \\ c \end{bmatrix} = e^1 \begin{bmatrix} a \\ b \\ c \end{bmatrix} + 2e^3 \begin{bmatrix} a \\ b \\ c \end{bmatrix} = a + 2c = \begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix} \cdot \begin{bmatrix} a \\ b \\ c \end{bmatrix},$$

and so we see that α is given by taking the dot product with the vector $\begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix}$. Of course,

we could write $\vec{e}_1 + 2\vec{e}_3$ instead of $\begin{bmatrix} 1 \\ 0 \\ 2 \end{bmatrix}$ and $a\vec{e}_1 + b\vec{e}_2 + c\vec{e}_3$ instead of $\begin{bmatrix} a \\ b \\ c \end{bmatrix}$. In doing so,

we would be analyzing vectors in a *coordinate free* manner. That is, we never need to talk about the coordinates of a vector in order to do a computation. One of the selling points of differential forms is that they allow you to do multivariable calculus in a coordinate free way.

Stepping back for a moment, the machinery we've developed is almost comically tautological. Let $\vec{v} \in \mathbb{R}^n$ and let α be a covector. By definition,

$$\begin{aligned} \vec{v} &= a_1\vec{e}_1 + \cdots + a_n\vec{e}_n \\ \alpha &= b_1e^1 + \cdots + b_ne^n, \end{aligned}$$

and it follows that

$$\alpha(\vec{v}) = (b_1e^1 + \cdots + b_ne^n)(a_1\vec{e}_1 + \cdots + a_n\vec{e}_n) = b_1a_1 + \cdots + b_na_n = \begin{bmatrix} b_1 \\ \vdots \\ b_n \end{bmatrix} \cdot \begin{bmatrix} a_1 \\ \vdots \\ a_n \end{bmatrix}.$$

Thus, while by definition

$$\vec{v} = a_1\vec{e}_1 + \cdots + a_n\vec{e}_n = \begin{bmatrix} a_1 \\ \vdots \\ a_n \end{bmatrix},$$

if we wrote

$$\alpha = b_1e^1 + \cdots + b_ne^n = \begin{bmatrix} b_1 \\ \vdots \\ b_n \end{bmatrix},$$

we would have gone through a whole lot of work to reinvent the dot product! It will be worth it though⁴.

Covectors as Rulers

XXX Finish

⁴ Some would argue that dot and cross products confuse vectors and covectors—that quantities like *force* are naturally covectors, and that physics should be taught with covectors (instead of dot products) from the beginning.

7.2 1-forms and Line Integrals

Recall that a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ was called a vector field because to each point in \mathbb{R}^n it assigns a vector in \mathbb{R}^n . A 1-form is a type of *differential form* that is analogous to a vector field except that instead of assigning to each point a vector, it assigns to each point a *covector*. We will also write 1-forms using a bold font.

Definition 7.2.1 — 1-form. A 1-form on the space \mathbb{R}^n is a function $\alpha : \mathbb{R}^n \rightarrow (\mathbb{R}^n)^*$.

We will always assume our 1-forms are *smooth*—that is, every 1-form can be differentiated as much as we like⁵.

Since 1-forms output covectors which themselves measure vectors, it is worth introducing some notation. If α is a 1-form, then $\alpha(\vec{v})$ is a covector. If we wish to apply this covector, $\alpha(\vec{v})$, to measure the vector \vec{w} , we write $[\alpha(\vec{v})](\vec{w})$. We'll be doing this a lot, so we have the alternative notation

$$\alpha(\vec{v}; \vec{w}) \equiv [\alpha(\vec{v})](\vec{w}).$$

As we've seen, covectors are used to measure vectors, and vectors themselves come from displacements. 1-forms are a way to measure *oriented curves*. We'll explore this by first considering some special 1-forms on \mathbb{R}^3 .

Definition 7.2.2 — Standard Basis for 1-forms. The *standard basis for 1-forms* are the 1-forms $\mathrm{d}x$, $\mathrm{d}y$, and $\mathrm{d}z$ defined by

$$\mathrm{d}x(\vec{v}) = \mathbf{e}^1 \quad \mathrm{d}y(\vec{v}) = \mathbf{e}^2 \quad \mathrm{d}z(\vec{v}) = \mathbf{e}^3$$

for all $\vec{v} \in \mathbb{R}^3$.

The names $\mathrm{d}x$, $\mathrm{d}y$, and $\mathrm{d}z$ are certainly evocative. When you see “ d ,” you should think “displacement.” After all $\mathrm{d}x(\vec{v})$ is the covector that measures the x displacement of a vector. Soon, you will also think “derivative.”

Measuring Curves with 1-forms

Let $\mathcal{C} \subset \mathbb{R}^3$ be a smooth curve that starts at $\vec{a} = (a_x, a_y, a_z)$ and ends at $\vec{b} = (b_x, b_y, b_z)$. The 1-form $\mathrm{d}x$ is meant to measure the displacement of \mathcal{C} from its starting point to its ending point. However, $\mathrm{d}x$ cannot *directly* interact with \mathcal{C} . Recall, $\mathrm{d}x : \mathbb{R}^3 \rightarrow (\mathbb{R}^3)^*$ is a function that inputs points and outputs covectors and covectors measure vectors. \mathcal{C} certainly has points, and if we squint our eyes, \mathcal{C} also has vectors—tangent vectors!

Let $\vec{r} : [0, 1] \rightarrow \mathcal{C}$ be a parameterization of \mathcal{C} defined by

$$\vec{r}(t) = \begin{bmatrix} r_x(t) \\ r_y(t) \\ r_z(t) \end{bmatrix}$$

⁵ We haven't defined what it means to differentiate a function whose codomain consists of covectors. The quick and dirty answer is, treat covectors as vectors and use the definition of differentiability for a vector field.

so that $\vec{r}(0) = \vec{a}$ and $\vec{r}(1) = \vec{b}$ (that is, it is oriented appropriately). Then,

$$\vec{r}'(t) = \begin{bmatrix} r'_x(t) \\ r'_y(t) \\ r'_z(t) \end{bmatrix}$$

is a tangent vector to \mathcal{C} at the point $\vec{r}(t)$.

XXX Figure including $\mathbb{d}x(r)$

Now, $\mathbb{d}x \circ \vec{r} = \mathbf{e}^1$ is the constant covector that measures the x displacement of a vector and so $\mathbb{d}x(\vec{r}(t); \vec{r}'(t)) = \mathbf{e}^1(\vec{r}'(t))$ is the x displacement of the tangent vector to \mathcal{C} coming from the parameterization \vec{r} at the point $\vec{r}(t)$. In other words,

$$\mathbb{d}x(\vec{r}(t); \vec{r}'(t)) = \mathbf{e}^1(\vec{r}'(t)) = r'_x(t).$$

What we're after is the *total* x displacement along the curve \mathcal{C} . That is $b_x - a_x$. But, we can get this quantity from the fundamental theorem of calculus! After some reverse engineering, we see

$$b_x - a_x = \int_0^1 r'_x(t) dt = \int_0^1 \mathbb{d}x(\vec{r}(t); \vec{r}'(t)) dt.$$

Further, this integral works out to the same value regardless of our choice of parameterization. If \vec{q} were another parameterization of \mathcal{C} starting at \vec{a} and ending at \vec{b} , then

$$\int_0^1 \mathbb{d}x(\vec{q}(t); \vec{q}'(t)) dt = b_x - a_x.$$

Based on this observation, we are ready to define what it means to integrate a 1-form.

Definition 7.2.3 — Integral of a 1-form. Let \mathcal{C} be an oriented curve that starts at \vec{a} and ends at \vec{b} and let α be a 1-form. We define the *integral of α over the curve \mathcal{C}* by

$$\int_{\mathcal{C}} \alpha = \int_a^b \alpha(\vec{q}(t); \vec{q}'(t)) dt = \int_a^b [\alpha \circ \vec{q}(t)](\vec{q}'(t)) dt$$

where $\vec{q} : [a, b] \rightarrow \mathcal{C}$ is a parameterization of \mathcal{C} satisfying $\vec{q}(a) = \vec{a}$ and $\vec{q}(b) = \vec{b}$.

With this definition, we have a powerful notation. For an oriented curve \mathcal{C} ,

$$\int_{\mathcal{C}} \mathbb{d}x$$

is the x displacement from the start of \mathcal{C} to the end of \mathcal{C} . Similarly, $\int_{\mathcal{C}} \mathbb{d}y$ gives the y displacement and $\int_{\mathcal{C}} \mathbb{d}z$ gives the z displacement.

There are many nice properties of the integral notation applied to 1-forms. For example, it behaves like a classical integral with respect to linear combinations. If α and β are 1-forms and t is a scalar, then

$$\int_{\mathcal{C}} (\alpha + t\beta) = \int_{\mathcal{C}} \alpha + t \int_{\mathcal{C}} \beta.$$

Exercise 7.4 Explain in words how to interpret the output of $\int_C (\mathrm{d}x + \mathrm{d}y + \mathrm{d}z)$.

Non-constant 1-forms

Let us seek the comfort of a one dimensional space as we explore non-constant 1-forms. In \mathbb{R}^1 we only have one standard basis 1-form: $\mathrm{d}x$. Let $f : \mathbb{R}^1 \rightarrow \mathbb{R}$ be a scalar-valued function. We will interpret \mathbb{R}^1 as a vector space rather than a set of scalars. Now, define

$$f \mathrm{d}x(\vec{v}) = f(\vec{v})\mathrm{d}x(\vec{v}) = f(\vec{v})\mathbf{e}^1.$$

That is, $f \mathrm{d}x$, when given a point, returns a scaled version of \mathbf{e}^1 , which will in turn measure a scaled version of the x displacement of a vector. Let $\mathcal{C} = [a, b]$ be the oriented interval starting at a and ending at b . \mathcal{C} is naturally parameterized by the function $\vec{r} : [a, b] \rightarrow [a, b]$ defined by $\vec{r}(t) = t$. Since \vec{r} is the identity function, $f \circ \vec{r}(t) = f(t)$ and $\vec{r}'(t) = 1$. Thus, we may compute

$$\int_{\mathcal{C}} f \mathrm{d}x = \int_a^b f \circ \vec{r}(t) \mathrm{d}x(\vec{r}(t); \vec{r}'(t)) \mathrm{d}t = \int_a^b f(t) \mathrm{d}t.$$

This shows that $\int_{\mathcal{C}} f \mathrm{d}x$ is the usual integral of f on the interval $[a, b]$.

Let's move now to \mathbb{R}^2 . Let $f_x : \mathbb{R}^2 \rightarrow \mathbb{R}$ and $f_y : \mathbb{R}^2 \rightarrow \mathbb{R}$ and consider the 1-form $\alpha = f_x \mathrm{d}x + f_y \mathrm{d}y$. Again, let \mathcal{C} be an oriented curve starting at \vec{a} , ending at \vec{b} , and parameterized by $\vec{r} : [0, 1] \rightarrow \mathcal{C}$ where $\vec{r}(t) = (r_x(t), r_y(t))$ and $\vec{r}(0) = \vec{a}$ and $\vec{r}(1) = \vec{b}$.

Now,

$$\begin{aligned} \int_{\mathcal{C}} \alpha &= \int_{\mathcal{C}} f_x \mathrm{d}x + \int_{\mathcal{C}} f_y \mathrm{d}y \\ &= \int_0^1 f_x \circ \vec{r}(t) \mathrm{d}x(\vec{r}(t); \vec{r}'(t)) \mathrm{d}t + \int_0^1 f_y \circ \vec{r}(t) \mathrm{d}y(\vec{r}(t); \vec{r}'(t)) \mathrm{d}t \\ &= \int_0^1 f_x \circ \vec{r}(t) r'_x(t) \mathrm{d}t + \int_0^1 f_y \circ \vec{r}(t) r'_y(t) \mathrm{d}t \end{aligned}$$

If we let $\vec{F} = (f_x, f_y)$, a familiar integral emerges.

$$\int_{\mathcal{C}} \alpha = \int_0^1 \vec{F} \circ \vec{r}(t) \cdot \vec{r}'(t) \mathrm{d}t = \int_{\mathcal{C}} \vec{F} \cdot \mathrm{d}\vec{r},$$

where the right-hand side is the vector line integral of \vec{F} along the curve \mathcal{C} . So, integrals of 1-forms along curves are actually vector line integrals.

■ **Example 7.2** Let $\alpha = -y \mathrm{d}x + x \mathrm{d}y$ and let \mathcal{C} be the circle of radius 2 centered at the origin. Compute $\int_{\mathcal{C}} \alpha$.

XXX Finish

■

7.3 2-forms and Surface Integrals

The Wedge Product

We can think of covectors as measuring oriented line segments. That is, any vector you can think of as a line segment with a start and an end. When the covector $\alpha = ae^1 + be^2 + ce^3$ is applied to a vector \vec{v} , it measures the “ (a, b, c) component” of \vec{v} . In other words, it measures how long \vec{v} appears to be from the perspective of the vector (a, b, c) . That is, $\alpha(\vec{v}) = (a, b, c) \cdot \vec{v}$.

We can invent a similar object that measures oriented parallelograms. Any pair of vectors (\vec{v}, \vec{w}) can be viewed as forming an oriented parallelogram $\mathcal{P}_{\vec{v}, \vec{w}}$. If (\vec{v}, \vec{w}) come in a right-handed order, we say $\mathcal{P}_{\vec{v}, \vec{w}}$ has a positive orientation. Otherwise it has a negative orientation. A function that takes in a pair of vectors and measures the oriented (signed) area of the parallelogram they determine is called a 2-covector.

Let us take a moment to discuss the *laws of area*⁶ that should apply to 2-covector. If α is a 2-covector and \vec{v}, \vec{w} are vectors, then

$$\alpha(\vec{v}, \vec{v}) = 0,$$

since the parallelogram $\mathcal{P}_{\vec{v}, \vec{v}}$ is degenerate⁷. Scaling one of the sides should affect the area accordingly:

$$\alpha(m\vec{v}, \vec{w}) = m\alpha(\vec{v}, \vec{w}) = \alpha(\vec{v}, m\vec{w}).$$

Further, if we slide one vector parallel to the other, it shouldn't affect the area. So,

$$\alpha(\vec{v}, \vec{w}) = \alpha(\vec{v} + m\vec{w}, \vec{w}) = \alpha(\vec{v}, \vec{w} + m\vec{v}).$$

Lastly, though not a direct consequence of the laws of area, we will require that α is linear in each argument. That is, that α is *multi-linear*.

Definition 7.3.1 — Multi-linear. A function $f : \underbrace{V \times \cdots \times V}_n \rightarrow W$ is *multi-linear* if it is linear in each argument. That is, for all $\vec{v}_1, \dots, \vec{v}_n, \vec{w} \in V$ and all scalars α ,

$$f(\vec{v}_1, \dots, \vec{v}_i + \vec{w}, \dots, \vec{v}_n) = f(\vec{v}_1, \dots, \vec{v}_i, \dots, \vec{v}_n) + f(\vec{v}_1, \dots, \vec{w}, \dots, \vec{v}_n)$$

and

$$f(\vec{v}_1, \dots, \alpha\vec{v}_i, \dots, \vec{v}_n) = \alpha f(\vec{v}_1, \dots, \vec{v}_i, \dots, \vec{v}_n).$$

Definition 7.3.2 — 2-covector. A 2-covector on the vector space V is a multi-linear function $\alpha : V \times V \rightarrow \mathbb{R}$ that obeys the laws of area.

The laws of area should look very familiar from our derivation of the cross product. And, now we can see that orientation is a direct consequence of multi-linearity and the fact that $\alpha(\vec{v}, \vec{v}) = 0$. Observe,

$$\begin{aligned} 0 &= \alpha(\vec{v} + \vec{w}, \vec{v} + \vec{w}) = \alpha(\vec{v}, \vec{v}) + \alpha(\vec{v}, \vec{w}) + \alpha(\vec{w}, \vec{v}) + \alpha(\vec{w}, \vec{w}) \\ &= \alpha(\vec{v}, \vec{w}) + \alpha(\vec{w}, \vec{v}), \end{aligned}$$

⁶ We'll be calling them *laws of volume* later on.

⁷ It isn't bad; it just has no area.

and so $\alpha(\vec{v}, \vec{w}) = -\alpha(\vec{w}, \vec{v})$. This property is called *alternating*.

We've generated a lot of theory, but let's actually see a 2-vector in action.

■ **Example 7.3** For a pair of vectors (\vec{v}, \vec{w}) , let $\mathcal{P}_{\vec{v}, \vec{w}}$ be the oriented parallelogram with sides \vec{v} and \vec{w} . Let α be the 2-covector so that $\alpha(\vec{v}, \vec{w})$ is the oriented area of $\mathcal{P}_{\vec{v}, \vec{w}}$ projected onto the xy -plane.

XXX Figure

Let $\vec{v} = (v_x, v_y, v_z)$ and $\vec{w} = (w_x, w_y, w_z)$ be the components of \vec{v} and \vec{w} . Using components, the projection of each vector onto the xy -plane is given by $(v_x, v_y, 0)$ and $(w_x, w_y, 0)$. We can compute the oriented area of this parallelogram in the xy -plane using the cross product, which results in the formula

$$\alpha(\vec{v}, \vec{w}) = v_x w_y - v_y w_x.$$

It is left to the reader to verify that α satisfies the area laws and is multi-linear. ■

The 2-covector $\alpha(\vec{v}, \vec{w}) = v_x w_y - v_y w_x$ from Example 7.3 is interesting because it only relies on the x and y components of its input vectors. And, the covectors e^1 and e^2 measure the x and y components of their input! Thus, we could write

$$\alpha(\vec{v}, \vec{w}) = e^1(\vec{v})e^2(\vec{w}) - e^2(\vec{v})e^1(\vec{w}).$$

There's a beautiful symmetry to that formula—so beautiful that it is encapsulated in a concept called the *wedge product*.

Definition 7.3.3 — Wedge Product. Let α and β be covectors. The *wedge product* (also known as the *exterior product*) is notated $\alpha \wedge \beta$ and is defined by

$$\alpha \wedge \beta(\vec{v}, \vec{w}) = \alpha(\vec{v})\beta(\vec{w}) - \beta(\vec{v})\alpha(\vec{w}).$$

Theorem 7.3.1 If α and β be covectors then the wedge product $\alpha \wedge \beta$ is a 2-covector.

Proof. XXX Finish ■

We now see that our area-in-the- xy -plane 2-covector from before could be written as $e^1 \wedge e^2$. But, it gets better.

Theorem 7.3.2 Every 2-covector can be represented as a linear combination of vectors in the set $\{e^i \wedge e^j : i \neq j\}$.

Note that in general, it is not true that if α is a 2-covector then α can be written as the wedge product of two covectors⁸. We really do need to allow linear combinations to get all 2-covectors. For example, the 2-covector $e^1 \wedge e^2 + e^3 \wedge e^4$ in \mathbb{R}^4 cannot be written as the wedge product of two covectors.

Exploring the consequences of the definition of the wedge product, we see that the wedge product is anti-symmetric:

$$\alpha \wedge \beta = -\beta \wedge \alpha.$$

⁸ This *does* happen to be true for 2-covectors on \mathbb{R}^3 , but that's just a happy coincidence.

And, that it is multi-linear:

$$(m\alpha) \wedge \beta = m(\alpha \wedge \beta) = \alpha \wedge (m\beta)$$

$$(\alpha + \gamma) \wedge \beta = \alpha \wedge \beta + \gamma \wedge \beta$$

$$\alpha \wedge (\gamma + \beta) = \alpha \wedge \gamma + \alpha \wedge \beta.$$

Strangely, 2-covectors seem to have all the properties of their inputs.

Geometric Interpretation of 2-covectors

XXX Finish. Some pictures about projections and oriented areas, again.

2-forms

Now that we have 2-covectors, we can define 2-forms in a very similar way to how we defined 1-forms.

■ **Definition 7.3.4 — 2-form.** A 2-form on the space \mathbb{R}^n is a function $\alpha : \mathbb{R}^n \rightarrow \{\text{2-covectors for } \mathbb{R}^n\}$.

Just as the wedge product of two covectors produces a 2-covectors, the wedge product of two 1-forms produces a 2-form. 2-forms in \mathbb{R}^3 also have a standard basis.

■ **Definition 7.3.5 — Standard Basis for 2-forms.** The standard basis for 2-forms on \mathbb{R}^3 are the 2-forms $dx \wedge dy$, $dy \wedge dz$, and $dz \wedge dx$.

Notice the order that dx , dy , and dz appear in the standard basis for 2-forms on \mathbb{R}^3 . In particular, the third basis element is $dz \wedge dx$ and not $dx \wedge dz$. This has to do with having a consistent orientation. When we use the basis 2-forms to measure area, we'd like them to have the property that either $dx \wedge dy(\mathcal{P}), dy \wedge dz(\mathcal{P}), dz \wedge dx(\mathcal{P}) \geq 0$ or $dx \wedge dy(\mathcal{P}), dy \wedge dz(\mathcal{P}), dz \wedge dx(\mathcal{P}) \leq 0$, but we never want a mix of signs.

We call $dx \wedge dy$, $dy \wedge dz$, and $dz \wedge dx$ the standard basis of 2-forms for \mathbb{R}^3 because any 2-form in \mathbb{R}^3 can be expressed as a linear combination of $dx \wedge dy$, $dy \wedge dz$, and $dz \wedge dx$.

■ **Example 7.4** Let $\alpha = xdx + 3y^2dz$ and $\beta = -dx + 2dy - 3dz$. Express $\alpha \wedge \beta$ as a linear combination of $dx \wedge dy$, $dy \wedge dz$, and $dz \wedge dx$.

XXX Finish

■

Integrating 2-forms

1-forms are integrated over oriented curves. 2-forms are integrated over oriented surfaces, and the definition of the integral of a 2-form is nearly identical to that of a 1-form—cumbersome notation and all.

Recall, 2-forms are functions that input pairs of points and output 2-covectors. In turn, the 2-covectors input pairs of vectors and output scalars. If α is a 2-form, we will define

$$\alpha(\vec{x}_1, \vec{x}_2; \vec{v}, \vec{w}) \equiv \left[\alpha(\vec{x}_1, \vec{x}_2) \right](\vec{v}, \vec{w})$$

for points \vec{x}_1, \vec{x}_2 and vectors \vec{v}, \vec{w} .

Definition 7.3.6 — integral of a 2-form. Let S be an oriented surface. The integral of the 2-form α over S is defined to be

$$\int_S \alpha = \int_a^b \int_c^d \alpha(\vec{p}(t,s); \frac{\partial \vec{p}}{\partial x}(t,s), \frac{\partial \vec{p}}{\partial y}(t,s)) ds dt,$$

where $\vec{p} : [a, b] \times [c, d] \rightarrow S$ is an orientation-preserving parameterization of S .

The integral of a 2-form over a surface S doesn't depend on your choice of how you parameterize S .

■ **Example 7.5** Let S be the region in the xy -plane below the curve $y = 1$ and above the curve $y = x^2$ oriented upward. Compute $\int_S dx \wedge dy$.

First, we need to parameterize this region. Let $\vec{p} : [-1, 1] \times [0, 1] \rightarrow S$ by

$$\begin{bmatrix} a \\ b \end{bmatrix} \mapsto \begin{bmatrix} a \\ (1-b)a^2 + b \end{bmatrix}.$$

To convince ourselves that \vec{p} actually is a parameterization, let's graph some curves $\vec{p}(t, b)$ while b is held constant and $\vec{p}(a, s)$ where a is held constant.

XXX Figure of parameterization

Computing,

$$\frac{\partial \vec{p}}{\partial x}(a, b) = \begin{bmatrix} 1 \\ 2(1-b)a \end{bmatrix} \quad \frac{\partial \vec{p}}{\partial y}(a, b) = \begin{bmatrix} 0 \\ 1 - a^2 \end{bmatrix}.$$

From the definition of $dx \wedge dy$,

$$dx \wedge dy(\vec{x}, \vec{y}; \vec{u}, \vec{v}) = \mathbf{e}^1 \wedge \mathbf{e}^2(\vec{u}, \vec{v}) = \mathbf{e}^1(\vec{u})\mathbf{e}^2(\vec{v}) - \mathbf{e}^1(\vec{v})\mathbf{e}^2(\vec{u}),$$

and so

$$\begin{aligned} dx \wedge dy(\vec{p}(a, b); \frac{\partial \vec{p}}{\partial x}(a, b), \frac{\partial \vec{p}}{\partial y}(a, b)) &= (1)(1 - a^2) - (0)(2(1 - b)a^2) \\ &= 1 - a^2. \end{aligned}$$

Thus,

$$\int_S dx \wedge dy = \int_{-1}^1 \int_0^1 1 - t^2 ds dt = \frac{4}{3},$$

which is the area of S . ■

You should notice some similarities between integrating a 2-form and surface integrals. When computing surface integrals, once we had a parameterization of a surface, we exploited the canonical parameterization of the tangent plane at a point on the surface to find a volume form (surface area element). That's actually what happened in Example 7.5. Let $\vec{p}(a, b) = (a, (1 - b)a^2 + b)$ be the parameterization from Example 7.5. Then, the canonical parameterization of the tangent plane to S at the point (a, b) is

$$\vec{L}_{\vec{p}}(t, s) = t \frac{\partial \vec{p}}{\partial x}(a, b) + s \frac{\partial \vec{p}}{\partial y}(a, b) + \vec{p}(a, b).$$

If we compute the area of $\vec{L}_{\vec{p}}([a, a + \Delta a] \times [b, b + \Delta b])$, we find

$$\text{area of } \vec{L}_{\vec{p}}([a, a + \Delta a] \times [b, b + \Delta b]) = (1 - a^2)\Delta a\Delta b,$$

and so the volume form associated with this parameterization would be

$$dV(t, s) = (1 - t^2)dt ds.$$

This is not a coincidence. The word *form* in the term “volume form” is the same as in the term “2-form.”

The techniques used in Example 7.5 will work in any dimension to integrate a 2-form over any surface, whether situated in the xy -plane or not.

Let \mathcal{R} be a surface in the xy -plane. The integral $\int_{\mathcal{R}} dx \wedge dy$ gives the signed area of \mathcal{R} . We can compute this integral using the techniques from Example 7.5, but in general, it's very hard to come up with parameterizations for strange surfaces. If all we wanted to do was find the signed area of \mathcal{R} , we could just do a double integral with clever bounds of integration. It turns out that this is the same integrating with respect to identity parameterization on a \mathcal{R} -shaped domain instead of a rectangularly-shaped one. You can integrate 2-forms in this way too.

■ **Example 7.6** Let S be the region below the curve $y = 1$ and above the curve $y = x^2$. Compute $\int_S dx \wedge dy$.

Let $\vec{p} : S \rightarrow S$ be the identity parameterization given by $\vec{p}(a, b) = (a, b)$. Now,

$$\frac{\partial \vec{p}}{\partial x}(a, b) = \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \hat{x} \quad \frac{\partial \vec{p}}{\partial y}(a, b) = \begin{bmatrix} 0 \\ 1 \end{bmatrix} = \hat{y}.$$

A little computation shows that $dx \wedge dy(\vec{p}(a, b); \frac{\partial \vec{p}}{\partial x}(a, b), \frac{\partial \vec{p}}{\partial y}(a, b)) = 1$. Since we can explicitly define S as

$$S = \{(x, y) : -1 \leq x \leq 1 \text{ and } x^2 \leq y \leq 1\},$$

we may write

$$\int_S dx \wedge dy = \int_{-1}^1 \int_{x^2}^1 1 \, ds \, dt = \frac{4}{3}.$$

■

Non-constant 2-forms

Just like integrating dx over a curve \mathcal{C} gave the change in x -coordinate from the start to the end of \mathcal{C} , Integrating the 2-form $dx \wedge dy$ over the surface S will give the signed area of S projected onto the xy -plane. Further, the integral of a non-constant 1-form α over a curve \mathcal{C} was equivalent to a vector line integral. Analogously, the integral of a non-constant 2-form β over a surface S turns out to be a flux integral.

Theorem 7.3.3 Let $S \subset \mathbb{R}^3$ be an oriented surface in \mathbb{R}^3 and let $\vec{F} = (f_x, f_y, f_z)$ be a vector

field in \mathbb{R}^3 . Then

$$\text{Flux}_S \vec{F} = \int_S f_x dy \wedge dz + f_y dz \wedge dx + f_z dx \wedge dy.$$

The proof of Theorem 7.3.3 amounts to careful bookkeeping, but it reveals an interesting fact. There is a correspondence between $a dy \wedge dz + b dz \wedge dx + c dx \wedge dy$ and the vector (a, b, c) . That is,

$$\begin{array}{lll} dy \wedge dz & \text{acts like} & \hat{x} \\ dz \wedge dx & \text{acts like} & \hat{y} \\ dx \wedge dy & \text{acts like} & \hat{z} \end{array}$$

at least when thinking about flux integrals. We will explore this connection further when discussing the *Hodge dual* operator.

7.4 Higher-order Forms

We will start by defining n -covectors after a minor expansion of the laws of area. A 3-covector is an object that measures volume, a 4-covector measures volume in \mathbb{R}^4 , etc. and so an n -covector must input n vectors and return a scalar. Further, every pair of those vectors must satisfy the laws of area. We call this extension of the laws of area to collections of n vectors the *laws of volume*. The laws of volume together with multi-linearity provide a sufficient basis to define n -covectors.

Definition 7.4.1 — n -covector. Let V be a vector space. A function $\alpha : \underbrace{V \times \cdots \times V}_n \rightarrow \mathbb{R}$ is called an n -covector if it multi-linear and obeys the laws of volume.

With the wedge product, we can create n -covectors (and in turn n -forms). Let α and β be covectors and recall

$$\alpha \wedge \beta(\vec{u}, \vec{v}) = \alpha(\vec{u})\beta(\vec{v}) - \beta(\vec{u})\alpha(\vec{v}).$$

If γ is another covector, we will recursively define the 3-covector $\alpha \wedge \beta \wedge \gamma$ by

$$\alpha \wedge \beta \wedge \gamma(\vec{u}, \vec{v}, \vec{w}) = \alpha(\vec{u})(\beta \wedge \gamma(\vec{v}, \vec{w})) - \beta(\vec{u})(\alpha \wedge \gamma(\vec{v}, \vec{w})) + \gamma(\vec{u})(\alpha \wedge \beta(\vec{v}, \vec{w})).$$

More succinctly, we might write $\alpha \wedge \beta \wedge \gamma = \alpha \cdot \beta \wedge \gamma - \beta \cdot \alpha \wedge \gamma + \gamma \cdot \alpha \wedge \beta$ where “ \cdot ” means multiplication, not a dot product⁹. Generalizing this, we can define the wedge product of any number of covectors.

Definition 7.4.2 — Wedge Product. Let $\alpha^1, \dots, \alpha^n$ be covectors. Then the n -covector $\alpha^1 \wedge$

⁹ You might recognize this formula as related to the cofactors expansion of the determinant of a matrix.

$\cdots \wedge \alpha^n$ is defined by the recursive formula

$$\alpha^1 \wedge \cdots \wedge \alpha^n = \sum_{i=1}^n (-1)^{i+1} \alpha^i \cdot (\alpha^1 \wedge \cdots \wedge \alpha^{i-1} \wedge \alpha^{i+1} \wedge \cdots \wedge \alpha^n).$$

This definition merely says: to find $\alpha^1 \wedge \cdots \wedge \alpha^n$, take each α^i out, one by one; multiply α^i by the wedge product of the remaining terms in order; sum the result together with a $+$ if i is odd and a $-$ if i is even.

It can be cumbersome to evaluate an n -covector directly from the definition. Fortunately the theories of linear algebra provide efficient ways to evaluate an n -fold wedge product. We won't go through them here, but we will explore some consequences of the definition.

First off, one must verify that $\alpha^1 \wedge \cdots \wedge \alpha^n$ follows the laws of volume. A direct consequence is that the wedge product of any number of covectors is *alternating*. That is,

$$\alpha^1 \wedge \cdots \wedge \alpha^i \wedge \alpha^{i+1} \wedge \cdots \wedge \alpha^n = -\alpha^1 \wedge \cdots \wedge \alpha^{i+1} \wedge \alpha^i \wedge \cdots \wedge \alpha^n.$$

In other words, switching adjacent pairs in a wedge product negates the product—this generalizes anti-symmetry for 2-covectors. By using the alternating property of the wedge product, we can rearrange large wedge products. We can now show if the expression $\beta \wedge \beta$ ever occurs in a wedge product, the wedge product must be zero¹⁰, and, by extension, if two β 's ever occur in a wedge product, the product must be zero.

n -forms

We define n -forms just like we defined 1-forms and 2-forms. Further, we define integration of n -forms.

Definition 7.4.3 — n -form. An n -form on the space \mathbb{R}^n is a function $\alpha : \mathbb{R}^n \rightarrow \{n\text{-covectors for } \mathbb{R}^n\}$.

Definition 7.4.4 — Integration of n -forms. Let \mathcal{R} be an n -dimensional oriented region and let α be an n -form. We define

$$\int_{\mathcal{R}} \alpha = \int_{a_1}^{b_1} \cdots \int_{a_n}^{b_n} \alpha(\vec{p}(\vec{t}); \frac{\partial \vec{p}}{\partial x_1}(\vec{t}), \dots, \frac{\partial \vec{p}}{\partial x_n}(\vec{t})) dt_n \cdots dt_1$$

where $\vec{p} : [a_1, b_1] \times \cdots \times [a_n, b_n] \rightarrow \mathcal{R}$ is an orientation-preserving parameterization of \mathcal{R} and $\vec{t} = (t_1, \dots, t_n)$.

3-forms

Of particular interest to us are 3-forms. 3-forms measure oriented volume, and in three dimensions, there's only one "direction" a volume can point. This is reflected in the fact that the standard basis for 3-forms consists of one element.

¹⁰ Move $\beta \wedge \beta$ to the back of the wedge product. Eventually Eventually you'll end up with $\sum a_i \beta \wedge \beta$ and since $\beta \wedge \beta = 0$, the whole sum is zero.

Definition 7.4.5 — Standard Basis for 3-forms. The standard basis for 3-forms on \mathbb{R}^3 is $dx \wedge dy \wedge dz$.

Because every 3-form on \mathbb{R}^3 is a multiple of $dx \wedge dy \wedge dz$, integrals of 3-forms in \mathbb{R}^3 are particularly easy. In fact, if $\mathcal{R} \subseteq \mathbb{R}^3$ is an oriented region,

$$\int_{\mathcal{R}} dx \wedge dy \wedge dz$$

is just the oriented volume of \mathcal{R} .

7.5 The Exterior Derivative

In calculus, we use the symbol “d” to notate both the variable of integration as well as differentiation. This notation helps us remember the relationship between integrals and derivatives—if we integrate a derivative we get back to (basically) the same thing. So far, our use of “d” seems to coincide with traditional integral notation, but it also means something in terms of derivatives. Namely, “d” is the symbol for the *exterior derivative*.

Before we define the exterior derivative, we need one bit of notation.

Definition 7.5.1 For a point $\vec{x} = (x_1, \dots, x_n) \in \mathbb{R}^n$, vectors $\vec{v}_1, \dots, \vec{v}_n \in \mathbb{R}^n$ and scalars $t_1, \dots, t_n \in \mathbb{R}$, define $B_{(\vec{x}; \vec{v}_1, \dots, \vec{v}_n)}(t_1, \dots, t_n)$ to be the oriented parallelotope with based at \vec{x} with sides $t_1 \vec{v}_1, t_2 \vec{v}_2, \dots, t_n \vec{v}_n$.

So as not to get bogged down with the intricacies of orientation, let’s go through some examples. If $\vec{x} \in \mathbb{R}^1$, then $B_{(\vec{x}; \vec{x})}(t)$ is a line segment of length t oriented left to right if $t > 0$ and right to left if $t < 0$. If $\vec{x} \in \mathbb{R}^2$, then $B_{(\vec{x}; \vec{x}, \vec{y})}(a, b)$ is a rectangle with side lengths a and b oriented upwards (i.e., in a right-handed way) if $a, b > 0$. Similarly, if $\vec{x} \in \mathbb{R}^3$, then $B_{(\vec{x}; \vec{x}, \vec{y}, \vec{z})}(a, b, c)$ is a box with side lengths a, b , and c oriented in a right-handed way (assuming $a, b, c > 0$), whereas $B_{(\vec{x}; \vec{y}, \vec{x}, \vec{z})}(a, b, c)$ is a box with side lengths a, b, c oriented in a left-handed way (assuming $a, b, c > 0$).

B should be a familiar object and can be used to tile n -dimensional space. For example, if we use the shorthand $B_{\vec{x}} = B_{(\vec{x}; \vec{x}, \vec{y})}(1, 1)$, then $B_{(0,0)}$, $B_{(1,0)}$, $B_{(0,1)}$, and $B_{(1,1)}$ perfectly fills out a 2×2 square in \mathbb{R}^2 .

XXX Figure

The fact that B is oriented is also very important. If B is oriented, then ∂B , the boundary of B , is also oriented. If we consider the partition of $B_{(\vec{0}; \vec{x}, \vec{y})}(2, 2) = [0, 2] \times [0, 2]$ by $B_{(0,0)}$, $B_{(1,0)}$, $B_{(0,1)}$, and $B_{(1,1)}$, we see that

$$B_{(\vec{0}; \vec{x}, \vec{y})}(2, 2) = B_{(0,0)} \cup B_{(1,0)} \cup B_{(0,1)} \cup B_{(1,1)}$$

and that

$$\partial B_{(\vec{0}; \vec{x}, \vec{y})}(2, 2) = \partial B_{(0,0)} \cup \partial B_{(1,0)} \cup \partial B_{(0,1)} \cup \partial B_{(1,1)},$$

provided we interpret the union of two overlapping and oppositely-oriented sets as empty.

XXX Figure showing boundary cancellation

This picture should be just like the one we used when deriving the divergence theorem and Stokes's theorem.

Now, we can define the *exterior derivative*.

Definition 7.5.2 — Exterior Derivative. The *exterior derivative*, notated d , is a map that takes n -forms to $(n + 1)$ -forms. If α is an n -form, then

$$d\alpha(\vec{x}; \vec{v}_1, \dots, \vec{v}_{n+1}) = \lim_{t_1, \dots, t_{n+1} \rightarrow 0} \frac{1}{t_1 t_2 \cdots t_{n+1}} \int_{\partial B} \alpha$$

where $B = B_{(\vec{x}; \vec{v}_1, \dots, \vec{v}_{n+1})}(t_1, \dots, t_{n+1})$.

A quick sanity check shows that since $B = B_{(\vec{x}; \vec{v}_1, \dots, \vec{v}_{n+1})}(t_1, \dots, t_{n+1})$ is an $n+1$ dimensional object, ∂B is an n -dimensional object over which an n -form α could be integrated. One must still prove that the result of $d\alpha$ is an $(n + 1)$ -form in that $d\alpha(\vec{x})$ outputs an honest $(n + 1)$ -covector (multi-linear and obeying the laws of volume). Showing this relies on smoothness of α as an n -form, which is beyond the scope of this book¹¹.

The Exterior Derivative on 0-forms

We haven't yet defined 0-forms, but we've already been using them. Since an n -covector inputs n vectors and outputs a scalar, by extension, a 0-covector must input zero vectors and return a scalar. A 0-form must input a point and output a 0-covector, which can be represented by a scalar. Therefore, we can think of scalar fields as 0-forms.

Now, 0-forms are to be integrated over oriented zero dimensional objects. These are just points with a positive or negative orientation. There's really not much choice when defining the integral over a point. Thus, if α is a 0-form,

$$\int_{\vec{x}} \alpha = \alpha(\vec{x})$$

if \vec{x} is oriented positively and

$$\int_{\vec{x}} \alpha = -\alpha(\vec{x})$$

if \vec{x} is oriented negatively¹². Further, integrals should be additive over disjoint domains, so

$$\int_{\{\vec{x}_1, \dots, \vec{x}_n\}} \alpha = \alpha(\vec{x}_1) + \cdots + \alpha(\vec{x}_n).$$

¹¹ Though thinking in high-dimensions and covectors can be hard, the idea is familiar—at every point there needs to exist a co-tangent plane!

¹² I hear you complaining—something doesn't smell right! $\alpha(\vec{x})$ should be a 0-covector which you can identify with a scalar but isn't truly a scalar. Further, we never had to add “−” signs to the *outside* of the integral for other n -forms. What gives?

We'll, it turns out, a 0-covector should really be a function that inputs no vectors *and* orientation, and returns a scalar. This is true for n -covectors in general. However, for $n > 0$, we could always deduce orientation from how we ordered the input-vectors. No input vectors means we have to specify the orientation manually. If it makes you feel better, write $\alpha(\vec{x}; +)$ and $\alpha(\vec{x}; -)$ when using 0-forms.

Finally, we're ready for action. Let α be a 0-form, and let's compute $\mathsf{d}\alpha$. By definition

$$\mathsf{d}\alpha(\vec{x}; \vec{v}) = \lim_{t \rightarrow 0} \frac{1}{t} \int_{\partial B} \alpha$$

where $B = B_{(\vec{x}; \vec{v})}(t)$. But, B is just a line segment and so ∂B consist of two points: $\vec{x} + t\vec{v}$ oriented positively, and \vec{x} oriented negatively.

XXX Figure

Thus,

$$\mathsf{d}\alpha(\vec{x}; \vec{v}) = \lim_{t \rightarrow 0} \frac{\alpha(\vec{x} + t\vec{v}) - \alpha(\vec{x})}{t},$$

which is just the definition of the directional derivative of α at \vec{x} in the direction \vec{v} . Further, $\mathsf{d}\alpha$ is a 1-form, which means to each point it assigns a covector. We know that if β is a covector, $\beta(\vec{v}) = \vec{\beta} \cdot \vec{v}$ for some vector $\vec{\beta}$. But, if we recall, the gradient of a function was exactly defined to be the vector that when dot producted with a direction gave the directional derivative. Thus,

$$\mathsf{d}\alpha(\vec{x}; \vec{v}) = \nabla \alpha(\vec{x}) \cdot \vec{v},$$

and we'd be justified in writing " $\mathsf{d}\alpha(\vec{x}) \approx \nabla \alpha(\vec{x})$."

The Exterior Derivative on 1-forms

We're on a roll with 0-forms. The exterior derivative replaces the need for the gradient operator. What does it do for 1-forms? Well, let α be a 1-form. Then, $\mathsf{d}\alpha$ inputs a point and outputs a 2-covector and is defined by

$$\mathsf{d}\alpha(\vec{x}; \vec{u}, \vec{v}) = \lim_{\Delta a, \Delta b \rightarrow 0} \frac{1}{\Delta a \Delta b} \int_{\partial B} \alpha$$

where $B = B_{(\vec{x}; \vec{u}, \vec{v})}(\Delta a, \Delta b)$. We're using Δa and Δb on purpose because ∂B is the oriented perimeter of a parallelogram.

XXX Figure

This limit looks exactly like the limit used to compute curl. In fact, if we refer back to the definition of the \vec{w} -component of the curl of a vector field, the definitions align flawlessly. Let $\vec{F} = (f_x, f_y, f_z)$ and consider the corresponding 1-form $\alpha = f_x \mathsf{d}x + f_y \mathsf{d}y + f_z \mathsf{d}z$. Then,

$$\begin{aligned} \text{curl}_{\hat{x}} \vec{F}(\vec{x}) &= \mathsf{d}\alpha(\vec{x}; \hat{y}, \hat{z}) \\ \text{curl}_{\hat{y}} \vec{F}(\vec{x}) &= \mathsf{d}\alpha(\vec{x}; \hat{z}, \hat{x}) \\ \text{curl}_{\hat{z}} \vec{F}(\vec{x}) &= \mathsf{d}\alpha(\vec{x}; \hat{x}, \hat{y}). \end{aligned}$$

Note the order we had to list $\hat{x}, \hat{y}, \hat{z}$ in to ensure $\mathsf{d}\alpha$ matched the curl. We've seen this order before—in the standard basis for 2-forms on \mathbb{R}^3 . Since we know we can find the curl of a vector field on \mathbb{R}^2 in any direction using just the \hat{x}, \hat{y} , and \hat{z} components, we'd be justified in writing " $\mathsf{d}\alpha(\vec{x}) \approx \nabla \times \alpha(\vec{x})$ " when α is a 1-form on three-dimensional space¹³.

¹³ Can you begin to see why curl could not be a "vector" in \mathbb{R}^4 ? The standard basis for 2-forms in \mathbb{R}^4 has six elements. Therefore, you would need to know the curl in six directions to describe the component of curl in an arbitrary direction.

The Exterior Derivative on 2-forms

Can this luck streak keep going? Let α is a 2-form. We know

$$\mathbb{d}\alpha(\vec{x}; \vec{u}, \vec{v}, \vec{w}) = \lim_{\Delta a, \Delta b, \Delta c \rightarrow 0} \frac{1}{\Delta a \Delta b \Delta c} \int_{\partial B} \alpha$$

where $B = B(\vec{x}; \vec{u}, \vec{v}, \vec{w})(\Delta a, \Delta b, \Delta c)$. Sure enough, this is exactly the limit of the flux through a little box as the size of the box goes to zero. After some careful analysis, if α is a 2-form on \mathbb{R}^3 , we might write “ $\mathbb{d}\alpha(\vec{x}) \approx \nabla \cdot \alpha(\vec{x})$,” and so the exterior derivative captures divergence, curl, and gradient in a single definition.

Practical Computing with \mathbb{d}

We’ve seen that \mathbb{d} can replace gradient, curl, and divergence in \mathbb{R}^3 and without changing any definitions, we can also bring these concepts to \mathbb{R}^n . But, using limits to compute derivatives can be cumbersome—especially when we know derivative rules from single-variable calculus.

The exterior derivative comes with its own set of rules.

Definition 7.5.3 — Exterior Derivative Rules. The exterior derivative, \mathbb{d} , obeys the following rules.

1. If $f(a, b, c, \dots)$ is a 0-form, then

$$\mathbb{d}f = \frac{\partial f}{\partial a} \mathbb{d}a + \frac{\partial f}{\partial b} \mathbb{d}b + \frac{\partial f}{\partial c} \mathbb{d}c + \dots$$

2. If α is an n -form, then

$$\mathbb{d}(\mathbb{d}\alpha) = 0.$$

3. If α is an n -form and β is an m -form, then \mathbb{d} obeys the following product rule:

$$\mathbb{d}(\alpha \wedge \beta) = (\mathbb{d}\alpha) \wedge \beta + (-1)^n \alpha \wedge (\mathbb{d}\beta).$$

The exterior derivative rules could actually be taken as axioms—the exterior derivative is the *only* non-trivial operator that satisfies all the rules¹⁴. With the exterior derivative rules in hand, we’re ready for action.

■ **Example 7.7** Compute $\mathbb{d}f$ where $f(x, y, z) = xy - z$ is a 0-form.

XXX Finish ■

■ **Example 7.8** Compute $\mathbb{d}f$ where $f(x, y, z) = -y\mathbb{d}x + x\mathbb{d}y + \mathbb{d}z$ is a 1-form.

XXX Finish ■

■ **Example 7.9** Compute $\mathbb{d}g$ where $g(x, y, z) = 2yx\mathbb{d}y \wedge \mathbb{d}z + 4\mathbb{d}z \wedge \mathbb{d}x + z^2\mathbb{d}x \wedge \mathbb{d}y$ is a 2-form.

XXX Finish ■

¹⁴ Many textbooks forgo limits and define the exterior derivative axiomatically. This makes the exterior derivative simpler in some respects, but it also makes it harder to see how it connects to physical concepts like work and curl.

Changing Coordinates

Differential forms make doing integrals in other coordinate systems almost trivial—it computes the volume form automatically.

Polar coordinates are related to rectangular coordinates by the equations

$$x = r \cos \theta \quad \text{and} \quad y = r \sin \theta.$$

Suppose $S \subset \mathbb{R}^2$ is a region in the xy -plane. If we wish to find the area of S , we need only to compute

$$\int_S dx \wedge dy.$$

However, using the exterior derivative, we see

$$\begin{aligned} dx &= d(r \cos \theta) = \cos \theta dr - r \sin \theta d\theta \\ dy &= d(r \sin \theta) = \sin \theta dr + r \cos \theta d\theta. \end{aligned}$$

Therefore,

$$\begin{aligned} dx \wedge dy &= (\cos \theta dr - r \sin \theta d\theta) \wedge (\sin \theta dr + r \cos \theta d\theta) \\ &= \cos \theta \sin \theta dr \wedge dr + r \cos^2 \theta dr \wedge d\theta - r \sin^2 \theta d\theta \wedge dr + r^2 \cos \theta \sin \theta d\theta \wedge d\theta \\ &= r \cos^2 \theta dr \wedge d\theta - r \sin^2 \theta d\theta \wedge dr \\ &= r \cos^2 \theta dr \wedge d\theta + r \sin^2 \theta dr \wedge d\theta \\ &= r dr \wedge d\theta, \end{aligned}$$

because $dr \wedge dr$ and $d\theta \wedge d\theta$ are 0 and $dr \wedge d\theta = -d\theta \wedge dr$. We conclude that

$$\int_S dx \wedge dy = \int_S r dr \wedge d\theta,$$

and have just used differential forms to find the volume form for polar coordinates¹⁵. Further, if we instead wanted to integrate a function $f : S \rightarrow \mathbb{R}$ over S , we could do the exact same thing since

$$\int_S f dx \wedge dy = \int_S f r dr \wedge d\theta.$$

As long as we can determine what the value of f is at a point in polar coordinates, we can compute the integral with ease.

■ **Example 7.10** XXX Polar coordinates example. ■

¹⁵ Notice that the terms that got thrown away because $dr \wedge dr = d\theta \wedge d\theta = 0$ were very similar to the terms that got thrown away from the pre-volume form when we used a limit process.

7.6 The Hodge Dual

There is a certain symmetry to n -covectors and n -forms. We haven't delved deeply into the linear algebra reasoning behind it, but (at least for \mathbb{R}^3) the standard basis for n -covectors and n -forms always consisted of a linear combinations of the wedge product of n covectors/1-forms. This is true in general.

Theorem 7.6.1 Let $B^1 = \{e^1, \dots, e^n\}$ be the standard basis of covectors for \mathbb{R}^n and let

$$B^k = \{\alpha : \alpha \text{ is the wedge product of } k \text{ items from } B_1\}.$$

Then, every k -covector can be expressed as a linear combination of k -covectors from B^k .

Theorem 7.6.1 applies equally well to k -forms and has an immediate consequence.

For the time being, assume we are always taking about k -covectors on \mathbb{R}^n . Since $\alpha \wedge \alpha = 0$ if α is a covector¹⁶, every non-zero element in B^k must be a wedge product of k different covectors from B^1 . Further, if $\alpha, \beta \in B^k$ are wedge products of the same k covectors from B^1 , by applying the alternating property, we know either $\alpha = \beta$ or $\alpha = -\beta$, and so β is a linear combination of α . This means, to create all k -covectors as linear combinations, all we really need a set consisting of one wedge product for each distinct set k covectors from B^1 .

Given n items, $\binom{n}{k} = \frac{n!}{k!(n-k)!}$ is the number of ways to choose k things from a set of n things with no repeats provided $0 \leq k \leq n$. The number $\binom{n}{k}$ is called a *binomial coefficient*¹⁷. If $k < 0$ or $k > n$, there are no ways to pick k items from a set of size n without repeats. We will extend notation to say that $\binom{n}{k} = 0$ if $k < 0$ or $k > n$.

This alone is significant. Our argument now says that if $k > n$, every k -covector can be represented as a linear combination of k -covectors from a set of size $\binom{n}{k} = 0$. In other words, every k -covector for $k > n$ is trivial.

Definition 7.6.1 — Basis & Dimension. Let V be a vector space. A set $B \subseteq V$ is called a *basis* for V if (i) every element of V can be represented as a linear combination of elements in B and (ii) if any element is removed from B , property (i) fails.

The *dimension* of V , notated $\dim(V)$, is the size of a basis for V .

You can think of a basis for V as a minimal set of vectors for which every element in V can be written as a linear combination. Every element in a basis is essential, since removing any element means you can no longer write everything in V as a linear combination.

Our argument about the number of essential elements in B^k along with the language of dimension allows us to state the following theorem.

Theorem 7.6.2 Let $\Lambda^k \mathbb{R}^n$ be the set of all k -covectors on \mathbb{R}^n . The dimension of $\Lambda^k \mathbb{R}^n$ is $\binom{n}{k}$.

Again, Theorem 7.6.2 applies equally well to k -forms.

Let's be concrete for a moment and consider k -forms on \mathbb{R}^3 . Theorem 7.6.2 says that the dimension of the set 0-forms should be 1, the dimension of the set of 1-forms should be 3, the dimension of the set of 2-forms should be 3, the dimension of the set of 3-forms should be 1,

¹⁶ Remember, this may not be true for n -covectors with $n > 1$, and certainly isn't true for 0-covectors.

¹⁷ Note that since $0! = 1$, we have $\binom{n}{0} = 1$, like it should. There's only one way to choose nothing.

and the dimension of the set of all higher forms should be 0. This is reflected in the standard bases we've already introduced.

	Standard Basis	Dimension
0-forms	the function $f(x, y, z) = 1$	1
1-forms	dx, dy, dz	3
2-forms	$dx \wedge dy, dy \wedge dz, dz \wedge dx$	3
3-forms	$dx \wedge dy \wedge dz$	1

Now we can directly see the symmetry. The dimension of the set of k -forms is $\binom{n}{k}$ which is equal to the dimension of the set of $(n - k)$ -forms¹⁸. In particular, the dimension of the set of 0-covectors is 1 and the dimension of the set of n -covectors is 1.

The Hodge Star Operator

A k -covector (or k -form) on \mathbb{R}^n is a linear combination of the wedge product of k covectors. An n -covector (or n -form) is a linear combination of the wedge product of n covectors. Since we have a basis of n covectors, every n -covector is a multiple of this element. Namely, every n -covector on \mathbb{R}^n is a multiple of $e^1 \wedge \cdots \wedge e^n$ and every k -covector could be thought as “missing” some items in its wedge product, to top it up to be an n -covector. The *Hodge star operator* sends a k -covector to its missing half.

Definition 7.6.2 — Hodge Star Operator. Let $\Lambda^k \mathbb{R}^n$ be the set of k -covectors on \mathbb{R}^n . The *Hodge star operator* (or *Hodge dual operator*), denoted “ \star ” is defined on $\Lambda^k \mathbb{R}^n$ as follows. Let $\alpha \in \Lambda^k \mathbb{R}^n$ be represented as $\alpha = \sum t_i b^i$ where b^i are standard basis elements for $\Lambda^k \mathbb{R}^n$. Then $\star \alpha$ is defined to be the $(n - k)$ -covector so that

$$\alpha \wedge \star \alpha = \left(\sum t_i^2 \right) \omega$$

where $\omega = e^1 \wedge \cdots \wedge e^n$. We call $\star \alpha$ the *Hodge dual* of α .

The Hodge dual of a k -covector α is a kind of compliment that fills it out to be a full n -covector. But why the $\sum t_i^2$ terms? We'll, that should make you think of norms!

Recall that if a vector $\vec{v} = \sum t_i \vec{e}_i$ where \vec{e}_i are the standard basis vectors for \mathbb{R}^n , then

$$\vec{v} \cdot \vec{v} = \|\vec{v}\|^2 = \sum t_i^2.$$

In the same way, if α is a k -covector, we can think of $\alpha \wedge \star \alpha$ as $\|\alpha\|^2$. If we spent some time to make this rigorous, we would see the same $\sum t_i^2$ for $\|\alpha\|^2$ as we do for $\|\vec{v}\|^2$.

We will now start working with differential forms.

■ **Example 7.11** Compute $\star dx$ and $\star(dy \wedge dz)$ on \mathbb{R}^3 .

Since dx is a 1-form on \mathbb{R}^3 , we know $\star dx$ must be a 2-form. Further, $dx = 1dx$ is already written in terms of the standard basis. Therefore, $\star dx$ satisfies

$$dx \wedge \star dx = dx \wedge dy \wedge dz,$$

¹⁸ $\binom{n}{k} = \frac{n!}{k!(n-k)!} = \frac{n!}{(n-k)!k!} = \binom{n}{n-k}$

and so $\star dx = dy \wedge dz$.

Similarly, $\star(dy \wedge dz)$ is a 1-form so that

$$(dy \wedge dz) \wedge \star(dy \wedge dz) = dx \wedge dy \wedge dz.$$

We conclude that $\star(dy \wedge dz) = dx$ since

$$dy \wedge dz \wedge dx = -dy \wedge dx \wedge dz = dx \wedge dy \wedge dz.$$

■

As in the example, it turns out applying “ \star ” twice always brings you back to the same form.

Theorem 7.6.3 If α is a k -form, then $\star\star\alpha = \alpha$.

This is the reason we call “ \star ” a *dual operator*¹⁹.

Since \mathbb{R}^2 and \mathbb{R}^3 are so important, we will make a table of basis forms and their Hodge duals.

Form	Hodge Dual (\mathbb{R}^2)
1	$dx \wedge dy$
dx	dy
dy	$-dx$

Form	Hodge Dual (\mathbb{R}^3)
1	$dx \wedge dy \wedge dz$
dx	$dy \wedge dz$
dy	$dz \wedge dx$
dz	$dx \wedge dy$

Exercise 7.5 Verify that corresponding elements in the table are indeed Hodge duals.

The Hodge Dual and Vector Fields

With the Hodge dual in hand, we can start nailing down some relationships.

Notice that 1-forms directly correspond to vector fields since covectors directly correspond to vectors. We make this correspondence explicit with an operator.

Definition 7.6.3 — Transpose. If $\alpha = \sum t_i e^i$ is a covector, then the *transpose* of α , denoted $\downarrow\alpha$, is the vector

$$\downarrow\alpha = \sum t_i \vec{e}_i.$$

If $\vec{a} = \sum t_i \vec{e}_i$ is a vector, then the *transpose* of \vec{a} , denoted $\uparrow\vec{a}$, is the covector

$$\uparrow\vec{a} = \sum t_i e^i.$$

¹⁹ Covectors are also called dual vectors for a similar reason.

If you're familiar with the transpose operator from linear algebra, this is the very same one²⁰. You might think of it as transposing superscripts into subscripts (since we use superscripts for covectors and subscripts for vectors).

The transpose is a little like the Hodge dual since $\alpha(\downarrow\alpha) = \sum t_i^2 = \|\downarrow\alpha\|^2$ and $\uparrow\vec{\alpha}(\vec{\alpha}) = \sum t_i^2 = \|\vec{\alpha}\|^2$.

The transpose directly extends to 1-forms and vector fields. If α is a 1-form, then $\downarrow\alpha$ is a vector field and if \vec{F} is a vector field, $\uparrow\vec{F}$ is a 1-form. Further note that $\uparrow\downarrow\alpha = \alpha$ and $\downarrow\uparrow\vec{\alpha} = \vec{\alpha}$.

It's time to recast some familiar operations in terms of differential forms and the Hodge dual.

Operation	Notation	Equivalency
Cross Product	$\vec{u} \times \vec{v}$	$\downarrow\star(\uparrow\vec{u} \wedge \uparrow\vec{v})$
Gradient	∇f	$\downarrow d(\uparrow f)$
Curl	$\nabla \times \vec{F}$	$\downarrow\star d(\uparrow\vec{F})$
Divergence	$\nabla \cdot \vec{F}$	$\downarrow\star d(\star\uparrow\vec{F})$

Or, if we use bold face for corresponding covectors and forgo some pedanticness, we have a simplified table.

Operation	Notation	Equivalency
Cross Product	$\vec{u} \times \vec{v}$	$\star(\mathbf{u} \wedge \mathbf{v})$
Gradient	∇f	$d\mathbf{f}$
Curl	$\nabla \times \vec{F}$	$\star d\mathbf{F}$
Divergence	$\nabla \cdot \vec{F}$	$\star d\star\mathbf{F}$

Although using “ \star ” and “ d ” we can recover the exact formulations of multivariable calculus using vector fields, by now there is a strong argument for using differential forms from the beginning. Somehow, the fact that covectors and 2-covectors are duals of each other in \mathbb{R}^3 allows us to get away with calling curl a *vector* rather than a 2-covector. Further, the cross product of two vectors in \mathbb{R}^3 is another vector. We can take the wedge product in any dimensional space, but only in \mathbb{R}^3 is $\star(\mathbf{u} \wedge \mathbf{v})$ a covector. The strongest argument yet for formulating multivariable calculus in terms of differential forms comes from the generalized Stokes's theorem.

7.7 Stokes's Theorem

We've now reached the centerpiece of differential forms—the generalized Stokes's theorem. And, its statement couldn't be easier.

Theorem 7.7.1 — Generalized Stokes's Theorem. Let α be a k -form and let \mathcal{R} be an ori-

²⁰ If all true vectors are column vectors, then covectors must be row vectors and $\alpha(\vec{v})$ is just $\alpha\vec{v}$ where α is a row matrix and \vec{v} is a column matrix.

ented k -dimensional region. Then,

$$\int_{\mathcal{R}} \mathbb{d}\alpha = \int_{\partial\mathcal{R}} \alpha.$$

This theorem is equivalent to the fundamental theorem of calculus for 0-forms, Stokes's curl theorem for 1-forms, and the divergence theorem for 2-forms. It also generalizes to all dimensions, and the proof is easy! We've already seen the main ideas of the proof.

1. Partition \mathcal{R} by oriented polytopes. Then, shared faces have opposite orientations and so cancel. All that is left are contributions from the boundary.
2. The definition of $\mathbb{d}\alpha$ is exactly taking the limit of $\int_B \alpha$ over smaller and smaller polytopes. Since these are the same polytopes in our partition of \mathcal{R} , we get a $\mathbb{d}\alpha$ on the left and a $\partial\mathcal{R}$ on the right.

The only hard part of the proof is bookkeeping—making sure the orientation of adjacent sides really to cancel—and exchanging a limit and a sum—this requires some assumptions about the differentiability of α and the smoothness of $\partial\mathcal{R}$.

7.8 de Rham Cohomology

Recall that a vector field \vec{F} is called *conservative* if there is a scalar field f so that $\nabla f = \vec{F}$. Further, if \vec{F} is conservative, then $\nabla \times \vec{F} = 0$ —this was called the screening test. However, it may be that \vec{F} is not conservative but even still $\nabla \times \vec{F} = 0$. Poincaré's lemma states that if $\nabla \times \vec{F} = 0$ on the interior of a ball (equivalently, the interior of a box), then \vec{F} is conservative. In other words, if you want a non-conservative vector field to satisfy $\nabla \times \vec{F} = 0$, you have to make the domain strange. For example,

$$\vec{F}(x, y, z) = \frac{1}{x^2 + y^2} \begin{bmatrix} -y \\ x \\ 0 \end{bmatrix}$$

has zero curl, but the z -axis must be excluded from the domain. The key insight here is that the difference between a conservative and a closed vector field only shows up when the domain has certain topologies²¹. Using differential forms, we get a supped up topology-recognizing machine called *cohomology*.

Before we begin, we need some definitions.

Definition 7.8.1 — Closed & Exact. An n -form α is called *exact* if there is an $(n-1)$ -form β so that $\alpha = \mathbb{d}\beta$. An n -form α is called *closed* if $\mathbb{d}\alpha = 0$.

Exact is the differential forms term for a conservative vector field.

We can use Stokes's theorem to prove that every exact form is closed.

²¹ A topology describes how a set is connected. A plane and the surface of a doughnut are both two dimensional objects, but they have very different topologies.

Theorem 7.8.1 Every exact differential form is closed. That is, $\mathfrak{d}(\mathfrak{d}\alpha) = 0$ for all differential forms α .

Proof. The proof uses Stokes's theorem and topology. Let \mathcal{R} be an oriented region and notice that $\partial(\partial\mathcal{R})$ is empty—a boundary has no boundary²². From Stokes's theorem we know

$$\int_{\mathcal{R}} \mathfrak{d}\mathfrak{d}\alpha = \int_{\partial\mathcal{R}} \mathfrak{d}\alpha = \int_{\partial(\partial\mathcal{R})} \alpha = 0$$

since an integral over an empty set is zero. Any form that integrates to zero over all regions must be zero²³. ■

Definition 7.8.2 — Exterior Power. Let \mathcal{M} be a domain (more specifically a *manifold*). The set $\Lambda^k(\mathcal{M})$ is the set of all k -forms on \mathcal{M} and is called the k^{th} exterior power of \mathcal{M} .

Using our new notation, we can say things like $\mathfrak{d} : \Lambda^k(\mathcal{M}) \rightarrow \Lambda^{k+1}(\mathcal{M})$. Further, since $\mathfrak{d}\mathfrak{d}\alpha = 0$, we know exact forms are in the kernel of \mathfrak{d} . This allows us to write what (in the abstract algebra) is called an *exact sequence*. If \mathcal{M} is n -dimensional, then we have the following exact sequence.

$$\Lambda^0(\mathcal{M}) \xrightarrow{\mathfrak{d}} \Lambda^1(\mathcal{M}) \xrightarrow{\mathfrak{d}} \dots \xrightarrow{\mathfrak{d}} \Lambda^{n-1}(\mathcal{M}) \xrightarrow{\mathfrak{d}} \Lambda^n(\mathcal{M}) \xrightarrow{\mathfrak{d}} 0$$

This notation means that the map \mathfrak{d} moves objects from a space on the left to a space on the right *and* the range of \mathfrak{d} is always contained in the kernel of \mathfrak{d} . In other words, a chain is called *exact* if two steps down the chain guarantees you become zero.

If we take $\mathcal{M} = \mathbb{R}^3$ then we have the explicit exact chain

$$\Lambda^0(\mathbb{R}^3) \xrightarrow{\mathfrak{d}} \Lambda^1(\mathbb{R}^3) \xrightarrow{\mathfrak{d}} \Lambda^2(\mathbb{R}^3) \xrightarrow{\mathfrak{d}} \Lambda^3(\mathbb{R}^3) \xrightarrow{\mathfrak{d}} 0.$$

Further, by Poincaré's lemma, the range of \mathfrak{d} is exactly the kernel of the next \mathfrak{d} . This sets our baseline. Now, consider the space $H = \mathbb{R}^3 \setminus \{\alpha\vec{x} : \alpha \in \mathbb{R}\}$ consisting of \mathbb{R}^3 without the z -axis. We have a very similar chain.

$$\Lambda^0(H) \xrightarrow{\mathfrak{d}} \Lambda^1(H) \xrightarrow{\mathfrak{d}} \Lambda^2(H) \xrightarrow{\mathfrak{d}} \Lambda^3(H) \xrightarrow{\mathfrak{d}} 0$$

Except, this time the range of \mathfrak{d} applied to 0-forms does not equal the kernel of \mathfrak{d} applied to 1-forms. This is because there are now covector fields that circulate around the z -axis but still have “zero curl.” We won't prove it, but the image of 1-forms under \mathfrak{d} , in this case, is kernel of \mathfrak{d} applied to 2-forms. This somehow means that some interesting curves fit into H , but no interesting surfaces. The interesting curves are those that circle around the z -axis!

We could play this game again. Let $S = \mathbb{R}^3 \setminus \{\vec{0}\}$. Now we have a chain

$$\Lambda^0(S) \xrightarrow{\mathfrak{d}} \Lambda^1(S) \xrightarrow{\mathfrak{d}} \Lambda^2(S) \xrightarrow{\mathfrak{d}} \Lambda^3(S) \xrightarrow{\mathfrak{d}} 0.$$

²² To formally prove this, we would need to appeal to the definition of boundary. Since $\partial\mathcal{R} \subseteq \mathcal{R}$, we know any point $\vec{x} \in \partial\mathcal{R}$ is also in \mathcal{R} . But, the tangent space to \mathcal{R} is either n -dimension (at points on \mathcal{R} 's interior) or $n - 1$ -dimensional when restricted to $\partial\mathcal{R}$. There's just no room for an $n - 2$ -dimensional tangent space to fit, so $\vec{x} \notin \mathcal{R}$. We conclude $\partial\partial\mathcal{R}$ is empty.

²³ This conclusion requires $\mathfrak{d}\mathfrak{d}\alpha$ to be continuous, and so α needs to be twice differentiable.

In this chain the image of \mathfrak{d} applied to 0-forms is the kernel of \mathfrak{d} applied to 1-forms. However, the image of \mathfrak{d} applied to 1-forms is *not* the kernel of \mathfrak{d} applied to 2-forms. For example, consider

$$\vec{F}(x, y, z) = \frac{1}{(x^2 + y^2 + z^2)^{3/2}} \begin{bmatrix} x \\ y \\ z \end{bmatrix},$$

which you may recognize as related to the force due to gravity of a point-mass centered at the origin acting on a particle at position (x, y, z) . $\nabla \cdot \vec{F} = 0$, but \vec{F} cannot be the curl of any vector field. Thus, $\mathfrak{f} = \star \uparrow \vec{F}$ is a 2-form in the kernel of \mathfrak{d} but not the image of any 1-form under \mathfrak{d} . This would be very hard to show directly from equations, but easy to show with Stokes's theorem. After all, $\mathfrak{d}\mathfrak{f} = 0$ is an easy computation²⁴, but the integral of \mathfrak{f} around a sphere centered at the origin and oriented outwards is clearly positive—all vectors coming from \vec{F} point radially and symmetrically outwards!

Based on these observations, we conclude that S cannot contain any loops that are more interesting than those in \mathbb{R}^3 , but it can contain surfaces that are more interesting. What are the interesting surfaces? The ones that enclose the origin!

²⁴ If you've ever done the computation to show that gravity acting on a sphere is equivalent to gravity acting on a point-mass at the center of the sphere, then you've done this computation.

Appendix A

Proofs

Below are some guidelines to help you write proofs. The following rules apply whenever you write a proof¹.

1. **The burden of communication lies on you, not on your reader.** It is your job to explain your thoughts; it is not your reader's job to guess them from a few hints. You are trying to convince a skeptical reader who doesn't believe you, so you need to argue with airtight logic in crystal clear language; otherwise the reader will continue to doubt. If you didn't write something on the paper, then (a) you didn't communicate it, (b) the reader didn't learn it, and (c) the grader has to assume you didn't know it in the first place.
2. **Tell the reader what you're proving.** The reader doesn't necessarily know or remember what "Theorem 2.13" is. Even a professor grading a stack of papers might lose track from time to time. Therefore, the statement you are proving should be on the same page as the beginning of your proof. For an exam this won't be a problem, of course, but on your homework, recopy the claim you are proving. This has the additional advantage that when you study for exams by reviewing your homework, you won't have to flip back in the notes/textbook to know what you were proving.
3. **Use English words.** Although there will usually be equations or mathematical statements in your proofs, use English sentences to connect them and display their logical relationships. If you look in your notes/textbook, you'll see that each proof consists mostly of English words.
4. **Use complete sentences.** If you wrote a history essay in sentence fragments, the reader would not understand what you meant; likewise, in mathematics you must use complete sentences with verbs to convey your logical train of thought.

Some complete sentences can be written purely in mathematical symbols, such as equations (e.g., $a^3 = b^{-1}$), inequalities (e.g., $x < 5$), and other relations (like $5 \mid 10$ or $7 \in \mathbb{Z}$). These statements usually express a relationship between two mathematical *objects*, like

¹ This list is an adaptation of *The Elements of Style for Proofs* written by Anders Hendrickson of St. Norbert College and modified by Dana Ernst of Northern Arizona University.

numbers or sets. However, it is considered bad style to begin a sentence with symbols. A common phrase to use to avoid starting a sentence with mathematical symbols is “We see that...”

5. **Show the logical connections among your sentences.** Use phrases like “Therefore” or “because” or “if... , then...” or “if and only if” to connect your sentences.
6. **Know the difference between statements and objects.** A mathematical object is a *thing*, a noun, such as a group, an element, a vector space, a number, an ordered pair, etc. Objects either exist or don’t exist. Statements, on the other hand, are mathematical *sentences*: they can be true or false.

When you see or write a cluster of math symbols, be sure you know whether it’s an object (e.g., “ $x^2 + 3$ ”) or a statement (e.g., “ $x^2 + 3 < 7$ ”). One way to tell is that every mathematical statement includes a verb, such as $=$, \leq , “divides”, etc.
7. **“=” means equals.** Don’t write $A = B$ unless you mean that A actually equals B . This rule seems obvious, but there is a great temptation to be sloppy. In calculus, for example, some people might write $f(x) = x^2 = 2x$ (which is false), when they really mean that “if $f(x) = x^2$, then $f'(x) = 2x$.”
8. **Don’t interchange $=$ and \implies .** The equals sign connects two *objects*, as in “ $x^2 = b$ ”; the symbol “ \implies ” is an abbreviation for “implies” and connects two *statements*, as in “ $ab = a \implies b = 1$.” You should avoid using \implies in your formal write-ups.
9. **Say exactly what you mean.** Just as the $=$ is sometimes abused, so too people sometimes write $A \in B$ when they mean $A \subseteq B$, or write $a_{ij} \in A$ when they mean that a_{ij} is an entry in matrix A . Mathematics is a very precise language, and there is a way to say exactly what you mean; find it and use it.
10. **Don’t write anything unproven.** Every statement on your paper should be something you *know* to be true. The reader expects your proof to be a series of statements, each proven by the statements that came before it. If you ever need to write something you don’t yet know is true, you *must* preface it with words like “assume,” “suppose,” or “if” (if you are temporarily assuming it), or with words like “we need to show that” or “we claim that” (if it is your goal). Otherwise the reader will think they have missed part of your proof.
11. **Write strings of equalities (or inequalities) in the proper order.** When your reader sees something like

$$A = B \leq C = D,$$

he/she expects to understand easily why $A = B$, why $B \leq C$, and why $C = D$, and he/she expects the *point* of the entire line to be the more complicated fact that $A \leq D$. For example, if you were computing the distance d of the point $(12, 5)$ from the origin, you could write

$$d = \sqrt{12^2 + 5^2} = 13.$$

In this string of equalities, the first equals sign is true by the Pythagorean theorem, the second is just arithmetic, and the *point* is that the first item equals the last item: $d = 13$.

A common error is to write strings of equations in the wrong order. For example, if you were to write “ $\sqrt{12^2 + 5^2} = 13 = d$ ”, your reader would understand the first equals sign, would be baffled as to how we know $d = 13$, and would be utterly perplexed as to why you wanted or needed to go through 13 to prove that $\sqrt{12^2 + 5^2} = d$.

12. **Avoid circularity.** Be sure that no step in your proof makes use of the conclusion!
13. **Don’t write the proof backwards.** Beginning students often attempt to write “proofs” like the following, which attempts to prove that $\tan^2(x) = \sec^2(x) - 1$:

$$\begin{aligned}\tan^2(x) &= \sec^2(x) - 1 \\ \left(\frac{\sin(x)}{\cos(x)}\right)^2 &= \frac{1}{\cos^2(x)} - 1 \\ \frac{\sin^2(x)}{\cos^2(x)} &= \frac{1 - \cos^2(x)}{\cos^2(x)} \\ \sin^2(x) &= 1 - \cos^2(x) \\ \sin^2(x) + \cos^2(x) &= 1 \\ 1 &= 1\end{aligned}$$

Notice what has happened here: the student *started* with the conclusion, and deduced the true statement “ $1 = 1$.” In other words, he/she has proved “If $\tan^2(x) = \sec^2(x) - 1$, then $1 = 1$,” which is true but highly uninteresting.

Now this isn’t a bad way of *finding* a proof. Working backwards from your goal often is a good strategy *on your scratch paper*, but when it’s time to *write* your proof, you have to start with the hypotheses and work to the conclusion.

14. **Be concise.** Most students err by writing their proofs too short, so that the reader can’t understand their logic. It is nevertheless quite possible to be too wordy, and if you find yourself writing a full-page essay, it’s probably because you don’t have a proof, but just an intuition. When you find a way to turn that intuition into a formal proof, it will be much shorter.
15. **Introduce every symbol you use.** If you use the letter “ k ,” the reader should know exactly what k is. Good phrases for introducing symbols include “Let $n \in \mathbb{N}$,” “Let k be the least integer such that. . .,” “For every real number a . . .,” and “Suppose that X is a counterexample.”
16. **Use appropriate quantifiers (once).** When you introduce a variable $x \in S$, it must be clear to your reader whether you mean “for all $x \in S$ ” or just “for some $x \in S$.” If you just say something like “ $y = x^2$ where $x \in S$,” the word “where” doesn’t indicate whether you mean “for all” or “some.”

Phrases indicating the quantifier “for all” include “Let $x \in S$ ”; “for all $x \in S$ ”; “for every $x \in S$ ”; “for each $x \in S$ ”; etc. Phrases indicating the quantifier “some” (or “there exists”) include “for some $x \in S$ ”; “there exists an $x \in S$ ”; “for a suitable choice of $x \in S$ ”; etc.

On the other hand, don’t introduce a variable more than once! Once you have said “Let $x \in S$,” the letter x has its meaning defined. You don’t *need* to say “for all $x \in S$ ” again, and you definitely should *not* say “let $x \in S$ ” again.

17. **Use a symbol to mean only one thing.** Once you use the letter x once, its meaning is fixed for the duration of your proof. You cannot use x to mean anything else.
18. **Don’t “prove by example.”** Most problems ask you to prove that something is true “for all”—You *cannot* prove this by giving a single example, or even a hundred. Your answer will need to be a logical argument that holds for *every example there possibly could be*.
19. **Write “Let $x = \dots$,” not “Let $\dots = x$.”** When you have an existing expression, say a^2 , and you want to give it a new, simpler name like b , you should write “Let $b = a^2$,” which means, “Let the new symbol b mean a^2 .” This convention makes it clear to the reader that b is the brand-new symbol and a^2 is the old expression he/she already understands. If you were to write it backwards, saying “Let $a^2 = b$,” then your startled reader would ask, “What if $a^2 \neq b$?”
20. **Make your counterexamples concrete and specific.** Proofs need to be entirely general, but counterexamples should be concrete. When you provide an example or counterexample, make it as specific as possible. For a set, for example, you must name its elements, and for a function, you must give its rule. Do not say things like “ θ could be one-to-one but not onto”; instead, provide an actual function θ that *is* one-to-one but not onto.
21. **Don’t include examples in proofs.** Including an example very rarely adds anything to your proof. If your logic is sound, then it doesn’t need an example to back it up. If your logic is bad, a dozen examples won’t help it (see rule 18). There are only two valid reasons to include an example in a proof: if it is a *counterexample* disproving something, or if you are performing complicated manipulations in a general setting and the example is just to help the reader understand what you are saying.
22. **Use scratch paper.** Finding your proof will be a long, potentially messy process, full of false starts and dead ends. Do all that on scratch paper until you find a real proof, and only then break out your clean paper to write your final proof carefully. *Do not hand in your scratch work!*

Only sentences that actually contribute to your proof should be part of the proof. Do not just perform a “brain dump,” throwing everything you know onto the paper before showing the logical steps that prove the conclusion. *That is what scratch paper is for.*

Appendix B

Creative Commons License Fulltext

Creative Commons Legal Code

Attribution-ShareAlike 4.0 International

Official translations of this license are available in other languages.

Creative Commons Corporation (“Creative Commons”) is not a law firm and does not provide legal services or legal advice. Distribution of Creative Commons public licenses does not create a lawyer-client or other relationship. Creative Commons makes its licenses and related information available on an as-is basis. Creative Commons gives no warranties regarding its licenses, any material licensed under their terms and conditions, or any related information. Creative Commons disclaims all liability for damages resulting from their use to the fullest extent possible.

Using Creative Commons Public Licenses

Creative Commons public licenses provide a standard set of terms and conditions that creators and other rights holders may use to share original works of authorship and other material subject to copyright and certain other rights specified in the public license below. The following considerations are for informational purposes only, are not exhaustive, and do not form part of our licenses.

Considerations for licensors: Our public licenses are intended for use by those authorized to give the public permission to use material in ways otherwise restricted by copyright and certain other rights. Our licenses are irrevocable. Licensors should read and understand the terms and conditions of the license they choose before applying it. Licensors should also secure all rights necessary before applying our licenses so that the public can reuse the material as expected. Licensors should clearly mark any material not subject to the license. This includes other CC-licensed material, or material used under an exception or limitation to copyright. More considerations for licensors.

Considerations for the public: By using one of our public licenses, a licensor grants the public permission to use the licensed material under specified terms and conditions. If the licensors permission is not necessary for any reason—for example, because of any applicable exception or limitation to copyright—then that use is

not regulated by the license. Our licenses grant only permissions under copyright and certain other rights that a licensor has authority to grant. Use of the licensed material may still be restricted for other reasons, including because others have copyright or other rights in the material. A licensor may make special requests, such as asking that all changes be marked or described. Although not required by our licenses, you are encouraged to respect those requests where reasonable. More considerations for the public.

Creative Commons Attribution-ShareAlike 4.0 International Public License

By exercising the Licensed Rights (defined below), You accept and agree to be bound by the terms and conditions of this Creative Commons Attribution-ShareAlike 4.0 International Public License (“Public License”). To the extent this Public License may be interpreted as a contract, You are granted the Licensed Rights in consideration of Your acceptance of these terms and conditions, and the Licensor grants You such rights in consideration of benefits the Licensor receives from making the Licensed Material available under these terms and conditions.

Section 1 Definitions.

- a. **Adapted Material** means material subject to Copyright and Similar Rights that is derived from or based upon the Licensed Material and in which the Licensed Material is translated, altered, arranged, transformed, or otherwise modified in a manner requiring permission under the Copyright and Similar Rights held by the Licensor. For purposes of this Public License, where the Licensed Material is a musical work, performance, or sound recording, Adapted Material is always produced where the Licensed Material is synched in timed relation with a moving image.
- b. **Adapter’s License** means the license You apply to Your Copyright and Similar Rights in Your contributions to Adapted Material in accordance with the terms and conditions of this Public License.
- c. **BY-SA Compatible License** means a license listed at <http://creativecommons.org/compatiblelicenses>, approved by Creative Commons as essentially the equivalent of this Public License.
- d. **Copyright and Similar Rights** means copyright and/or similar rights closely related to copyright including, without limitation, performance, broadcast, sound recording, and Sui Generis Database Rights, without regard to how the rights are labeled or categorized. For purposes of this Public License, the rights specified in Section 2(b)(1)-(2) are not Copyright and Similar Rights.
- e. **Effective Technological Measures** means those measures that, in the absence of proper authority, may not be circumvented under laws fulfilling obligations under Article 11 of the WIPO Copyright Treaty adopted on December 20, 1996, and/or similar international agreements.
- f. **Exceptions and Limitations** means fair use, fair dealing, and/or any other exception or limitation to Copyright and Similar Rights that applies to Your use of the Licensed Material.

- g. **License Elements** means the license attributes listed in the name of a Creative Commons Public License. The License Elements of this Public License are Attribution and ShareAlike.
- h. **Licensed Material** means the artistic or literary work, database, or other material to which the Licensor applied this Public License.
- i. **Licensed Rights** means the rights granted to You subject to the terms and conditions of this Public License, which are limited to all Copyright and Similar Rights that apply to Your use of the Licensed Material and that the Licensor has authority to license.
- j. **Licensor** means the individual(s) or entity(ies) granting rights under this Public License.
- k. **Share** means to provide material to the public by any means or process that requires permission under the Licensed Rights, such as reproduction, public display, public performance, distribution, dissemination, communication, or importation, and to make material available to the public including in ways that members of the public may access the material from a place and at a time individually chosen by them.
- l. **Sui Generis Database Rights** means rights other than copyright resulting from Directive 96/9/EC of the European Parliament and of the Council of 11 March 1996 on the legal protection of databases, as amended and/or succeeded, as well as other essentially equivalent rights anywhere in the world.
- m. **You** means the individual or entity exercising the Licensed Rights under this Public License. **Your** has a corresponding meaning.

Section 2 Scope.

a. License grant.

1. Subject to the terms and conditions of this Public License, the Licensor hereby grants You a worldwide, royalty-free, non-sublicensable, non-exclusive, irrevocable license to exercise the Licensed Rights in the Licensed Material to:
 - A. reproduce and Share the Licensed Material, in whole or in part; and
 - B. produce, reproduce, and Share Adapted Material.
2. Exceptions and Limitations. For the avoidance of doubt, where Exceptions and Limitations apply to Your use, this Public License does not apply, and You do not need to comply with its terms and conditions.
3. Term. The term of this Public License is specified in Section 6(a).
4. Media and formats; technical modifications allowed. The Licensor authorizes You to exercise the Licensed Rights in all media and formats whether now known or hereafter created, and to make technical modifications necessary to do so. The Licensor waives and/or agrees not to assert any right or authority to forbid You from making technical modifications necessary to exercise the Licensed Rights, including technical modifications necessary to circumvent Effective Technological Measures. For purposes of this Public License, simply making modifications authorized by this Section 2(a)(4) never produces Adapted Material.

5. Downstream recipients.

- A. Offer from the Licensor Licensed Material. Every recipient of the Licensed Material automatically receives an offer from the Licensor to exercise the Licensed Rights under the terms and conditions of this Public License.
- B. Additional offer from the Licensor Adapted Material. Every recipient of Adapted Material from You automatically receives an offer from the Licensor to exercise the Licensed Rights in the Adapted Material under the conditions of the Adapters License You apply.
- C. No downstream restrictions. You may not offer or impose any additional or different terms or conditions on, or apply any Effective Technological Measures to, the Licensed Material if doing so restricts exercise of the Licensed Rights by any recipient of the Licensed Material.

6. No endorsement. Nothing in this Public License constitutes or may be construed as permission to assert or imply that You are, or that Your use of the Licensed Material is, connected with, or sponsored, endorsed, or granted official status by, the Licensor or others designated to receive attribution as provided in Section 3(a)(1)(A)(i).

Other rights.

- b. 1. Moral rights, such as the right of integrity, are not licensed under this Public License, nor are publicity, privacy, and/or other similar personality rights; however, to the extent possible, the Licensor waives and/or agrees not to assert any such rights held by the Licensor to the limited extent necessary to allow You to exercise the Licensed Rights, but not otherwise.
- 2. Patent and trademark rights are not licensed under this Public License.
- 3. To the extent possible, the Licensor waives any right to collect royalties from You for the exercise of the Licensed Rights, whether directly or through a collecting society under any voluntary or waivable statutory or compulsory licensing scheme. In all other cases the Licensor expressly reserves any right to collect such royalties.

Section 3 License Conditions.

Your exercise of the Licensed Rights is expressly made subject to the following conditions.

a. Attribution.

- 1. If You Share the Licensed Material (including in modified form), You must:
 - A. retain the following if it is supplied by the Licensor with the Licensed Material:
 - i. identification of the creator(s) of the Licensed Material and any others designated to receive attribution, in any reasonable manner requested by the Licensor (including by pseudonym if designated);
 - ii. a copyright notice;
 - iii. a notice that refers to this Public License;
 - iv. a notice that refers to the disclaimer of warranties;

- v. a URI or hyperlink to the Licensed Material to the extent reasonably practicable;
 - B. indicate if You modified the Licensed Material and retain an indication of any previous modifications; and
 - C. indicate the Licensed Material is licensed under this Public License, and include the text of, or the URI or hyperlink to, this Public License.
- 2. You may satisfy the conditions in Section 3(a)(1) in any reasonable manner based on the medium, means, and context in which You Share the Licensed Material. For example, it may be reasonable to satisfy the conditions by providing a URI or hyperlink to a resource that includes the required information.
- 3. If requested by the Licensor, You must remove any of the information required by Section 3(a)(1)(A) to the extent reasonably practicable.

b. ShareAlike.

In addition to the conditions in Section 3(a), if You Share Adapted Material You produce, the following conditions also apply.

- 1. The Adapters License You apply must be a Creative Commons license with the same License Elements, this version or later, or a BY-SA Compatible License.
- 2. You must include the text of, or the URI or hyperlink to, the Adapter's License You apply. You may satisfy this condition in any reasonable manner based on the medium, means, and context in which You Share Adapted Material.
- 3. You may not offer or impose any additional or different terms or conditions on, or apply any Effective Technological Measures to, Adapted Material that restrict exercise of the rights granted under the Adapter's License You apply.

Section 4 Sui Generis Database Rights.

Where the Licensed Rights include Sui Generis Database Rights that apply to Your use of the Licensed Material:

- a. for the avoidance of doubt, Section 2(a)(1) grants You the right to extract, reuse, reproduce, and Share all or a substantial portion of the contents of the database;
- b. if You include all or a substantial portion of the database contents in a database in which You have Sui Generis Database Rights, then the database in which You have Sui Generis Database Rights (but not its individual contents) is Adapted Material, including for purposes of Section 3(b); and
- c. You must comply with the conditions in Section 3(a) if You Share all or a substantial portion of the contents of the database.

For the avoidance of doubt, this Section 4 supplements and does not replace Your obligations under this Public License where the Licensed Rights include other Copyright and Similar Rights.

Section 5 Disclaimer of Warranties and Limitation of Liability.

- a. Unless otherwise separately undertaken by the Licensor, to the extent possible, the Licensor offers the Licensed Material as-is and as-available, and makes no representations or warranties of any kind concerning the Licensed Material, whether express, implied, statutory, or other. This includes, without limitation, warranties of title, merchantability, fitness for a particular purpose, non-infringement, absence of latent or other defects, accuracy, or the presence or absence of errors, whether or not known or discoverable. Where disclaimers of warranties are not allowed in full or in part, this disclaimer may not apply to You.
- b. To the extent possible, in no event will the Licensor be liable to You on any legal theory (including, without limitation, negligence) or otherwise for any direct, special, indirect, incidental, consequential, punitive, exemplary, or other losses, costs, expenses, or damages arising out of this Public License or use of the Licensed Material, even if the Licensor has been advised of the possibility of such losses, costs, expenses, or damages. Where a limitation of liability is not allowed in full or in part, this limitation may not apply to You.
- a. The disclaimer of warranties and limitation of liability provided above shall be interpreted in a manner that, to the extent possible, most closely approximates an absolute disclaimer and waiver of all liability.

Section 6 Term and Termination.

- a. This Public License applies for the term of the Copyright and Similar Rights licensed here. However, if You fail to comply with this Public License, then Your rights under this Public License terminate automatically.
- b. Where Your right to use the Licensed Material has terminated under Section 6(a), it reinstates:
 - 1. automatically as of the date the violation is cured, provided it is cured within 30 days of Your discovery of the violation; or
 - 2. upon express reinstatement by the Licensor.

For the avoidance of doubt, this Section 6(b) does not affect any right the Licensor may have to seek remedies for Your violations of this Public License.

- c. For the avoidance of doubt, the Licensor may also offer the Licensed Material under separate terms or conditions or stop distributing the Licensed Material at any time; however, doing so will not terminate this Public License.
- d. Sections 1, 5, 6, 7, and 8 survive termination of this Public License.

Section 7 Other Terms and Conditions.

- a. The Licensor shall not be bound by any additional or different terms or conditions communicated by You unless expressly agreed.

- b. Any arrangements, understandings, or agreements regarding the Licensed Material not stated herein are separate from and independent of the terms and conditions of this Public License.

Section 8 Interpretation.

- a. For the avoidance of doubt, this Public License does not, and shall not be interpreted to, reduce, limit, restrict, or impose conditions on any use of the Licensed Material that could lawfully be made without permission under this Public License.
- b. To the extent possible, if any provision of this Public License is deemed unenforceable, it shall be automatically reformed to the minimum extent necessary to make it enforceable. If the provision cannot be reformed, it shall be severed from this Public License without affecting the enforceability of the remaining terms and conditions.
- c. No term or condition of this Public License will be waived and no failure to comply consented to unless expressly agreed to by the Licensor.
- d. Nothing in this Public License constitutes or may be interpreted as a limitation upon, or waiver of, any privileges and immunities that apply to the Licensor or You, including from the legal processes of any jurisdiction or authority.

Creative Commons is not a party to its public licenses. Notwithstanding, Creative Commons may elect to apply one of its public licenses to material it publishes and in those instances will be considered the Licensor. The text of the Creative Commons public licenses is dedicated to the public domain under the CC0 Public Domain Dedication. Except for the limited purpose of indicating that material is shared under a Creative Commons public license or as otherwise permitted by the Creative Commons policies published at creativecommons.org/policies, Creative Commons does not authorize the use of the trademark Creative Commons or any other trademark or logo of Creative Commons without its prior written consent including, without limitation, in connection with any unauthorized modifications to any of its public licenses or any other arrangements, understandings, or agreements concerning use of licensed material. For the avoidance of doubt, this paragraph does not form part of the public licenses.

Creative Commons may be contacted at <http://creativecommons.org>.

